

TFPT — Red Team / Stress Test layer

Five adversarial targets (A–E): try to *break* the reductions, not confirm them

Stefan Hamann

Alessandro Rizzo

June 22, 2026 · v5.2 (rev 245)

The TFPT document set — what is treated where

Plain language: TFPT (Topological Fixed-Point Theory) is a small discrete compiler. Two inputs — the seam constant $c_3 = \frac{1}{8\pi}$ and the carrier rank $g_{\text{car}} = 5$ — build $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ and read off the Standard-Model skeleton, the constants and the scale grammar. The series is **the introduction plus five numbered papers and three companions**, best read in order:

0. **introduction** — reading guide, compiler closure, predictions, the dependency DAG and the single proof ledger (`verification/status_ledger.csv`). (*entry point*)
1. **tfpt_1_architecture_e8** — the two axioms, the derivation map, the EM fixed point α^{-1} , the $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ construction, the QBL theorem chain.
2. **tfpt_2_standard_model** — the SM in one φ_0 -ladder formula, flavor from parabolic transport (sheet diamond, dual anchor, Q_+ cohomology), the worked closures.
3. **tfpt_3_e8_audit_bootstrap** — the seven E_8 slices as an audit raster, the cascade bridge, the Möbius bootstrap.
4. **tfpt_4_frontier** — honest status of η_B , m_p/m_e , Koide, dark matter, full QG.
5. **tfpt_5_redteam** — the adversarial audit: declared attacks, what survives, what each reduces to, and the kill tests.

Companions: **R.** **tfpt_research_contracts** — the open interfaces $(v_{\text{geo}}, G_{\text{net}}, F_{\text{transfer}})$ as numbered contracts; **H.** **tfpt_horizon_readouts** — Appendix H (unit-system reframe): one seam constant as the universal horizon thermal code, the Nariai/anchor surface, the resummed seam clock; **O.** **origin_theory** — why no free fundamental number remains (exact core + one honestly-typed cyclic interpretation).

You are reading document 5 (Red Team).

TFPT status at a glance — read this card the same way in every document

Two inputs, one compiler. From the boundary constant $c_3 = \frac{1}{8\pi}$ (axiom P1) and the five-slot carrier $g_{\text{car}} = 5$ (axiom P2) — jointly the elementary symmetric data of one anchor $a = (1, 1, 2)$ — the discrete compiler $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ is *built*, and the Standard-Model skeleton (gauge group, three families, hypercharges, the flavor matrix), the dimensionless constants ($\alpha^{-1} = 137.0359992$) and the scale grammar are read off as fixed points. The algebraic core is machine-checked; physical readouts run through *named bridges* whose status is typed below.

What is closed, and what genuinely remains. The discrete/algebraic layer is closed ($[E]$); the everything-open residual reduces to *three named interface problems*:

Interface	Question	Status
v_{geo}	the one metrology unit ($= 1/\sqrt{G} = m/\mu$); No-Unit Thm: no compiler scale	primitive $[O]$
G_{net}	simple-current extension $(D_5)_1 \otimes (A_3)_1 \rtimes \langle (1, 1) \rangle \cong (E_8)_1$	$[E]/[C]$
F_{transfer}	one functor, four typed interfaces (Koide, η_B , axion, m_p/m_e)	$[C]$

The single open *structural theorem* is G_{net} (= SEAM.EQUIV.01: the raw seam *is* the holomorphic $(E_8)_1$ net), a constructive-AQFT theorem whose only open input is the continuum Osterwalder–Schrader reconstruction of the gapped quasi-free collar. QGEO.SYM.01 (the μ_4 deck acting geometrically) is now its *corollary*: a conformal net’s vacuum is rotation-invariant by axiom, so the deck follows with no extra premise (v335, Lean FORM.QGEO.BW.01). The full nonperturbative quantum-gravity measure QG.AMB.01 is the *general* Euclidean-QG conformal-factor problem (Gibbons–Hawking–Perry 1978), from which TFPT is rigorously *gap-decoupled* (margin $1.648 > 0$, v76/v330) — an inherited, decoupled problem, *not* a TFPT-specific open item. So the honest residual is *one* solvable theorem (G_{net}) plus the typed v_{geo} and F_{transfer} interfaces. Everything carries a claim ID resolving to a row of **verification/status_ledger.csv** (the single source of truth); **if the text and the ledger ever disagree, the ledger wins**. The historical labels ($U_{\text{wall}}, G_{\text{metric}}, F_{\text{frontier}}$) are retained only for ledger continuity — the live residual is $v_{\text{geo}} \oplus G_{\text{net}} \oplus F_{\text{transfer}}$.


Marker key: $[E]$ exact / machine-proven · $[C]$ conditional (holds under named hypotheses) · $[O]$ open / axiom · $[X]$ falsifiable kill test. Per-claim fine type in **verification/status_ledger.csv**.

$[E]$ exact / machine-proven

$[C]$ conditional

$[O]$ open / axiom

$[X]$ kill test



What this note is

This is the deliberately *adversarial* layer requested in the v78 review. Its purpose is **not** to confirm TFPT but to attack the five load-bearing reductions at their weakest logical transitions. Each target is run through one fixed protocol — minimal statement, assumptions, logical chain, validity conditions, counterexample search, limiting/degenerate cases, alternative structures, provisional verdict — and produces a deliberately narrow output: where the statement works, where it could fail, and the residual risk. The layer

is machine-backed by `[verification/redteam/rt_A_e8net.py]` ... `[verification/redteam/rt_E_vgeo.py]`, `[verification/redteam/rt_F_qft4d.py]` and `[verification/redteam/run_redteam.py]`; a red-team check asserts an *adversarial* fact (a counterexample really exists, a hidden assumption is really needed, a firewall really holds), and the honest outcome lives in the **status** of each target, never in a green pass.

Target	Minimal claim	Status	Residual risk
A	seam–Calderón measure = $(E8)_1$ net	reduced, not closed	<i>one</i> statement: boundary-net holomorphy + $c=8$ ($\Leftrightarrow$ the index-4 inclusion); E_8 and bulk uniqueness then automatic (v83/v87/v89; Lie-level realisation v143). Since reduced to <i>zero new</i> open content: the free-bulk premise is a fixed-point theorem, the cone is eliminated (cone gap = one-particle gap), and the residual factors into the already-open A_2 (net assembly) + GATE.QGEO (v160–v165). The three-residual form below is the historical development.
B	$g_{\text{car}} = 5$ forced by the Pascal rule	survives (narrowed)	boundary derivation of half-spinor exhaustion (degree-2 truncation)
C	$k = c_3/2 = \frac{1}{16\pi}$, $S/A = \frac{1}{4}$	survives	absolute $1/G$ (UV-sensitive) is the one anchor
D	ratios+products $\Rightarrow$ one scale v_{geo}	survives (narrowed)	CP phases ($\delta_{\text{CKM/PMNS}}$) not covered by v_{geo}
E	one scale v_{geo} carries the theory	survives (narrowed)	EW/reheating/leptogenesis scales + seam=Planck identification

Method

The red team treats each reduction as a hostile witness. A confirmatory script that always passes is worthless here: the layer is built so that it *may* downgrade a claim. Three outcomes are allowed — **survives** (stands as worded), **survives (narrowed)** (stands only after a silent assumption is made explicit), and **reduced, not closed** (the conservative wording is the correct one). A fourth, **broken**, is reserved for an actual failure; none occurred.

Target A — the $(E8)_1$ boundary-net identification

Minimal claim. The ambient metric/projective measure is reduced to identifying the seam–Calderón boundary measure with the $(E8)_1$ lattice net, carrier subnet $(D5)_1 \times (A3)_1$. `[verification/redteam/rt_A_e8net.py]`

Logical chain (established). The Sugawara level-1 central charges are $c(E8)_1 = \frac{248}{31} = 8$, $c(D5)_1 = 5 = g_{\text{car}}$, $c(A3)_1 = 3 = N_{\text{fam}}$, so the conformal-embedding criterion $c_{\text{coset}} = c(E8) - c(D5) - c(A3) = 0$ holds [E]. This is *compatibility*.

Where it breaks: central charge underdetermines the net

Counterexample: $(D8)_1 = SO(16)_1$ also has $c = \frac{120}{15} = 8$, but is *not* holomorphic (four primaries 1, v , s , c) — a distinct $c = 8$ net. Hence equal central charge is *necessary, not sufficient*: $(E8)_1$ is the unique *holomorphic* $c = 8$ net (even unimodular rank-8 lattice = $E8$), so **holomorphy (single vacuum sector, μ -index 1) is the load-bearing extra assumption**. Dropping chirality is worse: the $c = 8$ Narain

family is positive-dimensional. The constructive map seam–Calderón kernel $\rightarrow (E8)_1$ and the uniqueness of bulk reconstruction are not established; the gap $\Delta_{\text{eff}} > 0$ *supports* tightness (clustering) but does not prove it.

Verdict. Keep “reduced to a rigorous boundary-net identification problem”; do *not* write “metric sector closed”. The red team confirms the conservative wording. Status [C]. Residual: holomorphy proof + constructive boundary map + bulk-reconstruction uniqueness.

Target B — carrier rank / the Pascal condition

Minimal formal claim. $2^g = g^2 + g + 2$ has the unique solution $g = 5$ (Lean-verified [E]). This arithmetic is *not* attacked. The attack is upstream: why must the carrier obey the Pascal half-spinor exhaustion $2^{g-1} = \binom{g}{0} + \binom{g}{1} + \binom{g}{2}$? [verification/redteam/rt_B_pascal.py]

Where it narrows: the rule, not the arithmetic, selects 5

Perturbing the constant breaks it: $2^g = g^2 + g + c$ gives $\{5\}$ only at $c = 2$; neighbours give other or empty solution sets. The truncation *degree* is the real choice: with $2^{g-1} = \sum_{k \leq K} \binom{g}{k}$ one finds $K=0 \rightarrow g=1$, $K=1 \rightarrow 3$, $K=2 \rightarrow 5$, $K=3 \rightarrow 7$, i.e. $g = 2K + 1$. So choosing $K = 2$ (to reach $g_{\text{car}} = 5$) is a physical postulate. It is corroborated — the $D5$ half-spinor $\dim S^+ = 2^{g-1} = 16 \Leftrightarrow g = 5$ and the family closure $N_{\text{fam}} = (2^{g-1} - 1)/g$, $g + N_{\text{fam}} = 8$ both pick 5 — but each is itself a postulate, not a theorem.

Verdict. Theorem A’s arithmetic core stands; the Pascal *selection* must be typed [O]/[C], not [E]. Residual: a boundary/seam derivation of half-spinor exhaustion.

Target C — $k = c_3/2$ and the seam-area coefficient

Minimal claim. In reduced units $k = c_3/2 = \frac{1}{16\pi}$ and Fursaev–Solodukhin gives $S/A = \frac{1}{4}$. [verification/redteam/rt_C_kc3.py]

The dimensional firewall holds

The only legal chain is

$$k_{\text{red}} = \frac{c_3}{2} = \frac{1}{16\pi} \text{ (dimensionless)}, \quad k_{\text{phys}} = \frac{k_{\text{red}}}{G}, \quad S = 4\pi k_{\text{phys}} A_{\text{phys}} = \frac{A_{\text{phys}}}{4G}.$$

The forbidden form $k_{\text{phys}} = c_3/2$ is dimensionally false (legal only at $G = 1$). A scan of v67/v68/v73 finds every $c_3/2$ used as the *reduced* coefficient and **zero** naked dimensionful identities; the firewall’s teeth are verified against a synthetic violation. Internally consistent with $1/G$ being UV-sensitive (Sakharov/Connes, v68).

Verdict. Safe as worded (reduced coefficient). Status [E] (structure) + [O] for the residual. Residual: the absolute $4D$ Newton scale $1/G$ ($\propto \Lambda^2 f_2$) stays the one dimensionful anchor; nothing *dimensionless* is open.

Target D — $U_{\text{point}} \rightarrow v_{\text{geo}}$

Minimal claim. Ratios plus sector products determine the dimensionless amplitudes, leaving one common scale v_{geo} . [verification/redteam/rt_D_upoint.py]

Where it narrows: the bijection needs explicit hypotheses

For positive reals, (ratios, product) $\Leftrightarrow$ (individuals) is a bijection, and the TFPT sectors satisfy it (positive rational c , odd size $3 = N_{\text{fam}}$). But it *fails* off-assumption: even sector size leaves a sign ambiguity ($\{2, 3\}$ vs $\{-2, -3\}$ share ratio $3/2$ and product 6); complex amplitudes leave an n -th-root branch ambiguity ($\{1, 1, 1\}$, $\{\omega, \omega, \omega\}$, $\{\omega^2, \omega^2, \omega^2\}$ share all ratios and product 1); a zero amplitude is degenerate. **Genuine residual:** the bijection fixes amplitude *magnitudes* only — CP phases $\delta_{\text{CKM}}, \delta_{\text{PMNS}}$ are complex and

are *not* folded into v_{geo} .

Verdict. Holds for TFPT once four hypotheses (real, positive, fixed branch, no zeros / odd sector size) are made explicit; they are silent in v75 and must be named. Status **[E]** (reduction) + **[O]** (v_{geo}). Residual: the CP-phase sector.

Target E — v_{geo} as the unique dimensional floor

Minimal claim. No pure-number theory can derive an absolute dimensionful scale; all scales are ratios to one unit v_{geo} . [\[verification/redteam/rt_E_vgeo.py\]](#)

Single-scale audit: exact for the certified tiers, conditional beyond

The closed compiler + protected IR readouts all reduce to (pure number) $\times v_{\text{geo}}$: $M_{\text{scal}} = c_3^{7/2} \bar{M}_{\text{Pl}}$, $f_a = (c_3/|\mu_4|)^{7/2} \bar{M}_{\text{Pl}}$, masses $= (\pi/\sqrt{2}) c_f \varphi_0^k v_{\text{geo}}$, and $1/G$ shares the same anchor (v68/v75). The audit *surfaces* candidates that are not pure reductions: the seam cutoff Λ (only an *identification* seam=Planck collapses it to v_{geo}), the electroweak matching scale, the reheating temperature, and the leptogenesis scale — all frontier **[C]**/**[O]** inputs that must not be silently absorbed into v_{geo} .

Verdict. v_{geo} is the unique dimensional floor of the certified tiers (true by dimensional analysis). Full single-scale-ness is conditional on the flagged identifications staying honestly typed. Status **[E]** (ratios) + **[O]** (one scale). Residual: explicit reduction (or honest frontier typing) of the EW / reheating / leptogenesis scales + the seam=Planck cutoff.

Target F — the perturbative 4D-QFT + scale round (v269–v275)

Minimal claim. The round published in release 5.2 — the Epstein–Glaser S_{pert} layer (v269/v271/v273), the PMNS Jarlskog CP strength (v270), the absolute ν -scale typing (v272), the scale over-determination (v274) and the QG.AMB.01 roadmap (v275). [\[verification/redteam/rt_F_qft4d.py\]](#)

Two attacks land; the round survives once they are named

(1) **The gravity S_{pert} is non-unitary.** Power-counting renormalizability is *not* unitarity: the R^2/Weyl^2 sector gives a four-derivative propagator $1/(p^2(p^2 + M^2))$ with a negative-residue massive spin-2 mode (the Stelle ghost). So S_{pert} is unitary for the matter+gauge sector only; the gravity perturbative S-matrix carries the ghost — which is exactly why the *nonperturbative* QG.AMB.01 is a *decoupled general* QG problem (the Euclidean-QG conformal-factor problem; gap-decoupled from every TFPT readout, not a TFPT structural item, v335), not a new gap. (2) **The over-determination is conditional.** v274's Route-2 inverts the $\rho_\Lambda/\bar{M}_{\text{Pl}}^4 = (3/4\pi^2)e^{-2\alpha^{-1}}$ prediction, which is itself LAMBDA.BRANCH.01 **[C]**; so the 0.11% gravity-vs-cosmology agreement is conditional on the Λ -branch (the gravity route is unconditional). **Already-honest:** J_{PMNS} inherits δ 's **[C]** provenance (assembled μ_6 phase, not an M_ν output) and the ν -scale is **[O]** with the NO floor $\Sigma m_\nu = 0.0586$ eV as the one kill test.

Verdict. SURVIVES (NARROWED): the round stands as typed; the two narrowings above are now folded into v269 and v274. Residual: the perturbative gravity S-matrix is non-unitary (the Stelle ghost) — this is the *decoupled general* QG.AMB.01 problem (gap-decoupled from every TFPT readout, not a TFPT structural item, v335); and the anchor over-determination is conditional on the Λ -branch.

Under examination: an infinite-derivative narrowing ([C], does *not* close the gate). The Stelle ghost is a feature of the *local* $R+R^2/\text{Weyl}^2$ *truncation*: the four-derivative propagator $1/(p^2(p^2 + M^2)) = (1/M^2)[1/p^2 - 1/(p^2 + M^2)]$ carries the negative-residue spin-2 mode. But that truncation is the a_4 Seeley–DeWitt order of the spectral action $\text{Tr } f(D^2/\Lambda^2)$, whose cutoff is *not* a free scheme — it is the $\beta=1$ seam KMS weight $f(u) = e^{-u}$ (v259, no new parameter). Keeping that exponential cutoff *un-truncated* points to a nonlocal, infinite-derivative graviton form factor $p^2 a(p^2)$; for an *entire* factor $a(p^2) = e^{p^2/M^2}$ (nowhere zero) the dressed propagator

$1/(p^2 a)$ keeps its *only* pole at $p^2=0$ — no ghost (Tomboulis 1997; Biswas–Mazumdar–Siegel 2006) — whereas any *polynomial* (finite-derivative) truncation has a real zero = the ghost pole. So *if* the resummed form factor is entire, the perturbative ghost is a truncation artefact: a conditional [C] narrowing of this attack. It does *not* close QG.AMB.01 — the trace regulator f is not yet shown to resum to an entire kinetic form factor $a(\square)$, and a ghost-free *perturbative* graviton is still not the nonperturbative ambient measure. (QGAMB.IDG.01, [C]/[O]; consistent with the SEAM.EQUIV.01 firewall.) [verification/v304_idg_nonlocal_ghost.py]

Outcome

What the red team changed. No target broke. The hardest gate (A) is confirmed *open* with the conservative wording; four targets (B, D, E, F) *survive once a silent assumption is named*; one (C) survives as worded. The net effect is not new physics but sharper typing: the package now states *where* each reduction would fail and *which* assumptions are truly necessary, exactly at the weakest logical transitions. Reproduce with `cd verification/redteam && python run_redteam.py`.

Follow-up — the verdicts were then used as a worklist [E]
([verification/v83_e8net_holomorphic_uniqueness.py])

Two residuals were attacked directly after the red-team pass; the result is recorded in ledger rows GATE.METRIC.04 and CAR.PASCAL.01.

Target A, residual 1 — closed [E]. At level 1 the number of primaries equals $\det(\text{Cartan}) = |\text{fundamental group}|$: $A_3=4$, $D_5=4$, $D_8=4$, $\mathbf{E}_8=1$. So holomorphy (single vacuum sector, μ -index 1) is *necessary* — it excludes the same- c rival $(D_8)_1 = SO(16)_1$ ($c = 120/15 = 8$ but four primaries $1, v, s, c$) — *and sufficient*: a holomorphic $c = 8$ chiral CFT is the lattice theory of an even unimodular rank-8 lattice, of which there is exactly one, E_8 , by the Minkowski–Siegel mass formula (mass = $1/|W(E_8)| = 1/696729600$, and $|\text{Aut}(E_8)|$ saturates it). The $(D_5)_1 \times (A_3)_1$ embedding $c = 5 + 3 = 8$ stays *compatibility* (16 primaries vs 1). **Consequence:** the constructive map need only show the seam–Calderón boundary net is holomorphic with $c = 8$ (then E_8 is automatic), so Target A drops from *three* residuals to *two*: (i) boundary-net holomorphy + $c = 8$ ($\Delta_{\text{eff}} > 0$ supports, not proves), and (ii) bulk-reconstruction uniqueness — both still [C]/[O].

Target B — reduced [E]. The “ $K = 2$ Pascal truncation” is not free: $K = (g - 1)/2$ is the Pascal-row midpoint $\sum_{k \leq (g-1)/2} \binom{g}{k} = 2^{g-1}$ (odd g), i.e. the half-spinor split; the carrier is the even Clifford half-spinor $\Lambda^{\text{even}}(\mathbb{C}^5) = 1 + 10 + 5 = 16 = \dim S^+$. So B’s residual reduces to the single standard input “carrier = half-spinor of $\text{Spin}(10)$ ” (the Lean arithmetic core stays [E]).

Target B — a backward certificate ([verification/v219_icosahedral_mckay.py]). The choice $(|\mathbb{Z}_2|, N_{\text{fam}}, g_{\text{car}}) = (2, 3, 5)$ that B attacks is, *given* the E_8 closure, the icosahedron: by the McKay correspondence E_8 is the exceptional top of the finite- $SU(2)$ tower ($2T \rightarrow \hat{E}_6$, $2O \rightarrow \hat{E}_7$, $2I \rightarrow \hat{E}_8$), the binary icosahedral group $2I$ (order $120 = |R^+(E_8)|$) has irrep degrees $\{1, 2, 2, 3, 3, 4, 4, 5, 6\}$ equal to the affine- E_8 Kac marks ($\sum = 30 = h(E_8)$), and it contains the 2, 3, 5 rotation-axis orders. This does *not* break B’s verdict — it is a *backward* certificate (it presupposes E_8), so $g_{\text{car}} = 5$ stays the P2 interface [O]; what it removes is the impression that $(2, 3, 5)$ is an arbitrary triple, since no spherical McKay node sits above the icosahedron.

Still open, typed honestly: A (i)+(ii); Target D’s CP-phase residual (at $\theta_{13} = 0$ the Dirac phase is undefined, so this needs a texture correction); Target E’s reheating/leptogenesis scales and the seam=Planck identification. Target C is a dimensional *floor*, not a gap. **No target was promoted.**

Second follow-up round [E]
([verification/v87_bulk_uniqueness_reduction.py], [verification/v88_cp_phase_audit.py], [verification/v86_nstar_reheating.py])

Target A, residual (ii) — conditionally closed [E]: Target A = ONE residual. Bulk-reconstruction uniqueness is *not* independent of residual (i). By the algebraic-QFT classification of full 2D CFTs over a chiral net (Longo–Rehren; Kawahigashi–Longo–Müger; Bischoff–Kawahigashi–Longo–Rehren), the possible bulks are the haploid commutative Q-systems $\simeq$ physical modular invariants of $\text{Rep}(A)$. For a *holomorphic* net $\text{Rep}(A) = \text{Vect}$, so the bulk pairing is *unique*. The contrast is machine-checked: the same- c rival $SO(16)_1$ (discriminant category $\mathbb{Z}_2 \times \mathbb{Z}_2$, μ -index 4) admits *six*

modular invariants — diagonal, $s \leftrightarrow c$ swap, both E_8 -extension invariants $|\chi_0 + \chi_{s/c}|^2$ (the lattice glue $E_8 = D_8 \cup (D_8 + s)$, $h_s=1$ bosonic), and two twisted pairings — so without holomorphy the $c=8$ bulk is *multiply* ambiguous; with it, trivially unique. Hence the entire Target A reduces to the single theorem “seam–Calderón boundary net is holomorphic with $c=8$ ” ([C]/[O]; ledger GATE.METRIC.05). The theorem also has an equivalent, more tractable *index* form ([verification/v89_carrier_index_lemma.py], ledger GATE.METRIC.06): by Kawahigashi–Longo–Müger, $\mu_A = [B:A]^2 \mu_B$, the carrier inclusion has Jones index $[(E_8)_1 : (D_5)_1 \times (A_3)_1] = \sqrt{16} = 4 = |\mu_4|$ — the glue-group order *is* the inclusion index — and the extension is the μ_4 simple-current extension (all three glue sectors are $h=1$ currents; $248 = 45+15+64+64+60$). Since an index-4 simple-current extension of the carrier has $\mu = 16/4^2 = 1$, **holomorphy follows from the index statement**: Target A $\Leftrightarrow$ “the seam net contains the carrier net with Jones index 4 as its μ_4 simple-current extension” — both ends constructed objects, the index controllable by Longo theory, the gap supporting finiteness. And *which* extension carries no freedom ([verification/v92_glue_uniqueness.py], ledger GATE.METRIC.07): exhaustive classification of the carrier discriminant form $(\mathbb{Z}_4 \times \mathbb{Z}_4, q=(5x^2+3y^2)/8)$ shows the extension tower is *completely rigid* — exactly two Lagrangian glues (the two chiralities, identified by the sheet $\mathbb{Z}_2$; the same two objects as the E_8 -extension invariants above), and exactly one halfway extension, whose induced form $q = \{0, 0, 0, \frac{1}{2}\}$ is the D_8 discriminant form: the $SO(16)_1$ rival is the unique intermediate of the tower, not an arbitrary counter-model. Tower: carrier ($\mu=16$) $\rightarrow SO(16)_1$ ($\mu=4$) $\rightarrow (E_8)_1$ ($\mu=1$, sheet pair) — nothing else exists. Target A is thereby the *bare* index statement.

Target A, final form — the Simple-Current Extension Theorem [E]/[C] ([verification/v154_simple_current_theorem.py]). Assembled into one statement: the carrier net $A = (D_5)_1 \otimes (A_3)_1$ extended by the isotropic glue $L = \langle(1,1)\rangle$ has $[B:A]=|L|=4$, $c=5+3=8$, $\mu(B)=\mu(A)/|L|^2=16/16=1 \Rightarrow$ holomorphic $\Rightarrow B \cong (E_8)_1$. The *entire* adversarial residual of Target A is now the single seam-side identification “the RP seam–Calderón net *is* A ” — the algebra $(B \cong (E_8)_1)$ is exact and machine-checked (v89/v92/v125/v143/v148). That residual is now sharply located: the analytic core (“the kernel is quasi-free”) *follows* from a free RP bulk — the boundary marginal of a Gaussian bulk is the compression of a contraction (v155), the *same* premise the area law and the EH mechanism use — and the net is *explicitly constructed and verified*: DtN = $|k|$, 16 free Majoranas, currents $248=120+128$, holomorphic character $E_4/\eta^8 = j^{1/3}$ with 248 at level 1 (v156). *Final reduction* (v160–v165): the free-bulk premise is recast as a fixed-point theorem (quasi-free $\Rightarrow$ all higher cumulants $\kappa_{2n}=0$, v160); the infinite Schwinger cone is *eliminated* — by Bogoliubov second quantisation the cone gap equals the one-particle gap $(2/3)^6$ (v161), whose value is forced by μ_4 -equivariance (v162); and the remaining 16-dim identification factors into the two *already-open* items — A_2 net existence (an assembly of established net constructions, [C]) and $R_1 = \text{GATE.QGEO}$ (“seam boundary = $\mathbb{P}^1 \setminus \mu_4$ ”) — with the irreducible core $\{\pi, v_{\text{geo}}\}$ a theorem (v165). So Target A carries *no new independent open content*: it is a unification, not an unconditional proof. The two structural residuals are then discharged to a theorem (v175): the CAR second-quantisation functor $\Gamma(t) = \bigoplus_m \Lambda^m(t)$ makes full-cone reflection positivity reduce *for every* m to the one-particle contraction (verified on the complete 2^{16} -dim Fock space), and A_2 is an assembled, verified $(E_8)_1$ certificate; net existence and full-cone RP become [E]. Target A’s single residual is therefore the seam realisation, now the named keystone SEAM.EQUIV.01 — *the raw RP seam is the holomorphic $(E_8)_1$ net at $\tau=i$* — whose conformal-deck premise QGEO.SYM.01 is its *downstream* definitional face, Lean-pinned (FORM.SEAMEQUIV.01, v308) to named cited steps plus the derived Recovery gap. The sprint-by-sprint reduction (v176–v302) is recorded in the changelog, not replayed here.

Target D — the CP-phase residual quantified [E] (not closed). With the frozen leading reading $\delta = \pi/3 + 3\lambda_C^2 = 68.65^\circ$: pull vs $\gamma_{\text{PDG}} = 65.7^\circ \pm 3.0^\circ$ is $+0.98\sigma$ (survives). Audit, *not* promoted: the data central value sits 0.07° from the alternative coefficient reading $\pi/3 + 2\lambda_C^2$ ($|\mathbb{Z}_2|$ instead of N_{fam}) — a textbook look-elsewhere trap; the registry keeps coefficient 3 (REG.FREEZE.01). Decision threshold: the readings differ by $\lambda_C^2 = 2.88^\circ$, i.e. 3σ -distinguishable once $\sigma_\gamma \leq 0.96^\circ$. The J -inversion is flagged magnitude-contaminated (source-scale s_{23} vs M_Z). Ledger FLAV.CP.01.

Target E — one flagged scale computed [C]. The reheating temperature is no longer a silent input: $M_{\text{scal}} = c_3^{7/2} \bar{M}_{\text{Pl}}$ [E] + standard physics give $T_{\text{reh}} = 9.6 \times 10^9$ GeV and $N_*(0.05/\text{Mpc}) = 51.4$ [C] (ledger COSMO.NSTAR.01); honestly recorded: this slow Higgs-channel point is A_s -disfavoured (-11.4σ ; the measured A_s requires near-instantaneous reheating), so the point is conditional on the decay channel and the frozen band [50, 60] stays the surface of record; the leptogenesis scale and the seam=Planck identification remain typed inputs. **Again: no target was promoted; one residual was merged (A), one was quantified (D), one input was computed (E).**

Target D — the CP residual geometrically *located* (not closed) [E]/[C]

The follow-up round only *quantified* the CP residual; this round *locates* it geometrically, in three consistent readings, without promoting it. **(i) A channel typing** ([`verification/v227_degree_exponent_channel_split.py`]). The E_8 split $248 = 120 + 128$ is a magnitude/phase split: the exponents sum to $120 = |R^+(E_8)|$ (the *magnitude* channel, where the (ratios, product) $\Rightarrow v_{\text{geo}}$ bijection lives), the fundamental degrees sum to $128 = \text{rank } E_8 \cdot \dim S^+$ (the *phase/glue* channel). The un-covered CP phases of Target D are exactly the 128 phase channel — a clean home, not a new gate; whether 128 also hosts a dark sector is Frontier, *not* asserted (this replaces the over-strong “ S^- is dark matter” reading). **(ii) A hexagonal modulus** ([`verification/v220_cp_hexagonal_modulus.py`]). The seam deck is the *square* modulus ($j=1728, \mathbb{Z}/4$ clock; the kill-test selects it), which is real and carries no phase; its exceptional partner is the *hexagonal* modulus $j=0$ ($\tau=\rho=e^{i\pi/3}, \mathbb{Z}/6$, CM by $\mathbb{Z}[\omega]$). The frozen leading reading $\delta = \pi/3 + 3\lambda_C^2 = 68.65^\circ$ has leading term $\pi/3 = \arg \rho$, the hexagonal CM phase; CP lives in the $\mathbb{Z}/6$ phase *fiber* over the $\mathbb{Z}/4$ seam, not in the deck. **(iii) An orientation** ([`verification/v225_dual_normal_frame.py`]). The dual normal frame $d = a^\top R^{-1} = -\frac{1}{2}(1, 1, -2)$ (Nariai/horizon normal) and $n = (5, -9, 6)$ (first-generation selector, $n^\top R = (8, 0, 0)$) has oriented volume $\det(\mathbf{1}, d, n) = 21 = N_{\text{fam}} \cdot 7$, and rotating by the hexagonal unit gives $\text{Im} \det(\mathbf{1}, d, \rho n) = 21 \sin(\pi/3) \neq 0$ — CP sits precisely in the *orientation* of the normal frame, exactly where the magnitude bijection fails. All three are [E] geometry; the physical CP *derivation* stays the Target-D [C] residual (the seam deck is not re-selected, the registry δ is untouched). **(iv) One hexagonal unit, two sheets** ([`verification/v231_cp_mu6_phases.py`]). The reduction is sharper still: *both* CP phases are μ_6 powers of the single hexagonal CM unit $\rho = e^{i\pi/3}$ — $\delta_{\text{CKM}}^{\text{lead}} = \arg(\rho^1) = \pi/3$ (quark) and $\delta_{\text{PMNS}} = \arg(\rho^4) = 4\pi/3$ (lepton, the neutrino phase lattice) — with $\rho^4 = -\rho$, so $\delta_{\text{PMNS}} = \delta_{\text{CKM}}^{\text{lead}} + \pi$ is the quark phase times the $\mathbb{Z}_2$ sheet ($\rho^3 = -1$); the frame orientations of (iii) are correspondingly sheet-flipped ($\pm 21 \sin(\pi/3)$). So Target D’s CP sector is *not* two independent phase inputs but *one* hexagonal unit read on the two sheets — the magnitude residual $3\lambda_C^2$ (quark, v88) and the physical derivation stay [C], but one of the two phase inputs is removed. **(v) The universal triality phase** ([`verification/v233_cp_triality_phase.py`]). In the canonical $\mu_6 = \mu_3(\text{family/triality}) \times \mu_2(\text{sheet})$ factorisation, $\rho = \omega^2(-1)$ with $\omega = e^{2\pi i/3}$ the $\mathbb{Z}_3$ triality centre (the $2/3$ cusp weight); since $\rho^4 = \rho^1 \rho^3$ with $\rho^3 \in \mu_2$, the quark (ρ^1) and lepton (ρ^4) phases have the *same* family/triality class and differ *only* by the sheet. So the CP phase *is* the universal triality phase, sheet-split — not even a free choice of μ_6 power: the power choice of (iv) is removed, and ω is the μ_3 factor of the order-30 (2, 3, 5) Coxeter/Milnor monodromy that builds the hull (v232). The physical CP derivation still stays [C]. **(vi) The lock as a pre-registered kill test** ([`verification/v320_galois_cp_relation.py`]). Because $\delta_{\text{PMNS}} = \delta_{\text{CKM}}^{\text{lead}} + \pi$ is now *forced* (one unit $\rho = \zeta_6$, the μ_4 -power ρ^4 , the $\mathbb{Z}_2$ sheet $\rho^3 = -1$), it is a falsifiable cross-prediction, not a free input: a measured δ_{PMNS} robustly away from 240° at DUNE/Hyper-K (beyond the sub-leading budget, $> 3\sigma$) kills the Galois-CP organisation. The relation is [E], the physical phase reading stays [C], the kill test is [X].

Global look-elsewhere closure [E] ([`verification/v100_numerology_null_mc.py`])

The per-target look-elsewhere traps recorded above (the $\pi/3 + 2\lambda_C^2$ alternative in Target D, the $\ln(46080)$ scale reading destroyed in [`verification/v99_koide_flow_time.py`]) are local instances of the global adversarial question: *could the whole scorecard be formula-fishing?* The numerology null test answers it with an explicit, declared null model run in adversarial mode: the *entire* complexity-matched formula grammar (which provably contains every scored TFPT formula — the null space includes the theory) is enumerated against conservative data windows. Joint formula-fishing probability $\prod_i p_i = 10^{-25.8}$ over the 13 scored observables; all 94 500 $F_{U(1)}$ variants solved, exactly one hits CODATA ($p_\alpha \leq 1.1 \cdot 10^{-5}$); total $\leq 10^{-30.7}$. The test is falsifiable in both directions: the negative controls (seed perturbation 1%/10% collapses TFPT’s own hit count $13 \rightarrow 9/6$; data shuffling $\rightarrow 1.2$) prove it has power, and the census is stable under budget/window variation (< 2 orders). Honest limit, stated as such everywhere: this is a null-model rejection *conditional on the declared grammar* — no statistical test certifies a theory; the tension observables keep their large windows and contribute nothing to the number. Details and the grammar definition: the audit note (tfpt_3) and the module docstring.

Structural firewall + out-of-sample seed cross-validation [E]
([verification/v305_witness_independence.py], [verification/v306_seed_crossval.py])
The global null test answers “is the <i>scorecard</i> fishing?”; two companions answer the deeper adversarial worries “is the same integer counted many times as if independent?” and “is one shared seed dressed up as many predictions?”. Witness economy ([verification/v305_witness_independence.py]): every headline integer ($\mu_4 =4$, $g_{\text{car}}=5$, $\mathbb{Z}_2 =2$, $N_{\text{fam}}=3$, $\dim S^+=16$, $\text{rank } E_8=8$, 240, 248, $2I =120$, $h=30$, $W(D_5) =1920$, $\Omega_{\text{adm}}=48$, the gap exponent 6) is a documented image of the <i>single</i> anchor $a=(1,1,2)$ with π the only transcendental primitive, so the recurring $2/3= \mathbb{Z}_2 /N_{\text{fam}}$ and $(2/3)^6$ are <i>one</i> ratio, not many — 13 atoms from 2 free primitives ($6.5\times$ compression made explicit). The E_8 spine is genuinely overdetermined (≥ 3 structurally distinct derivations sharing only leaf inputs), while the pure-seed readouts are single-witness, so their joint surprise is scored jointly by v100, never multiplied; the negative control shows α^{-1}’s “137” is <i>foreign</i> to the anchor family (an independent witness via the $F_{U(1)}$ cubic, v3). Seed cross-validation ([verification/v306_seed_crossval.py]): treating the axiom seed φ_0 adversarially as a free dial and back-solving it from each of the five seed-grammar observables gives an error-weighted mean reproducing $\varphi_0 = 1/(6\pi) + 48c_3^4$ to 0.04%; leave-one-out predicts every held-out observable within 3σ <i>out of sample</i> (the most-tensioned hold-out is $\sin^2 \theta_{13}$, $\sim 2\sigma$), and a data$\leftrightarrow$formula shuffle explodes the cross-validation χ^2 by $> 100\times$ (the test has power). Both stay [E] as structural/statistical checks, conditional on the same declared grammar. No-golden-dynamics rule ([verification/v305_witness_independence.py]): the golden ratio $\phi = 2\cos(\pi/5)$ lives in the <i>static</i> carrier field $\mathbb{Q}(\sqrt{5})$ (the $\mathbb{Z}/5$ hand; minimal polynomial $x^2 - x - 1$, discriminant 5 non-square), while every <i>dynamic</i> rate is rational ($\mathbb{Q}$, the $\mathbb{Z}/6$ = sheet$\times$family hand, e.g. the recovery factor $(2/3)^6$) — so a dynamic rate “explained” by ϕ is a cross-field false friend, and CP phases and recovery rates must run on the order-6 hand (the field firewall of v314/v315).




Figure 1: The seed-hyperplane test ([verification/v306_seed_crossval.py]): each scorecard observable is *independently* inverted to the latent seed φ_0 (with its measurement error). All five project onto the one axiom value $\varphi_0 = \frac{1}{6\pi} + 48c_3^4 = 0.053172$ (green), and their error-weighted mean reproduces it to 0.04% — one real seed, not five fits. The honest pressure point is $\sin^2 \theta_{13}$ ($\sim 2\sigma$ low, a *separate* carrier-trace channel, **v268**); a data $\leftrightarrow$ formula shuffle explodes the cross-validation χ^2 by $> 100\times$. [E].

External-review residual map [E]/[C] ([verification/v182_reviewer_residual_map.py])
A careful external review (2026-06-14) raised four concerns at the [C]/[O] seams: (A) the mapping $2D \rightarrow 4D$ (“why does the dynamics read the masses as the Plücker norm 11?”), (B) the abstractness of UV gravity, (C) the meta-look-elsewhere of the grammar choice, and (D) the Koide Möbius flow (non-integer flow time $t \sim 2.84$, missing QFT generator). None adds an <i>independent</i> open gap. A [E]/[C]: the readout integer is exact ($11 = \sum_{k \leq 2} \binom{4}{k} = \dim S^+ - g_{\text{car}}$, the QBL boundary count, v174) and the $2D \rightarrow 4D$ step

is the holographic reduction, so $A = \text{QGEO.SYM.01} + F_{\text{transfer}}$. **B [E]**: “gravity is the geometric readout of the boundary” *is* **QGEO.SYM.01**. **D [E]/[C]**: $\text{Aut}(\mathbb{P}) = \text{PSL}(2, \mathbb{C})$ is the Möbius group and the cusp flow lives on the deck $\mathbb{P} \setminus \mu_4$, so the Möbius *nature* is forced by the geometry; the non-integer $t \sim 2.84$ only locates the physical point and the missing generator *is* F_{transfer} , so $D = \text{QGEO.SYM.01} + F_{\text{transfer}}$. So A, B, D collapse onto the *two already-named* residuals; only **C** is genuinely *experiment-only*: the grammar is frozen (**v84**), minimal and over-determined (the anchor $a = (1, 1, 2)$ lands in ≥ 6 mutually-independent, not-selected-for structures), a look-elsewhere-resistant signature whose residual — correctly — falls only to out-of-sample data (JUNO, CMB-S4), never to argument. The four concerns are a unification of the residual *structure*, not a closure of the underlying premises, which stay **[O]/[C]**.