

TFPT — Research Contracts for the Remaining Interfaces:

$$v_{\text{geo}}, G_{\text{net}}, F_{\text{transfer}}$$

(U_{wall}): the parabolic flavor wall-selection and (G_{metric}): the full quantum-gravity measure

Stefan Hamann

Alessandro Rizzo

July 15, 2026 · v5.4 (rev 370)

The TFPT document set — what is treated where

Plain language: TFPT (Topological Fixed-Point Theory) is a small discrete compiler. Two inputs — the seam constant $c_3 = \frac{1}{8\pi}$ and the carrier rank $g_{\text{car}} = 5$ — build $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ and read off the Standard-Model skeleton, the constants and the scale grammar. The series is **the introduction plus five numbered papers and four companions**, best read in order:

0. **introduction** — reading guide, compiler closure, predictions, the dependency DAG and the single proof ledger (**verification/status_ledger.csv**). (*entry point*)
1. **tfpt_1_architecture_e8** — the two axioms, the derivation map, the EM fixed point α^{-1} , the $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ construction, the QBL theorem chain.
2. **tfpt_2_standard_model** — the SM in one φ_0 -ladder formula, flavor from parabolic transport (sheet diamond, dual anchor, Q_+ cohomology), the worked closures.
3. **tfpt_3_e8_audit_bootstrap** — the seven E_8 slices as an audit raster, the cascade bridge, the Möbius bootstrap.
4. **tfpt_4_frontier** — honest status of η_B , m_p/m_e , Koide, dark matter, full QG.
5. **tfpt_5_redteam** — the adversarial audit: declared attacks, what survives, what each reduces to, and the kill tests.

Companions: **R. tfpt_research_contracts** — the open interfaces ($v_{\text{geo}}, G_{\text{net}}, F_{\text{transfer}}$) as numbered contracts; **H. tfpt_horizon_readouts** — Appendix H (unit-system reframe): one seam constant as the universal horizon thermal code, the Nariai/anchor surface, the resummed seam clock; **O. origin_theory** — why no free fundamental number remains (exact core + one honestly-typed cyclic interpretation); **S. tfpt_safeguards** — the verification discipline: every mechanism that defends a claim against coincidence and numerology (status calculus, no-free-pattern, the over-determination map, the firewall, frozen predictions + null model, the independent Wolfram and Lean paths, the red team).

You are reading companion R (Research Contracts).

TFPT status at a glance — read this card the same way in every document

Two inputs, one compiler. From the boundary constant $c_3 = \frac{1}{8\pi}$ (axiom P1) and the five-slot carrier $g_{\text{car}} = 5$ (axiom P2) — jointly the elementary symmetric data of one anchor $a = (1, 1, 2)$ — the discrete compiler $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ is *built*, and the Standard-Model skeleton (gauge group, three families, hypercharges, the flavor matrix), the dimensionless constants ($\alpha^{-1} = 137.0359992$) and the scale grammar are read off as fixed points. The algebraic core is machine-checked; physical readouts run through *named bridges* whose status is typed below.

What is closed, and what genuinely remains. The discrete/algebraic layer is closed ([E]); the everything-open residual reduces to *three named interface problems*:

Interface	Question	Status
v_{geo}	the one metrology unit ($= 1/\sqrt{G} = m/\mu$); No-Unit Thm: no compiler scale	primitive [O]
G_{net}	simple-current extension $(D_5)_1 \otimes (A_3)_1 \rtimes \langle (1, 1) \rangle \cong (E_8)_1$	[C]
F_{transfer}	one functor, four typed interfaces (Koide, η_B , axion, m_p/m_e)	[C]

Gravity is parameter-free. The classical metric-sector field equation is no longer only an $R+R^2$ readout: the fixed-volume entanglement equilibrium (Jacobson; Faulkner et al.), run with TFPT's atoms, gives the *full* covariant Einstein equation $G_{ab} + \Lambda g_{ab} = c_3^{-1} T_{ab}$ with *both* coefficients fixed — $c_3^{-1} = 8\pi$ (no free dimensionless Newton dial; the thermodynamic origin $2\pi/\eta$ *coincides* with the geometric one $|\mathbb{Z}_2|2\pi\chi$ via $|\mu_4| = |\mathbb{Z}_2|\chi(S^2) = 4$, so c_3 is triply over-determined — anchor, geometry, thermodynamics, v358) and Λ from α ($\rho_\Lambda = (3/4\pi^2)e^{-2\alpha^{-1}}$, v60); the Einstein tensor (not Ricci) is forced by Lovelock, so matter conservation is an output (v359). The residual here is only the equation-of-state status (an external candidate action — Bianconi's entropic action, arXiv:2408.14391 — is quantified in v473 without changing the typing) and the one unit v_{geo} ; the global measure (C7 = QG.AMB.01) is now a [C] redundancy (v369), a certification object rather than missing dynamics, conditional on SEAM.EQUIV.01 and Bisognano–Wichmann.

The seam-net structural theorem G_{net} (= SEAM.EQUIV.01: the raw seam *is* the holomorphic $(E_8)_1$ net) is now *closed modulo a cited theorem*. The explicit lattice model (v367/v368) and the S_3 closure stack pin the target at every computable level — central charge $c = 8$ (v376), the $(E_8)_1$ character with 248 currents and one primary (v377), genus-one torus $\text{GSD} = 1$ (v378) and reflection positivity (v379) — and it is Lean-formalised as FORM.SEAM.MMST.01: the collar's MMST hypotheses are kernel-proved, the MMST scaling-limit and Adamo–Moriwaki–Tanimoto OS-reconstruction theorems enter as named cited axioms, and the #print axioms check is clean. The *only* residual that stays [O] is the abstract continuum existence of the scaling limit (v336) — its 128-spinor extension leg since certified at net level by the peer-reviewed crossed-product package ($h_s=16/16=1 \in \mathbb{Z}$, the Longo–Rehren locality criterion; Böckenhauer; KLM $\mu=4/2^2=1$; AGT/AMT the independent second witness) with the realisation input reduced to invariant level and the parallel Lean route `seamResidualClosed'` (v469); so this is *closed modulo a cited theorem*, not solved. QGEO.SYM.01 (the μ_4 deck acting geometrically) is its *corollary*: a conformal net's vacuum is rotation-invariant by axiom, so the deck follows with no extra premise (v335, Lean FORM.QGEO.BW.01). The nonperturbative quantum-gravity measure QG.AMB.01 is no longer carried as an open item: it is a [C] *redundancy* (v369), a certification object rather than missing dynamics, conditional on SEAM.EQUIV.01 and Bisognano–Wichmann (and TFPT is moreover rigorously gap-decoupled from the general Euclidean-QG conformal-factor problem, margin $1.648 > 0$, v76/v330). So the honest residual is the one unit v_{geo} plus the typed F_{transfer} interfaces, above the cited-theorem ceiling on the seam and with QG.AMB.01 a [C] redundancy. Everything carries a claim ID resolving to a row of `verification/status_ledger.csv` (the single source of truth); **if the text and the ledger ever disagree, the ledger wins**. The historical labels ($U_{\text{wall}}, G_{\text{metric}}, F_{\text{frontier}}$) are retained only for ledger continuity — the live residual is $v_{\text{geo}} \oplus F_{\text{transfer}}$, with the seam G_{net} closed modulo a cited theorem and QG.AMB.01 a [C] redundancy.


Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in `verification/status_ledger.csv`.

[E] exact / machine-proven

[C] conditional

[O] open / axiom

[X] kill test



What this document is about

After the compiler closure the entire residual is

$$\text{Rest} = (U_{\text{wall}}) \oplus (G_{\text{metric}}) \oplus (F_{\text{frontier}})$$

This note turns the two genuine research gates into *contracts*: a numbered chain of lemmas, the single theorem that closes each gate, and — for every step — whether it is machine-certifiable today. (F_{frontier}) (Koide, η_B , axion relic scale, m_p/m_e) is *not* a gate: those are deliberately typed QFT/cosmology interfaces (see `tfpt_4_frontier`). **Priority:** (U_{wall}) *first* (finite, algebraic, falsifiable), then (G_{metric}) (deep analytic programme).

Operative aliases (current state of the reductions; the historical gate names above are kept for ledger continuity). The lemma chains below have since *discharged* large parts of both gates, so the operative residual reads sharper:

$$\text{Rest}_{\text{operativ}} = v_{\text{geo}} \oplus G_{\text{net}} \oplus F_{\text{transfer}}$$

— (U_{wall}) $\rightarrow v_{\text{geo}}$: the quark *ratios* are closed combinatorially (Readout Rigidity + Plücker, U2/v49/v71), U_{point} collapses to the *one* dimensionful unit v_{geo} shared with $1/G$ (v75/v68) — what remains is a scale anchor, not a flavor gap (the *full* amplitude matrix beyond the ratios is the only place where U_{point} still appears); (G_{metric}) $\rightarrow G_{\text{net}}$: holographically reduced to the single boundary-net statement “*the seam-Calderón inclusion has Jones index 4 = $|\mu_4|$* ” — holomorphy, $(E_8)_1$ and the unique 2D bulk then follow (v89/v92/GATE.METRIC.05-07); and the carrier-side QBL input now merges into the *same* statement: the tower carrier $\rightarrow SO(16)_1 \rightarrow (E_8)_1$ carries one and the same 16-fermion content, whose single quasi-free kernel (rank = c ; Wick/Pfaffian exact) is the certified Calderón datum — closing G_{net} therefore also closes the carrier choice [E] ([`verification/v113_quasifree_kernel.py`]); (F_{frontier}) $\rightarrow F_{\text{transfer}}$: named QFT/cosmology *transfer* interfaces, never compiler powers — and *not* four unrelated interfaces but *one* discrete geodesic on the Sheet Diamond (the $K \rightarrow C \rightarrow F$ sheet axis $M(1, t)$, with the Koide source $\rightarrow$ pole as its 1-D Möbius projection), along which the external continuous solver merely evaluates ([`verification/v224_diamond_ftransfer_path.py`]).

The residual is certification, not construction ([`verification/v384_residual_certification.py`]).

A ledger-CI audit classifies every remaining item into exactly three *external* kinds, with *zero* open *internal* TFPT mechanisms: [E] (A) an external math proof — SEAM.EQUIV.01 continuum existence (the cited MMST theorem plus the crossed-product extension package, v336/v469, with AGT/AMT as second witness), ALPHA.QUILLEN.EXACT.01 (v382, its level and normalisation sharpened by v470, the determinant-line/moduli bridge lemma exhibited at the finite level by v472 — the continuum ζ -det identification stays the open part) and the constructive QG.AMB.01 measure (a [C] redundancy, v369); [E] (B) theorem-forbidden — v_{geo} , the one unit, by the No-Unit theorem (v153/v364); [E] (C) external physics — F_{transfer} (v371–v374, firewalled v187). The two load-bearing facts are checked (No-Unit dimensionlessness; the gap-decoupling margin $1.648 > 0$), so the only hard things left are theorems and external physics, not missing dynamics (RESIDUAL.CERTIFICATION.01).

The Coxeter-clock facet symmetry, resolved as a structural lemma (RES.COXETER.SYMMETRY.01, [`verification/v394_coxeter_atoms.py`]/[`verification/v409_coxeter_prime2_lemma.py`]). The order-30

clock $30 = 2 \cdot 3 \cdot 5 = h(E_8)$ has all three primes facet-typed, and they are exactly the three anchor atoms ($2=|Z_2|$, $3=N_{\text{fam}}$, $5=g_{\text{car}}$) wearing number-field hats ($\mathbb{Q}(i)/\mathbb{Q}(\sqrt{-3})/\mathbb{Q}(\sqrt{5})$, [v390](#)). The table is *not* symmetric: the seam (prime-2) appears as the *common factor* of both dynamic rates ($2/3 = |Z_2|/N_{\text{fam}}$ and $(\varphi+2)/4 = (\varphi+2)/|\mu_4|$, $|\mu_4|=2^2$), not as an independent attractor. *The structural lemma* ([v409](#)) *closes the falsifiable corollary*: the sheet involution $\sigma = \text{diag}(1, -1, -1)$ has spectrum $\{+1, -1\}$ and projector $(I+\sigma)/2$ spectrum $\{0, 1\}$ — none in the open interval $(0, 1)$, so prime-2 supplies a *parity/projection*, never a nonzero contraction rate; the genuine rates $2/3$ and $(1/3)^6$ lie in $(0, 1)$ and carry prime-2 only as a *factor*. Hence prime-2 is *necessary* (admissible boundary transport is sheet-paired) but *not autonomous*: **no prime-2-only attractor exists** — dynamics begins only after coupling to family prime-3 or carrier prime-5 ($|\mu_4|=2^2$ the Galois bridge). This closes the corollary [\[E\]](#) and types the necessity [\[C\]](#), sharpening [v383](#)’s “every sector is the same gapped operator” towards “the same *seam-driven* operator”; the full “appears in every gap” universality stays the [v383](#) reading. The corollary is **machine-proved in Lean 4** (`TfptCarrier/CoxeterPrime2.lean`, `FORM.COXETER.PRIME2.01`: axiom-pure, `lake build clean`). [\[E\]](#)/[\[C\]](#).

QFT layer (this round). The boundary QFT is now *one relative object* $\text{TFPT}_{\text{QFT}} = (\mathcal{A}_\Sigma, \omega_\Sigma, \Delta_\Sigma, \rho, A_F, H_F, D_F, J, \gamma, S_{\text{rel}})$, closed modulo the *single keystone* `SEAM.EQUIV.01` (whose downstream conformal-deck face is `QGEO.SYM.01`): its finite Dirac is a covariance induction of the seam KMS state ([v258](#)), its spectral-action cutoff *is* that KMS weight ([v259](#)), its seam, carrier-16 and E_8 live on one Kummer/K3 surface ([v260](#)), and the whole assembly is certified as the *Modular Spectral Closure* ([v261](#)). The ambient quantum-gravity measure (G_{metric}) is gap-decoupled and kept separate by design — the QFT layer does not wait on it.

Research Contract 1 — (U_{wall})

Goal. Prove

$$\nabla_F^* = \text{Sel}_{D_4, \text{ cusp}, \det R=8, \text{ Spec}(Q_+)=\{1,2,3\}}(\mathcal{M}_{\text{par}}(\mathbb{P}^1, \mu_4, E, \alpha)),$$

with $E = \mathcal{O}(-2) \oplus \mathcal{O}(-1) \oplus \mathcal{O}(-1)$ and $\alpha = \{0, \frac{1}{3}, \frac{2}{3}\}$ at the four points of μ_4 . The substrate is standard: degree-zero stable parabolic bundles correspond to unitary representations with prescribed local monodromies (Mehta–Seshadri); tame non-abelian Hodge gives the Higgs / flat-connection side. The TFPT content is to pin the *one* D_4 -symmetric, rank-3, four-point point that realises the selectors.

Lemma 1 (U_0 — normalisation). *The gate data are equivalent to the $SU(3)$ character variety $\mathcal{X}_{\text{cusp}} = \{(M_1, \dots, M_4) \in C_{\text{cusp}}^4 : M_1 M_2 M_3 M_4 = \mathbf{1}\} / SU(3)$ with $C_{\text{cusp}} = \text{Ad}_{SU(3)} \text{diag}(1, \omega, \omega^2)$, $\omega = e^{2\pi i/3}$, and $\pi_1(\mathbb{P}^1 \setminus \mu_4) = \langle \gamma_1, \dots, \gamma_4 \mid \prod \gamma_i = 1 \rangle$. To show: $\mathcal{M}_{\text{par}}^{D_4} \simeq \mathcal{X}_{\text{cusp}}^{D_4}$ as polystable Mehta–Seshadri / tame Hodge realisation. Machine: the group presentation and D_4 -action are finite and codable (Sage/Wolfram).*

Finite rigidity — now proven [\[E\]](#). The finite half of U_0 is closed exactly: μ_4 has cross-ratio 2 and Möbius stabiliser $\langle z \mapsto iz, z \mapsto 1/z \rangle$ of order $8 = |D_4|$ (a 4-cycle of the marks + a reflection, faithful); $H^1(\mathbb{P}^1 \setminus \mu_4)$ has rank $4 - 1 = N_{\text{fam}} = 3$; and the forms $\omega_k = z^{k-1} dz / (z^4 - 1)$, $k = 1, 2, 3$, are μ_4 -eigenforms ($z \mapsto iz : \omega_k \mapsto i^k \omega_k$) carrying the three nontrivial characters χ_1, χ_2, χ_3 with weights $(1, 2, 3) = \text{Spec}(Q_+)$. Hence a genus-zero boundary with faithful D_4 , four parabolic marks and this μ_4 -decomposition is Möbius-equivalent to $\mathbb{P}^1 \setminus \mu_4$ with the $(1, 2, 3)$ grading (`QGEO.RIGID.01`, [\[E\]](#), [\[verification/v168_qgeo_rigidity.py\]](#)). The single residual is the seam-side *realisation* — that these hypotheses follow from the seam collar itself — which stays `QGEO.REALIZE.01` ([\[C\]](#)).

Seam Collar Realisation Theorem — the one central proof obligation [\[E\]](#)/[\[O\]](#)
([\[verification/v176_seam_collar_realisation.py\]](#))

Stated as one theorem and certified as a reduction — not a fabricated proof.

Theorem (target). *Given* (1) an oriented one-sided RP seam collar Σ ; (2) a sheet-odd Calderón involution; (3) quasi-free one-particle seam data with gap $(2/3)^6$; (4) a faithful D_4 action on the boundary marks; (5) four parabolic marked points with μ_4 -character decomposition; (6) an admissible $c=8$ boundary net — *then* the conformal boundary of Σ is Möbius-equivalent to $\mathbb{P}^1 \setminus \mu_4$, its H^1 character

basis is canonically the generation basis with grading $(1, 2, 3)$, the seam Calderón contraction is the rank-8 K_Σ polarisation, and the index-4 simple-current extension is $(E_8)_1$.

The proof decomposes into six steps, five already [E]:

- 1–2 *Geometry / uniformisation* [E] (v168): μ_4 has cross-ratio 2 and a faithful Möbius D_4 stabiliser of order 8; four genus-0 marks with faithful D_4 force the μ_4 square up to Möbius; $H^1(\mathbb{P}^1 \setminus \mu_4)$ has rank $4 - 1 = N_{\text{fam}} = 3$.
- 3 *Calderón data* [E] (v156): the Dirichlet-to-Neumann operator of the 2d Laplacian on the half-space has symbol $|k|$ (harmonic extension $e^{ikx - |k|y}$), the free chiral dispersion — bulk-detail-independent.
- 4 *H^1 character = generation grading* [E] (v137/v168): $\omega_k = z^{k-1}dz/(z^4 - 1)$ carry μ_4 characters i^k with weights $(1, 2, 3)$ = the A_3 exponents = $\text{Spec}(Q_+)$ — natural, not merely isomorphic.
- 5 *One-particle contraction* [E] (v162): a μ_4 -equivariant operator is diagonal in the cusp basis $(\frac{1}{|\mu_4|} \sum_j U^j X U^{-j} = \text{diag } X \text{ for } U = \text{diag}(1, i, -1, -i))$; its spectrum is the cusp weights $\{0, \frac{1}{3}, \frac{2}{3}\}$, so the transfer eigenvalues $(1 - \alpha)^{p_2}$ ($p_2=6$) are $\{1, (2/3)^6, (1/3)^6\}$, gap $(2/3)^6$ — forced, no free knob.
- 6 *AQFT closure* [E] (v154/v175): $(D_5)_1 \otimes (A_3)_1$ with the isotropic μ_4 glue has index $4 = |\mu_4|$, $c = 5+3 = 8$, μ -index $16/16 = 1 \Rightarrow$ holomorphic $\Rightarrow (E_8)_1$ ($248 = 120+128$); full-cone RP holds for all m via the CAR second-quantisation functor.

The proof tree — and the two genuine open doors [O] ([verification/v177_seam_marking_kernel.py]). Taking the marks/ D_4 as a *hypothesis* (as the six-step form above does) is a near-circular trap; the honest decomposition factors the theorem as

$$\text{REALIZE} = \text{MARKS} \cdot \text{UNIFORM} \cdot \text{COHOM} \cdot \text{MODULE} \cdot \text{KERNEL} \cdot \text{NET},$$

of which *four nodes are now closed constructively*: UNIFORM [E] (an order-4 Möbius map conjugates to $z \mapsto iz$, a non-fixed orbit scales to μ_4 , and $\sigma\rho\sigma = \rho^{-1}$ for $\sigma : z \mapsto 1/z$ gives a faithful D_4 — so a genus-0 four-marked faithful- D_4 boundary *is* $(\mathbb{P}^1, \mu_4)$, strengthening v168's cross-ratio to the full uniformisation — and now **Lean-formalised** as FORM.QGEO.02: $\rho^4 = \text{id}$, $\rho^2 \neq \text{id}$, $\sigma^2 = \text{id}$, $\sigma\rho\sigma = \rho^{-1}$, the orbit scaling to μ_4 , the μ_4 permutation, and the multiplier classification “order-4 $\Leftrightarrow \zeta = \pm i$ ” are machine-proved in TftptCarrier/MobiusUniformisation.lean, AUDIT: PASS); COHOM [E] ($\rho^* \omega_k = i^k \omega_k$, grading $(1, 2, 3)$) — also **Lean-formalised** as FORM.QGEO.03: the pullback characters i^1, i^2, i^3 and the reflection parity $\sigma^* \omega_1 = \omega_3$, $\sigma^* \omega_2 = \omega_2$, $\sigma^* \omega_3 = \omega_1$ are machine-proved in TftptCarrier/CohomologyGrading.lean, AUDIT: PASS); MODULE [E] (a *unique* μ_4 -equivariant iso — multiplicity-free + residue normalisation, reflection $\omega_1 \leftrightarrow \omega_3$, ω_2 fixed; the parity is the Lean FORM.QGEO.03 above, the *uniqueness* stays symbolic); and NET [E] (v154/v175). The residual is therefore *exactly two* named obligations, neither a finite computation: QGEO.MARKS.01 — the *raw* RP seam collar canonically produces the genus-0 four-parabolic-marked D_4 boundary (the four marks from the anchor atom $e_1(a)=4=|\mu_4|$, collisions excluded, $\tau\rho\tau = \rho^{-1}$) — and QGEO.KERNEL.01 — the *raw* seam Calderón operator equals the μ_4 -equivariant free $c=8$ gapped one-particle contraction *as an operator* ($C_\Sigma = U^{-1} C_{\mu_4} U$), not merely the same spectral list. **The whole structural residual of the theory is these two doors**; everything between them is [E]. They need a human constructive-geometry/AQFT argument (or a proof assistant), not more arithmetic.

How far each door opens by computation ([verification/v178_marks_kernel_attempt.py]). Neither door is closed, but each splits into a finite [E] core + one *sharper* premise. *For MARKS*: given a genus-0 double with an order-4 clock ρ (fixed points $\{0, \infty\}$) and reflection τ , the marks are forced to be a generic ρ -orbit $\{1, i, -1, -i\} = \mu_4$ (not the fixed points), distinct iff $b_1 = 3 = N_{\text{fam}}$ (any collision drops b_1), with $\tau\rho\tau = \rho^{-1}$ giving faithful D_4 — so MARKS reduces to the *single* topological premise “the seam double is genus-0 with an order-4 clock whose marks form one orbit”. *For KERNEL*: on the 3-dim multiplicity-free H^1 the μ_4 characters $(1, 2, 3)$ are distinct, so by Schur a μ_4 -equivariant operator is forced *diagonal* and is *completely determined by its eigenvalue per character* — two such with the forced transfer spectrum $\{1, (2/3)^6, (1/3)^6\}$ and the unique residue-normalised basis are *equal as operators*, not merely isospectral; the “spectrum vs operator” distinction *vanishes* on the finite block. So KERNEL reduces to the *single* premise “the full DtN on L^2 reduces to this 3-dim gapped transfer block + the fixed rank-8 polarisation” (principal symbol $|k|$, Lee–Uhlmann v156; no drift v158). The two doors are thereby *milder*, but still open.

Marks without a geometric posit: a Lefschetz/character argument ([verification/v195_marks_lefschetz_character.py]). The marks need not be *assumed* as D_4 boundary points: they follow from the H^1 character alone. The order-4 clock $\rho: z \mapsto iz$ fixes exactly

$\{0, \infty\}$, so a puncture at a fixed point would contribute a *trivial* character to H^1 ; but the carrier datum is $H^1 = \chi_1 \oplus \chi_2 \oplus \chi_3$ (the A_3 exponents $(1, 2, 3)$, with $\text{Tr}(\rho|H^1) = i + i^2 + i^3 = -1$), which has *no* trivial component — so no puncture sits at a fixed point. With $\text{rank } H^1 = n - 1 = 3 \Rightarrow n = 4$ punctures, all off the fixed locus and each in a free order-4 orbit, the only possibility is a *single* free μ_4 orbit (cross-ratio 2), and the closed-surface Lefschetz number $L(\rho) = 1 + 1 = 2 = |\{0, \infty\}|$ is consistent. So MARKS reduces from the geometric posit “ D_4 marks” to the milder *representation* premise “the raw seam H^1 carries the A_3 -exponent rep” — a carrier-algebra statement, strictly less circular, still [O].

And the two doors are in fact ONE ([verification/v179_conformal_realisation.py]). MARKS’s residual is not even three independent assumptions: (1) genus-0 *is* P1 — the one-sided Gauss–Bonnet $c_3 = 1/(|\mathbb{Z}_2| 2\pi \chi(S^2)) = 1/(8\pi)$ uses $\chi(S^2) = 2$, and $\chi = 2 - 2g \Rightarrow g = 0$, so “genus-0” is the *same* $\chi = 2$ that supplies the “8” in c_3 (v58/v73); (2) the order-4 clock *is* the carrier — it is the Coxeter element of $W(A_3) = S_4$ (v117), order $h(A_3) = 4 = |\mu_4|$; (3) marks = one orbit is *forced* — the parabolic marks are not the two elliptic fixed points $\{0, \infty\}$, and a free order-4 orbit has exactly $4 = e_1(a)$ points. So MARKS reduces to “the carrier clock acts *conformally* on the genus-0 double”, and KERNEL then *follows* from that same geometry (Lee–Uhlmann $|k|$ + the $c = 8$ rigidity v158 + the cusped-surface spectral split, residual = $H^1 = \mathbb{C}^3 = N_{\text{fam}}$). **Hence the entire structural residual of the theory is a single geometric premise** QGEO.CONF.01: *the raw RP seam normal double is conformally* ($\mathbb{P}^1, \mu_4$) *with the carrier clock as the order-4 Möbius automorphism* $z \mapsto iz$. Given it, everything — genus-0, the four D_4 marks, the $(1, 2, 3)$ grading, the DtN block, the rank-8 polarisation and the $(E_8)_1$ net — follows. This single premise is a constructive-geometry statement (the conformal class of the raw seam *is* the cusped sphere), *not* a finite computation, and is left honestly [O].

And even that conformal premise reduces to a milder isometry premise ([verification/v180_clock_is_mobius.py]). QGEO.CONF.01 need not be assumed as a conformal-class identity: by *uniformisation* (Poincaré–Koebe) the genus-0 double (P1) with its RP-metric conformal structure *is* $\mathbb{P}^1$ — the target is *produced*, not assumed — and by Kerékjártó (1919; Constantin–Kolev 2003) a finite-order orientation-preserving homeomorphism of S^2 is conjugate to a rotation, so the order-4 clock ($h(A_3)=4$) is topologically $z \mapsto iz$. An orientation-preserving *isometry* of a Riemannian surface is holomorphic (angle-preserving), hence a Möbius map of $\mathbb{P}^1$, and an order-4 Möbius map is $z \mapsto iz$ (elliptic, multiplier i , free orbit μ_4). So QGEO.CONF.01 is *implied* by the strictly milder, non-circular premise QGEO.ISO.01: “the carrier clock is an order-4 orientation-preserving *isometry* of the RP seam metric” — which does *not* name $\mathbb{P}^1/\mu_4$ and is the natural statement that the seam transport is metric-compatible (a holonomy preserving the RP inner product). **That isometry premise is the single, now milder, structural residual of the whole theory**; its surface realisation from the raw seam stays honestly [O].

And one step further reaches bedrock ([verification/v181_clock_is_conformal_symmetry.py]). By *equivariant uniformisation* (Nielsen realisation; for cyclic groups Kerékjártó + uniformisation) a finite-order orientation-preserving action on a genus-0 surface preserves a complex structure for which it acts by Möbius maps — so a *conformal automorphism* (strictly weaker than an isometry: every isometry is conformal, conformal allows any positive scale) already suffices to put the order-4 clock in the form $z \mapsto iz$. Since the clock is the Coxeter element of $W(A_3) = S_4$ (v117), i.e. the μ_4 deck, and a deck transformation preserves the pulled-back conformal structure *by definition*, the residual collapses to the *definitional* identification QGEO.SYM.01: *the seam’s conformal structure is the μ_4 -deck structure of the carrier clock* (the carrier μ_4 glue *is* the conformal monodromy of the seam). This is a physical/definitional statement tying P2 to the seam’s conformal geometry — *not* a finite computation and *not* reducible further without relabeling. **This is bedrock**: the chain $\text{REALIZE} \rightarrow \text{theorem} \rightarrow \text{MARKS} + \text{KERNEL} \rightarrow \text{CONF} \rightarrow \text{ISO} \rightarrow \text{SYM}$ (v175–v181) ends at a single definitional premise, left honestly [O]; everything above it is a theorem or an established citation.

A sharper, operator-checkable restatement: QGEO.ENERGY.01 ([verification/v192_qgeo_energy_clock.py]). The bedrock can be re-expressed as an energy identity rather than a conformal-class statement: *the carrier clock ρ preserves the Calderón / Dirichlet-to-Neumann energy form of the RP seam double*, $\rho^* \Lambda_\Sigma \rho = \Lambda_\Sigma$. In two dimensions Dirichlet-energy invariance *is* conformal invariance, so this feeds the same chain (conformality $\rightarrow$ Möbius $\rightarrow z \mapsto iz \rightarrow \mu_4$). Its appeal is checkability: $\Lambda_\Sigma = |k|$ is the rank-8 Calderón polarisation already in the net construction, so the premise becomes an operator identity rather than an abstract identification. **Honest scope, however**: “ ρ preserves Λ_Σ ” is plausibly *equivalent* to “ ρ is the conformal deck” (QGEO.SYM.01), so QGEO.ENERGY.01 *relocates* the bedrock to a more falsifiable surface but does *not* eliminate it — it becomes a theorem only if ρ is defined *independently* (from the carrier algebra)

and the invariance is then *proven*, which is open. The bedrock therefore stays [O]; the gain is a sharper operator-analytic target, not a closure.

The non-circular sharpening: QGEO.ENERGY.02 ([verification/v193_qgeo_energy_commutator.py]). The way to make “ ρ defined independently” precise is the *energy commutator* $[\rho, \Lambda_\Sigma] = 0$ on the rank-8 Calderón polarisation, with three sub-claims: (1) ρ is the Coxeter element of $W(A_3) = S_4$ (order $4 = |\mu_4|$, v117), defined from the *carrier* algebra, not the seam; (2) *the non-circularity hinge* — Λ_Σ must be the DtN of the *raw* RP seam (definable from the RP data alone), not the μ_4 normal-form operator; (3) $[\rho, \Lambda_\Sigma] = 0$ forces the μ_4 character grading $(1, 2, 3)$ on H^1 , already established (QGEO.COHO.01). An exact consistency rider pins the direction: the induced EH coefficient $k = (\ln m)/(12\pi)$ (v150) reproduces $c_3/2$ iff $\ln m = q(A_3) = \frac{3}{4}$ (the *family* glue norm), while $q(D_5) = \frac{5}{4}$ (the carrier norm) misses by $g_{\text{car}}/N_{\text{fam}} = \frac{5}{3}$ — so the seam energy form must reproduce the *family* norm (gravity is family-geometry-induced). If sub-claim (2) holds, $[\rho, \Lambda_\Sigma] = 0$ turns QGEO.SYM.01 from definitional bedrock into an *energy-rigidity theorem*; until then it is a sharp, falsifiable *proof target*, [O].

Subclaim (2) tackled — one notch closer, not closed ([verification/v194_raw_seam_dtn_rp.py]). The non-circularity hinge is essentially *met* at the canonical level: Osterwalder–Schrader reconstruction yields a canonical transfer operator $T = e^{-H}$ from the RP state and reflection alone — no coordinates (v54; §(4) of tfpt_2) — on a *quasi-free* seam state (v155: the boundary marginal of a Gaussian bulk is the compression PGP, a contraction). So Λ_Σ is RP-definable without ever naming $\mathbb{P}^1 \setminus \mu_4$. The commutator $[\rho, \Lambda_\Sigma] = 0$ then *closes on the finite cohomology* H^1 [E]: $\rho: z \mapsto iz$ acts on $\langle w_1, w_2, w_3 \rangle$ ($w_k = z^{k-1} dz / (z^4 - 1)$) as $\rho^* w_k = i^k w_k$, distinct characters $i, -1, -i$, so any ρ -equivariant operator is diagonal there (v177). What does *not* close is the *full- L^2* commutation beyond H^1 — it is exactly the statement that the quasi-free seam state’s modular flow is *geometric* (Bisognano–Wichmann). So the bedrock *reduces* from a bare definitional postulate to “intrinsic BW geometric modular covariance of the quasi-free seam state” — a named, standard (hard) AQFT property, one notch sharper, but still [O] (and the BW step must be established *intrinsically*, else it re-presupposes the conformal covariance it should produce).

A concrete variational handle on the residual ([verification/v196_seam_energy_functional.py]). The open BW statement can be made a discretisable extremum problem via the failure functional $E_{\text{fail}}(\Sigma, \rho) = \|[\rho, \Lambda_\Sigma]\|_{HS}^2 + \|\rho^4 - 1\|^2 + \|\Theta\rho\Theta - \rho^{-1}\|^2$, whose zero locus is exactly the QGEO.SYM.01 conditions (the energy commutator, the order-4 clock, and the RP-reflection twist). On the finite H^1 block $E_{\text{fail}} = 0$ at the μ_4 configuration ($\rho = \text{diag}(i, -1, -i)$, Λ μ_4 -equivariant, $\Theta = \text{conjugation}$) and > 0 for any μ_4 -breaking perturbation, so μ_4 is a genuine minimiser [E]; minimising E_{fail} over the *full* raw-seam DtN is precisely the BW residual above. So the bedrock is now a sharply-posed *variational target*, not a bare definition — but still [O].

The decisive crack: principal symbol exact + Tomita–Takesaki ([verification/v198_modular_commutator_reduction.py]). Two observations collapse the residual to a single non-circular state-invariance. *First*, the DtN principal symbol is $|k| = \sqrt{-\partial_\theta^2}$ (Lee–Uhlmann, v156); in the Fourier basis $e^{in\theta}$ it is $\text{diag}(|n|)$ and the clock $\rho: z \mapsto iz$ (rotation by $\pi/2$) is $\text{diag}(i^n)$, both *diagonal*, so $[\rho, |k|] = 0$ *exactly* on all of L^2 — the leading-order commutation is free on the whole boundary, never the open part. *Second*, by Tomita–Takesaki the modular operator Δ_ω is intrinsic to (M, Ω) , so a state-preserving symmetry ($\rho\Omega = \Omega$) commutes with Δ and hence with $\log \Delta \sim \Lambda_\Sigma$: thus $[\rho, \Lambda_\Sigma] = 0$ *follows from* $\omega \circ \rho = \omega$ — using only the state and reflection (which RP supplies) and *no* conformal covariance, which **removes the BW circularity worry above**. Since the finite H^1 block is already μ_4 -invariant (v177), the entire residual is now the single *state-invariance* “the raw quasi-free seam state is μ_4 -invariant” ($\omega \circ \rho = \omega$, equivalently the Gaussian bulk measure is deck-invariant) — the maximally-operational, checkable form of QGEO.SYM.01. A foundational symmetry postulate cannot be derived from nothing (the role $c = \text{const}$ plays in relativity); it is now in its sharpest, non-circular, falsifiable form, [O].

The residual, fully localised and numerically probed ([verification/v199_seam_state_invariance.py], [verification/v200_seam_variational_scan.py]). Attacking $\omega \circ \rho = \omega$ on the *raw* seam (not the μ_4 normal form) confirms it is non-circular — ρ is the Coxeter element of $W(A_3) = S_4$ ($\rho^4 = 1$), an automorphism of the raw CAR algebra, defined entirely carrier-side — and reduces it, via the quasi-free covariance $C_\Sigma = \frac{1}{2}(1 + \text{sgn } H_1)$, to $[\rho, H_1] = 0$, i.e. H_1 *is block-diagonal in the four μ_4 -character classes* $\{n \equiv r \pmod{4}\}$. The principal symbol $|k| = \text{diag}(|n|)$ passes automatically, so the entire residual is the single bounded statement “the sub-principal symbol of the raw DtN has no off-character matrix elements”. The companion variational scan minimises $E_{\text{fail}}(\theta) = \|\rho(\theta)^4 - 1\|^2 + (\text{faithfulness})$ over a free clock angle and finds its *unique* minimum at $\theta = \pi/2$ —

the μ_4 clock $z \mapsto iz$ (the only faithful order-4 grading $i, -1, -i$) — a numerical sign that the variational principle selects μ_4 . This is the “next executable theorem” a working mathematician can take up; it does not close the postulate, [O].

The bounded sub-principal residual is the image of the μ_4 marks, not a free postulate ([verification/v201_seam_subprincipal_marks.py]). One genuine step further: the seam DtN is $\Lambda = |k| + M_f$ with M_f the bounded sub-principal piece = multiplication by a curvature function $f(\theta)$ on the seam circle (the standard Steklov/DtN structure), so $\langle n|M_f|n' \rangle = f_{n-n'}$. Then M_f is μ_4 -character-block-diagonal iff f has Fourier support only on modes $\equiv 0 \pmod{4}$, i.e. iff f is $\mathbb{Z}_4$ -invariant (the principal $|k|$ already commutes, v198). But a curvature sourced by the four marks, $f(\theta) = \sum_{j=0}^3 g(\theta - 2\pi j/4)$ for any profile g , has $f_m = 4 g_m [m \equiv 0 \pmod{4}]$ — automatically $\mathbb{Z}_4$ -invariant. So the μ_4 -mark orbit (forced by v195) forces the sub-principal symbol block-diagonal; 3 equally-spaced marks ($\mathbb{Z}_3$) or 4 generic marks both break it (negative controls). The residual of $\omega \circ \rho = \omega$ thus collapses to “the DtN sub-principal symbol is *mark-local* (the seam is flat away from the μ_4 marks = the conformal-deck structure, v177)” — a locality/flatness property, structurally definitional, the narrowest form yet of QGEO.SYM.01, [O].

Certified on realistic Steklov profiles, and at the state level ([verification/v210_mark_local_dtn.py]). The mode-level argument above is carried to concrete, smooth boundary curvature: a real von-Mises bump $g(\theta)$ localised at a puncture, summed over the four marks $f(\theta) = \sum_{j=0}^3 g(\theta - j\pi/2)$, has every off-(mod 4) Fourier coefficient vanishing to $< 10^{-16}$ (exact 4-fold cancellation), the assembled DtN $\Lambda = |D_\theta| + M_f$ has $\|[\rho, \Lambda]\| < 10^{-15}$, and — the genuinely new step — the *quasi-free state* it defines, with covariance $C = \frac{1}{2}(1 + \text{sgn } H_1)$, is clock-invariant: $\rho C \rho^\dagger = C$ to $< 10^{-13}$, i.e. the actual $\omega \circ \rho = \omega$, not merely $[\rho, \Lambda] = 0$. A single off-centre bump, a $\mathbb{Z}_3$ (three-mark) source and four unequally-spaced marks each break both by $O(1)$, and the mark-local commutator stays $< 10^{-10}$ as the truncation grows ($N \in \{16, 32, 64\}$) — not an artefact. This *numerically certifies the implication* “mark-local DtN $\Rightarrow \omega \circ \rho = \omega$ ” for real profiles; the premise that the physical seam *is* mark-local stays [O] (it does not close QGEO.SYM.01).

The same residual, read dynamically: modular flow + a recovery Lindblad generator ([verification/v238_modular_lindblad_dynamics.py]). The whole QGEO.SYM.01 discussion above is *static* — a commutator condition $[\rho, \Lambda_\Sigma] = 0$. Read as a *dynamics* it is the statement that the modular flow $\sigma_t = \Delta^{it}$ (the Tomita–Takesaki one-parameter group generated by the modular Hamiltonian $\Lambda_\Sigma = \log \Delta_\omega$) *conserves* the μ_4 clock: $\rho \sigma_t \rho^{-1} = \sigma_t \forall t \Leftrightarrow [\rho, \Lambda_\Sigma] = 0$ — so the open premise is exactly “modular time conserves the carrier charge”. The dissipative half is the recovery channel T of v221 read as a generator: $L = \log T$ is a doubly-stochastic (CPTP-classical) Markov generator ($e^L = T$, $e^{L\tau} \rightarrow$ the democratic projector = the arrow of time), and its dissipative gap is $-\log((2/3)^6) = 6 \log(3/2) = \Delta$, the OS mass gap of v64 — the static recovery bound and the dynamical mass gap are *one* number, one exponentiated. Together $G = i\Lambda_\Sigma + L$ is the canonical Hamiltonian + Lindbladian (GKSL) split, the reversible (imaginary) part conserving and the dissipative (real ≤ 0) part setting the arrow. This is a finite-block exhibit of where the static compiler becomes a dynamics; the residual it leaves — that σ_t is the *geometric* modular flow (Connes–Rovelli “thermal time = physical time”) on the full seam algebra — is the same QGEO.SYM.01 bedrock, [O].

From the dynamics to a QFT skeleton on the seam ([verification/v239_kms_thermal_time.py], [verification/v240_gns_os_reconstruction.py], [verification/v241_dhr_sectors.py], [verification/v242_spin_statistics_modular.py], [verification/v243_haag_ruelle_braiding.py], [verification/v244_spectral_action_contract.py]). With the static $\rightarrow$ dynamic flip in hand, the standard AQFT machinery assembles a finite-block QFT skeleton, each step typed honestly. *Thermal time* (v239): the modular flow is KMS at $\beta = 1$, and the seam unit $2\pi = 1/(4c_3)$ fixes the modular temperature as the horizon temperature $T_H = c_3/M$. *Reconstruction* (v240): GNS gives the Hilbert space and a cyclic vector from $(\mathcal{M}, \omega)$, and the OS Hamiltonian $H_{\text{OS}} = -\log T = -L$ is positive with gap Δ — the v64 “ $H \geq 0$ + mass gap” is the recovery generator read as Euclidean time. *Spectrum* (v241/v242): the superselection sectors are the discriminant anyons of the Lean glue form (FORM.GLUE.01); fusion is recovered from the modular S by Verlinde, the Gauss–Milgram sum returns $c = 8$, and the condensation tower $16 \rightarrow 4 \rightarrow 1$ (v235/v237) selects the holomorphic $(E_8)_1$. *Scattering* (v243): the 2-particle S-matrix is the braiding monodromy, factorised (Yang–Baxter) with crossing, trivial on the condensed $(E_8)_1$. The one genuinely 4d step is typed as a contract, CONTRACT.QFT4D.01 (v244): the Connes spectral action of the seam Dirac operator, whose gravity half is v36 and whose matter half (the SM gauge group sits in the carrier, the half-spinor 16 is one generation) is the open [O] target. Its representation-theory core is now discharged ([verification/v245_spectral_matter_content.py]): the **16** is exactly one *anomaly-free* SM generation ($\sum Y = \sum Y^3 = [SU(2)]^2 U(1) = [SU(3)]^2 U(1) = 0$) with

the forced SM hypercharges and the spectral-action value $\sin^2 \theta_W = 3/8$ ($g_3 = g_2 = \sqrt{5/3} g_1$), so only the seam→Dirac realisation and the run-down couplings stay open. So the residual to a full 4d QFT is *one premise* (QGEO.SYM.01) plus *one contract* (CONTRACT.QFT4D.01), not a diffuse list.

The perturbative 4d S-matrix is constructible ([[verification/v269_spert_paft_skeleton.py](#)]). The 4d scattering layer is typed precisely as three *distinct* objects: S_{top} (the 2d DHR braiding monodromy, $|M|=1$ statistics, v243), S_{pert} (the 4d Epstein–Glaser / BV S-matrix of the spectral action), and S_{phys} (LSZ on the OS-reconstructed Wightman functions, v240) — the 2d braiding is *not* the 4d cross-section S-matrix. [E] the spectral-action a_4 interaction is power-counting renormalizable (every operator mass dimension ≤ 4 : 7 marginal + 3 relevant, a *finite* counterterm basis), so [C] by the Epstein–Glaser theorem (causal perturbation theory; Brunetti–Fredenhagen pAQFT) the perturbative S-matrix $S(g) = T \exp(i \int g L_{\text{int}})$ is built order by order by causal factorisation + finite local extension — no path integral, *no* UV divergences. [C] the gap $\Delta = 6 \log(3/2)$ gives an IR-safe adiabatic limit. [C] this is *perturbative* only: the nonperturbative ambient measure QG.AMB.01 is now a [C] redundancy (v369), a certification object rather than missing dynamics, and the canonical 4d reading remains boundary-only (v265); S_{pert} is the perturbative cross-section layer on top (QFT4D.SPERT.01).

A concrete one-loop EG number ([[verification/v271_eg_oneloop_quartic.py](#)]). The skeleton is taken from “exists” to an actual Epstein–Glaser renormalization: the order-1 product $T_1 = :L_{\text{int}}:$ needs no extension, and the first one is the order-2 T_2 , whose UV behaviour is the Steinmann *scaling degree*. [E] the massless propagator has $\text{sd} = 2$ in $d = 4$, the one-loop bubble (two propagators) has $\text{sd} = 4 = d$, so the singular order $\omega = \text{sd} - d = 0 \Rightarrow$ *exactly one* logarithmic local counterterm ($C(\omega+d, d) = 1$ — finite, renormalizable, no infinite tower); the one-loop factor is $\Omega_3/(2(2\pi)^4) = 1/(16\pi^2)$ exactly, the *same* $1/(16\pi^2)$ that normalises the spectral-action a_4 geometric quartic. [C] the ϕ^4 one-loop running is then $\beta_\lambda = 3\lambda^2/(16\pi^2)$ (symmetry factor 3 = the s, t, u channels), with the EG initial condition fixed by the spectral action, $\lambda_\Phi(\text{UV}) = 1/(16\pi^2)$. [O] this is order-2/one-loop, one diagram class; the all-order perturbative construction stays open (QFT4D.SPERT.02), while the nonperturbative measure is discharged as a [C] redundancy (below).

Red-team caveat, now discharged ([[verification/redteam/rt_F_qft4d.py](#)]). Power-counting renormalizability is *not* unitarity for the gravity sub-sector: the local R^2/Weyl^2 *truncation* gives a four-derivative propagator $1/(p^2(p^2 + M^2))$ whose massive spin-2 mode has a negative residue (the Stelle ghost). But that ghost is a Seeley–DeWitt *truncation artefact*: the entire seam KMS form factor $a(p^2) = e^{p^2/M^2}$ keeps its only pole at $p^2=0$ (v304), the spin-2 sector is ghost-free via the Barnes–Rivers decomposition (v370), and the nearest truncation-zero modulus runs to infinity (v380), so *perturbative* graviton unitarity is established. The matter+gauge S_{pert} is unitary outright; the *nonperturbative* QG.AMB.01 (asymptotic safety v207 / the boundary $(E_8)_1$ net) is then a [C] *redundancy* (v369; gap-decoupled from every TFPT readout, not a TFPT structural item, v335), a certification object rather than a new gap.

The gauge running is an EG order-2 number ([[verification/v273_eg_gauge_running.py](#)]). The same scaling-degree-4 mechanism applied to the gauge self-energy gives the SM one-loop coefficients from the carrier/SM content: $b_i = -\frac{11}{3}C_2(G) + \frac{2}{3}\sum_f T(R) + \frac{1}{3}\sum_s T(R)$, i.e. [E] $b_3 = -7$, $b_2 = -\frac{19}{6}$, $b_1 = \frac{41}{10}$ (GUT-normalised; the same 41 as the carrier algebra $10 b_1 = g_{\text{car}} 2^{g_{\text{car}}-2} + 1$, v159). [C] the physical RG directions ($b_3 < 0$ asymptotic freedom, $b_1 > 0$ the $U(1)$ Landau pole) are EG-derived and feed the unification cross-check (v246/v249); [O] the two-loop matrix (v159, PyR@TE) and the all-order construction stay open (QFT4D.SPERT.03).

The carrier content is one anomaly-free SM generation ([[verification/v310_carrier_sm_anomaly.py](#)]). Those one-loop coefficients sit on top of an exact spectrum/anomaly closure: the carrier half-spinor $16 = \Lambda^{\text{even}}(\mathbb{C}^5) = 1+10+5$ decomposes under $SU(5) \subset SO(10)$ as $10+\bar{5}+1$ = exactly one Standard-Model generation ($Q, u^c, d^c, L, e^c, \nu^c$), and that content is [E] gauge-anomaly-free ($\sum Y = 0$, $\sum Y^3 = 0$, $[SU(2)]^2 U(1) = [SU(3)]^2 U(1) = 0$, exact) while reproducing $(b_1, b_2, b_3) = (\frac{41}{10}, -\frac{19}{6}, -7)$; a fourth family would shift b_1 off $41/10$, so $N_{\text{fam}}=3$ *pins* the RG. [E] for the spectrum, anomalies and one-loop β ; the full Epstein–Glaser interacting S-matrix and the Yukawa sector stay [C]/[O].

The bridge to the physical S-matrix and one-loop unitarity ([[verification/v278_lszy_bridge_unitarity.py](#)]). The perturbative layer connects to the physical asymptotic S-matrix: [E] the EG time-ordered products of S_{pert} are the OS Wick-rotated correlators (v240), so amputating external legs and taking the on-shell residue gives $S_{\text{pert}} \xrightarrow{\text{LSZ}} S_{\text{phys}}$ (the three-layer S-matrix $S_{\text{top}}|S_{\text{pert}} \xrightarrow{\text{LSZ}} S_{\text{phys}}$), valid because the gap $\Delta=6 \log(3/2)$ gives Haag–Ruelle asymptotic states. [E] *one-loop unitarity* is the optical theorem: the s -channel bubble discontinuity

$\text{Im } B(s)/\pi = x_+ - x_- = \sqrt{1 - 4m^2/s}$ is *exactly* the two-body phase space (turning on at the physical threshold $s=4m^2$), so $2 \text{Im } M = \sum \int d\Pi |M|^2$ holds and the matter+gauge sector is unitary order by order. [C] the gravity sector's Stelle ghost is a truncation artefact (perturbative spin-2 unitarity established, v304/v370/v380); the nonperturbative QG.AMB.01 is a [C] redundancy (v369) (QFT4D.SPERT.04).

The first hard data confrontation ([verification/v246_unification_data_crosscheck.py]). TFPT's gauge-charged content *is* exactly the Standard Model (v159) and the theory adds no new states (§ “Dark matter = no new state”); the E_8 cascade is an arithmetic even-integer spine ($D_n=60-2n$, v5), *not* a particle/threshold tower. So a “TFPT-adapted” RGE run *is* the SM run — and the SM does not unify. Tested against the *measured* couplings (v159 betas, plus the standard two-loop gauge matrix): at one loop the pairwise crossings spread over $\sim 10^{13}-10^{17}$ GeV with residual $\Delta\alpha^{-1} \approx 9$ at 2×10^{16} GeV; at two loops (RK4) the minimal spread is $\Delta\alpha^{-1} \approx 3.2$ at $\mu \approx 1.7 \times 10^{14}$ GeV — the three couplings never meet (Fig. 2). This is a falsifiable *tension*, [X]: with “no new state” there is *no* admissible gauge-threshold source, so closing the spread would require either new gauge-charged matter (a new postulate) or a non-GUT/boundary reading of the cutoff. The spectral-action 4d-GUT route is thus in genuine tension — the same status as the Connes–Chamseddine NCG Standard Model — and is likely falsified unless the theory is extended; the AQFT skeleton, DHR structure and SM rep-core survive regardless as an emergent boundary QFT. Recorded as a kill-test, never as a confirmation.

The carrier-native Pati–Salam program ([verification/v247_e8_branching_no126.py], [verification/v248_ps_algebra.py], [verification/v249_ps_unification.py], [verification/v250_dirac_triple_contract.py]). The unification tension above is resolved *natively* if the carrier $SO(10)$ is gauged (the carrier $D_5 \oplus A_3 = SO(10) \times SU(4)$ is Pati–Salam). *First*, the E_8 audit hull *forbids* the **126**: the **248** branches as $(45, 1) + (1, 15) + (10, 6) + (16, 4) + (\overline{16}, \overline{4})$, so the available $SO(10)$ reps are exactly $\{1, 10, 16, 45\}$ — the renormalisable 126_H is algebraically absent, making the “ 16_H over 126_H ” selection a theorem of the hull, not a fit (v247). *Second*, the Pati–Salam NCG algebra $A_F = \mathbb{H}_L \oplus \mathbb{H}_R \oplus M_4(\mathbb{C})$ gives the unimodular group $SU(2)_L \times SU(2)_R \times SU(4)_c$, the exact SM hypercharges $Y = T_{3R} + (B-L)/2$ and the matching $\alpha_1^{-1} = \frac{3}{5}\alpha_{2R}^{-1} + \frac{2}{5}\alpha_4^{-1}$, $\sin^2 \theta_W = 3/8$ (v248). *Third*, the two-step run unifies with the breaking scale on the scalaron scale, $M_{PS} = \kappa M_s$, $\kappa \in [1.0, 1.7]$ at one loop ($\rightarrow [1.0, 1.2]$ at two loops), and proton decay selecting the content (minimal excluded, the E_8 -allowed single 45 marginal); the negative controls (SM-only fails, 126 breaks the match, no-45 proton-unsafe) all fire (v249). *Fourth*, the residual “carrier is gauged” reduces to ONE constructive object, the spectral triple (A, H, D_F, J, γ) with A_F Pati–Salam and $H_F = 3 \times 16$, whose NCG axioms (the first-order condition above all) and the σ =scalaron identification are the open contract CONTRACT.QFT4D.DIRAC.01 (v250). *Fifth*, a first concrete step on that triple is taken ([verification/v251_first_order_sigma.py]): on the lepton block the first-order condition is verified to *hold* for the Yukawa operator and to be *violated exactly* by the Majorana/ σ term — the Chamseddine–Connes–van Suijlekom “beyond first order” mechanism by which the inner fluctuations promote the SM to Pati–Salam and generate σ (honest correction: σ does *not* rescue the first-order condition; its controlled violation *is* the mechanism). *Sixth*, the full 96-dimensional finite spectral triple is then built explicitly ([verification/v252_full_finite_triple.py]) — H_F the particle+antiparticle doubling of three $SO(10)$ **16**-plets — and the NCG axioms are verified by direct computation: reality J , grading γ , KO-dimension 6, order-zero, the first-order condition (*holding* for the SM algebra including the Majorana, *violated* for the Pati–Salam algebra exactly by the Majorana on $\{\nu_R, e_R\}$), and the spectrum of D_F reproducing the fermion masses and the type-I seesaw $m_{\text{light}} = m_D^2/M_R$ — thereby discharging five of the seven obligations of CONTRACT.QFT4D.DIRAC.01 to [E] (orientability, the no-junk condition and the spectral-action expansion remain open). *Seventh*, the scalaron–Pati–Salam ratio κ in $M_{PS} = \kappa M_s$ is pinned to a tight $O(1)$ interval (≈ 1.3) by the RG and shown to be a scale-free ratio of heat-kernel coefficients, $\kappa = \sqrt{(f_2/f_0) c_{PS}/c_{\text{grav}}}$, over the same 96-dim H_F ([verification/v253_kappa_scalaron_ps.py]). *Eighth*, the remaining NCG obligations are then discharged or honestly typed: the inner fluctuations Ω_D^1 give exactly the one SM Higgs doublet $(1, 2)_{1/2}$ with no junk and no **126**, while σ (the ν_R Majorana scalar, $B-L=2$) appears only as the beyond-first-order Pati–Salam field ([verification/v254_inner_fluctuations.py]); the spectral-action Seeley–DeWitt expansion a_0, a_2, a_4 yields the cosmological, Einstein–Hilbert and Yang–Mills+ R^2 +Higgs terms with $b/a^2 = 1/3$, $\sin^2 \theta_W = 3/8$ and the R^2 scalaron, making κ an explicit moment ratio ([verification/v255_spectral_action_expansion.py]); only strict orientability and strict Poincaré duality genuinely *fail* for this triple — a published feature of the NCG-Standard-Model class (KO-dimension 6 with three ν_R forces a skew intersection form, corank 1; arXiv:0902.2068, hep-th/0601192), not a TFPT defect; the physically-correct weakened versions,

quasi-orientability (left/right irrep boxes disjoint by chirality) and *modified Poincaré duality*, do hold ([`verification/v256_orientability_poincare.py`], [`verification/v257_quasiorient_modpd.py`]). *Ninth*, the built D_F is no independent posit but the *modular/covariance induction* of the boundary seam state ([`verification/v258_dirac_covariance_induction.py`]): a quasi-free (KMS, $\beta=1$) state is fixed by its covariance C (a positive contraction), with one-particle modular Hamiltonian $H = \log((1-C)C^{-1})$ (Araki; Casini–Huerta) inverse to the KMS occupation $C = (1+e^H)^{-1}$; inducing the carrier covariance $C_F = P_F C_\Sigma P_F$ returns the `v252` operator *exactly* — same spectrum, chirality-odd, the Majorana an off-diagonal covariance entry with the seesaw m_D^2/M_R recovered — so D_F is the Kasparov shadow $[D_\Sigma] \hat{\otimes}_{(E_8)_1} [K_{\text{car}}]$ of the boundary geometry and the `v18` Yukawas are the *readout* of C_Σ , [`E`] modulo the one [`O`] seam realisation. *Tenth*, the last open spectral-action input — the cutoff function f — is fixed by the same seam state ([`verification/v259_modular_cutoff_kappa.py`]): by Tomita–Takesaki/Bisognano–Wichmann the seam is KMS at $\beta=1$ (`v239`) and the seam unit $2\pi=1/(4c_3)$ (`v58`) fixes that β , so the heat-kernel cutoff $f(u) = e^{-u}$ is *derived*, not chosen, with moments $f_2/f_0 = f_4/f_2 = 1$ *exactly* (a generic Gaussian cutoff gives $\sqrt{\pi}/2 \neq 1$, so the choice is non-vacuous); hence $\kappa = \sqrt{(f_2/f_0) c_{PS}/c_{\text{grav}}} = \sqrt{c_{PS}/c_{\text{grav}}}$ loses its scheme factor and the `v253/v255` open [`O`] cutoff freedom becomes a [`C`] finite trace ratio over the 96-dim H_F (the RG $\kappa \approx 1.3$ pinning $c_{PS}/c_{\text{grav}} \approx 1.7$). So unification is now *one* theory decision (gauge $SO(10)$?) plus *one* constructive object, both falsifiable — not a fog. The colour $SU(4)_c$ here sits inside $SO(10)=D_5$ and is distinct from the carrier family factor A_3 .

The implication is now Lean-formalised (FORM.QGEO.01). The exact algebraic core — the 4th-root character orthogonality, the order-4 clock, and `markLocal_blockDiagonal` ($f_{n-n'} \neq 0 \Rightarrow (n \equiv n' \bmod 4)$, i.e. $[\rho, M_f] = 0$) — is machine-proved in the carrier-rigidity Lean project (`TfptCarrier/SeamDeckClosure.lean`; AUDIT: PASS, axioms `propext`, `Classical.choice`, `Quot.sound` only). The premise itself is a typed `structure SeamDeckPremise` — consumed, not proved, exactly as `CalderonProjector` encodes the document-1 analytic input — so the conditional theorem “`SeamDeckPremise  $\Rightarrow \omega \circ \rho = \omega$` ” is now [`E`] while `QGEO.SYM.01` stays the one open [`O`] postulate.

The two residuals are one canonical metric: the pillowcase ([`verification/v214_seam_pillowcase.py`]). The isometry residual `QGEO.ISO.01` (`v180`) and the mark-locality residual `QGEO.SUBPRIN.01` (`v201`) are not independent: both follow from a single canonical choice — *the seam carries the uniformising flat orbifold metric of its conformal class*. The four marks make the orbifold sphere $S^2(2, 2, 2, 2)$, whose orbifold Euler characteristic $\chi_{\text{orb}} = 2 - 4(1 - \frac{1}{2}) = 0$ is *Euclidean*: by Troyanov’s cone-metric uniformisation it carries a flat metric, unique up to scale, with curvature concentrated at the four cone points (Gauss–Bonnet: four deficits π sum to $4\pi = 2\pi\chi(S^2)$) — this *is* the `v201` “flat away from the marks”. And the order-4 clock is no longer an extra assumption but a *consequence* of the cross-ratio 2 that `v168` already proved: $j(\lambda) = 256(\lambda^2 - \lambda + 1)^3/(\lambda^2(\lambda - 1)^2)$ gives $j(2) = 1728$, so the modulus is the *square* torus $\tau = i$, whose complex multiplication by $\mathbb{Z}[i]$ supplies the order-4 isometry $z \mapsto iz$ (explicitly $(x, y) \mapsto (-x, iy)$ on $y^2 = x^3 - x$, order 4); a generic cross-ratio gives $j \notin \{0, 1728\}$ and only $\mathbb{Z}/2$, so the clock is specific to μ_4 . Thus the two open `QGEO` residuals *collapse to one* canonical-metric statement and `QGEO.ISO.01` ceases to be an imposed symmetry. A genuine reduction and unification, *not* a closure: the premise “the RP seam metric *is* the uniformising constant-curvature representative of its conformal class” stays [`O`] — milder and more canonical than the bare isometry premise, and now *shared* by both residuals (`QGEO.PILLOW.01`).

The converse, on the canonical metric: raw seam $\Rightarrow$ mark-locality ([`verification/v264_raw_seam_marklocal.py`]). Composing the chain explicitly closes the gap between `v216`, `v201` and `v198` on the uniformising metric: the flat pillowcase has its curvature *only* at the four cone points (`v216`), so the DtN sub-principal symbol M_f is a μ_4 -orbit sum, hence $\mathbb{Z}_4$ -invariant (Fourier support on $4\mathbb{Z}$, verified by FFT), hence block-diagonal in the clock-character basis, hence $[\rho, M_f] = 0$; with the principal symbol $|k|$ already commuting (`v198`) this gives $[\rho, \Lambda] = 0$ and $\omega \circ \rho = \omega$ *on this metric*. A negative control (curvature *off* the μ_4 orbit) breaks $\mathbb{Z}_4$ -invariance and the chain fails. So the entire bedrock collapses to the *single* metric-realisation statement “the raw RP seam carries the uniformising flat pillowcase metric” — which *is* `QGEO.SYM.01`. A genuine reduction/sharpening (`QGEO.ENERGY.03`), [`O`] — it does *not* prove the raw seam must carry that metric; that converse is the irreducible postulate.

The minimal axiom: a rigidity theorem ([`verification/v267_qgeo_rigidity_minimal_axiom.py`]). One can sharpen the bedrock to its irreducible kernel. [`E`] the *single* bare statement “the seam admits the carrier’s order-4 clock as a conformal symmetry cyclically permuting its four marks” forces everything: an order-4 Möbius orbit $\{a, ia, -a, -ia\}$ has cross-ratio 2 *independent of a* ; cross-ratio 2 $\Leftrightarrow j=1728$ (only

$\lambda \in \{-1, \frac{1}{2}, 2\} \Leftrightarrow$ the order-4 CM automorphism $(x, y) \mapsto (-x, iy)$ on $y^2 = x^3 - x$; and from the square configuration the unique flat pillowcase metric (Trojanov), mark-locality, $[\rho, \Lambda] = 0$ and $\omega \circ \rho = \omega$ all follow (v214/v201/v264/v198). [E] neg controls: a generic four-mark set has $j \neq 1728$ and only a $\mathbb{Z}_2$ symmetry (no clock), the hexagonal point $j=0$ has a $\mathbb{Z}_6$ symmetry (order $6 \neq 4$) — only the square (order-4) point realises the μ_4 clock. [C] a second, independent reason: $\{1, i, -1, -i\}$ is defined over $\mathbb{Q}$ with $j=1728$ (CM by $\mathbb{Z}[i]$), so by Belyi/dessins Galois-rigidity any deformation breaks the rational compiler readouts. [O] So QGEO.SYM.01 is now in its minimal sharpest form — the bare existence of that order-4 conformal symmetry, which (like $c=\text{const}$ in relativity) cannot be derived from nothing. The *reduction* is closed (axiom $\Rightarrow$ everything); the axiom itself stays the one foundational postulate (QGEO.SYM.02).

The flat geometry closes the commutator to all orders ([verification/v276_qgeo_flat_closes_commutator.py]). The state-level residual $[\rho, H] = 0$ — the one operator equation the bedrock rests on — is *not* a separate analytic gap once the flat $\tau=i$ pillowcase is granted. [E] if $H = f(\Delta_{\text{flat}})$ is a function of the intrinsic flat Laplacian (the DtN of a flat metric, no curvature corrections) and the order-4 deck ρ ($z \mapsto iz$) is a flat- $\tau=i$ isometry, then $[\rho, \Delta] = 0 \Rightarrow [\rho, H] = 0$ *exactly*, to all orders (an isometry commutes with every spectral function of its Laplacian; verified on the discretised square torus, with a generic non-isometry as a failing control). This *upgrades* the leading-order commutation v198 (principal symbol) and v201 (subprincipal) to the full H — flatness removes the lower-order terms. [O] so QGEO.SYM.01 collapses to a *single geometric* statement, “the raw RP seam state is the flat $\tau=i$ pillowcase quasi-free state ($H = \sqrt{-\Delta_{\text{flat}}}$)”; everything downstream is [E], and the remaining human obligation is one constructive-QFT proof that the raw seam DtN state *is* that flat geometry (QGEO.SYM.03).

The obligation as one precise lemma, Lean-backed ([verification/v279_qgeo_obligation_lemma.py]). The bedrock is now written as a single constructive-QFT *lemma*: *given* the premise “the raw seam carries the flat $\tau=i$ pillowcase metric g_{flat} (Trojanov-unique, cone angles π), so $\Lambda_{\Sigma} = \sqrt{-\Delta_{g_{\text{flat}}}}$ ”, the conclusion $\omega_{\Sigma} \circ \rho = \omega_{\Sigma}$ (and hence mark-locality, the metric inclusion, E_8) follows. [E] the reduction DAG from “free RP seam” to this premise has *exactly one* un-discharged leaf — “the raw seam state *is* the flat state” — every other node being closed; and the geometric side is fixed (the flat DtN is Fourier-diagonal), so the obligation is a *state-identification*, not a geometric ambiguity, with two candidate routes ([C]): RP-state uniqueness on the flat $\tau=i$ orbifold, or Trojanov + OS reconstruction. [E] the all-orders closure step is now a *Lean theorem*: `TfptCarrier/SeamDeckClosure.lean (flat_all_orders_clock)` proves that any spectral function $g(\Delta_{\text{flat}})$ of the Fourier-diagonal flat Laplacian commutes with the carrier clock $\rho = \text{diag}(i^n)$ — upgrading v198/v201 to the full operator with only the standard axioms, no *sorry* (FORM.QGEO.03). [O] the one premise stays the honest axiom (QGEO.OBLIG.01).

Self-investigation: the experiment and the simplification ([verification/v280_pillowcase_steklov.py], [verification/v282_e8_tau_i_unification.py]). Two things can be done *ourselves*. (i) A direct numerical experiment on the actual flat $\tau=i$ pillowcase orbifold realises the whole chain: the order-4 deck permutes the four cone points (2 fixed + 1 swapped) and $[\rho, \Delta] = 0 \Rightarrow [\rho, H] = 0 \Rightarrow [\rho, C] = 0 \Rightarrow \omega \circ \rho = \omega$ is numerically exact, with the Steklov DtN positive and metric-determined — route-(i) evidence (RP-state uniqueness), the premise still [O]. (ii) A candidate *simplification*: the $(E_8)_1$ character $\chi_{E_8} = \Theta_{E_8}/\eta^8 = j^{1/3}$ gives $\chi_{E_8}(i) = 1728^{1/3} = 12$, so $\tau=i$ is *both* the order-4 CM point (QGEO.SYM.01, flat geometry) *and* the $(E_8)_1$ partition-function modulus (QG.AMB.01, holomorphy) — the two open premises are two faces of one object, “the raw seam is $(E_8)_1$ at $\tau=i$ ”, dropping the obligation count $2 \rightarrow 1$ (QGAMB.UNIFY.01).

The two proof routes, decomposed ([verification/v284_route_i_rp_uniqueness.py], [verification/v285_route_ii_seam_condensation.py]). The keystone — now *closed modulo a cited theorem* via the lattice/ S_3 stack (below) — can be attacked from two sides, and both are now lemma chains with the dischargeable steps discharged. *Route (i)* (RP-state uniqueness) is a six-lemma chain — [E] five are already [E]/[O] (RP $\rightarrow$ DtN v194; marks $\rightarrow$ pillowcase v195/v216; the flat Steklov operator $= \sqrt{-\Delta_{\text{flat}}}$, principal symbol $|\xi|$ with no curvature term; covariance from the DtN v258; μ_4 -invariance v276), and [C] Trojanov ($\chi_{\text{orb}}=0$) reduces the residual to “the raw seam is the constant-curvature geometric state”. *Route (ii)* (seam-net condensation) is a four-lemma chain — [E] three are discharged (no abelian sector $\Leftrightarrow \det K=1 \Leftrightarrow$ holomorphic $\Leftrightarrow E_8$ v234/v277/v281; the tower $16 \rightarrow 4 \rightarrow 1$ v92), leaving “the free RP seam condenses the μ_4 -Lagrangian glue” = QGEO.SYM.01. [E] **the two open lemmas coincide** — both say “the raw seam is the $(E_8)_1$ at $\tau=i$ ” (v282), so the canonical equivalence between the raw seam KMS/DtN state and the holomorphic $(E_8)_1$ net at $\tau=i$ is the *single* [O] statement that closes both (QGEO.ROUTEI.01/QGAMB.ROUTEII.01).

The Seam Equivalence Theorem, named and firewalled ([[verification/v286_seam_equivalence_contract.py](#)]). That single statement is now the named target **SEAM.EQUIV.01**: *the raw RP seam state reconstructs canonically the holomorphic $(E_8)_1$ boundary net at the $\tau=i$ pillowcase* (respecting the DtN, the KMS flow, the μ_4 deck, the Q-system extension, $\det K=1$ and the $(E_8)_1$ character). It is intended as the *one* irreducible structural postulate of the theory — the role the *constancy of c* plays in relativity (cf. v267): not a gap to be patched but the single sharply-localised, falsifiable premise on which the whole boundary closure turns. [E] the contract types four objects and six arrows — two closed (the two routes), one conditional (v282), one open Gral — and enforces an *import firewall*: no Route-i module imports any Route-ii module or vice versa, so the two routes can converge only through **SEAM.EQUIV.01** (this kills the “ E_8 smuggled into the geometry and pulled back out” attack). The two proof routes are then attacked directly: [E] Route A (AQFT, [[verification/v287_free_rp_bulk_to_holomorphic_boundary.py](#)]) reduces the Gral to *one* standard theorem, “invertible (SRE) Gaussian bulk $\Rightarrow$ single-sector boundary” (4/5 lemmas discharged); and [E] Route B (DtN, [[verification/v288_full_l2_subprincipal_z4.py](#)]) *proves* the full- L^2 lift of the sub-principal $\mathbb{Z}_4$ block-diagonality (a mark-sourced curvature has Fourier support only on $m \equiv 0 \pmod{4}$), so $[\rho, \Lambda]=0$ on all of L^2), shrinking the residual to the single sharper question “why is the raw seam sub-principal term mark-local?” (**SEAM.EQUIV.A01/SEAM.EQUIV.B01**).

Route B sharpened to one analytic lemma; both routes precisely scoped. That last question is decomposed ([[verification/v289_marklocal_raw.py](#)], **MARKLOCAL.RAW.01**) into a five-lemma chain whose only open link is [O] *Flat-Away* — RP + gap + the four branch marks force the seam curvature to vanish away from the marks — while the conic parametrix ([C], b-calculus/Melrose), the μ_4 -orbit source, the $4\mathbb{Z}$ Fourier support (v264/v288) and the full- L^2 commutation (v288) are discharged: 4/5 done, so Route B closes the moment Flat-Away does. The two claims are firewalled against over-reading: Route B does *not* assert “full- L^2 QGEO” — its exact-label split is Fourier-lemma (Formal/Exact), operator test (Numerical Exact), transfer to the raw conic DtN (Conditional, v264), mark-locality from the raw seam (Open Selection, v289); and Route A’s “one standard theorem” is itself a composition of citable results (Kitaev/Freed–Hopkins, Muger/Kawahigashi–Longo–Muger, Kawahigashi–Longo, [[verification/v287_free_rp_bulk_to_holomorphic_boundary.py](#)]) whose only open hypothesis — seam-bulk invertibility — coincides with Route B’s Flat-Away. Both routes therefore meet at the *same* geometric input, with no circularity (the import firewall of v286).

Red-team: $\mathbb{Z}_4$ -invariance is not mark-locality. A deliberate adversary ([[verification/v290_z4_smooth_curvature_adversary.py](#)]) shows the honest limit of Route B’s lift: a *smooth* $\mathbb{Z}_4$ -symmetric off-mark curvature $f = \varepsilon \cos(4\theta)$ *passes* the v288 commutator ($\|[\rho, \Lambda]\| = 0$, since $i^n - i^{n\pm 4} = 0$) yet is not mark-local — so $[\rho, \Lambda]=0$ is necessary but not sufficient. The modulus is instead excluded by the spectral data: it shifts the DtN spectrum off the flat ladder and the heat trace ($\max|\Delta\lambda| = 0.155$, $\Delta\text{Tr} = 0.019$ at $\varepsilon = 0.3$), i.e. the KMS weight and $(E_8)_1$ character that **SEAM.EQUIV.01** must match. Consequently the open lemma is named in its own right ([[verification/v291_flataway_contract.py](#)], **FLATAWAY.RP.01**): *the smooth curvature density in the DtN sub-principal symbol vanishes away from the four branch marks*, with three independent attack routes — spectral-rigidity and heat-kernel (the flat seam is the *unique* spectral/heat match of the smooth $\mathbb{Z}_4$ family) and RP-energy/Troyanov (the constant-curvature minimiser) — closing any one of which closes Route B.

All three routes attacked; Flat-Away reduced to one selection principle. Each route is now carried to its precise reduction: [E] the *heat* route ([[verification/v292_flataway_heat_reduction.py](#)]) shows the heat-trace deviation $\Delta\text{Tr}(e^{-t\Lambda}) = c(t)\varepsilon^2 + O(\varepsilon^4)$ is a positive-definite quadratic form in the off-mark curvature (leading invariant $\|f\|_{L^2}$), so Flat-Away $\Leftarrow$ “the seam fixes the a_2 heat coefficient”; [E] the *spectral* route ([[verification/v293_flataway_spectral_hessian.py](#)]) lifts this to the full smooth- $\mathbb{Z}_4$ space — mode $4k$ splits the degenerate Steklov level $|2k|$, the mismatch Hessian over $\{4, 8, 12\}$ is positive-definite ($\{0.497, 0.500, 0.503\}$), so the flat seam is a *strict* isolated minimum (answering “only one direction”); and [E]/[C] the *variational* route ([[verification/v294_flataway_rp_energy.py](#)]) gives, with the rigid mark divisor and Troyanov uniqueness, the flat cone metric as the unique strictly-convex curvature-energy minimiser. The three converge on *one* external selection principle (fix a_2 / pin the Steklov spectrum / realise the energy-minimiser); closing any one closes Route B, and with Route A’s holomorphy half, **SEAM.EQUIV.01**.

The heat route made exact and Lean-formalised. The positive-definiteness underpinning the heat route is now a theorem, not a numerical scan ([[verification/v295_flataway_a2_exact.py](#)]): the first variation of $\text{Tr } e^{-t\Lambda}$ vanishes (the smooth off-mark curvature has zero mean, Gauss-Bonnet), and since $\Lambda \mapsto \text{Tr } e^{-t\Lambda}$ is convex (Klein/Peierls), $\varepsilon=0$ is the *global* minimum — so $\Delta\text{Tr} \geq 0$ for every deformation, $= 0$ iff

$f = 0$; the exact second-order coefficient is the divided-difference kernel of e^{-tx} (validated to 10^{-6}). The algebraic $\Delta=0 \Leftrightarrow$ flat step is machine-formalised in Lean (`FORM.FLATAWAY.01`, `SeamDeckClosure.lean`: a positive-weighted sum of squares is positive-definite; axiom-clean). So Flat-Away’s heat route is reduced to the single external invariant “the seam fixes a_2 ”, with its positivity both proved and formalised. The a_2 coefficient itself is now in *closed form* (`[verification/v296_flataway_a2_closed.py]`): the operator tails telescope geometrically to $W_k^{\text{tail}} = t(1 - e^{-4kt})/(8k(1 - e^{-t}))$, leaving a finite $(4k-1)$ -term middle, e.g. $W_1(t) = t(2te^t + e^{3t} + 3e^{2t} + e^t - 1)e^{-3t}/8$ (manifestly positive; small- t $W_k = t/2 + O(t^3)$). Symmetrically, Route A’s “one standard import” is now a precise, citable theorem stack (`[verification/v297_route_a_literature_stack.py]`): LIT-A (Kitaev; Freed–Hopkins) $\rightarrow$ LIT-B (Müger; Kawahigashi–Longo–Müger) $\rightarrow$ LIT-C (Conway–Sloane; Dong–Xu) compose into “invertible bulk $\Rightarrow (E_8)_1$ ”, all three legs established, the one shared hypothesis (seam-bulk invertibility) coinciding with Route B’s Flat-Away — so both routes are reduced to citable/established machinery plus the *one* shared geometric input, which the lattice/ S_3 stack (below) now pins at every computable level, making the keystone *closed modulo a cited theorem*.

Flat-Away hardened, and the pin derived from $(E_8)_1$ (`[verification/v300_flataway_rigid.py]`, `FLATAWAY.RIGID.01`). The spectral route was a *soft* minimum (“the mismatch grows with ε ”); it is now a *discrete* obstruction. [E] the flat fingerprint is the integer Steklov ladder with every level $n \geq 1$ of multiplicity exactly 2; [E] the resonant mode $4k$ splits level $2k$ by gap $= g_k$ (the exact 2×2 block $[2k, g/2], [g/2, 2k] \Rightarrow 2k \pm g/2$, validated against the full Toeplitz operator), so [E] any $g_k \neq 0$ lifts the degeneracy ($2 \rightarrow 1+1$) and changes the spectral *multiset* — preserving the flat ladder *with its multiplicities* forces every $g_k = 0$, a degeneracy obstruction robust to rescaling and strictly stronger than the v293 “deviation grows” scan; with the v295 convexity, flat is isolated analytically *and* discretely. Moreover the external pin is now *derived*: [C] the $(E_8)_1$ vacuum character $E_4/\eta^8 = q^{-1/3}(1 + 248q + 4124q^2 + \dots)$ has integer q -coefficients (E_8 even unimodular $\Rightarrow$ integer conformal weights), so by 2d Steklov conformal rigidity the integer ladder $\Leftrightarrow$ flat seam — i.e. “the seam fixes its spectrum” *is* the $(E_8)_1$ rationality of Route A. [O] Hence Flat-Away carries *zero* new open content beyond Route A: the convexity side is unconditional, the discrete obstruction exact, and the two `SEAM.EQUIV.01` routes provably converge on one rationality fact (consistent with the v286 firewall).

Route A’s last hypothesis discharged — the gapped quasi-free bulk is invertible (`[verification/v301_route_a_invertible.py]`, `SEAM.EQUIV.A.INV.01`). Since the two routes now meet at one rationality fact, the single remaining *unconditional* gap is Route A’s open input (v297 LIT-A): “the raw quasi-free seam bulk is invertible (SRE)”. A *gapped* free-fermion (Gaussian) 2+1D bulk cannot carry intrinsic topological order — free fermions realise only the integer (IQHE/Chern) phases, all *invertible*; intrinsic topological order requires interactions (Kitaev’s periodic table, class D, the only invariant being the chiral central charge). [E] the carrier is 16 Majoranas, so $c = g_{\text{car}} + N_{\text{fam}} = 8 = 16/2$ (v148); [E] $SO(16)_1 \supset (E_8)_1$ ($\dim \mathfrak{so}(16) = 120$ currents + 128 spinor = 248, v154); [C] hence a gapped quasi-free bulk is invertible automatically, and [E] the invariant $|\det K| = \# \text{anyons} = 1$ for E_8 (invertible) against 4 for the gauged $D_8=SO(16)$ (anyons), which free fermions cannot produce. [O] So LIT-A’s ‘invertible bulk’ *follows* from ‘gapped 16-Majorana quasi-free bulk’ (v148/v160 + the mass gap): the residual changes from a topological-order statement to a spectral one. *Bottom line*: together with v300, the entire open content of `SEAM.EQUIV.01` reduces to the single spectral statement “the quasi-free seam bulk is gapped” — the topological-order obstruction (hidden anyons) is removed, and the gap itself is not manufactured.

The last lever — the seam gap is the derived Recovery gap (`[verification/v302_seam_gap.py]`, `SEAM.EQUIV.GAP.01`). That single remaining input is not free: the seam transfer operator has the frozen spectrum $\text{spec}(T) = \{1, (2/3)^6, (1/3)^6\}$ (v160/v162), a simple Perron eigenvalue 1 separated from the rest, and the separation is the *Recovery gap* $\Delta = -\ln((2/3)^6) = 6 \ln(3/2) = 2N_{\text{fam}} \ln(3/2) \approx 2.4328 > 0$ — a number forced by the carrier $((2/3)^6$ the μ_4 -deck transfer eigenvalue), with the decoupling margin $\Delta - 31/(4\pi^2) \approx 1.648 > 0$ (v76) making it robust. [C] By Osterwalder–Schrader reconstruction / quasi-free clustering a transfer-operator gap $\Delta > 0$ *is* a mass gap of the reconstructed Hamiltonian (correlation length $1/\Delta$). *Bottom line*: with v300/v301, the whole residual of `SEAM.EQUIV.01` is now a composition of standard cited theorems (OS/quasi-free clustering + Kitaev free-fermion invertibility + the v297 AQFT stack) over established TFPT facts (16 quasi-free Majoranas; transfer gap $6 \ln(3/2)$): *no undischarged TFPT-internal assumption remains*. It stays [O] (not machine-proved end-to-end), but the open front is reduced to citable machinery over a spectral gap that is provably positive.

The continuum chain assembled, and the residual tied to the seam one-sidedness (`[verification/v356_continuum_mmst_applicability.py]`, `SEAM.EQUIV.CONTINUUM.03`). Writing the

chain out link-by-link: **S1** the gapped (v302) quasi-free 16-Majorana collar $\rightarrow$ **S2** invertible (Kitaev free fermions, v301) $\rightarrow$ **S3** the *nontrivial* invertible phase ($c_- = 8 \neq 0$) $\rightarrow$ **S4** a non-gappable anomalous chiral edge (anomaly inflow / bulk-edge) $\rightarrow$ **S5** its chiral-CFT scaling limit (MMST, v336; range $\text{rank} \leq c \leq D$, i.e. $8 \leq 8 \leq 16$) $\rightarrow$ **S6** the order-4 μ_4 clock + the AQFT stack pin $(E_8)_1$ (v351/v297). Every link is a theorem or an established TFPT fact *except* S3, and S3 is precisely the seam being *one-sided* (chiral, $c_- \neq 0$) — the *same* one-sidedness that makes $c_3 = 1/(8\pi)$ a one-sided Gauss–Bonnet ($|\mathbb{Z}_2| \nmid K$, v58). So the lone realization input is tied to the constant that *defines* the theory, not a free dial; the residual is now “the abstract collar is realized as a genuine lattice chiral free-fermion invertible phase”, [O] but theorem-bounded on both sides.

MMST in-class verified for the collar, so the scaling limit is $(E_8)_1$ ([verification/v366_mmst_seam_collar.py], SEAM.MMST.INCLASS.01). The continuum leg is now checked *hypothesis-by-hypothesis* for the explicit seam collar: it is a $D = \dim S^+ = 16$ Majorana quasi-free (CAR) state (v155/v160), gapped with $\Delta = 6 \log \frac{3}{2} > 0$ *derived* (v302), and chiral with $c_- = D/2 = 8 = g_{\text{car}} + N_{\text{fam}}$, so it sits squarely in the MMST free-lattice-fermion class with $\text{rank} \leq c \leq D$ reading $8 \leq 8 \leq 16$ (*in* range, saturating $\text{rank} = c = 8$). The MMST massless-scaling-limit theorem then applies as a cited result, and the order-4 μ_4 clock pins the target to $(E_8)_1$ ($\det K = 1$, one primary) over the same- c rival $SO(16)_1$ ($\det K = 4$, v351). [E] the class arithmetic, the derived gap, the chirality and the $(E_8)_1$ pinning; [C] the CAR-class typing and the cited MMST theorem; [O] the lone residual is exactly S3 above (the lattice realization). The continuum side of SEAM.EQUIV.01 is thereby reduced to a single, named, theorem-bounded input. The lattice $\rightarrow$ conformal-net step itself is now a *rigorous AQFT construction*: Adamo–Giorgetti–Tanimoto (arXiv:2506.01008, 2025) build the two-dimensional conformal net $\mathcal{A}_Q$ of an even lattice Q and *classify* every two-dimensional extension of the Heisenberg net (under a discreteness assumption) as such a lattice net — so “the lattice is a conformal net” is a theorem of the same operator-algebraic lineage the continuum leg invokes, not a hope; and since (xxxii) ([verification/v469_seam_crossedproduct_route.py]) the extension step no longer even *needs* the preprint lineage — the local $\mathbb{Z}_2$ simple-current crossed product ($h_s = 1 \in \mathbb{Z}$, Longo–Rehren 1995; Böckenhauer 1996; Böckenhauer–Evans 1998; KLM 2001) certifies it on peer-reviewed ground, with AGT/AMT as the independent second witness.

S3 made concrete: an explicit gapped chiral-Majorana lattice model ([verification/v367_seam_s3_lattice.py], [verification/v368_seam_s3_inflow.py]). The S3 input is no longer abstract. The $p+ip$ Bloch model $h(k) = \sin k_x \sigma_x + \sin k_y \sigma_y + (M - \cos k_x - \cos k_y) \sigma_z$ is, at $M=1$, a gapped free-fermion phase with numerically-computed Chern number $|C|=1$ ($C=0$ trivially at $M=3$), so $N_{\text{Maj}} = \dim S^+ = 16$ copies give a chiral edge of $c_- = 16/2 = 8 = g_{\text{car}} + N_{\text{fam}}$, with the edge K-matrix $SO(16)_1$ ($\det 4$) condensed to $(E_8)_1$ ($\det 1$) by the order-4 μ_4 clock (v351/v369). A finite STRIP of the same model carries an edge-localised in-gap chiral branch (absent in the trivial phase) — so by anomaly inflow the chiral edge *exists* as a consequence of the invertible bulk (the S4 leg), discharging the “does the edge exist?” question on the lattice. [E] the gapped Chern model, the $c_- = 8$ counting and the edge existence; [C] the μ_4 condensation to $(E_8)_1$; [O] only the genuine *continuum scaling limit* (MMST, v336) stays open — S3 is a runnable lattice model that pins the target at every computable level, so the keystone is *closed modulo a cited theorem*.

The S3 closure stack: $c = 8$, the $(E_8)_1$ character, the torus count and reflection positivity — all computed on the lattice ([verification/v376_seam_s3_centralcharge.py], [verification/v377_seam_s3_e8character.py], [verification/v378_seam_s3_modular.py], [verification/v379_seam_s3_rp.py]). On that explicit model the target of the scaling limit is now pinned at every level a computation can reach. *Central charge* (v376): the Calabrese–Cardy finite-size entanglement scaling gives $c = 1.0000$ per critical complex mode, so the 16-Majorana ($= 8$ complex) collar has $c = 8 = g_{\text{car}} + N_{\text{fam}}$, and it is chiral ($c_- = 8$). *Convergence* ([verification/v392_seam_s3_scalinglimit.py]): across five sizes $L = 66 \dots 514$ the slope $c_1(L)$ is a monotone Cauchy-like sequence approaching 1 and $1/L$ -Richardson-extrapolates to $c_1(\infty) \approx 1.000$ ($c \approx 8.00$ in the limit) — the existence-relevant *convergence* the two-size value did not test (evidence, not a proof of existence; the residual stays [O]). *The net* (v377): the exact affine character $\chi = E_4/\eta^8 = q^{-1/3}(1 + 248q + 4124q^2 + \dots)$ gives $c = 8$ and 248 level-1 currents (one primary), and the μ_4 /GSO promotion $248 = 240 + 8 = 120 + 128$ ($240 = 16 \cdot 5 \cdot 3$) pins $(E_8)_1$ over the same- c rival $SO(16)_1$ (120 currents, four primaries). *Genus-1* (v378): the KLM condensation $\mu(\text{carrier}) = 16 \rightarrow \mu(E_8) = 16/4^2 = 1$ and the single character’s modular T -eigenvalue $e^{-2\pi i/3}$ give torus ground-state degeneracy 1 — the holomorphy / $\det K = 1$ statement that the genus-0 arguments could not reach (v344). *Reflection positivity* (v379): the gapped collar’s Euclidean two-point function has a positive spectral measure, so the OS reflection Gram matrix is positive semidefinite (a ghost mode breaks it). [E] all four (numerical

c , exact character, genus-1 count, RP); [O] the one thing a computation cannot supply is the abstract *continuum existence* of the limit (the cited MMST theorem, v336). With RP in hand the ambient measure C7 is then discharged by the redundancy theorem (v369), so S3 is the master key: the target is pinned exactly, only the existence proof is external.

And that external step is now Lean-formalised as “closed modulo a cited theorem” (FORM.SEAM.MMST.01, TftptCarrier/SeamScalingLimit.lean). The continuum leg is machine-formalised in the audit-contract style of SeamEquivChain.lean: the collar’s MMST applicability hypotheses are *provable* arithmetic facts by Lean kernel `decide` (no axioms) — $D=16$ Majoranas, $\text{rank}=c=8$, the range $8 \leq 8 \leq 16$, the chirality $c_- = 8 \neq 0$, and the holomorphy discriminator $\det K=1$ vs $SO(16)$ ’s 4; the MMST scaling-limit theorem and the Adamo–Moriwaki–Tanimoto OS reconstruction are declared as *named cited axioms*; and the conclusion $(E_8)_1$ follows by a well-typed composition whose `#print axioms` pins the residual to exactly those two published theorems (plus the carrier gap), with no hidden assumption and no `sorry`. Every TFPT-internal hypothesis is thus kernel-checked; only the two cited theorems are assumed — the honest closure ceiling for SEAM.EQUIV.01. *Two of the seven convergent routes below are now kernel-hardened the same way:* FORM.SEAM.BW.HSMI.01 (TftptCarrier/SeamStandardPair.lean) proves the four Borchers/Wiesbrock *standard-pair relations* exactly over $\mathbb{Z}$ on the centred ladder ($[K, P]=P$, $J^2=1$, $JKJ=-K$, $JPJ=P_+$, all by kernel `decide`, no domain axiom), with the cited Borchers step and the one continuum input as named axioms; and FORM.SEAM.APPLICABILITY.01 (TftptCarrier/SeamApplicabilityLedger.lean) kernel-checks the audit *count* itself — of the 12 cited-theorem hypotheses exactly 11 are internal and 1 external (and $8 \leq 8 \leq 16$, $\det K=1 \neq 4$), all axiom-free — so the very bookkeeping the keystone’s honesty rests on is machine-verified, leaving only the one named continuum fact open.

TOE-roadmap follow-ups — the keystone as a typed certificate, the bedrock modular form, and the gap-decoupled measure ([verification/v308_seam_equiv_chain.py], [verification/v309_modular_bedrock.py], [verification/v311_gap_decoupled_measure.py]). Three machine-checked steps push the keystone arc further. (i) *Composition certificate* ([verification/v308_seam_equiv_chain.py]): the closing implication is one well-typed logical chain $\text{gap}>0 \rightarrow \text{invertible} \rightarrow \text{holomorphic } c=8 \rightarrow (E_8)_1$ whose arithmetic discriminators are exact — $|\det \text{Cartan}(E_8)|=1$ against $|\det \text{Cartan}(D_8)|=4$, the carrier tower $4 \cdot 4=16 \rightarrow 4 \rightarrow 1$, $c=g_{\text{car}}+N_{\text{fam}}=8$ — so the only undischarged premise is the cited OS step; [E] the discriminators, [O] the keystone. (ii) *Modular bedrock* ([verification/v309_modular_bedrock.py]): with the four marks derived from Gauss–Bonnet (v216) and the order-4 clock from cross-ratio 2 (v214), the quasi-free modular Hamiltonian $K = \log((1-C)/C)$ is built explicitly and $[\rho, K]=0 \Leftrightarrow \omega \circ \rho = \omega$ (a $\mathbb{Z}_2$ profile breaks it), so the bedrock reduces *non-circularly* to the Bisognano–Wichmann intrinsicity of the raw collar state; [E] the model, [O] the intrinsicity. (iii) *Gap-decoupled measure* ([verification/v311_gap_decoupled_measure.py]): the admissible sector clusters exponentially at rate $(2/3)^6$ (correlation length $1/\Delta$, finite susceptibility 729/665) and the v76 margin $\Delta - 31/(4\pi^2) > 0$ decouples it from the ambient, so the physical sector is a convergent object while QG.AMB.01 is a [C] redundancy (v369). Each pins the residual to a single named step; with the lattice/ S_3 stack the keystone is now *closed modulo a cited theorem*. (iv) *The intrinsic-BW residual as a cited reduction* ([verification/v424_seam_lto_rp_reduction.py]): the open bedrock lemma of (ii) — the raw collar’s intrinsic Bisognano–Wichmann property — is reduced to the recent commuting-projector theorem of Naaijken–Penneys–Wallick (arXiv:2605.10693), whose reflection-positivity axiom (LTO-RP) is exactly $u_\Theta=J$ (the reflection *is* the modular conjugation), i.e. the BW condition $\theta K \theta = -K$ — *not* $\omega \circ \rho = \omega$. [E] the μ_4 clock is a modular *symmetry* ($[\rho, K]=0$), not the flow ($\sigma^\psi = \rho$ is type-wrong, a one-parameter group versus a finite order-4 element); a positive character-block-diagonal K has $[\rho, K]=0$ yet, being positive definite, never $\theta K \theta = -K$, so clock-invariance does *not* imply (LTO-RP). [O] make-or-break: NPW26 prove (LTO-RP) for *topologically ordered* commuting-projector models (Toric Code via a trace, Levin–Wen anyonic), whereas the seam is the opposite corner — invertible ($|\det \text{Cartan}(E_8)|=1$, no anyons) and a $\beta=1$ KMS state (not a trace) — so two sub-steps stay open: realising the raw seam as a $\mathbb{Z}_2$ -reflection commuting-projector LTO, and extending (LTO-RP) to the invertible KMS case. An MMST-style reduction (cf. v308/v336), *not* a closure. *Sub-step (ii) is now constructively exhibited* ([verification/v426_seam_lto_rp_kms.py]): a gapped reflection-symmetric μ_4 -even collar at $\beta=1$ KMS satisfies $\Theta K \Theta^{-1} = -K$ ($u_\Theta=J$) and $[\rho, K]=0$, so (LTO-RP) holds there by direct Tomita–Takesaki — the corner NPW26 leave open — with $|\det \text{Cartan}(E_8)|=1$ (invertible, no anyons) and the state a genuine KMS, not a trace; only the continuum sub-step (i) remains [O]. On the covariance side, the recent theorem of Adamo–Giorgetti–Tanimoto (arXiv:2508.07109, 2025) — every DHR representation of a conformal net extending a loop-group/Virasoro net is *automatically* positive-energy and diffeomorphism-covariant — defuses the old “Bisognano–Wichmann presupposes the

very covariance it should produce” circularity (v215): once the seam net is in hand, geometric covariance is a *consequence*, not an input.

Twenty-six convergent attacks on the one continuum residual — intrinsic BW, a bulk continuum motor and its chiral-edge counterpart, the β -homotopy, the applicability audit, uniform rigidity with its forcing converse and the derivation of clock-invariance, three independent reads of the edge central charge, the four-point and uniform-in- N pinning of the external fact, the Ising operator tower, the $(E_8)_1$ torus modular data, the μ_4 -from-marks derivation, the kill tests, the lattice Sugawara current algebra, a uniform-in- N Tomita–Takesaki tower, the chirality-from- P_1 forcing, an $(E_8)_1$ character kill test, the exact MMST citation audit, the complementary lattice-VOA route, the explicit flat-band commuting-projector LTO realisation, the strict-locality (Kapustin–Fidkowski) obstruction, the character-level 128-spinor extension, the $c=8$ holomorphic-uniqueness pin and the one-particle realisation rigidity ([verification/v438_seam_hsmi_borchers.py], [verification/v439_seam_lieb_robinson.py], [verification/v440_seam_lto_rp_beta.py], [verification/v441_seam_applicability_ledger.py], [verification/v442_seam_rigidity_uniform.py], [verification/v443_seam_e8_discriminator.py], [verification/v444_seam_edge_scaling.py], [verification/v445_seam_rigidity_forcing.py], [verification/v446_seam_clock_invariance.py], [verification/v447_seam_edge_chern.py], [verification/v448_seam_edge_fourpoint.py], [verification/v449_seam_mmst_uniform.py], [verification/v450_seam_edge_entanglement.py], [verification/v451_seam_edge_cardy_tower.py], [verification/v452_seam_e8_modular.py], [verification/v453_seam_mu4_from_marks.py], [verification/v454_seam_edge_virasoro.py], [verification/v455_seam_bw_uniform.py], [verification/v456_seam_chirality_from_c3.py], [verification/v457_seam_character_killtest.py], [verification/v458_seam_mmst_citation_audit.py], [verification/v459_seam_lattice_voa_route.py], [verification/v460_seam_s3_lto.py], [verification/v461_seam_strict_locality.py], [verification/v462_seam_spinor_continuum.py], [verification/v463_seam_c8_holomorphic_uniqueness.py], [verification/v464_seam_oneparticle_rigidity.py], [verification/v469_seam_crossedproduct_route.py]).

None closes SEAM.EQUIV.01; together they ground the one continuum input on both its halves, reduce it to a single analytic fact and make the claim falsifiable. (v) *The intrinsic Bisognano–Wichmann route* ([verification/v438_seam_hsmi_borchers.py]): the seam supplies a Longo *standard pair* / +half-sided modular inclusion — a positive lightlike translation $P \geq 0$, the boost K with $[K, P]=P$ (the Borchers scaling $\Delta^{it}P\Delta^{-it}=e^{-t}P$), and a reflection J with $JKJ=-K$, $JPJ=P_-$ — so by Borchers’ theorem the modular flow is geometric *intrinsically*, *without* presupposing the rotation-covariance of the vacuum; [E] the ladder algebra and the $sl(2, \mathbb{R})$ positivity (min eig > 0 , with the indefinite boost as control), [C] the synthesis, [O] assembling the finite algebra and the positive-energy rep into one continuum standard pair (the MMST/NPW26 wall). (vi) *The continuum-existence motor* ([verification/v439_seam_lieb_robinson.py]): the explicit $p+ip$ collar (v367) has a uniform-in- N gap, exponential clustering with N -independent ξ , and a finite Lieb–Robinson light cone, so by Nachtergaele–Sims its *bulk* thermodynamic limit exists — grounding the bulk half of the cited MMST existence in our model and narrowing the open step to the chiral massless *edge* scaling limit; [E] gap/clustering/light-cone, [C] the lift, [O] the edge CFT. (vii) *The β -homotopy* ([verification/v440_seam_lto_rp_beta.py]): the (LTO-RP)/BW data $(\Theta K \Theta = -K, [\rho, K]=0)$ are β -independent, and along the whole path $C_\beta = 1/(1+e^{\beta K})$ from the NPW26-proven *trace* state ($\beta=0$) to the seam KMS state ($\beta=1$) every covariance is reflection-related, μ_4 -invariant and valid — a norm-continuous, obstruction-free homotopy, upgrading the single point of v426 to a path argument for sub-step (ii). (viii) *The applicability audit* ([verification/v441_seam_applicability_ledger.py]): a hypothesis-by-hypothesis ledger of MMST/NPW26/Adamo finds 11 of 12 hypotheses are TFPT-internal computed facts ([E]) and exactly one is external — the continuum existence of the chiral scaling limit — covariance being an Adamo consequence, not an input; the keystone’s external dependence is pinned to that one analytic fact. (ix) *Uniform rigidity* ([verification/v442_seam_rigidity_uniform.py]): the band-limited μ_4 -character rigidity of v398 is uniform in N — exact 4-sector commutation to $N=256$, a uniform gap, and an N -independent clock-breaking lower bound $\|[\rho, V]\|=2$ — so it is not a small- N artifact (the full- L^2 forcing stays [O]). (x) *The kill tests* ([verification/v443_seam_e8_discriminator.py]): mode D of the closure-mode classification (v347) as concrete *falsifiable* discriminators against the nearest $c=8$ rival $SO(16)_1$ — torus GSD 1 vs 4, level-1 currents 248 vs 120, lattice theta 240 vs 112 norm-2 roots — with $c=8$ alone unable to decide (v277) and only $\det K=1$ (the genus-1 statement, v90) selecting E_8 ; a seam with GSD >1 would falsify $(E_8)_1$. [X] the tests, [O] the realisation. (xi) *The chiral-edge scaling motor* ([verification/v444_seam_edge_scaling.py]): the counterpart of (vi) on the *edge*, attacking exactly the chiral massless scaling limit that the bulk motor left open. On the same $p+ip$ collar

(v367/v368) the lattice edge already carries the conformal data of a chiral $c=\frac{1}{2}$ Majorana CFT — a linear, chiral in-gap dispersion (velocity $v\approx 0.98$, $R^2=0.9999$, with *opposite* per-edge velocities), an *algebraic* edge two-point $\sim r^{-\eta}$ with $\eta\approx 1$ (versus the *exponential* bulk row and the exponential trivial-phase control), and a Cauchy-convergent scaling-collapse under refinement ($\eta=1.001\pm 0.007$, successive spread $0.028\rightarrow 0.013$) — so the 16-copy edge is a chiral $c_-=8$ $(E_8)_1$ CFT; [E] the kinematics, correlator and collapse, [C] the central-charge reading, [O] the rigorous uniform-convergence theorem (cited MMST). With (vi) for the bulk and (xi) for the edge, *both halves* of the one continuum residual are numerically grounded and pinned to a single cited theorem. (xii) *The rigidity forcing converse* ([verification/v445_seam_rigidity_forcing.py]): the converse of (ix), upgrading rigidity from “block-diagonal *permits* commuting” to “commuting with the order-4 clock *forces* block-diagonal”. For $\rho=\text{diag}(i^n)$, $[\rho, K]=0$ iff $K_{ij}=0$ whenever $i\not\equiv j \pmod 4$ (verified entrywise, swept $N=4..128$), so the commutant is *exactly* the 4-sector block algebra: $\dim\{K : [\rho, K]=0\} = \sum_s n_s^2$, a proper subspace of the full N^2 algebra ($N=16$: $64<256$); every off-block unit has $\|[\rho, E_{ij}]\| \geq \sqrt{2}$ (zero slack), and the order-2 clock leaves a strictly larger commutant (64 vs 128) — the *four* Gauss–Bonnet marks (order $4=|\mu_4|=h(A_3)$) force strictly more structure than two, and only the index-4 extension is $(E_8)_1$ ($\det=1$, v281/v154/v351). [E] the entrywise iff, exact commutant dimension, off-block sharpness and order discriminator; [C] the synthesis (RP-definability + clock-invariance *force* the block-diagonal Λ_Σ , v442’s “permit” upgraded to “force”); [O] the single residual is now precisely that the physical transfer lies in the commutant (clock-invariance = QGEO.SYM.01, v323/v335). (xiii) *Clock-invariance derived* ([verification/v446_seam_clock_invariance.py]): closing (xii)’s residual one notch further — the operator condition $[\rho, \Lambda_\Sigma]=0$ is itself a *consequence* of a structural symmetry, not an extra assumption. On explicit finite OS data, the canonical transfer $\Lambda=C(1)C(0)^{-1}=e^{-h}$ (the transfer fixed by RP+reflection, v54) of a μ_4 -invariant covariance satisfies $[\rho, \Lambda]=0$ (swept $N=4..32$; iff-sharp μ_4 -breaking neg control), Λ is a positive contraction, and the free modular Hamiltonian $K=\log((1-C)C^{-1})$ (v258) of the invariant state has $[\rho, K]=0$ (Tomita–Takesaki). [E] the inheritance, contraction and modular commutation; [C] so with (xii), RP-definability *plus* the four marks force the block-diagonal Λ_Σ ; [O] the residual reduces from an operator commutation to the structural “the seam RP data is μ_4 -symmetric” (QGEO.SYM.01, v216/v323). (xiv) *The independent topological edge reading* ([verification/v447_seam_edge_chern.py]): a second, fit-free route to the edge central charge of (xi). On the same $p+ip$ collar (v367/v368) the bulk Chern number is the integer $C=+1$ (Fukui–Hatsugai–Suzuki link variables; $C=0$ in the trivial phase, $C=-1$ at the opposite mass, matching (xi)’s opposite per-edge velocities), the open strip carries exactly $|C|=1$ chiral edge branch per boundary, and the Li–Haldane entanglement spectrum has in-gap modes at $\frac{1}{2}$ only in the topological phase; [E] the Chern integer, channel count and entanglement signature, [C] the bulk–edge reading $c_-=N_{\text{Maj}} \cdot \frac{1}{2} \cdot |C|=16 \cdot \frac{1}{2} \cdot 1=8=g_{\text{car}}+N_{\text{fam}}$ — the *same* $c_-=8$ as (xi) from a logically independent (topological) observable. (xv) *The four-point pinning of the external fact* ([verification/v448_seam_edge_fourpoint.py]): extending (xi)’s two-point scaling limit to the next order to support exactly the one external fact (viii) isolates. The edge state is quasi-free (the ground-state correlation matrix is an exact projector $G^2=G \sim 10^{-13}$, the MMST CAR hypothesis), and the edge four-point cross-ratio $Q=[G_{01}G_{23}]/[G_{02}G_{13}]$ at commensurate positions is Cauchy-convergent under refinement to the free-fermion conformal value $\sin(5\pi/12)/\sin(\pi/12)=2+\sqrt{3}\approx 3.732$ (the bulk control off-scale, exponential); [E] the quasi-free property, four-point convergence and conformal form, [C] so the one external MMST fact is empirically supported at *both* two- and four-point order. (xvi) *The uniform-in-N pinning* ([verification/v449_seam_mmst_uniform.py]): strengthening (xv) from convergence at a single cross-ratio to the *uniformity* the cited theorem asserts — two *independent* edge four-point cross-ratios converge to *distinct* conformal values ($Q(1, 2, 4, 7)/12 \rightarrow 2+\sqrt{3}$ and $Q(1, 3, 5, 9)/12 \rightarrow 2$) at one N -independent $1/N$ rate ($|Q(N)-Q_\infty| \cdot N \leq 4$ for both, $N_x \in \{36, 48, 60\}$), the bulk control off-scale; [E] the two limits and the uniform rate, [C] so the MMST uniform-convergence *hypothesis* is empirically met. (xvii) *The entanglement reading of c_-* ([verification/v450_seam_edge_entanglement.py]): a *third*, correlator-independent route to the edge central charge — the Calabrese–Cardy block-entropy law gives $c_{\text{Dirac}}=1.000$ for the critical free-fermion chain (vs a saturating, ℓ -independent gapped control), so c per Majorana $=\frac{1}{2}$ and $c_-=16 \cdot \frac{1}{2}=8=g_{\text{car}}+N_{\text{fam}}$; [E] the law and the control, [C] the assembly. (xviii) *The Ising operator tower* ([verification/v451_seam_edge_cardy_tower.py]): naming the edge CFT by its *operator content* — the Cardy finite-size spectrum reproduces the exact Ising chiral primaries $h_\sigma=\frac{1}{16}$ (the periodic–antiperiodic ground split) and $h_\epsilon=\frac{1}{2}$ (the lowest NS pair gap), both L -independent, so $\{c, h_\sigma, h_\epsilon\}=\{\frac{1}{2}, \frac{1}{16}, \frac{1}{2}\}$ is the free-Majorana minimal model $M(3, 4)$; [E] the dimensions and conformal scaling, [C] the identification. (xix) *The $(E_8)_1$ torus modular data* ([verification/v452_seam_e8_modular.py]): pinning the $(E_8)_1$ identity on the *torus*, complementing the edge reads (xi)/(xiv)/(xvii)/(xviii) — the character

$\chi=E_4/\eta^8$ is S -invariant ($\chi(-1/\tau)=\chi(\tau)$, one primary), T -multiplies by $e^{-2\pi i/3}=e^{-2\pi ic/24}$ ($c=8$) and is holomorphic ($c \equiv 0 \pmod 8$, T -phase order 3); [E] the full modular signature. (xx) *The μ_4 -from-marks derivation* ([verification/v453_seam_mu4_from_marks.py]): folding the last structural premise into the existing one — the four Gauss–Bonnet marks are μ_4 (roots of z^4-1), the form basis ω_k is μ_4 -graded ($\rightarrow i^k \omega_k$, weights $(1,2,3)=A_3$ exponents) and the order-4 clock $z \rightarrow iz$ fixes their cross-ratio $\lambda=2$ (the D_4 stabiliser, v168), so the residual QGEO.SYM.01 of (xiii) follows from marks = μ_4 plus the existing QGEO.REALIZE.01 — [E] the three mark facts, [C] no new premise: the chain (xx) $\rightarrow$ (xiii) $\rightarrow$ (xii) closes modulo realisation. The integer backbone of (xiv)–(xix) ($c_-=8$, the Chern integers, the order-3 T -phase) is kernel-hardened in Lean axiom-free as FORM.SEAM.EDGE.CHERN.01. (xxi) *The lattice Sugawara current algebra* ([verification/v454_seam_edge_virasoro.py]): testing the cited MMST “Koo–Saleur generators $\rightarrow$ continuum Virasoro with the correct c ” fact on the edge algebra, not just its central charge. The critical free-fermion ring has the Affleck–Cardy Casimir energy $E_0(L)=e_\infty L - \pi v c / (6L)$ fitting $c_{\text{Dirac}}=1.000$ ($\frac{1}{2}$ per Majorana, $L=64..1024$), the lattice $U(1)$ current closes its affine level exactly ($\langle 0|J_n J_{-n}|0\rangle=n$, $n=1..8$; vanishing for a gapped control), and the level-1 Sugawara central charge $c=\dim G/(1+h^\vee)=8$ for both $SO(16)_1$ (120/15) and $(E_8)_1$ (248/31); [E] Casimir, Kac–Moody level and Sugawara, [C] $16 \cdot \frac{1}{2}=8$ is the level-1 current-algebra c_- . (xxii) *The uniform-in- N Tomita–Takesaki tower* ([verification/v455_seam_bw_uniform.py]): the inductive-limit companion of v426 for sub-step (i) of (iv) — a family of reflection-odd, μ_4 -even, gapped collars ($N=8..128$) on which the intrinsic Bisognano–Wichmann condition $u_\Theta=J$ ($\|\Theta K \Theta + K\|=0$) holds exactly at every N , with an N -independent gap ($6 \log \frac{3}{2}$) and an N -independent neg-control margin ($\|\Theta K \Theta + K\|=2\|K\|$ for a side-even reflection); [E] the uniform BW, gap, clock symmetry and teeth, [C][O] so the inductive limit inherits $u_\Theta=J$ (the abstract continuum theorem stays the cited backbone). (xxiii) *The chirality-from- P_1 forcing* ([verification/v456_seam_chirality_from_c3.py]): the S3 input of the continuum chain (v356) is removed as a free choice — c_3 ’s “8” is the one-sided count $8=|\mathbb{Z}_2| \cdot (\int K/\pi)=2 \cdot 4$ ($8\pi=|\mathbb{Z}_2| \cdot \int K=2 \cdot 2\pi\chi(S^2)$, v73), and a reflection sends the edge Chern integer $C \rightarrow -C$ (verified $C=+1 \rightarrow -1$ on the explicit $p+ip$ model), so any two-sided (reflection-symmetric) gapped boundary forces $C=-C=0$ while the one-sided seam — the same $|\mathbb{Z}_2|$ quotient that gives c_3 its 8π — has no such reflection and forces $C \neq 0$; [E] the one-sided count and the $C \rightarrow -C$ flip, [C] chirality (S3) is a consequence of P_1 , magnitude $c_-=8$. (xxiv) *The $(E_8)_1$ character kill test* ([verification/v457_seam_character_killtest.py]): a falsifiable sharpening of the (x)/(x”) discriminators — the $(E_8)_1$ tower $\chi=E_4/\eta^8=q^{-1/3}(1+248q+4124q^2+34752q^3+\dots)$, the weight-1 count $248=120+128$ (vs the free 120) and the sector count $|\det \text{Cartan } E_8|=1$ (vs $|\det \text{Cartan } D_8|=4$); a measured 120, four sectors, or a wrong tower coefficient would falsify $(E_8)_1$, and the data give $248/1/\{1, 248, 4124\}$ (passed). [X] the kill criterion, [E] the exact discriminators. (xxv) *The exact MMST citation audit* ([verification/v458_seam_mmst_citation_audit.py]): a hypothesis-by-hypothesis reading of the cited MMST theorem (arXiv:2107.13834) against the collar — H1 CAR class, H2 massless limit, H3 the per-copy $c=\frac{1}{2}$ Thm 4.11, H4 decoupled additivity to $c=8$, and H5 the WZW range $8 \leq 8 \leq 16$ with its 120 bilinear currents all match, isolating the single residual H6 to the 128-spinor extension $SO(16)_1 \rightarrow (E_8)_1$ (the $\det K 4 \rightarrow 1$ holomorphy step), outside MMST’s bilinear method; [E] the per-copy theorem, [C] the additivity and range, [O] the one open spinor extension — since (xxxii) certified at net level by the peer-reviewed crossed product. (xxvi) *The complementary lattice-VOA route* ([verification/v459_seam_lattice_voa_route.py]): the second, independent construction that supplies exactly H6 — for the even unimodular E_8 the lattice net A_{E_8} (Adamo–Giorgetti–Tanimoto, arXiv:2506.01008; OS arXiv:2407.18222) has weight-1 space $248=8+240$ with the 240 roots splitting 112 integer + 128 half-integer (the spinor, $248=120+128$), so the lattice route builds the 128-spinor currents MMST (xxv) leaves open; [E] the even unimodularity and root split, [C] AGT constructs the holomorphic extension — since (xxxii) an independent second witness, the primary net-level certification being the peer-reviewed crossed product — [O] the routes leave one shared realisation input (the collar’s scaling limit = A_{E_8}), tied to P_1 by (xxiii). The G-block arithmetic (the one-sided $8=|\mathbb{Z}_2| \int K/\pi$, the reflection $C \rightarrow -C$, $c_-=8$, the range $8 \leq 8 \leq 16$, $248=8+240=120+128$, $240=112+128$, $\det K 1$ vs 4) is kernel-hardened in Lean (FORM.SEAM.RESIDUAL.01, TfptCarrier/SeamResidualAxiom.lean), reducing the residual to ONE named TFPT-internal realisation axiom (CollarRealizedAsLatticePhase) plus the two cited existence theorems. (xxvii) *The explicit flat-band commuting-projector (LTO) realisation* ([verification/v460_seam_s3_lto.py]): NPW26 prove the intrinsic Bisognano–Wichmann lemma of (iv) for commuting-projector (local-topological-order) models, so sub-step (i) asks for the seam realised as such a $\mathbb{Z}_2$ -reflection LTO. On the same v367 $p+ip$ collar the flattened Bloch $Q=h/|h|$ is a frustration-free flat-band parent — $Q^2=I$ and the occupied $P=(I-Q)/2$ is an exact projector EQUAL to the ground-state projector — whose real-space kernel is gap-protected quasi-local ($\xi \approx 1.14$ at $M=1$,

$|P(20,0)| < 10^{-6}$; a clean power law $|R|^{-2.1}$ at the gapless point, so locality needs the gap), and whose flattened modular Hamiltonian $K_{\text{flat}} = \text{sign}(K)$ inherits the $\mathbb{Z}_2$ reflection $\Theta K_{\text{flat}} \Theta = -K_{\text{flat}}$ and the μ_4 clock $[\rho, K_{\text{flat}}] = 0$ EXACTLY; [E] the flat-band/projector and reflection algebra, [C] the quasi-locality, [O] the chiral $c_- = 8 \neq 0$ obstructs *strictly* finite-range commuting-projector locality (Kapustin–Fidkowski, arXiv:1810.07756), so the realisation is exactly the quasi-local LTO net NPW26 use — advancing sub-step (i) of (iv) from abstract existence to an explicit gap-protected net, strict locality the cited residual. (xxviii) *The strict-locality obstruction made explicit* ([verification/v461_seam_strict_locality.py]): upgrading (xxvii)’s “strict locality is the cited residual” from a citation to an EXHIBITED topological obstruction. On the same v367 collar the FHS Chern is $C = +1$ (chiral) / 0 (trivial) / -1 (opposite mass), and the hybrid Wannier-centre (Wilson-loop) phase $\theta(k_y)$ winds by $|C|=1$ (chiral) and 0 (trivial) over the BZ — exponentially-localised Wannier functions exist *iff* that winding is 0 (Brouder–Panati et al.; Thouless), so the chiral collar has none; both phases are gapped, yet only the trivial one (winding 0) admits a strict finite-range commuting projector, so it is the *chirality* $c_- = 8 \neq 0$, not the gap, that obstructs. [E] the Chern integer, the winding = $|C|$ identity and the gapped neg-control, [C][O] Kapustin–Fidkowski (arXiv:1810.07756): a strictly finite-range commuting projector would force winding 0 hence $C=0$, so for $C=1 \neq 0$ none exists — sub-step (i)’s realisation is *provably* the quasi-local NPW26 LTO net, strict locality topologically forbidden rather than a missing premise. (xxix) *The character-level 128-spinor extension* ([verification/v462_seam_spinor_continuum.py]): the constructive companion of (xxv)/(xxvi), exhibiting the 128-spinor extension three ways — the exact Jacobi/ E_8 identity $\theta_2^8 + \theta_3^8 + \theta_4^8 = 2E_4$ makes the holomorphic character the sum of the $SO(16)_1$ vacuum and spinor sectors, $\chi_{(E_8)_1} = E_4/\eta^8 = \chi_o^{SO16} + \chi_s^{SO16}$ (weight-1 counts $120+128=248$); the finite 16-Majorana Fock space already carries them ($C(16,2)=120$ currents and $2^{16/2}=256=128+128$ Ramond states); and a critical free-fermion ring converges to the chiral CFT (Casimir $c_- \rightarrow 8$, lowest scaled gap $x_1 \rightarrow 1$ with $|x_1-1| \sim 1/N^2$). [E] the θ identity, the $120+128=248$ split and the finite- L convergence, [C][O] the rigorous holomorphic extension is certified at net level by the crossed product (xxxii), with AGT/AMT (xxvi) as the independent second witness. (xxx) *The $c=8$ holomorphic-uniqueness pin* ([verification/v463_seam_c8_holomorphic_uniqueness.py]): the positive lever that makes the *identification* half classification-forced rather than assumed. At level 1 the central charge $c = \dim \mathfrak{g}/(1+h^\vee) = 8$ has *three* simple solutions — A_8 (dim 80), $D_8 = \mathfrak{so}(16)$ (dim 120) and E_8 (dim 248) — so $c=8$ is necessary but *not* sufficient, and the $\det K=4$ $SO(16)$ rival of (xxiii) is a genuine same- c competitor, not a strawman. But a holomorphic (single-sector) $c=8$ theory has an $SL(2, \mathbb{Z})$ -covariant character, and the only weight-0 form with leading $q^{-1/3}$ and integer coefficients is $j^{1/3} = E_4/\eta^8$, whose q^1 coefficient (the one-dimensional candidate space, vacuum $a_0=1$) is 248 — so $\dim V_1=248$ is *forced*, selecting E_8 uniquely and excluding A_8/D_8 . [E] the three candidates, the forced 248, the tower $\{1, 248, 4124, 34752, 213126\} = E_4/\eta^8 = j^{1/3}$; [C] Dong–Mason/ Schellekens give the holomorphic $c=8$ VOA unique = V_{E_8} (the even unimodular rank-8 lattice unique = E_8 , Mordell–Witt), so “ $c=8$ holomorphic limit = $(E_8)_1$ ” is classification-forced. (xxxi) *The one-particle realisation rigidity* ([verification/v464_seam_oneparticle_rigidity.py]): attacking the *realisation* input R_1 itself. Since the seam is quasi-free (v113/v160/v161), by Araki’s self-dual CAR the whole net is the second quantisation of its one-particle symbol P , so R_1 collapses to a finite-per-mode statement about P : the gapped ground state is the *unique* quasi-free state with $P =$ the spectral projection onto $K < 0$ ($P^2 = P$ to $< 10^{-12}$, K fixed by gap + chirality + reflection v426); its real-space kernel is Cauchy on every fixed window ($\max |P_{2N} - P_N| \sim 1/N \rightarrow 0$, so P_∞ exists), and the limit carries $c=1$ per Dirac (entanglement $S(L) = (c/3) \ln L$, 16 copies $\Rightarrow c_- = 8$, tying v450). [E] idempotency, Cauchy convergence and the c -fit; [C] Araki/Shale–Stinespring give the state $\leftrightarrow$ symbol bijection and implementable Bogoliubov map, so R_1 is upgraded from a bald axiom to “the unique quasi-free realisation, its one-particle limit exhibited, modulo the named CAR functor”. (xxxii) *The net-level crossed-product route + the invariant-level realisation* ([verification/v469_seam_crossedproduct_route.py]): both halves of the residual re-founded on harder-to-reject ground. The 128-spinor extension needs no preprint: the $SO(16)_1$ spinor has $h_s = N/16 = 16/16 = 1 \in \mathbb{Z}$, so its statistics phase is $+1$ — *exactly* the Longo–Rehren locality criterion — and the $\mathbb{Z}_2$ simple-current crossed product $SO(16)_1 \rtimes \{1, s\}$ is a *local* index-2 extension by 1995–2001 *peer-reviewed* subfactor theory (Longo–Rehren, Rev. Math. Phys. 7 (1995); the level-1 $so(N)$ net with its DHR sectors from even-CAR: Böckenhauer, Rev. Math. Phys. 8 (1996); simple-current extensions of nets: Böckenhauer–Evans, Comm. Math. Phys. 197 (1998)); KLM (Comm. Math. Phys. 219 (2001)) gives $\mu(B) = 4/2^2 = 1 \Rightarrow$ holomorphic $\Rightarrow (E_8)_1$ by the (xxx) pin — so the AGT/AMT route (xxvi) is *demoted to an independent second witness*. The index-4 μ_4 glue runs on the same integer, $h(J^k) = \{1, 1, 1\}$ for $k=1, 2, 3$ (the v125 isotropy $q(k(1,1)) = k^2$ at net level, $\mu = 16/4^2 = 1$). And the realisation input R_1 is reduced from model fiat to *invariants* (R_1'): quasi-free [C] (v155/v160) + gap

[E] (v302) + class D + $c_- = 8$ [E] ((xxiii), from P_1) — computed on the collar as FHS Chern $|C|=1$ (control 0), $\nu=16 \times |C|=16$, $c_- = \nu/2=8$; and $\nu \equiv 0 \pmod{16}$ is exactly Kitaev’s class-D phase whose edge admits a purely bosonic description = the $(E_8)_1$ level-1 state (the sixteen-fold way; c_- a *phase* invariant, Kapustin–Spodyneiko; Kim et al.), so any realisation with the $R1'$ invariants is phase-equivalent to the v367 stack. [E] the locality integers, the KLM μ -arithmetic, the computed invariants and the composed chain; [C] the re-founded citation package; [O] the cited theorems stay cited. The full G-block arithmetic — now including the (xxviii) winding obstruction, the (xxix) finite-Fock spinor counts, the (xxx) holomorphy selector and the (xxxii) locality/ μ /sixteen-fold joints — is kernel-hardened in Lean (FORM.SEAM.RESIDUAL.01), the two formerly separate cited theorems are *merged* into one combined external axiom `seam_realisation_theorem`, so the residual reduces to ONE realisation axiom plus ONE cited theorem (`#print axioms` drops from six to four); and the (xxxii) route is Lean-carried as the *parallel* derivation `seamResidualClosed'` (the invariant-level axiom `collar_invariants` plus the single combined package `crossedproduct_route_theorem`, its joints — `lr_locality_integer`, `glue_locality_integers`, `klm_holomorphic`, `sixteenfold_pin` — kernel facts with no axioms). Across (v)–(xxxii), SEAM.EQUIV.01 stays [O]: with (xxx) pinning the identification, (xxxi) reducing the realisation to its one-particle data and (xxxii) re-founding the extension leg on peer-reviewed ground, the residual is now *entirely certification* — a named, hypothesis-audited package of published theorems (MMST, Longo–Rehren/Böckenhauer/Böckenhauer–Evans/KLM, Dong–Mason, Araki, Shale–Stinespring; AGT/AMT as second witness) with no open internal mechanism — yet still rests on cited continuum-existence theorems we do not re-prove.

The TOE/QFT closure contracts as three citable certificates (`[verification/v398_seam_state_rigidity.py]`, `[verification/v399_gravity_complete.py]`, `[verification/v401_metrology_closure.py]`). Three companion contracts package the closure status as named targets without fabricating any closure. (i) *Seam State Rigidity* (SEAM.RIGIDITY.01): the scattered state-invariance reductions (v177/v199/v308/v309/v335) as ONE theorem target — [E] the RP-definability hinge (OS canonical transfer from RP + reflection alone, v54, quasi-free v155), [E] the band-limited μ_4 -character-block-diagonal core swept $N=4.64$ (beyond the 3-dim H^1 , made uniform in N to $N=256$ by v442), [E] holomorphy $\Rightarrow (E_8)_1$; and [E] the forcing converse (v445) that commuting with the order-4 clock is *exactly* μ_4 -block-diagonality, so the residual is no longer “does block-diagonal permit” but only $[\rho, \Lambda_\Sigma]=0$ on the full L^2 — the physical transfer lying in that commutant (intrinsic Bisognano–Wichmann). [E] Moreover (v446) that operator commutation is itself *derived*: the OS canonical transfer of μ_4 -symmetric RP data inherits the clock (with the modular Hamiltonian commuting, v258), so the residual reduces further to the *structural* statement that the seam RP data is μ_4 -symmetric (the four marks, QGEO.SYM.01), leaving [O] only that structural symmetry on the raw continuum collar — for which the four marks now have an *established* continuum frame: the four-interval multilocal modular geometry (v480, keybox below), where the invariance is manifest at the state level in the exactly solvable class. (ii) *Gravity-complete* = *Boundary-complete* + *Ambient-redundancy* (GRAVITY.COMPLETE.01): the local equation parameter-free (v358/v359) plus the redundancy discriminators (v369), with $\mathcal{O}_{\text{phys}} \subset \mathcal{A}_\Sigma$ checked over five gravitational readout classes. (iii) *Metrology closure* (METROLOGY.CLOSURE.01): the final input set $\{a, \pi, v_{\text{geo}}\}$ — 0 dimensionless dials + 1 unit, the dimensional bookkeeping that every dimensionful observable is a number $\times v_{\text{geo}}^d$ (No-Unit, v153), and v_{geo} theorem-forbidden (TOE_{strict}=TOE_{dimensionless}+ $[v_{\text{geo}}]$, v384).

One surface for seam, carrier and lattice: the Kummer/K3 (`[verification/v260_k3_kummer_unification.py]`). The pillowcase is the elliptic-involution quotient $T_{\tau=i}^2/(z \mapsto -z)$, its four marks the four 2-torsion fixed points; squaring the construction to complex dimension two gives the *Kummer surface* $\text{Km}(T^2 \times T^2)$, a K3, whose three facets reproduce three TFPT objects at once. [E] the 16 nodes (the 2-torsion $|A[2]|=2^4=16$ of the abelian surface) resolve to 16 exceptional $\mathbb{P}^1$ ’s = $\dim S^+ = \Lambda^{\text{even}}(\mathbb{C}^5) = 16$ (v197), the carrier-16 as the *square* of the seam-4; [E] the cup-product form on $H^2(K3, \mathbb{Z})$ is the even unimodular lattice $U^3 \oplus E_8(-1)^2$ of signature (3, 19), so the same E_8 the carrier glues to (v1) is the surface’s intersection form. (Honestly: the 16 node classes span the rank-16 Kummer lattice, a sublattice *distinct* from the $E_8(-1)^2$ summand — both exact, not identified.) Thus the v1 “glue $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ ” is the lattice shadow of a single smooth object that already carries the seam and the carrier; E_8 appears both *locally* as the du Val $\mathbb{C}^2/2I$ singularity (the Poincaré-sphere link, v232/v236; by Kronheimer that local model *is* the hyper-Kähler quotient of the marks quiver, v479) and *globally* as this H^2 summand. A unification (ARCH.K3.01), *not* a closure: the lattice arithmetic is [E], the carrier $\leftrightarrow$ node identification [C].

Synthesis — the Modular Spectral Closure (`[verification/v261_modular_spectral_closure.py]`).

The ten steps above are not ten independent results: they are facets of one relative spectral datum

$$\text{TFPT}_{\text{QFT}} = (\mathcal{A}_\Sigma, \omega_\Sigma, \Delta_\Sigma, \rho, A_F, H_F, D_F, J, \gamma, S_{\text{rel}})$$

(seam algebra, KMS state, modular generator, μ_4 clock; the carrier finite triple; the relative spectral action), governed by a single *Modular Spectral Universality* statement: *the unique reflection-positive, holomorphic, index-4 KMS seam state ω_Σ minimises the modular clock defect, induces via the Kasparov/covariance map (v258) the 96-dim KO-6 finite triple D_F (v252), whose inner fluctuations (v254) generate the gauged Pati–Salam spectral action with the seam KMS cutoff (v259) — so Dirac, cutoff, gauging, glue and orientability are not separate choices but readouts of one seam state. The certificate v261 pins the cross-consistency: one number $4 = [B:A] = |\mu_4| = 2\chi = |(\mathbb{Z}/2)^2|$ (index = marks = fixed points = glue order), one carrier-16 ($\dim S^+ = \Lambda^{\text{even}}(\mathbb{C}^5) = \text{Kummer nodes}, H_F = 96$), and one gap $\Delta = 6 \log(3/2)$ (OS mass gap = recovery rate = Lindblad gap, static and dynamic). The honest account of completeness is therefore:*

- **[E]/[C] QFT-complete (modulo one premise).** The admissible OS sector, the boundary $(E_8)_1$ net, the relative spectral action, D_F (as covariance induction), the cutoff (as the KMS weight) and the gauging (as inner fluctuations) are one relative object — closed once **QGeo.SYM.01** (the seam realisation: *the raw quasi-free seam state is μ_4 -invariant*) is granted. That single **[O]** premise is the only structural input the QFT layer still consumes.
- **[C] Ambient QG — a redundancy, separate by design.** The full covariant metric-sector / ambient quantum-gravity measure (**QG.AMB.01**) is gap-decoupled ($\Delta_{\text{eff}} = 1.648 > 0$) and is now a **[C] redundancy (v369)**, not a blocker for the QFT layer: a certification object, not a missing piece of the boundary QFT.

So “solving the QFT” here means exactly this: the boundary QFT is a single relative object reduced to one foundational symmetry postulate, with ambient QG held apart — not a diffuse list of open interfaces. The MSC closes the boundary relative QFT object modulo **QGeo.SYM.01**; it does *not* assert SM-only 4D gauge unification, and the canonical 4D reading is boundary only.

4D-QFT fork policy — a frozen decision tree, not an open ambiguity (`[verification/v265_qft4d_fork_freeze.py]`)

The 4D reading is a *branch policy*, machine-frozen in `qft4d_fork_freeze.json`:

- **A (canonical) — boundary only.** **[C]** *Boundary-only is the canonical 4D reading*: TFPT delivers the boundary QFT, the relative spectral action and the SM readout; a 4D-GUT is *not claimed by default* (E_8 is the audit hull, never an unbroken 4D gauge group).
- **SM-only 4D-GUT unification is killed.** **[X]** the plain SM couplings do not meet (at the $\alpha_1=\alpha_2$ scale α_3^{-1} is off by ~ 5.6 ; spread ~ 8.8 at 2×10^{16} GeV, v246). This is the canonical *expectation*, not a failure.
- **B (conditional) — carrier-native Pati–Salam UV branch.** **[C]/[O]** permitted *only* with the E_8 -allowed content $\{1, 10, 16, 45\}$ (no **126**, v247), $\sin^2 \theta_W = 3/8$ (v248), $\kappa = M_{PS}/M_s$ in the frozen $O(1)$ band (v249/v253), and a threshold model with a proton-decay kill test; carrier gauging itself stays the contract **CONTRACT.QFT4D.DIRAC.01**.
- **Proton-decay rule, now applied** (`[verification/v266_ps_threshold_proton.py]`). **[X]** the explicit two-step thresholds give $d=6$ $\tau_p(p \rightarrow e^+ \pi^0)$: the *minimal* 16-Higgs content ($M_{GUT} \sim 2.4 \times 10^{15}$) gives $\tau_p \sim 4.4 \times 10^{33}$ yr < Super-K (2.4×10^{34}) and is *killed*; the E_8 -allowed (15, 1, 1) [the single 45] raises $M_{GUT} \sim 5.9 \times 10^{15} \Rightarrow \tau_p \sim 1.6 \times 10^{35}$ yr, proton-safe *now* but sitting at the Hyper-K reach ($\sim 1.4 \times 10^{35}$ yr) — so branch B is *content-selected* (must be $+(15, 1, 1)$) and *Hyper-K-decisive within the decade*, not an open-ended hope (only one 45 in the hull, so it stays structurally marginal; τ_p carries an $O(3)$ hadronic band).

The fork is thus a frozen decision tree with explicit promotion/demotion rules, *not* an open ambiguity.

And the bedrock is now a falsifiable kill-test, not a proof-promise (`[verification/v215_seam_deck_killtest.py]`). A foundational symmetry cannot be derived from nothing (Bisognano–Wichmann would presuppose the very covariance it should produce); the honest move is to recast **QGeo.SYM.01** as four independently-falsifiable predictions about the raw seam DtN. The non-circular reduction (v198/v199/v201) gives $\omega \circ \rho = \omega \Leftrightarrow [\rho, C] = 0$ for the quasi-free covariance, with the principal symbol $|k|$ commuting exactly, so the whole residual sits in the bounded

sub-principal M_f (equivalently: the raw Gaussian bulk form A is μ_4 -deck-invariant). The four levers are **K1** exactly four marks ($b_1 = \# \text{marks} - 1 = 3 = N_{\text{fam}} \Rightarrow \# \text{marks} = 4 = |\mu_4|$); **K2** the *square* configuration (cross-ratio $2 \Rightarrow j=1728$, $\text{Aut}=\mathbb{Z}/4$; the clock order $h(A_3)=4$ selects $\mathbb{Z}/4$, not the $j=0$ $\mathbb{Z}/6$, and a generic mark set has only $\mathbb{Z}/2$); **K3** the mod-4 sub-principal support ($[\rho, M_f]=0 \Leftrightarrow f$ is $\mathbb{Z}_4$ -invariant, automatic for a μ_4 -mark sum); and **K4** the transfer gap $(2/3)^6=64/729$. These are frozen in `freeze_file.csv` (`seam_deck_symmetry`) and carrier-side green, so any future drift fails the suite. The *decisive* kill is deferred to QGEO.REALIZE.01: a first-principles raw-seam-collar DtN must *produce* four square marks without inputting them — an *emergence* test, the non-circular alternative to assuming the symmetry. [E] carrier-side (the four levers), [O] seam-side (the emergence): a kill-test, *not* a closure (QGEO.KILL.01).

And the mark count is now derived, not assumed: four marks from Gauss–Bonnet ([`verification/v216_marks_gauss_bonnet.py`], [`verification/v217_marks_emergence_scan.py`]). The K1 lever (“exactly four marks”) is the one that can be *executed* today. If the marks are $\mathbb{Z}_2$ branch points (cone angle π , deficit π) on the flat seam sphere ($\chi=2$ — the same χ that supplies the 8 in c_3), Gauss–Bonnet *forces* their number: $n\pi = 2\pi\chi = 4\pi \Rightarrow n = 2\chi = 4 = |\mu_4| = N_{\text{fam}}+1$. The closed Euclidean sphere 2-orbifolds are exactly $(2, 3, 6)$, $(2, 4, 4)$, $(3, 3, 3)$, $(2, 2, 2, 2)$, and *three independent* criteria — all cone orders $= 2$ (the $|\mathbb{Z}_2|$ sheet branch), $\text{rank } H^1 = \# \text{marks} - 1 = 3 = N_{\text{fam}}$ (the three generations pick the 4-mark square over the 3-mark hexagonal $(3, 3, 3)$), and a free order-4 deck — converge on the pillowcase $(2, 2, 2, 2)$. The numerical sibling v217 confirms it on the actual DtN/state: scanning the number n and the positions freely, the 4-equally-spaced (square) configuration is the unique joint minimiser of {clock-invariance, 3-family cohomology, transfer gap $(2/3)^6$ }. So the mark *count* and *equality* move from input to derived; only the square *modulus* (cross-ratio 2, the order-4 clock) stays the one carrier input, and the from-zero raw-seam emergence QGEO.REALIZE.01 stays [O]. [E] (count + equality), [O] (the square modulus and the raw-seam realisation) — QGEO.MARKS.03 / QGEO.EMERGE.01.

And the four marks are a four-interval modular geometry ([`verification/v480_multilocal_four_interval.py`]). The seam circle minus its four marks is the $n=4$ disjoint-interval region of a chiral free fermion — the one continuum class whose *multi-interval* modular theory is exactly known (Casini–Huerta mixing term; Longo–Xu multi-interval subfactors; Rehren–Tedesco multilocal fermions; and Kawahigashi–Longo–Müger: *holomorphic* $\Leftrightarrow$ trivial multi-interval inclusion, $\mu=1$ — the same one bit as $\det K=1$). Until now the suite used only the single-ball CHM boost (v358); v480 exhibits the multilocal mechanism on the exactly solvable antiperiodic critical ring: [E] the quarter-turn clock is a signed permutation with $\rho^4=-1$ *exactly* (the fermionic μ_4 clock is the order-8 double-cover lift — the same $\mathbb{Z}_2$ -lift pattern as $2I$ over I); [E] its primitive-8th-root character projectors decouple the 4-interval covariance at 10^{-15} (the diagonal 4-fold-cover isomorphism), each sector a *single-interval* quarter-ring kernel with boundary twist (10^{-14} spectra); [E] $\rho C \rho^{-1} = C$ at 10^{-14} , so $\omega \circ \rho = \omega$ holds *manifestly* at the 4-interval state level — and a one-site displacement of a single interval re-couples the sectors at 10^{-2} (the invariance is the μ_4 -symmetric configuration, not an artifact); and the cross-interval modular coupling peaks on the conjugate-point diagonal, the lattice shadow of the Casini–Huerta mixing term (numerical). [C] the continuum multi-interval theory is cited, not re-derived; [O] the raw-seam premise (QGEO.SYM.01) is untouched — this puts the bedrock and the AP2 continuum leg (v476/v478) in one established, exactly solvable frame; it does not close them.

QGEO.SYM.01 is *not* a theorem — it is the final seam-identification premise

The single load-bearing bolt of the whole construction, on one page.

- What is identified.** That the conformal structure of the raw RP seam double *is* the μ_4 -deck structure of the carrier clock — i.e. the carrier (P2) and the seam’s conformal geometry are the same object. Equivalently, in Riemann–Roch *index* language: the seam’s degree-4 mark divisor on the genus-0 double has $h^0 = \deg + 1 = 5 = g_{\text{car}}$ and $\text{rank } H_1 = \deg - 1 = 3 = N_{\text{fam}}$ (with $\Lambda^{\text{even}}(\mathbb{C}^5) = 1+10+5 = 16 = \dim S^+$), so P2 is the genus-0 mode count of the four-mark seam divisor ([`verification/v228_rr_index_gate.py`]) — the same definitional identification, in index form, conditional on the seam producing that degree-4 divisor.
- Why it is weaker than QGEO.CONF.01/ISO.01.** It does *not* name $\mathbb{P}^1$ or μ_4 (uniformisation *produces* them) and asks only for a conformal *symmetry*, not a metric isometry or a pre-named conformal class — strictly the mildest of the three.
- What follows automatically (the conditional theorem).** *Given* QGEO.SYM.01, equivariant uniformisation puts the order-4 clock in the form $z \mapsto iz$ on $\mathbb{P}^1$, the four parabolic marks are its free orbit μ_4 , H^1 carries the $(1, 2, 3)$ grading, the seam Calderón contraction is the rank-8 K_Σ polarisation,




Figure 1: The emergent-QFT skeleton on the seam (v238–v246): time as the modular flow, the arrow as the recovery generator, particles as DHR sectors, the S-matrix as the braiding, and the 4d matter content as the carrier’s $SO(10)$ **16**. Each step is machine-checked; the one residual — SEAM.EQUIV.01, closed modulo cited theorems (v367–v379) with the cited continuum scaling-limit existence (v336) as its **[O]** leg — is in red, the frozen 4D fork and the QG certification layer in gold. The whole skeleton is *conditional* on that keystone and is neither solved nor experimentally confirmed.

and the index-4 simple-current extension is $(E_8)_1$ — the *entire* AQFT closure of v175–v181.

4. **What would falsify or unseat it.** A seam double of genus > 0 (uniformisation to $\mathbb{P}^1$ fails); a carrier clock that is *not* a conformal symmetry of the seam geometry (a non-deck transport); or a fourth chiral generation ($N_{\text{fam}} \neq 3$ collapses $b_1 = N_{\text{fam}}$ and the four-mark μ_4 structure). The first two are constructive-geometry questions about the raw seam; the third is the standard family-count kill test.

Honest framing. The correct external statement is *conditional*: “Given QGEO.SYM.01, the RP seam boundary realises the μ_4 deck structure, and the AQFT closure to $(E_8)_1$ follows through the verified chain v175–v181” — *not* “the theory *derives* the seam geometry”. The premise stays **[O]**; everything downstream of it is **[E]**. Per v323 (Bisognano–Wichmann) QGEO.SYM.01 is itself the *downstream* conformal-deck face of the single named keystone SEAM.EQUIV.01 (the raw RP seam = the holomorphic $(E_8)_1$ net at $\tau=i$), which is now *closed modulo a cited theorem* (the lattice/ S_3 stack pins it at every computable level, Lean FORM.SEAM.MMST.01; only the cited continuum existence stays **[O]**).

This is not a gap to apologise for — it is TFPT’s one fundamental postulate. Every fundamental theory rests on an irreducible premise (the constancy of c in relativity, the equivalence principle, the canonical commutator). QGEO.SYM.01 is that premise for TFPT: a single geometric identification (seam conformal structure = carrier μ_4 deck) from which the entire algebraic and topological structure unfolds




Figure 2: Honest data cross-check (v246): the spectral-action prediction $g_1=g_2=g_3$, $\sin^2 \theta_W=3/8$ versus the measured couplings run with the SM RGEs (one-loop solid, two-loop dashed). The three lines do *not* meet at any scale (1-loop residual $\Delta\alpha^{-1} \approx 9$ at 2×10^{16} GeV; 2-loop minimal spread ≈ 3.2 at 1.7×10^{14} GeV) — a falsifiable tension. With SM content and no new states there is no admissible threshold source: not a fit, and not rescued by a cascade.

by theorem. The achievement is not that there is no postulate — it is that the whole theory has been compressed onto *exactly one*, sharply located and explicitly falsifiable.

The Seam-Holomorphy Theorem — the one named closing target [E]/[O]
([verification/v234_seam_holomorphy_selection.py])

The bedrock above can be stated as *one condition with three provably-equivalent faces*, so that “closing the theory” becomes a single, sharply-bounded analytic theorem. **The condition:** the seam carries *no nontrivial abelian sector*. Its three faces all force E_8 :

1. **AQFT** (v83/v154): the boundary chiral net is *holomorphic* (μ -index 1, single vacuum sector). Holomorphic + $c = 8 \Rightarrow$ the unique even unimodular rank-8 lattice net $(E_8)_1$, and $c = 8 = g_{\text{car}} + N_{\text{fam}}$ is itself structural.
2. **Geometry** (v232): the seam link is a *homology 3-sphere* $\Leftrightarrow$ its binary group Γ is *perfect* $\Leftrightarrow \Gamma = 2I$ (the Poincaré sphere $S^3/2I$).
3. **Representation theory** (v219): exactly *one* 1-dim irrep (the trivial one).

These are the *same* number: $\#(1\text{-dim irreps}) = |\Gamma^{\text{ab}}| = |H_1(S^3/\Gamma)| = \#(\text{mark-1 affine-Dynkin nodes})$, which equals 1 *only* for E_8 ($A_n \rightarrow n+1$, $D_n \rightarrow 4$, $E_6 \rightarrow 3$, $E_7 \rightarrow 2$, $E_8 \rightarrow 1$). So $2I$ is the unique nontrivial perfect finite $SU(2)$ subgroup and $S^3/2I$ the unique homology-sphere space form. **Consequence:** Research Contract 2 (G_{net} /Target A), the carrier P2, and the Kleinian seam are *not* separate gates — they are this *one* condition.

The closing chain (what a complete closure would be), [O] at one step:

$$\begin{array}{c} \text{RP} + \text{gap} + \text{chirality} \xrightarrow{\text{v160}} \text{quasi-free bulk} \xrightarrow{?} \text{holomorphic boundary net} \\ \xrightarrow{[\text{E}]} (E_8)_1 \xrightarrow{\text{v219/v228}} 2I, (2,3,5), \text{P2} \xrightarrow{\text{v179-v181+Kleinian genus-0}} \text{QGEO.SYM.01.} \end{array}$$

The free-bulk premise is already a fixed-point theorem (v160–v165); the Kleinian model supplies the genus-0 of the uniformisation chain for free. **The single remaining analytic step is “free Gaussian RP bulk $\Rightarrow$ holomorphic (single-sector) boundary net”** — the bulk–boundary statement that

an invertible (Gaussian) bulk has a holomorphic chiral boundary. That step is now *closed modulo a cited theorem*: the lattice/ S_3 stack pins it at every computable level (v367/v368, v376–v379, Lean FORM.SEAM.MMST.01), only the cited continuum existence staying [O] (v336); v_{geo} stays the No-Unit metrology primitive and F_{transfer} stays external physics. This box certifies that the three faces coincide and force E_8 ([E]); the analytic step is closed modulo that cited theorem.

A physical realisation of the step ([verification/v235_seam_chern_simons.py]). The step lands in standard physics. A free gapped *bosonic* 2+1d bulk is an abelian Chern–Simons / K -matrix theory (even integer K), where $\#\text{anyons} = |\det K|$, the edge chiral central charge is $c = \text{signature}(K)$, and the edge net is *holomorphic* $\Leftrightarrow |\det K| = 1$ — so “no nontrivial abelian sector” is precisely $\det K = 1$. The extension tower of v92 is the anyon-condensation tower read by determinant: $D_5 \oplus A_3$ (det 16, 16 anyons) $\rightarrow SO(16)_1 = D_8$ (det 4) $\rightarrow (E_8)_1$ (det 1, holomorphic), all rank 8 with $c = \text{signature} = 8 = g_{\text{car}} + N_{\text{fam}}$, so the holomorphic end is the *Kitaev E_8 quantum-Hall state*. Holomorphy $\Leftrightarrow$ condensing the order- $|\mu_4| = 4$ *Lagrangian* glue (det 16 $\rightarrow$ 16/4²=1); a partial order-2 isotropic only reaches det 4 (D_8). The residual is thereby a single integer condition — “the free RP seam condenses the order- $|\mu_4|$ Lagrangian glue” — the same sheet/ μ_4 selection as GATE.QGEO, but now with holomorphic $\Leftrightarrow$ det 1 a theorem (GATE.HOLO.02, [E]/[O]).

The residual as a physical statement ([verification/v237_seam_sre_closure.py]). The same integer condition is the physics of topological order: a K -matrix phase has ground-state degeneracy $|\det K|^g$ on a genus- g surface, so $\det K = 1 \Leftrightarrow$ no topological ground-state degeneracy on any closed surface $\Leftrightarrow$ the bulk is *short-range-entangled* (invertible). Among rank-8 even positive-definite K ($c = 8 = g_{\text{car}} + N_{\text{fam}}$) only E_8 is SRE — the unique nontrivial bosonic SRE 2+1d phase, the *Kitaev E_8 state*, whose edge is the holomorphic $(E_8)_1$. TFPT already has a unique vacuum on the plane (RP + gap + cluster + OS); the plane cannot see $\det K$ (no non-contractible cycles), the torus can. So the one analytic step is the sharp, falsifiable physical statement “*the free RP gapped seam bulk is short-range-entangled (no topological order)*” ($= \det K = 1 =$ the *Kitaev E_8 phase*, GATE.HOLO.03, [E]/[C]) — and the lattice/ S_3 stack now answers it at every computable level (torus $\det K=1$, v378), so the keystone is *closed modulo a cited theorem*, only the cited continuum existence (v336) staying [O].

Lemma 2 (U1 — wall landscape, corrected). *Parabolic degree zero forces the line weight-sums onto the wall (w_1, w_2, w_3) = (2, 1, 1). An explicit balanced representative (weights in $\frac{1}{3}$) is*

$$W_{\text{wall}} = \frac{1}{3} \begin{pmatrix} 2 & 1 & 2 & 1 \\ 1 & 0 & 0 & 2 \\ 0 & 2 & 1 & 0 \end{pmatrix},$$

*each column a permutation of (0, 1, 2), row sums (2, 1, 1). **Honest result:** the 6⁴=1296 column-permutation matrices give 144 walls; under the full group $D_4 \times S_3(\text{lines}) \times S_3(\text{values})$ they form five orbits, not one. So W_{wall} is not singled out by symmetry — the uniqueness must come from the selector (Lemma U_4), not from a symmetry quotient. The wall landscape is nonetheless a finite, explicitly enumerated set; W_{wall} is one of the five. Machine: fully certified $C_U^{(1)}$ ([verification/v29_research_contract_certs.py]).*

The splitting type is now algebraic (RH route, [verification/v32_rh_splitting.py]). *Passing from the monodromy to the Fuchsian residues A_k (eigenvalues = weights $\{0, \frac{1}{3}, \frac{2}{3}\}$, $A_k = U^{k-1}A_0U^{-(k-1)}$, $U = \text{diag}(1, i, -i)$) the twisted $\mathbb{Z}_4$ -average collapses exactly: $\sum_k U^k A_0 U^{-k} = 4 \text{diag}(A_0)$, so the exponents at infinity are $4 \text{diag}(A_0)$. A flat bundle with regular ∞ needs integer exponents summing to $4 = -\deg E$; with the cusp weights the only options are perms of (2, 1, 1) and (2, 2, 0), i.e.*

$$\mathcal{O}(-2) \oplus \mathcal{O}(-1)^2 \iff \text{diag}(A_0) = (\tfrac{1}{2}, \tfrac{1}{4}, \tfrac{1}{4})$$

(the sibling (2, 2, 0) $\rightarrow \mathcal{O}(0) \oplus \mathcal{O}(-2)^2$). Such a Hermitian A_0 (spectrum $\{0, \frac{1}{3}, \frac{2}{3}\}$, that diagonal) exists by Schur–Horn. So the splitting is an algebraic condition on $\text{diag}(A_0)$. Honest wall: the global $\prod M_k = \mathbf{1}$ is the path-ordered monodromy (not $\exp(2\pi i \sum A_k)$), constraining off-diag(A_0) via the genuine RH solve; and R-extraction still needs the $C_6 \leftrightarrow \mathbb{P}^1 \setminus \mu_4$ bridge. The RH route fixes the splitting algebraically but not c_u/c_d .

Lemma 3 (U2 — the kill switch, run and hardened [E]). *The selected polystable wall representative must not collapse to the diagonal abelian $U(1)^3$ point if it is to give non-trivial $SU(3)_F$ transport mixing: (A) ∇_F^* polystable but not simultaneously diagonal (the desired breakthrough), or (B) diagonal collapse, in which case c_u/c_d stays honestly [C]. **Result: the cheap test is run robustly and lands on A.** The kill-switch fires iff the D_4 -symmetric cusp representation is reducible (commutant non-scalar). Across 400 random cusp-class generators of the $\mathbb{Z}_4$ family $M_k = U^{k-1} M U^{-(k-1)}$ the commutant is scalar at every point (irreducible 400/400), the reducible locus is hit with probability 0 (measure-zero, codimension ≥ 1), and the mixing strength stays bounded away from zero (min off-diag = 0.35 > 0.05). So case A is the generic behaviour, not an accident of the one explicit point of [verification/v33_explicit_flat_bundle.py]: (U_{wall}) survives its own cheapest falsification test. [verification/v38_uwall_killswitch.py]*

Lemma 4 (U3 — D_4 fixed locus as an explicit variety, positive-dimensional). *With $M_k = U^{k-1} A U^{-(k-1)}$ ($U = \text{diag}(1, i, -i)$, the $\mathbb{Z}_4$ rotation) and the reflection $M_1 = V M_1^{-1} V^{-1}$, $M_3 = V M_3^{-1} V^{-1}$, $M_2 = V M_2^{-1} V^{-1}$ (with $V = -[[1, 0, 0], [0, 0, 1], [0, 1, 0]]$, $V^2 = 1$, $V U V^{-1} = U^{-1}$), the relations become polynomials in the entries of $A \in C_{\text{cusp}}$; $\mathcal{X}_{\text{cusp}}^{D_4}$ is cut out as $\bigcup_k C_k$, parametrised by trace coordinates $x = \text{tr}(M_1 M_2)$ (note $\text{tr} A = 1 + \omega + \omega^2 = 0$). **Result** ($C_U^{(2)}$, [verification/v30_d4_character_variety.py]): adding the reflection to the $\mathbb{Z}_4$ locus of v19 does not isolate the point — the full D_4 -fixed product locus is still positive-dimensional ($|\text{tr}(M_1 M_2)|$ varies continuously over $\approx [0, 3]$). So even the full D_4 symmetry cannot select ∇_F^* . Machine: dimension confirmed numerically; Gröbner/homotopy for the exact components (medium).*

Lemma 5 (U4 — selectors as equations). *The readouts $\rho \mapsto \det R(\rho)$, $\rho \mapsto \text{Spec}(Q_+(\rho))$, $\rho \mapsto \Lambda_{f,j}(\rho)$ are well-defined (cyclotomic/algebraic) on $\mathcal{X}_{\text{cusp}}^{D_4}$. **Acceptance criterion:***

$$\#\{\rho \in \mathcal{X}_{\text{cusp}}^{D_4} : \det R = 8, \text{Spec}(Q_+) = \{1, 2, 3\}\} = 1$$

*in the polystable quotient. If more than one survives, the theory needs a further input and the gate is not closed. **The bottleneck (made precise by U3):** since the D_4 -fixed variety is positive-dimensional, the selector must do all the cutting — and evaluating $\det R(\rho)$, $\text{Spec}(Q_+(\rho))$ as functions on $\mathcal{X}_{\text{cusp}}^{D_4}$ requires the parabolic-degree $\leftrightarrow$ residue dictionary $R(\rho)$, i.e. the H2 equivalence. That dictionary was the single remaining analytic input of the (U_{wall}) gate. Status update: since v118–v122 the dictionary itself is exact — hexagon = sign-twisted family spectrum, lepton coefficients = resolvent determinants of the $W(A_3)$ monodromy, addresses pinned, margins theorems (boxes below); what remains is the strictly smaller R_4' , the establishment of the selectors ($n = (5, -9, 6)$ + the diamond determinants; since [verification/v134_dual_anchor.py]: the pair (d, n)).*

What $R(\rho)$ is ([verification/v31_R_dictionary.py]). *Two facts pin its nature. (i) Case A is alive: every D_4 -fixed cusp monodromy found is irreducible, so the wall point carries non-trivial $SU(3)_F$ mixing — the U2 kill switch does not fire to the diagonal case B. (ii) $R(\rho)$ is not algebraic: R is integer-valued, but the algebraic invariants of ρ (e.g. $\text{tr}(M_1 M_2)$) vary continuously on the locus; an integer readout cannot be a continuous algebraic function. Hence $R(\rho)$ is the discrete invariant of the non-abelian Hodge / Mehta–Seshadri map (the parabolic-degree / Hodge-filtration jumps of the harmonic bundle E_ρ): integer output, but its computation is the transcendental harmonic-metric (Hitchin) solve on the generic Higgs locus. **Refinement (U4'',** [verification/v40_harmonic_metric.py]): *at the polystable point TFPT actually selects this pessimism does not apply — see the box below.**

Progress: the selector side of U4 is closed on the explicit point ([E]). While $R(\rho)$ as a function on the whole locus needs the Hitchin solve, the two U4 selectors are algebraically accessible on the explicit flat bundle of [verification/v33_explicit_flat_bundle.py] — no PDE. Computed exactly from A_0 : the exponents at infinity $\text{eig}(\sum_k A_k) = \{2, 1, 1\}$ give the splitting type $E = \mathcal{O}(-2) \oplus \mathcal{O}(-1)^2$, i.e. the parabolic anchor $a = (1, 1, 2)$; the parabolic degree is 0 ($\deg E = -4 = -|\mu_4|$ plus four cusp weight-sums of 1), so the point is *polystable* (Mehta–Seshadri

substrate). The two selectors then read off the bundle directly:

$$\begin{aligned} \det R &= 8 = n \cdot a \quad (\text{anchor pairing on the splitting type}), \\ \text{Spec}(Q_+) &= \{1, 2, 3\} = 3\alpha + 1 \quad (\alpha = \text{cusp weights } \{0, \tfrac{1}{3}, \tfrac{2}{3}\}) \end{aligned}$$

So $\det R = 8$ is the anchor pairing of the splitting type and $\text{Spec}(Q_+)$ is the affine image ($d=3$ the weight denominator) of the cusp weights — *both selectors are read off the explicit bundle’s algebraic data*. **Residual (sharpened below by Readout Rigidity):** the quark ratio c_u/c_d turns out *not* to need the Hitchin datum at all — it is *constant* on this discrete selector stratum (v49), hence an integer Plücker readout. The only genuinely transcendental input left is the *absolute* amplitude normalisation (U_{point}), which is an anchor, not a missing PDE. [verification/v39_uwall_selectors.py]

Breakthrough: the harmonic metric is finite linear algebra, not a PDE ([E]). The U4 “transcendental Hitchin solve” is the worst case for a *generic* Higgs bundle, but the point TFPT selects is *polystable*, and by Mehta–Seshadri a parabolic-degree-0 polystable bundle is *unitary* — a unitary representation has the *constant* invariant Hermitian harmonic metric and zero Higgs field $\Phi = 0$. So at the selected point the harmonic metric is finite linear algebra. On the explicit bundle of [verification/v33_explicit_flat_bundle.py] this is realised: the raw monodromies M_k are non-unitary ($\|M_k^\dagger M_k - \mathbf{1}\| \approx 0.83$), but there is a *unique* common invariant Hermitian form H ($\dim \ker\{M_k^\dagger H M_k = H\} = 1$), it is *positive-definite*, and the unitarised holonomy $\widetilde{M}_k = H^{1/2} M_k H^{-1/2}$ is unitary ($\sim 10^{-8}$), lies in the cusp class $\{1, \omega, \omega^2\}$ with $\prod \widetilde{M}_k = \mathbf{1}$, and has *clean* matrix elements $|\widetilde{M}_k|_{ij} \in \{0, \frac{1}{2}, \frac{1}{\sqrt{2}}\}$. **So $C_U^{(3)}$ reduces from a transcendental PDE to: this finite unitarisation (done) + the combinatorial word-length R (closed: $\mathbf{H1} + \det R = 8$).** *Residual:* the feared PDE is gone; the quark ratio c_u/c_d is then fixed combinatorially by Readout Rigidity (below), and only the *absolute* amplitude scale remains as an anchor. [verification/v40_harmonic_metric.py]

The final leg-assignment test (honest result) ([E]). Does the harmonic-frame holonomy $\{0, \frac{1}{2}, \frac{1}{\sqrt{2}}\}$ feed the $\delta = \frac{1}{2}$ resolvent to give the lepton amplitudes $\Lambda = (0.475, 1.107, 0.917)$? **(i) Ruled out as a literal identity:** $\Lambda_\mu = 1.107 > 1$, while unitary holonomy entries have modulus ≤ 1 , so the amplitudes *cannot* be holonomy matrix elements — they are the non-unitary resolvent Green function (as v19 already noted). **(ii) Suggestive positive:** the harmonic-frame holonomy diagonal modulus is exactly $|\text{diag } \widetilde{M}_0| = (0, \frac{1}{2}, \frac{1}{2})$, and that $\frac{1}{2}$ *equals* the distinguished lepton transport value $\delta = \frac{1}{2}$ that v20 used by hand — so the geometry plausibly *supplies* δ (an observation at the explicit point; genericity across the locus is future work). **(iii) Clean negative:** no natural holonomy construction (e.g. $(\mathbf{1} - \frac{1}{2}\widetilde{M})^{-1}$) maps the holonomy to the amplitude vector. **Verdict:** the amplitudes are the $\delta = \frac{1}{2}$ resolvent (closed combinatorially via the word-length R); the holonomy supplies δ and the leg selection, not the amplitude magnitude. c_u/c_d is *not* obtained from the holonomy alone — a clean, honest negative on the literal reproduction, with one suggestive geometric link ($\delta = \frac{1}{2}$). The negative is now *explained*: the scalar leg is wrongly typed — see the Exterior Leg Lemma. [verification/v41_leg_assignment.py]

The torsion normal form and the $\delta = \frac{1}{2}$ theorem ([E]). The v41 “genericity is future work” item is now executed — and yields a theorem. *(i) Flatness = μ_4 torsion [E]:* for any M (symbolic), $\prod_k U^k M U^{-k} = (MU)^4$ — the $\mathbb{Z}_4$ -family flatness constraint *is* the torsion statement “ $T = MU$ is a fourth root of unity”: the μ_4 atom appears as the *order* of the twisted cusp generator. *(ii) The $\delta = \frac{1}{2}$ theorem [E] (both directions):* on the *involution* branch ($T^2 = \mathbf{1}$) T is a Hermitian reflection $2vv^\dagger - \mathbf{1}$, and the cusp trace condition $2v^\dagger U^{-1}v = 1$ splits *exactly* into $|v_1|^2 = \frac{1}{2}$ (real part) and $|v_2| = |v_3|$ (imaginary part), forcing

$$\text{diag } M = (0, \tfrac{i}{2}, -\tfrac{i}{2}), \quad \text{spec } M = \{1, \omega, \omega^2\} \text{ automatic}$$

($\chi_M = \lambda^3 - 1$ from unitary + trace 0 + $\det 1$). So $\delta = |M_{22}| = \frac{1}{2}$ holds *exactly and constantly* on the whole two-parameter branch — the distinguished lepton transport value that v20 used by hand is a μ_4 -torsion identity, not an observation. *(iii) Reflection lemma [E]:* the D_4 reflection

forces $\text{diag } M = (r, z, \bar{z})$ with r real and $\text{Re } z = -r/2$ — the whole D_4 diagonal freedom is *one* real parameter, and $\delta = \frac{1}{2} \iff r = 0$. (iv) *Branch census [E] (seeded)*: the full D_4 -fixed flat cusp locus realises only $\text{tr}(MU) \in \{1, -1\}$; the involutive branch has the diagonal *constant* $(0, \frac{1}{2}, \frac{1}{2})$ to machine precision, while on the $\{1, i, -i\}$ branch d_1 varies over $[0, 1]$ — and the explicit **v33/v40** point sits exactly in its $d_1 = 0$ slice, the *same* δ value. *Residue [C] (recorded)*: whether the anchor splitting $\mathcal{O}(-2) \oplus \mathcal{O}(-1)^2$ (Mehta–Seshadri side) forces the $d_1 = 0$ slice — or selects the involutive branch — is the remaining question; $\delta = \frac{1}{2}$ itself is no longer the open item. [verification/v114_torsion_delta.py]

The anchor pins the residue: an exact normal form for the (U_{wall}) bundle ([E]). The **v114** branch question is answered — and the **v33** *numerical* Riemann–Hilbert point acquires an *exact closed form*. (i) μ_4 -average lemma [E]: $\sum_k U^k X U^{-k} = 4 \text{diag } X$ for any X — μ_4 conjugation-averaging *is* the diagonal projection, so the exponents at infinity are $4 \text{diag } A_0$ and the anchor splitting holds *iff* $\text{diag } A_0 = (2, 1, 1)/4 = a/|\mu_4|$: the residue diagonal **is** the anchor over μ_4 (**v33**’s diagonal was forced, not chosen). (ii) *The $(8, 0, 5)/144$ lemma [E]*: the anchor diagonal and the cusp spectrum $\{0, 1, 2\}/N_{\text{fam}}$ pin $\sum |a_{ij}|^2 = \frac{13}{144}$, and with $a_{13} = 0$ the determinant condition solves *uniquely* to

$$\boxed{(|a_{12}|^2, |a_{13}|^2, |a_{23}|^2) = \frac{(8, 0, 5)}{144}} \quad \begin{array}{l} 8 = \text{rank } E_8, \quad 5 = g_{\text{car}}, \\ 144 = (|\mu_4| N_{\text{fam}})^2, \quad 8+5=13=\Delta_Q \end{array}$$

— the off-diagonal weights are compiler atoms, and the total 13 is the *same* Δ_Q that denominates $c_u/c_d = 55/117 = 5 \cdot 11/(9 \cdot 13)$. (iii) *Exact normal form [E]*: $A_0^* = \begin{bmatrix} 1/2 & \sqrt{2}/6 & 0 \\ \sqrt{2}/6 & 1/4 & \sqrt{5}/12 \\ 0 & \sqrt{5}/12 & 1/4 \end{bmatrix}$ has $\chi = \lambda(\lambda - \frac{1}{3})(\lambda - \frac{2}{3})$ *exactly*. (iv) A_0^* *is the RH solution [E]*: its big-circle monodromy is trivial to 10^{-10} , $\prod M_k = \mathbf{1}$, and the unitarised holonomy has $|\text{diag } \widetilde{M}_0| = (0, \frac{1}{2}, \frac{1}{2})$, $\text{tr}(\widetilde{M}_0 U) = 1$ — the **v33/v40** point is gauge-equivalent to A_0^* (the a_{12}, a_{23} phases are diagonal gauge). (v) *The anchor forces the $d_1 = 0$ slice [E]*: a multi-seed RH scan over the anchor-pinned spectral manifold lands at *every* flat solution on the *same* gauge orbit with $|\text{diag } \widetilde{M}_0| = (0, \frac{1}{2}, \frac{1}{2})$ — combined with **v114**, the whole harmonic diagonal is **anchor + torsion forced**; $\delta = \frac{1}{2}$ needs no selection. *Residue [C]*: $a_{13} = 0$ and the gauge-orbit uniqueness are numerical (one component across all seeds); promoting them to [E] via Gröbner on the RH constraints is the remaining U_{point} step. [verification/v115_anchor_residue.py]

The resonance theorem: $M_\infty = \mathbf{1} \iff a_{13} = 0$ — **one gauge orbit [E]**
([verification/v116_resonance_uniqueness.py])

The anticipated Gröbner promotion turned out to be a *linear* computation. Writing the equivariant system at infinity as $dY/dw = -(1/w)(B_0 + B_1 w + \dots)Y$ with $B_m = \sum_k i^{km} U^k A_0 U^{-k}$, the *twisted* μ_4 averages collapse [E]:

$$B_0 = 4 \text{diag } A_0, \quad \boxed{B_1 = 4a_{12}E_{12} + 4\bar{a}_{13}E_{31}}$$

(only the eigenvalue-ratio $-i$ cells survive). The anchor exponents $\text{diag}(2, 1, 1)$ have resonance gap 1; the level-1 formal-gauge equation $(1 - \text{ad } R_0)H_1 = R_1$ is singular exactly on the cells $\{(2, 1), (3, 1)\}$, where B_1 has entries $(0, 4\bar{a}_{13})$; all higher levels are non-resonant and the formal monodromy is $e^{-2\pi i \text{diag}(2, 1, 1)} = \mathbf{1}$. Hence [E]:

$$\boxed{M_\infty = \mathbf{1} \iff a_{13} = 0}$$

— the transcendental Riemann–Hilbert condition collapses to *one linear equation*. *Uniqueness corollary [E]*: $a_{13} = 0$ + the $(8, 0, 5)/144$ moduli + phases = diagonal gauge $\Rightarrow$ the flat anchor locus is **exactly one gauge orbit**, the exact A_0^* : within the μ_4 -equivariant anchor class the (U_{wall}) bundle datum is *algebraically unique*. *Sharpness [E]*: $\|M_\infty - \mathbf{1}\| \sim 8.5 |a_{13}|$ (0.17 at $a_{13} = 0.02$) vs 10^{-10} at A_0^* . *Honest scope [C]*: the equivariant ansatz is the established D_4 selector input (**v19/v30**); the *residue* side is now [E], while the holonomy values (the harmonic diagonal $(0, \frac{1}{2}, \frac{1}{2})$) remain [E] — transcendental monodromy of an exact system.

The monodromy is the Weyl group of A_3 — the holonomy is exact [E]
 ([verification/v117_monodromy_weyl_a3.py])

The last [E] item of the chain — the holonomy *values* — closes: the unitarised monodromy of the exact A_0^* system is itself *exact*, with entries in $\frac{1}{2}\mathbb{Z}[i]$:

$$\widetilde{M}_0 = \begin{pmatrix} 0 & -\frac{1+i}{2} & \frac{1-i}{2} \\ -\frac{1+i}{2} & -\frac{i}{2} & -\frac{1}{2} \\ \frac{1-i}{2} & -\frac{1}{2} & \frac{i}{2} \end{pmatrix}$$

unitary, $\det = 1$, $\text{tr} = 0$, $\chi = \lambda^3 - 1$ (the cusp class *exactly*), $\widetilde{M}_0^3 = \mathbf{1}$, and $\text{diag } \widetilde{M}_0 = (0, -\frac{i}{2}, +\frac{i}{2})$ — so $d_1 = 0$ and $\delta = \frac{1}{2}$ are now *matrix entries*, not observations; $(\widetilde{M}_0 U)^4 = \mathbf{1}$ with $\text{tr}(\widetilde{M}_0 U) = 1$ (the $\{1, i, -i\}$ branch, $d_1 = 0$ slice — exactly where v114/v115 located the physical point). *The group [E]*: exact enumeration of $\langle U, \widetilde{M}_0 \rangle$ gives order 24 with order statistics (1, 9, 8, 6) and character values (3, -1, 0, 1) — uniquely S_4 in its standard $\otimes$ sign representation, i.e. the (twisted) reflection representation of the **Weyl group of the family lattice A_3** :

$$\text{monodromy}(U_{\text{wall}}) = W(A_3) = S_4, \quad |W(A_3)| = 4! = 24$$

— the flavor wall realises the A_3 factor of the glue $D_5 \oplus A_3 + \mu_4$ as its Weyl group (U = a 4-cycle, the μ_4 deck; $\widetilde{M}_0$ = a 3-cycle, the family rotation). *Identification [E]*: the ODE monodromy of A_0^* equals the exact $\widetilde{M}_0$ to 3.5×10^{-10} , and the invariant form H is U -invariant to 10^{-11} . *Status*: holonomy values [E]; only the analytic identification stays [E]. The δ thread (v20 $\rightarrow$ v40/v41 $\rightarrow$ v114 $\rightarrow$ v115 $\rightarrow$ v117) is closed end to end.

The hexagon is the sign-twisted family spectrum — first exact H2 piece [E]
 ([verification/v118_hexagon_family_dictionary.py])

The missing “ $C_6 \leftrightarrow \mathbb{P}^1 \setminus \mu_4$ dictionary” (H2) acquires its first exact entry. *Sign-twist lemma [E]*: $-\widetilde{M}_0$ has order 6 and

$$\text{spec}(\widetilde{M}_0) \cup \text{spec}(-\widetilde{M}_0) = \mu_6$$

— the six-site hypercharge hexagon spectrum *is* the family monodromy spectrum together with its sheet twist ($\mathbb{Z}_6 = \mathbb{Z}_2 \times \mathbb{Z}_3$ realised as $(-1) \times \langle \widetilde{M}_0 \rangle$); the v20 denominators $\frac{5}{4} - \cos(r\pi/3)$ are exactly the eigen-denominators $|1 - \zeta_6^r/2|^2$ of the two family resolvents $(\mathbf{1} \mp \frac{1}{2}\widetilde{M}_0)^{-1}$. *Cyclotomic determinant [E]*: $\det(\mathbf{1} - t\widetilde{M}_0) = 1 - t^3 = (1 - t)\Phi_3(t)$, so $\det(\mathbf{1} \mp \frac{1}{2}\widetilde{M}_0) = (\frac{7}{8}, \frac{9}{8})$ — and the lepton coefficients become resolvent determinants: $c_e = |\mu_4|/(2\det_-) = \frac{16}{7}$, $c_\mu = N_{\text{fam}}/(2\det_+) = \frac{4}{3}$, $c_\tau = |\mu_4|\det_-/N_{\text{fam}} = \frac{7}{6}$, product rule = $|\mu_4|/\det_+ = \frac{32}{9}$. The lepton 7 and 9 *are* the family-resolvent determinants; $\delta = \frac{1}{2} = |(\widetilde{M}_0)_{22}|$ is supplied by the matrix itself. *Sheet-extended group [E] (audit)*: $\langle U, \widetilde{M}_0, -\mathbf{1} \rangle$ has order $48 = \Omega_{\text{adm}} = N_{\text{fam}} \dim S^+$ — the seed’s quartic coefficient. *Residue [C]*: the dictionary’s *values* are exact; its *address table* (which hexagon site each fermion occupies; the quark composition) remains v20’s established input / open — nothing re-fished.

The address table: the lepton words are the compiler atoms [E]/[C]
 ([verification/v120_address_table.py])

The address table acquires its structure — on word-length integers *frozen* in v18/v20, long before today’s readings (the red-team firewall). (i) *The lepton words are the atoms [E]*:

$$(L_e, L_\mu, L_\tau) = (8, 5, 3) = (\text{rank } E_8, g_{\text{car}}, N_{\text{fam}}) = (p_0 + e_2, e_2, p_0)$$

in mass order (longest word = lightest lepton). (ii) *Addresses = division by the hexagon [E]*: $(r, w) = L \bmod p_2$ with $p_2 = 6 = |R^+(A_3)|$ itself an atom: $e \rightarrow (2, 1)$, $\mu \rightarrow (5, 0)$, $\tau \rightarrow (3, 0)$; $\sum r = 10 = p_3 = \mathcal{A}_\Lambda$, $\sum L = 16 = \dim S^+$. (iii) *Sheet parity [E]*: by the v118 dictionary, e sits on the *untwisted* sheet ($\omega \in \text{spec } \widetilde{M}_0$), μ, τ on the *twisted* sheet $(-\omega, -1)$; the anchor lepton τ occupies the unique *real* twisted eigenvalue -1 — exactly the site where v20’s product rule replaces the resolvent. (iv) *Quark sum rules [E] (frozen v18 words 7, 7, 5, 3, 2, 0)*: up-type sum = $10 = p_3$; down-type sum = $14 = p_1 + p_3$; **all quarks** = $24 = |W(A_3)|$ (the monodromy group order, v117); **all nine charged fermions** = $40 = p_1 p_3 = a^T Ra$

(v119); quark $\sum r = 2p_2$, $\sum w = |\mathbb{Z}_2|$; the top sits at the vacuum site $(0, 0)$ — zero word, no suppression, the established ladder anchor. *Residue [C]*: five independent closures $(16, 10, 14, 24, 40)$ constrain the address map; the per-fermion *assignment* remains the v18/v20 input — deriving it is the remaining H2 content.

The address table is pinned: R is the unique hexagon matrix with atom margins [E] ([verification/v121_address_pinning.py])

The per-fermion assignment closes at the *information* level. The established identity $L = R + 2U$ (v95) says the word-length table *is* $L[\text{sector}][\text{generation}]$ — rows (up; down; lepton) = $((7, 3, 0); (7, 5, 2); (8, 5, 3))$, the v18/v20 words verbatim — so the addresses are the entries of the *residue operator* R plus one hexagon turn for the first generation, and R 's first column is the **anchor** $a = (1, 1, 2)$. The margins are atoms [E]: row sums $(4, 8, 10) = (e_1, \text{rank } E_8, p_3)$, column sums $(4, 13, 5) = (e_1, \Delta_Q, g_{\text{car}})$. And the census closes it [E]:

among all 3×3 matrices with entries in the hexagon $\{0..5\}$,
exactly 17 have the atom margins — and exactly ONE has $\det = 8$

— and it is R (also the unique one with $\text{SNF} = (1, 1, 8)$; all 17 determinants are distinct, 8 occurs once — control computed explicitly). **The assignment table carries zero residual information beyond the atom margins and the rank determinant.** Within-sector ordering is the established transport reading ($m \sim \lambda_V^L$: longer word = lighter fermion). *Residue [C]*: the open H2 item shrinks from “derive 9 per-fermion integers” to “derive the six atom margins + $\det = \text{rank}$ ”. *Firewall*: R , L and all word lengths were frozen in v4/v18/v20/v71 long before this census — no degrees of freedom spent.

The margins are theorems ([verification/v122_margin_theorem.py]). Even the margins are not input. Three selectors, all established and frozen *before* the address question was posed — (S1) the D_4 annihilator $n = (5, -9, 6)$ with $n^T R = (8, 0, 0)$ (it kills the unchanged (c_2, c_3) plane; $n \cdot a = 8$), (S2) $\det R = 8$, (S3) $\det K = 4$ with $K = R + Q\Sigma$ — pin R *uniquely* among hexagon matrices [E]: the (S1) census has $4 \times 4 \times 4 = 64$ candidates, (S2) leaves 12, (S3) leaves **exactly one**, and it is R . Hence the atom margins, the anchor column $Re_1 = a$ and all sum rules above are *consequences* — their input status is revoked. *Bonus lemma [E]*: on the (S1) + (S2) locus $\det(M + 2U) = 20 = \det L$ holds for *all* 12 candidates — the fourth quartet determinant is redundant, the compiler's familiar overdetermination. *Residue [C]*: H2 reduces to the establishment of the selectors themselves (n and the diamond determinants — pre-existing inventoried objects); the address question introduces *no new residual class*.

The dual-normal selector: (d, n) **pins** R **columnwise** ([verification/v136_dual_normal_selector.py]). The selector list shrinks once more. With the dual anchor $d = a^\top R^{-1}$ ([verification/v134_dual_anchor.py]) and the torsion normal n , each *column* of R obeys $d \cdot c_j = a_j$ and $n \cdot c_j = 8\delta_{1j}$; the common kernel of (d, n) is the primitive vector $(6, 8, 7)$ (audit: p_2 , rank E_8 , scalaron-7), so within the address box $\{0..5\}^3$ each column is the *unique* lattice point (census: 1 of 216 each) — $(d, n) \Rightarrow R$, and the 6^9 census above collapses to three column problems [E]. *Honesty*: K is *not* pinned directly ($d^\top K = (1, \frac{1}{2}, 1)$); $\det K = 4$ comes via the μ_4 wall ([verification/v135_det_surface.py]). *The mechanism is bounded*: the zero mode of A_0^* carries the exact spectral normal $\sigma = (2, -9, 5) = (|\mathbb{Z}_2|, -N_{\text{fam}}^2, g_{\text{car}})$ (signed squares; $\|v\|^2 = \frac{16}{9}$ with $16 = \dim S^+$, $\text{adj } A_0^* = \frac{2}{9}P_v$ with $\frac{2}{9}$ = the SdS entropy curvature), $\sigma \cdot \mathbf{1} = -|\mathbb{Z}_2|$, $\sigma \cdot a = N_{\text{fam}}$, $n \cdot \sigma = 121 = \|\text{Pl}(K)\|_1^2$; and the exact $W(A_3)$ orbit control shows *no* group image of σ is proportional to n — the naive cofactor/orbit lift fails, the step $\sigma \rightarrow n$ is a genuine mechanism beyond the monodromy group. R_4' *restated [C]*: derive the *pair* (d, n) ; then R is forced columnwise.

The selector triangle ([verification/v139_selector_triangle.py]). The pair itself now collapses further. First, d is *pure anchor data*: $d = \frac{3}{2}a - 2 \cdot \mathbf{1}$ exactly — the closed form never mentions R , so the first selector is *derived*, not residual. Second, the frame $(\mathbf{1}, a, \sigma)$ is a basis of $\mathbb{Z}^3$ with $\det(\mathbf{1}|a|\sigma) = 11 = \|\text{Pl}(K)\|_1$ (the quark constant as the frame volume), and n is the *unique* integer covector with the three atom pairings

$$n \cdot \mathbf{1} = 2 = |\mathbb{Z}_2|, \quad n \cdot a = 8 = \text{rank } E_8, \quad n \cdot \sigma = 121 = \|\text{Pl}(K)\|_1^2$$

— the cofactor definition is replaced by frame data plus three atoms. On the $(\mathbf{1}, a)$ -constrained line the σ -pairing is $11(11+t)$: divisible by 11 for *every* t , a perfect square exactly at the true n ($t=0$). The

review's predicted lift is exact as an audit identity: $n = \text{reverse}(\sigma) + |\mu_4|e_3$ (generation flip + μ_4 lift at the third slot).

Frame integrality: three pairings → two ([`verification/v142_frame_integrity.py`]). The third pairing is *not* independent. The pairing triples $(m \cdot \mathbf{1}, m \cdot a, m \cdot \sigma)$ of *integer* covectors form an index-11 sublattice of $\mathbb{Z}^3$, cut out by the single congruence

$$x - 3y + z \equiv 0 \pmod{11} \quad (\text{index } 11 = \|\text{Pl}(K)\|_1 = \det(\mathbf{1}|a|\sigma)),$$

so with the two line pairings at their atoms $(x, y) = (|\mathbb{Z}_2|, \text{rank } E_8) = (2, 8)$, integrality *forces* $z \equiv 0 \pmod{11}$ — the $11(11+t)$ line *is* the integer line. Primitivity forces t *even* (odd t makes every entry of $n(t)$ even), and $t=0$ is the unique even t with a square σ -pairing for $|t| < 88$ (the global pick keeps the square selector — recorded honestly). Both line pairings are Cramer-reducible to inventoried identities ($n \cdot a = \det R$ by Laplace along the anchor column; $n \cdot \mathbf{1} = 8(R^{-1}\mathbf{1})_1$ with $R^{-1}\mathbf{1} = \frac{1}{4}(1, 1, -1)$, **v134**) — restructuring, not establishment.

The pairing values are atom identities ([`verification/v145_pairing_atoms.py`]). Even the two line-pairing *values* carry no independent information. Reading the **v139** lift structurally — $n = w_0(\sigma) + |\mu_4|e_3$ with w_0 the A_3 *exponent duality* $m \leftrightarrow h-m$ (an established Lie involution; lift slot = the top-exponent line, lift size = the glue order) — all three pairings become atom arithmetic:

$$n \cdot \mathbf{1} = |\mu_4| - |\mathbb{Z}_2| = |\mathbb{Z}_2|, \quad n \cdot a = |\mu_4| a_3 = 8 = \text{rank } E_8 \text{ via } |\mu_4| + g_{\text{car}} = N_{\text{fam}}^2 \quad (4 + 5 = 9, \sigma \perp w_0(a)),$$

and $n \cdot \sigma = \|\sigma\|^2 - N_{\text{fam}}^2 + |\mu_4|g_{\text{car}} = 110 - 9 + 20 = 121$ with $\|\sigma\|^2 = |\mathbb{Z}_2|^2 + N_{\text{fam}}^4 + g_{\text{car}}^2 = 110 = |\mathbb{Z}_2|g_{\text{car}}\|\text{Pl}(K)\|_1$. R_4' *residue [C]* (*final form*): derive the *one* map $n = w_0(\sigma) + |\mu_4|e_3$ — no independent number remains; everything else in the chain anchor $\rightarrow (d, n) \rightarrow R$ is exact.

Both dual normals are anchor data ([`verification/v149_cusp_normal.py`]). The remaining opacity dissolves in the canonical frame: the pairings of n with the *cusp eigenbasis* (the geometric basis of **v98**) are exactly

$$n \cdot e_1 = 6 = p_2(a), \quad n \cdot e_2 = 3 = p_0(a), \quad n \cdot e_3 = 5 = e_2(a)$$

— one anchor invariant per Q_+ degree, and the cusp matrix is unimodular, so these three pairings determine n uniquely. With $d = \frac{3}{2}a - 2 \cdot \mathbf{1}$ the dual normal *pair* (**v134**) is anchor data through and through — one normal per frame; the cusp data $(6, 3, 5)$, the frame pairings $(2, 8, 121)$ and the w_0 -lift are three presentations of one covector, all atomic (equivalences exact). *Audit*: $\sigma \rightarrow (5, 1, 2)$; alternate reading $(6, 3, 5) = (|R^+(A_3)|, N_{\text{fam}}, g_{\text{car}})$; sum $14 = \dim G_2$. R_4' *residue [C]* (*cusp form*): derive the *assignment* — why the Q_+ degrees $(1, 2, 3)$ carry (p_2, p_0, e_2) in that order; one discrete assignment, zero continuous freedom.

Exterior Leg Lemma — the quark u/d leg is a $\Lambda^2 F$ area, not a scalar [E]/[C]

The **v41** negative is a *category error*, not a dead end. u and d share the hexagon-radial data $(r, w) = (1, 1)$ ($L = 7$ for both), and the C_6 resolvent $c(r, w) = |\mu_4|^w / (\frac{5}{4} - \cos \frac{r\pi}{3})$ reads *only* (r, w) — so any *scalar* leg gives $c_u/c_d \equiv 1$. A scalar sees radius, not orientation. The u/d separation lives in the smallest *non-scalar* invariant of the $SU(3)_F$ composition: the exterior square $\Lambda^2 F$, i.e. the oriented anchor-plane area $\text{Pl}_{1,a}(K)$. The “11” is then *generated*, not read:

$$(K, \mathbf{1}, a) \rightarrow B_K = \begin{pmatrix} 13 & 8 & 4 \\ 18 & 11 & 6 \end{pmatrix} \rightarrow \text{Pl}(K) = (-1, 6, 4) = (-N_\Phi, |R^+(A_3)|, |\mu_4|) \rightarrow \|\text{Pl}(K)\|_1 = 11,$$

$$\frac{c_u}{c_d} = \frac{g_{\text{car}} \|\text{Pl}_{1,a}(K)\|_1}{N_{\text{fam}}^2 \Delta_Q} = \frac{5 \cdot 11}{9 \cdot 13} = \frac{55}{117} \quad (\text{anchor microcode} = \frac{e_2(a)(p_3(a)+1)}{p_0(a)^2(2p_2(a)+1)}).$$

Typing: [E] for the Plücker / microcode identities ($\text{Pl}(K) = (-1, 6, 4)$, $\|\cdot\|_1 = 11$, $5 \cdot 11 / (9 \cdot 13) = 55/117$); and — by Readout Rigidity (**v49**, below) — [E] also for the *physical* ratio $\mathcal{C}_{ud}(\rho^*) = \text{this } \Lambda^2 F \text{ readout}$, which is *constant* on the derived selector stratum and so does *not* need any point-selection on $\Lambda^2 F$. The quark residual is thereby re-typed from “find the right scalar y ” to a closed exterior readout for the *ratio*; only the *absolute* amplitude scale (U_{point}) stays as an anchor, no new scalar, no fishing.

Audit (the 121 lemma, [`verification/v119_review_validation_2.py`]): the “11” has a second, independent appearance [E]: with $h = \mathbf{1} + a = (2, 2, 3)$ in the canonical anchor ordering, $h^T R h = 121 = 11^2 = \|\text{Pl}(K)\|_1^2 = (p_3 + 1)^2$ — the quadratic *anchor norm* of the residue operator R reproduces the Plücker integer (sensitivity control: the three orderings of h give $\{105, 121, 135\}$; only

the canonical one yields 121). Bonus checksums: $\mathbf{1}^T R \mathbf{1} = 22 = 2 \times 11$ (the entry sum of R) and $a^T R a = 40 = p_1 p_3 = 10 b_1 - 1$. Typed *audit*, not load-bearing — the oriented Plücker mechanism above remains the derivation. [verification/v42_exterior_leg.py]

Bridge test ([verification/v43_exterior_bridge.py]). The typing is confirmed: for the $SU(3)$ holonomy the exterior square is the conjugate representation, $\Lambda^2(\widetilde{M}) = \widetilde{M}$ exactly — the exterior leg *is* the $\bar{\mathbf{3}}$. But the *continuous* action $\Lambda^2(\widetilde{M}) \cdot (\mathbf{1} \wedge a)$ does *not* equal the *integer* $\text{Pl}(K) = (-1, 6, 4)$: the integer Plücker is the combinatorial K datum. So the bridge $\rho^* \rightarrow \text{Pl}(K)$ is the *discrete* non-abelian-Hodge invariant (the integer R ; $\det R = 8$ and $\text{Spec}(Q_+) = \{1, 2, 3\}$ already confirmed on the bundle, Lemma U4 box), not a continuous exterior-square computation — exactly the H2 equivalence, now pinned from the $\Lambda^2 F$ side.

Lie-level grounding — no special math ([verification/v44_carrier_exterior.py]). The discrete invariant has a home already in TFPT: the carrier half-spinor is the *even exterior algebra* of the 5-slot carrier, $16 = \Lambda^{\text{even}}(5) = \binom{5}{0} + \binom{5}{2} + \binom{5}{4} = 1 + 10 + 5$, and under $SU(5) \subset SO(10)$ the up quark sits in $10 = \Lambda^2(5)$ while the down sits in $\bar{5} = \Lambda^4(5)$. So the *exterior degrees* are $\deg(u^c) = 2$, $\deg(d^c) = 4$, with $\deg(u) + \deg(d) = 6 = |R^+(A_3)|$ (the hexagon) and $\deg(d) - \deg(u) = 2 = |\mathbb{Z}_2|$ (the sheet). The “ Λ^2 ” of the exterior leg is therefore not exotic Hodge data — it is the carrier’s own Λ^2 grading (the up quark *lives* in $\Lambda^2(5)$). This re-types the residual once more: from “discrete non-abelian-Hodge invariant” to “the carrier exterior degree” — a standard [E] fact. *Honest scope*: this grounds the *type*; the Plücker *number* 55/117 stays the family $\text{Pl}(K)$ [E], and the carrier- $\Lambda^2 \leftrightarrow$ family- $\Lambda^2 F$ link (via family $\otimes$ carrier = 15 = 16–1) plus the normalisation remain [C].

The “11” is Pascal — a third, simplest origin ([verification/v45_family_exterior.py]). Following that link, $\text{Pl}(K)$ is the *exterior algebra of the family fundamental* $4 = \mu_4$ (the $A_3 = SU(4)$ fundamental): $\Lambda^k(4) = \binom{4}{k}$ is Pascal row 4 = (1, 4, 6, 4, 1) with full sum $2^4 = 16 = \dim S^+$, and $|\text{Pl}(K)| = (1, 4, 6) = \{\Lambda^0, \Lambda^1, \Lambda^2\}(4)$ are exactly its distinct exterior powers. Hence

$$\|\text{Pl}(K)\|_1 = \sum_{k=0}^2 \binom{4}{k} = 11 = 2^4 - \binom{4}{3} - \binom{4}{4} = 16 - g_{\text{car}}$$

(the cumulative $\Lambda^\bullet(\mu_4)$ up to degree 2 = $\deg(u^c)$), and family $\otimes$ carrier = 15 = $\dim \mathfrak{su}(4) = \Lambda^2(5) \oplus \Lambda^4(5) = 10 + 5 = N_{\text{fam}} g_{\text{car}}$ read three ways. So the “11” is generated a *third* way — the cumulative family exterior algebra up to the up-quark’s degree — pure Pascal/branching, *no special math*. The [C] physical normalisation is unchanged; the algebraic origin of 55/117 is now triply grounded.

Lemma 6 (U5 — the “11” worry, *discharged* by Readout Rigidity). *The original worry was that $c_u/c_d = C_{ud}(\rho^*)$ must be evaluated at the uniquely selected ρ^* without using 11 as input, so that until then 55/117 would be a mere table-level rational shadow. **This is now resolved (v49):** C_{ud} is constant on the discrete selector stratum $S_{8,Q}$, so the “11” is earned as the rigid Plücker readout $\|\text{Pl}(K)\|_1$ (v42) — independently of point-uniqueness. Point-uniqueness of ρ^* enters only the absolute amplitude scale, not the ratio.*

Lemma 7 (U6 — the four certificates). $C_U^{(1)}$ wall enumeration (done); $C_U^{(2)}$ D_4 -fixed character solver; $C_U^{(3)}$ selector uniqueness (one S -equivalence class); $C_U^{(4)}$ amplitude extraction — the ratios $R, Q, c_u/c_d$ are read off the (derived) selector stratum (closed, v49/v71); only the absolute U_f^*, c_q normalisation is the anchor.

Progress: an explicit valid flat bundle is realised (RH solve) ([E]). A numerical Riemann–Hilbert solve produced an *explicit* Hermitian residue A_0 (eigenvalues $\{0, \frac{1}{3}, \frac{2}{3}\}$, $\text{diag} = (\frac{1}{2}, \frac{1}{4}, \frac{1}{4})$) for which the Fuchsian system $A(z) = \sum_k U^{k-1} A_0 U^{-(k-1)} / (z - i^k)$ simultaneously satisfies: the cusp class at each puncture; the splitting $\mathcal{O}(-2) \oplus \mathcal{O}(-1)^2$ (exponents $\sum_k A_k = \text{diag}(2, 1, 1)$); the *path-ordered* monodromy at infinity is trivial, $\|M_\infty - \mathbf{1}\| \sim 10^{-9}$, i.e. $M_1 M_2 M_3 M_4 = \mathbf{1}$; and the representation is *irreducible* — **case A** (non-trivial $SU(3)_F$ mixing, not the diagonal case B). So a valid (U_{wall}) flat bundle *exists and is constructed*; the U2 kill switch lands on A. *Honest scope*: this is one realising point (existence + case A), likely one of a residual family; the unique ∇_F^* still needs the $\det R = 8$ selector and c_u/c_d still needs the $C_6 \leftrightarrow \mathbb{P}^1 \setminus \mu_4$ dictionary (H2). [verification/v33_explicit_flat_bundle.py]

H2 bridge probed ([verification/v34_h2_bridge_attempt.py]). The explicit per-puncture monodromies

M_k (cusp class, $\prod M_k = \mathbf{1}$ to 10^{-5}) were computed; but the natural resolvent-dressed diagonal extraction ($|\text{diag } M_k| = (0, \frac{1}{2}, \frac{1}{2})$ times the C_6 resolvent factor) does *not* reproduce the lepton amplitudes (0.475, 1.107, 0.917). So the precise $\Gamma_{ij}^{\min}$ geodesic-to-word dictionary (how the 6-site hypercharge hexagon combines with the 3-dim family ρ^*) was the genuine remaining analytic input — it is *not* fixed by a natural guess, and we do *not* fish for 55/117. *Resolved since (v118): the correct extraction is the resolvent determinant of the $W(A_3)$ monodromy, not the diagonal dressing — the probe's failure is thereby explained, and the dictionary is exact.*

Theorem U — the four-way split (sharper than one monster gate) [E]/[O]

D_4 fixes the admissible chamber, not the physical point (the D_4 -fixed locus is positive-dimensional, U3); the discrete selectors $\det R = 8$, $\text{SNF}(R) = (1, 1, 8)$, $\text{Spec}(Q_+) = \{1, 2, 3\}$ do the cutting. The gate splits into four pieces, three of them essentially closed:

- U_{unitary} [E]: parabolic degree 0 + polystable $\Rightarrow$ unitary (Mehta–Seshadri) $\Rightarrow \Phi = 0$; the harmonic metric is the *finite* invariant form H ($M_k^\dagger H M_k = H$), not a Hitchin PDE ([verification/v40_harmonic_metric.py]).
- U_{H_2} [E]: R, Q, K are read off the bundle (splitting type = a , $\det R = 8$, $\text{Spec}(Q_+)$) ([verification/v39_uwall_selectors.py]).
- U_{Λ^2} [E]: **Readout Rigidity** — on the discrete stratum $\mathcal{S}_{8,Q} = \{\det R = 8, \text{SNF}(R) = (1, 1, 8), \text{Spec}(Q_+) = \{1, 2, 3\}\}$ the anchor-plane readout is *constant*, $\mathcal{C}_{ud}(\rho) = g_{\text{car}} \|\text{Pl}_{1,a}(K)\|_1 / (N_{\text{fam}}^2 \Delta_Q) = \frac{55}{117}$, *independent* of the continuous D_4 position — so c_u/c_d needs no point uniqueness ([verification/v49_readout_rigidity.py], [verification/v42_exterior_leg.py]).
- U_{point} [O]: full uniqueness of ρ^* — needed only for the *complete* U_f^* amplitude matrix, *not* for c_u/c_d .

Headline: the remaining flavor bridge is not a PDE gate — it is a finite exterior *readout-rigidity* theorem; the only residual, U_{point} , is not an independent gate — it reduces to the single dimensionful anchor v_{geo} (Gate 1 below).

The simple closure ([verification/v71_simple_r_bridge.py]). The selector stratum $\mathcal{S}_{8,Q}$ is now *fully derived*: $\det R = 8 = \text{rank } E_8$ with $\text{SNF}(R) = (1, 1, 8)$ is the lattice invariant (cf. $\det Q = 3 = N_{\text{fam}}$, v70), and $\text{Spec}(Q_+) = \{1, 2, 3\}$ is the D_4 -equivariant result (v69). The four flavor operators are integer lattice maps with compiler-atom determinants ($\det Q, \det K, \det R, \det L$) = (3, 4, 8, 20), product 1920 = $|W(D_5)|$. So by Readout Rigidity the quark *ratios* ($c_u/c_d = g_{\text{car}} \cdot 11 / (N_{\text{fam}}^2 \cdot 13) = \frac{55}{117}$, etc.) are *pure integer Plücker readouts* — *no transcendental monodromy solve*. The earlier monodromy machinery was over-engineering for the ratios; the only transcendental piece, U_{point} , is the *absolute* amplitude normalisation — an anchor of the same nature as the one dimensionful scale, not a missing computation.

Gate 1 complete: $U_{\text{point}} \rightarrow v_{\text{geo}}$ ([verification/v75_upoint_to_vgeo.py]). The absolute amplitudes are not free either. The lepton amplitudes are exact ($c = (\frac{16}{7}, \frac{4}{3}, \frac{7}{6})$, v20), every cross-sector c -ratio is closed (Plücker), and each *sector product* is a clean φ_0^{ret} -power (Grand Mass Volume, v46: $\det M_{\text{sector}} \sim (\varphi_0^{\text{ret}})^{6,9,10}$; the lepton product is $\frac{16}{7} \cdot \frac{4}{3} \cdot \frac{7}{6} = \frac{32}{9} = 2^{g_{\text{car}}} / N_{\text{fam}}^2$). Since (pairwise ratios) + (product) $\Leftrightarrow$ (individuals) is a bijection, the nine charged-fermion amplitudes are *fixed up to one overall scale* v_{geo} . **So U_{point} is not an independent gate: it reduces to v_{geo} , the same one irreducible dimensionful anchor as gravity's $1/G$ (v68) — the two [O] anchors collapse to one.**

Research Contract 2 — (G_{metric})

Goal. Construct the TFPT metric measure

$$Z[J] = \lim_{\chi \rightarrow \infty} \int_{\mathcal{G}_\chi / \text{Diff}} \exp(-S_{\text{rel},\chi}[g, A, \psi] + J \cdot O) d\mu_{\text{Cal},\chi},$$

$$S_{\text{rel},\chi} = \text{Tr} f\left(\frac{D_{\text{rel}} + A_\Sigma}{\chi}\right) - \text{Tr} f\left(\frac{D_{\text{ref}}}{\chi}\right) + \frac{i\pi}{2} \Delta\eta_\Sigma,$$

the relative spectral action whose low-curvature shadow is the already-closed $R + R^2$ equation.

Lemma 8 (G0–G1 — configuration space & BRST/Hodge Fredholm). *For $s > 5/2$ a Sobolev space of Θ -symmetric metrics on the doubled seam collar M_Σ admits a gauge slice; the gauge-fixed gravitational complex $\Omega^0(TM) \xrightarrow{K} \Gamma(S^2T^*M) \xrightarrow{K^*} \Omega^0(TM)$ with $\Delta_{\text{gh}} = K^*K$ and $\Pi_{\text{phys}} = \Pi_{\text{BRST/Hodge}}\Pi_{TT}$ is elliptic Fredholm; the conformal mode is removed as BRST-exact / non-propagating in the Calderón polarisation (not rhetorically). Machine: principal symbols in xAct/-Cadabra; linear-algebra certificates in Lean (medium).*

Lemma 9 (G2 — finite relative spectral action, heat-kernel grounded [E]). *$S_{\text{rel},\chi}$ is finite, differentiable, with a local Seeley–DeWitt expansion $\text{Tr} f(D^2/\chi^2) \sim f_4\chi^4 a_0 + f_2\chi^2 a_2 + f_0 a_4 + \dots$. For the Lichnerowicz Dirac operator $D^2 = \nabla^2 + \frac{R}{4}$ (so $E = -\frac{R}{4}$, spinor trace $\text{tr} \mathbf{1} = 4$) the Gilkey coefficients give, per $(4\pi)^{-2}$,*

$$a_2 = \frac{1}{6} \text{tr}(6E + R \mathbf{1}) = -\frac{R}{3} \quad (\text{Einstein–Hilbert } \int R), \quad a_4|_{R^2} = \frac{1}{360} \text{tr}(180E^2 + 60RE + 5R^2 \mathbf{1}) = \frac{R^2}{72}$$

(the R^2 /Starobinsky term), so the spectral action structurally produces $S_{\text{eff}} = \int \sqrt{g} \frac{\bar{M}_{\text{Pl}}^2}{2} (R + \frac{R^2}{6\bar{M}^2}) + \dots$ with the scalaron mass ratio fixed by the cutoff moment, $M^2/\bar{M}_{\text{Pl}}^2 = 6(4\pi)^2/f_0$. The TFPT closure $M^2/\bar{M}_{\text{Pl}}^2 = c_3^7$ fixes $f_0 = 6(4\pi)^2/c_3^7$, i.e. $M = c_3^{7/2} \bar{M}_{\text{Pl}} = 3.06 \times 10^{13} \text{ GeV}$ — the boundary normalisation c_3 is the gravitational scale — and hence the Starobinsky slow-roll forms $n_s \simeq 1 - \frac{2}{N}$, $r \simeq \frac{12}{N^2}$. The $R + R^2$ structure is convention-independent (standard Gilkey); the precise rational $\frac{1}{72}$ and factor 6 are the standard Dirac conventions used here ([verification/v36_spectral_action_g2.py], [verification/v28_gravity_fR.py]).

The simpler resolution: “full QG” for physics is the gap-decoupled admissible sector ([E]). The infinite-dimensional ambient measure (G6) is the summit, but the admissible IR sector does not wait for it. The metric coupling is gap-dominated (G5): $2\|V_{\text{metric}}\| = 0.785 < \Delta = 6 \log \frac{3}{2} = 2.433$, so the admissible sector keeps a *strictly positive* effective gap $\Delta_{\text{eff}} = 1.648 > 0$ under the stated RP and metric-coupling assumptions. A positive gap dynamically *isolates* the admissible sector from the un-built ambient/deep-UV: under those assumptions the admissible (IR) measure is closed, and the ambient completion is decoupled from it (we do *not* assert it is physically irrelevant — that would itself be a physical interpretation). So the honest reading is

admissible IR sector = closed under stated RP; ambient metric measure (G6) = G_{metric} gate
--

G2 supplies the local action ($R + R^2$, heat-kernel); G5 supplies the decoupling of the admissible sector; together they reduce (G_{metric}) to the ambient projective limit (G6). **Precise status (Alessandro’s boundary):** the admissible IR sector is closed under the stated RP and gap-dominance assumptions; the ambient metric projective measure is now discharged as a [C] *redundancy* (v369, the holographic-route keybox below), a certification object rather than a missing piece. Its absence does *not* affect the bounded IR claim. [verification/v36_spectral_action_g2.py]

The glue Q-system: the index-4 object of the G_{net} statement exists explicitly ([E]/[C]). The G_{net} target — “the seam–Calderón inclusion has Jones index $4 = |\mu_4|$ ” — no longer refers to an unknown object. On the glue group algebra $A = \mathbb{C}[\mathbb{Z}_4]$ (multiplication $\delta_a \otimes \delta_b \mapsto \delta_{a+b}$, unit δ_0) all Longo Q-system axioms hold exactly [E]: associativity, unit law, the Frobenius identity $(m \otimes 1)(1 \otimes m^*) = m^*m = (1 \otimes m)(m^* \otimes 1)$, and specialness $mm^* = |\mathbb{Z}_4| \mathbf{1}$ — a special C^* Frobenius algebra with

$\dim \theta = 4 = \mu_4 = [B : A]$

exactly the Carrier Index Lemma value (KLM: $\mu_{\text{carrier}} = 16 = 4^2 \mu_{E_8}$ re-verified). *Locality = isotropy* [E]: on the Lagrangian glue $\langle (1, 1) \rangle$ the discriminant form gives $q(k(1, 1)) = k^2 \in \mathbb{Z}$ for every k — the braiding obstruction vanishes (v92 in Q-system language); the two Lagrangian glues (the sheet pair) are the only maximal isotropic choices. *Halfway step* [E]: $\{0, 2\} < \mathbb{Z}_4$ closes and gives the

unique intermediate Q-system $\mathbb{C}[\mathbb{Z}_2]$ = the $SO(16)_1$ tower step, with index factorisation $4 = 2 \times 2$. *R2 sharpened [C] (recorded, not claimed)*: what remains is the operator-algebraic *identification* “the seam–Calderón inclusion of the 16-Majorana net (v113) realises *this* Q-system” — matching two constructed objects, not a search; the three loads (metric + carrier + QBL, v123) hang on this one identification.

The hull is the graded module of the Q-system ([verification/v128_graded_hull.py]). The Q-system is not merely an abstract companion of E_8 — E_8 *realises* it [E]: the 240 roots, built as explicit norm-2 vectors in $D_5 \oplus A_3$ +glue coordinates, decompose as $240 = 52 + 64 + 60 + 64$ over the glue $\mathbb{Z}_4$ (the v89 branching as vectors), and for *all* 6720 root pairs whose sum is again a root, $\text{coset}(r_1+r_2) = \text{coset}(r_1) + \text{coset}(r_2) \bmod 4$:

E_8 is a $\mathbb{Z}_4$ -graded Lie algebra over the carrier; the grading group is the glue = the Q-system

— the checksum hull is the Q-system’s graded module, realised in 248 dimensions. The R2 identification thus matches the Calderón inclusion against a structure that E_8 already *carries*. [verification/v125_glue_qsystem.py]

Graded Frobenius ([verification/v143_graded_frobenius.py]). The identification holds at the finite Lie level: the graded module is in fact a graded *Frobenius* structure, i.e. everything the Q-system demands of its module is already exact inside E_8 [E]: (i) *coset duality* — $\text{coset}(-r) = -\text{coset}(r) \bmod 4$ for all 240 roots ($-C_1 = C_3$ bijectively, C_0, C_2 self-dual), so the Killing form restricts to a *nondegenerate* invariant pairing $g_k \times g_{-k} \rightarrow \mathbb{C}$; (ii) the glue character average $\frac{1}{4} \sum_k i^k \deg$ is *exactly* the carrier projector (the conditional-expectation shadow, quasi-basis index $|\mathbb{Z}_4| = 4$); (iii) each glue sector is *one* $W(D_5) \times W(A_3)$ orbit with multiplicity-one dimensions $64 = 16 \cdot 4 = (\mathbf{16}, \bar{\mathbf{4}})$, $60 = 10 \cdot 6 = (\mathbf{10}, \mathbf{6})$ — simple currents, sector ring $= \mathbb{Z}_4$ = the $\mathbb{C}[\mathbb{Z}_4]$ fusion, halfway subgroup $\{0, 2\} = SO(16)_1$. *R2 residue [C] (final form)*: exactly *one* conformal-net statement remains — the seam–Calderón inclusion of the 16-Majorana *net* carries this Q-system — with no remaining finite freedom.

Where that statement lives: the NS/R sector census ([verification/v148_fock_census.py]). The remaining work is scoped exactly. Every *untwisted* (Neveu–Schwarz) state of the 16 Majoranas has *integer* ($D_5, SO(6)$) weights, and integer weights lie in the even discriminant classes — so the NS module supports only $\{(0, 0), (2, 0), (0, 2), (2, 2)\}$: **the odd glue classes $k(1, 1)$ — the actual μ_4 simple currents — have zero NS support; they are twisted (Ramond-type) modules.** The Ramond zero-mode module ($2^8 = 256$ ground states) carries exactly the four odd class pairs, 64 each, and splits $128 + 128$ into the odd sectors of the *two* Lagrangian glues $\langle(1, 1)\rangle$ and $\langle(1, 3)\rangle$ (v92): *choosing the glue is choosing the Ramond projection* ($128 =$ one $SO(16)$ spinor chirality), and E_8 assembles as $248 = 120_{\text{NS}} + 128_{\text{R}}$. *Consequence [C]*: the index-4 extension cannot even be *stated* on the untwisted module alone; R2’s one remaining statement is irreducibly about the twisted-sector structure of the seam–Calderón net, with the sheet choice realised as the R-projection; the finite shadow is exhausted (v113/v125/v143 + this census).

Lemma 10 (G3–G4 — reflection positivity, fixed and integrated). *For each Θ -invariant g in the Calderón ball the physical boundary kernel $K_g = P_{\text{adm}} \Pi_{\text{phys}} \mathcal{N}_g^{TT} \Pi_{\text{phys}} P_{\text{adm}}$ is reflection positive, $\langle \Theta F, F \rangle_{K_g} \geq 0$ (G3); and $d\mu_{\text{Cal}, \chi}$ is Θ -invariant, supported where G3 holds uniformly, so $\int \langle \Theta F, F \rangle_{K_g} d\mu_{\text{Cal}, \chi} \geq 0$ (G4). **G4 is the decisive step:** RP after integrating over dynamical metrics, not merely at fixed background.*

Lemma 11 (G5 — gap dominance [E]). $2\|V_{\text{metric}}\|_{\text{rel}} < \Delta = 6 \log \frac{3}{2}$, with $\|V_{\text{metric}}\|_{\text{rel}} \leq 248c_3^2 = \frac{31}{8\pi^2} = 0.39262$, hence $2\|V\| = 0.7852 < 2.4328$ and $\Delta_{\text{eff}} \geq \Delta - 2\|V\| > 1.647 > 0$. *The numerical inequality is trivial* ([verification/v29_research_contract_certs.py]); *the work is the operator-norm bound (hard)*.

Sugawara reading + atomic OS moment ([verification/v171_os_moment_cluster.py]). The G5 budget is the E_8 Sugawara constant: $248c_3^2 = \frac{31}{8\pi^2}$ with $31 = 1 + h^\vee(E_8)$ and $c((E_8)_1) = 248/31 = 8$,

so gap dominance reads $|R^+(A_3)| \log(N_{\text{fam}}/|\mathbb{Z}_2|) > \frac{1+h^\vee(E_8)}{4\pi^2}$ — the A_3 transfer gap beats the E_8 Sugawara metric budget. Independently, the G3 reflection positivity is a *positive moment problem*: $\text{spec}(T) = \{1, (2/3)^6, (1/3)^6\}$ makes every two-point correlator $C(n) = A + B(64/729)^n + C(1/729)^n$ the moment sequence of a positive atomic measure, so each OS Hankel matrix factorises as a Vandermonde Gram and is positive semidefinite (a clean RP certificate; cluster constant $729/665$, correlation length $\xi = 1/(6 \log \frac{3}{2}) = 0.4111$).

Theorem 1 (Decoupling theorem — physical low-energy claims do *not* need the full G_{metric}). *If the admissible sector is reflection-positive (G4) and the metric coupling is gap-dominated, $2\|V_{\text{metric}}\| < \Delta = 6 \log \frac{3}{2}$ (G5), then the admissible IR measure is stable under the ambient completion: a strictly positive effective gap $\Delta_{\text{eff}} = \Delta - 2\|V\| > 0$ protects the admissible sector from the (un-built) ambient/deep-UV, so every physical low-energy read-out (masses, mixings, α^{-1} , the $R+R^2$ regime) is independent of the open ambient measure. Symbolically*

$$\boxed{\text{RP} + 2\|V\| < \Delta \implies \text{admissible IR measure stable under ambient completion .}}$$

Thus (G_{metric}) is not a diffuse “full quantum gravity missing” block: it is the sharply delimited ambient projective limit (G6), a strict-TOE certification target whose closure the falsifiable IR physics does not depend on. **Typing:** the inequality is machine-checked [E]; the operator-norm bound on $\|V_{\text{metric}}\|$ is the conditional input.

Lemma 12 (G6–G7 — projective limit & Ward identities). *A projective system μ_{χ_n} with $\pi_{m,n*}\mu_{\chi_m} = \mu_{\chi_n}$, uniform moment bounds $\sup_n \int \|g\|_{H^s}^p d\mu_{\chi_n} < \infty$, tightness and regulator-independence yields $\mu_{\text{QG}} = \lim \mu_{\chi_n}$ on $\mathcal{G}/\text{Diff}$ (G6, the hardest step); the limit obeys the BRST Ward identities $\langle s\mathcal{O} \rangle = 0$, whence $\nabla^\mu T_{\mu\nu} = 0$ and the η boundary phase cancels the APS anomalies (G7).*

Gate 2 reduction: G6 is a *boundary projective limit, not a bulk one* ([verification/v76_gmetric_reduction.py])

The exact decoupling margin is $\Delta_{\text{eff}} = \Delta - 2\|V\| = 6 \log \frac{3}{2} - \frac{31}{4\pi^2} = 1.648 > 0$, with $\|V_{\text{metric}}\|_{\text{rel}} \leq 248 c_3^2 = \dim(E_8) c_3^2 = \frac{31}{8\pi^2}$ ($31 = 2^{g_{\text{car}}} - 1$). **Holographic reduction:** because the seam is a *finite causal boundary*, the ambient *bulk* metric measure is reconstructed from a finite-codimension seam (Calderón) boundary measure — so G6 is a *boundary projective limit*. The conditional theorem is then

$$\boxed{\text{RP}(\text{seam-Calderón boundary kernel}) + \text{tightness} \implies G_{\text{metric}} \text{ closes}}$$

i.e. the dimension of the open problem drops from bulk QG to a finite seam-boundary measure.

Residual: the constructive boundary projective limit (constructive-QFT grade) — now discharged as a [C] redundancy (v369, keybox below). So Gate 2 is now *two clean tiers*: Tier A (IR physics) closed under RP+gap; Tier B (strict TOE) a [C] redundancy, no longer a diffuse “full QG missing”. [C]

The measure as one-loop fluctuations around the parameter-free saddle ([verification/v365_qg_oneloop_saddle.py], QGAMB.SADDLE.01)

With the saddle now parameter-free (v358/v359), the boundary measure of Tier B organises as a Gaussian one-loop fluctuation determinant whose ingredients are atom-fixed. **Fixed quadratic form [E]:** the fluctuation inverse-propagator stiffness is the saddle coefficient $1/c_3 = 8\pi$ — no free dimensionless dial, so the one-loop problem is *well-posed* (it was not before v358). **Gap-controlled convergence [E]:** the fluctuation operator obeys $M \geq \Delta_{\text{eff}} = 6 \log \frac{3}{2} - \frac{31}{4\pi^2} \approx 1.648 > 0$ (v76/v337), so the Gaussian integral converges mode-by-mode with no zero/negative modes besides the saddle directions; on the OS spectrum $\{0, 6 \log \frac{3}{2}, 6 \log 3\}$ the regularised one-loop log-det $\text{Tr}' \log M = 6 \log \frac{9}{2}$ is finite. **Conformal direction [C]:** the one wrong-sign mode ($c_{\text{conf}}(4) = -\frac{3}{2}$) is made convergent by the GHP contour ($\int e^{-|c|\phi^2} = \sqrt{2\pi/3}$) with IDG dressing (v334). **Tightness [E]:** $\chi = 1/(1 - (2/3)^6) = 729/665 < \infty$ (v330), so the one-loop measure has a projective limit on the admissible sector, with *zero* new dimensionless dials (v364). **Residual:** the full *non-perturbative* projective limit (G6 on the ambient sector) is now discharged as a [C] redundancy (v369, keybox below) — this sharpens C7 from “diffuse full QG” to a one-loop-controlled Gaussian problem around a parameter-free saddle and then to a certification object. [E]/[C]

The holographic route: the ambient measure is boundary-redundant, not missing dynamics
 ([verification/v369_qgamb_redundancy.py], QGAMB.REDUNDANCY.01)

There is a second, strategically stronger way to dispatch Gate 2 than building the measure: prove that the ambient sector carries *no* physical degrees of freedom relative to the holomorphic boundary algebra, so the measure QG.AMB.01 is a non-fundamental *certification* object — the role a coordinate choice plays in general relativity — rather than a missing piece of dynamics. The discriminators that make this *ambient-redundancy* statement well-posed are all in hand. **Trivial superselection [E]**: the number of DHR sectors of a chiral net is $|\det \text{Cartan}|$, which is 1 for E_8 (one primary, the vacuum; $\text{DHR}((E_8)_1) = \text{Vec}$) against 4 for the same- c rival $SO(16)_1$ — a holomorphic boundary net has *no* nontrivial superselection charge that a bulk could carry invisibly. **No torus degeneracy [E]**: the ground-state degeneracy on T^2 equals the number of anyons $= |\det K| = 1$ for E_8 , so there is no topologically protected bulk d.o.f. hidden from the boundary (the genus-1 obstruction of v344, in its redundancy reading). **Bounded reconstruction [E]**: the seam transfer deviation subspace contracts at the finite gapped Petz rate $(2/3)^6$ (v221), and the gap-decoupling margin $\Delta_{\text{eff}} = 6 \log \frac{3}{2} - \frac{31}{4\pi^2} \approx 1.648 > 0$ (v76/v337) separates the un-built ambient sector from the admissible one. **The statement [C]**: with the Bisognano–Wichmann modular reconstruction map (v258/v309/v323), under SEAM.EQUIV.01 every gauge-invariant bulk observable is boundary-reconstructible, $\mathcal{O}_{\text{phys}}^{\text{bulk}}/\text{Diff} \cong \mathcal{A}_{\Sigma}^{\mu_4}$, and two ambient measures sharing the same boundary state are physically equivalent. **Residual [O]**: the construction consumes exactly two premises — SEAM.EQUIV.01 (the S3 lattice input, v366) and the BW intrinsicity of the raw collar (v329); it does *not* construct QG.AMB.01 but supplies the *alternative* route, so that Gravity-complete = Boundary-complete + Ambient-redundancy. [E]/[C]/[O]

The classical field equation, parameter-free: entanglement equilibrium gives the linearised Einstein equation with G from c_3 ([verification/v358_grav_entropy_equilibrium.py])

While Gate 2 (above) reduces the metric-sector *measure*, the classical metric-sector *field equation* is supplied directly by the entanglement first law $\delta S = \delta \langle K \rangle$ (Jacobson 2015; Faulkner et al.), carried out with TFPT's atoms. **Parameter-free coefficient [E]**. Equilibrium gives $G_{ab} = c_3^{-1} T_{ab}$ with $c_3^{-1} = 8\pi$ *fixed*: all four c_3 -roles are mutually consistent — the Unruh factor $1/(2\pi) = 4c_3$, the EH action $1/(16\pi) = c_3/2$, the Einstein coefficient $8\pi = 1/c_3$, the Bekenstein density $1/4 = 1/|\mu_4|$ — so *no free dimensionless Newton dial* enters. **Thermo = geo coincidence [E] (new)**. The first law requires the coefficient $2\pi/\eta$ (η the entropy density); the geometry gives $1/c_3 = |\mathbb{Z}_2| 2\pi \chi(S^2)$ (one-sided Gauss–Bonnet, v58); these are the *same* equation iff $|\mu_4| = |\mathbb{Z}_2| \chi(S^2) = 4$ — a TFPT atom identity (v216). So the thermodynamic (first-law) and geometric (Gauss–Bonnet) origins of c_3 are *one*. **The matter side, assembled [C] (closing J3)**. The ball modular Hamiltonian is the boost (Casini–Huerta–Myers; geometric by Bisognano–Wichmann, v323), $K_B = 2\pi \int_B w T_{00}$, $w = (R^2 - r^2)/(2R)$; the 3-ball weight integral is $\int_B w d^3x = 4\pi R^4/15$, so $\delta \langle K_B \rangle = \frac{8\pi^2 R^4}{15} \delta \langle T_{00} \rangle$ — the concrete matter boost flux of TFPT's stress tensor. **The entropy density, atom-fixed [E]**. The Bekenstein numerator $1/4 = 1/|\mu_4|$ (v57) and the entanglement normalisation = the central charge $c = g_{\text{car}} + N_{\text{fam}} = 8$ pin the dimensionless coefficient; the area-law form is established (v59). **Residual**: the full covariant extension is now *closed* (v359, next keybox); what remains is only (i) the Jacobson equation-of-state *status* (a thermodynamic derivation, not a from-action quantisation) [O] and (ii) the absolute scale v_{geo} (the Planck-area unit) [O], while the global ambient measure (Gate 2/C7) is now a [C] redundancy (v369) — *no free dimensionless parameter survives in the gravity coupling*.

And that boundary measure is a rigorously constructed net ([verification/v77_e8_conformal_net.py]). The level-1 central charges $c = \dim \mathfrak{g}/(1 + h^\vee)$ give $c(E_8)_1 = \frac{248}{31} = 8 = \text{rank } E_8$, $c(D_5)_1 = \frac{45}{9} = 5$, $c(A_3)_1 = \frac{15}{5} = 3$, and $c(D_5) + c(A_3) = 5 + 3 = 8 = c(E_8)$ is a *conformal embedding* $(D_5)_1 \times (A_3)_1 \subset (E_8)_1$ (coset $c = 0$). Now $(E_8)_1$ is the holomorphic $c = 8$ E_8 -lattice VOA / conformal net — one of the most rigorously constructed objects in mathematical QFT (Frenkel–Kac–Segal; Kawahigashi–Longo / Carpi). **So the conditional theorem becomes concrete: identify the seam-Calderón boundary measure with the $(E_8)_1$ lattice net** (carrier $= (D_5)_1 \times (A_3)_1$ subnet), and the gap $\Delta_{\text{eff}} > 0$ supplies the clustering/tightness — so $G6$ is imported into existing RCFT/conformal-net rigor rather than built from scratch. Residual = the net identification + bulk reconstruction. [E] (charges, embedding) + [C] (identification)

The full covariant field equation, parameter-free: fixed-volume equilibrium gives the Einstein tensor with conservation as an output ([verification/v359_grav_nonlinear_einstein.py])

The linearised result above (v358) extends to the *full* covariant equation by Jacobson’s *fixed-volume* (maximal-vacuum-entanglement) equilibrium — no linearisation assumed. **Einstein tensor, not Ricci [E]**. Stationarity of the entanglement entropy at fixed volume forces the divergence-free combination: the 4D trace identity $g^{ab}G_{ab} = (1 - \frac{d}{2})R = -R$ (for $d = 4$) selects $G_{ab} = R_{ab} - \frac{1}{2}Rg_{ab}$, so

$$G_{ab} + \Lambda g_{ab} = c_3^{-1}T_{ab}, \quad c_3^{-1} = 8\pi$$

with *both* coefficients TFPT-fixed: the Newton coupling $8\pi = 1/c_3$ (v358) and the cosmological constant Λ from α , $\rho_\Lambda = (3/4\pi^2)e^{-2\alpha^{-1}}$ (v60). **Conservation is an output [E]**. By Lovelock’s theorem the only divergence-free, second-order, local symmetric tensor built from the metric is $aG_{ab} + bg_{ab}$, so $\nabla^a G_{ab} = 0$ forces $\nabla^a T_{ab} = 0$ — matter conservation is *derived*, not assumed. **So Gate B6 — the classical metric-sector field equation — is closed parameter-free at the local level. Residual:** the Jacobson equation-of-state *status* (a thermodynamic derivation, not a from-action quantisation) and the single unit v_{geo} stay [O], while the global ambient measure (Gate 2/C7, the constructive boundary projective limit) is now a [C] redundancy (v369) — no free dimensionless gravity dial remains. An *external candidate* for the missing action level is quantified in the next keybox (v473); it does not change this typing. [E]/[C]

The entropic-action bridge: an external candidate for the missing action level, quantified — nothing closes ([verification/v473_entropic_action_bridge.py])

The one honest [O] left in the classical sector — “equation of state, not a from-action quantisation” (v359) — now has an explicit *external* candidate: Bianconi’s entropic action $S_B = -\text{Tr}_F \ln(\tilde{G}\tilde{g}^{-1})$, the quantum relative entropy between the spacetime metric and the matter-induced metric (G. Bianconi, *Gravity from entropy*, Phys. Rev. D **111**, 066001 (2025); arXiv:2408.14391; her Appendix C identifies $\Delta_{\tilde{G},\tilde{g}}^{1/2} = \tilde{G}\tilde{g}^{-1}$ with an Araki-type modular operator) — a local covariant action whose variation yields the Einstein equations plus an emergent nonnegative Λ at low coupling: exactly the missing variational layer, *if* the bridge identities hold. What is machine-checked today (GRAV.ENTROPIC.ACTION.01): **carrier Hodge count [E]**. Her $\Omega^{0,1,2}$ matter fiber, placed on the five-slot carrier $\mathbb{C}^5$ (*not* on 4d spacetime, where the fiber is $1+4+6=11$), counts $1+5+10=16=2^{g_{\text{car}}-1} = \dim S_{D_5}^+$ and Hodge-folds ($\star: \Lambda^1 \cong \Lambda^4$) onto $\Lambda^{\text{even}}\mathbb{C}^5$ — the carrier half-spinor’s own Pascal row (v2/v44); a vector-space identity *only*. **Coefficient match [E] (the one quantitative pin)**. Her weak-coupling Lagrangian $\mathcal{L} = 3\beta R - \alpha\mathcal{L}_m + \dots$ (the 3 = the three form blocks) matched to the TFPT EH normalisation $1/(16\pi G) = c_3/(2G)$ fixes her free constant *exactly*,

$$\beta'_B = \frac{c_3}{6} = \frac{1}{48\pi}$$

(the $6 = 3 \times 2 = p_2 = |R^+(A_3)|$ is recorded as a structural echo, not a proof). **Entropy potential = quadratic Λ [E]**. Her $\Lambda_G = \frac{1}{2\beta} \text{Tr}_F(\mathcal{G} - I - \ln \mathcal{G})$ is nonnegative and quadratic at the fixed point; with the TFPT displacement $\varepsilon_\star = e^{-\alpha^{-1}}Q$ it reproduces the closed Λ branch $\rho_\Lambda/\bar{M}_{\text{Pl}}^4 = (3/4\pi^2)e^{-2\alpha^{-1}}$ (v60), and the prefactor match yields the *exact* normalisation target $\text{Tr}_F Q^2 = 32c_3^4 = 1/(128\pi^4) = \frac{2}{3}\delta_{\text{top}}$ — a closed algebraic target for the two-form (pair) block ($\binom{5}{2}=10$), [C] until an explicit Q is constructed. **R^2 kill test [E] (pre-registered)**. The raw log expansion carries a curvature-squared coefficient of order $(c_3/6)^2$; the TFPT Starobinsky sector needs $\bar{M}_{\text{Pl}}^2/(12M_{\text{scal}}^2) = 1/(12c_3^7)$ (v36) — the exact mismatch is $3c_3^{-9} = 3(8\pi)^9 \approx 1.2 \times 10^{13}$: a direct identification of the raw entropic action with the TFPT R^2 sector is *false* without a mechanism (the scalaron as a light trace mode of her $\mathcal{G}$ -field generating c_3^{-9} , or a KMS-spectral renormalisation of the log action). **The bridge proper [C] (recorded, not proven)**. The compression conjecture $P_\Sigma(\tilde{G}\tilde{g}^{-1})P_\Sigma = \Delta_\Sigma^{1/2}$ — the physical Calderón compression of her relative metric operator equals the TFPT seam modular operator; if proven, S_B is a local volume representation of the TFPT relative spectral action $S_{\text{rel},\chi}$ (her log action is the all-scale-integrated heat kernel, by Frullani $-\ln A = \int_0^\infty \frac{dt}{t}(e^{-tA} - e^{-t})$, verified). Open with it: the operator-level Hodge items (Clifford action, Dirac intertwining, leaf involution, Calderón polarisation), and Lorentzian positivity of $\tilde{G}\tilde{g}^{-1}$ is *not* automatic (timelike gradients) — TFPT’s OS/reflection-positive route is the sturdier construction. **Residual [O] (unchanged)**. Nothing closes: the v359 equation-of-state typing stays [O]; adopting S_B would close it *only* after the compression conjecture *and* an R^2 mechanism; zero TFPT knobs are added.

The operator level, executed ([`verification/v474_entropic_hodge_carrier.py`]). Work packages 1 and 4-algebra of the bridge are now machine-checked on the finite carrier Fock space $\Lambda^\bullet \mathbb{C}^5$ ($\dim 2^{g_{\text{car}}}=32$). **Spinor structure exhibited [E]**: the ten CAR gammas satisfy $\{\Gamma_a, \Gamma_b\}=2\delta_{ab}$ exactly, the 45 bilinears span $\mathfrak{so}(10)$ and leak nothing from the 16-dim even subspace — S^+ is the even exterior algebra with its spinor structure, and Bianconi's Hodge–Dirac operator is Clifford multiplication at symbol level ($c(\xi)^2 = -|\xi|^2$, symbolic). **The fold is the 5→5 conjugation [E]**: the traceless $\mathfrak{u}(5)$ weights of Λ^4 are the negatives of those of Λ^1 , so her fiber $1+5+10$ becomes the GUT half-spinor $16=1+5+10$ — the SM carrier decomposition, not a dimension match; honest obstruction: $\Omega^{\leq 2}$ is *not* Clifford-invariant, so the identification must go through the fold, never through degree truncation. **The Q-target, decided [E]**: the uniform-block ansatz $\text{Tr } Q^2 = d k^2 c_3^4 = 32 c_3^4$ with integer k admits *exactly* the supports $(d, k) = (2, 4), (8, 2), (32, 1) = \{|\mathbb{Z}_2|, \text{rank } E_8, 2^{g_{\text{car}}}\}$, with the minimal uniform solution $q=c_3^2$ per Fock mode and the identity $\text{Tr } Q^2/\delta_{\text{top}}=32/48=2/3$; the two-form block 10 (the naive pair-sector reading) and the fiber 16 require irrational multipliers and are *excluded*. The physical choice among the three supports stays [C]; AP2 and the R^2 mechanism stay [O].

The R^2 kill test, executed ([`verification/v475_entropic_scalaron_gate.py`]). On a maximally symmetric background the relative curvature operator has eigenvalues *exactly* $\{R; R/4 \times 4; R/6 \times 6\}$ (her Appendix-B flattening verified, incl. the 6×6 two-form block), so the exact vacuum action is $\mathcal{L}_B(R) = -[\ln(1-\beta R) + 4\ln(1-\beta R/4) + 6\ln(1-\beta R/6)] = 3\beta R + \frac{17}{24}\beta^2 R^2 + O(R^3)$ — the 3 re-derives her Eq. (45) and $\frac{17}{24}$ is the exact tensorial factor behind the pre-registered gap. With the pinned $\beta'=c_3/6$ *both* mass routes (Starobinsky matching and $f'/3f''$) give the raw entropic scalaron

$$m_{\text{raw}}^2 = \frac{4608\pi^2}{17} \bar{M}_{\text{Pl}}^2 \Rightarrow m_{\text{raw}} \approx 51.7 \bar{M}_{\text{Pl}}$$

— *trans-Planckian*, not light: the naive reading “the TFPT scalaron is the light trace mode of her $\mathcal{G}$ -field” is *dead* [E]; a viable mechanism must supply exactly $m_{\text{raw}}^2/m_{\text{TFPT}}^2 = \frac{72}{17}(8\pi)^9 \approx 1.69 \times 10^{13}$ in mass². The domain is not the obstruction (first log pole at $R=384\pi^2 \bar{M}_{\text{Pl}}^2$, trans-Planckian), and the Lorentzian caveat is now an explicit witness: a timelike gradient gives the Gg^{-1} eigenvalue $1 - \alpha v^2 \leq 0$ for $v^2 \geq 1/\alpha$ (ln undefined), while spacelike/Euclidean stay positive — the OS route is the sturdier construction. **Net verdict [C]**: of the three bridge interpretations, the entropic action covers the Einstein + Λ level (survives); the light-mode reading is dead; KMS-spectral renormalisation is the *only* surviving route to the R^2 sector. Nothing closes.

AP2, made well-posed ([`verification/v476_entropic_compression_toy.py`]). The compression conjecture itself is sharpened on an exactly solvable gapped chain. **The reading is decided**: on a *pure* bulk (the seam situation: pure bulk, thermal boundary) the covariance is a projection, so $f(x)=x/(1-x)$ is singular on its spectrum — the literal operator-side reading $P_\Sigma f(C)P_\Sigma$ is *ill-posed*, while the state-side $f(P_\Sigma C P_\Sigma)$ (Peschel) is finite: the identity can only mean *build Δ_Σ from the compressed relative metric*, which is exactly Bianconi's own local construction — AP2 retyped, not weakened. **The mismatch is exactly second order [E]**: symbolically, $[C^n]_{AA} - (C^n)_{AA}$ carries the cross-cut block twice (first order vanishes identically), so for any analytic f the two readings agree up to $O(\|C_{AB}\|^2)$ — the modular (entanglement) content *is* that quadratic correction. **Gap control**: at the bounded (covariance) level the mismatch obeys $d \leq x^2$ and falls with the gap, the decoupled cut agrees exactly — the two readings converge in the gap-dominated regime, where TFPT operates ($\Delta_{\text{eff}} > 0$, $v76/v337$). **Residual [O] (unchanged)**: the continuum seam statement itself. No marker moves.

The surviving R^2 route, typed as one moment ([`verification/v477_entropic_scaleflow.py`]). The Frullani identity upgrades to a *scale-flow representation*: $-\ln a = \int_0^\infty \frac{dt}{t} (e^{-ta} - e^{-t}) = 2 \int_0^\infty \frac{dx}{x} (e^{-a/x^2} - e^{-1/x^2})$ — Bianconi's log action *is* the flat scale-integral of relative heat-kernel actions, while TFPT's $S_{\text{rel},x}$ evaluates *one* KMS scale with a weight f : the bridge between the two action principles is a choice of *scale measure*. In the moment dictionary (weight $w(t) dt/t$, μ_k its e^{-t} -moments) the coefficients read $3\beta R \mu_1 + \frac{17}{24}\beta^2 R^2 \mu_2$; the flat weight has $\mu_1=\mu_2=1$ (= the raw v475 coefficients), and demanding $m^2=c_3^7 \bar{M}_{\text{Pl}}^2$ forces *exactly*

$$\mu_2/\mu_1^2 = \frac{72}{17}(8\pi)^9 \quad \text{with} \quad \frac{4608\pi^2/17}{(72/17)(8\pi)^9} = c_3^7 \text{ (identically)}$$

— the v475 kill-test number *is* a moment ratio, and the “13 orders of magnitude” are not a bridge mismatch but a scale-measure datum: the flat measure fixes it wrongly, the TFPT KMS measure ($f_0=6(4\pi)^2/c_3^7$, v36) fixes it correctly — one KMS moment, two parametrisations, zero new dials. [C] (consistency, not derivation: the measure is TFPT input); the R^2 gate and the v359 typing stay [O].

The two remaining legs, first computable steps ([[verification/v478_entropic_continuum_legs.py](#)]).

Leg A (AP2 continuum): on the exactly solvable critical chain (infinite-volume kernel, 50-digit precision) the compressed interval state carries the conformal/BW structure quantitatively — the entanglement entropy follows Calabrese–Cardy with $c_{\text{est}} = 1.00089/1.00038/1.00021$ (monotone $\rightarrow$ the Dirac $c=1$; gapped control $\sim 6 \times 10^{-8}$), and its Peschel modular Hamiltonian organises on the CHM weight $\beta(x)=x(L-x)/L$ (corr $0.87 \rightarrow 0.99$, even-range bands vanishing identically, odd tail 1.8%): the state-side Δ_Σ (the reading *forced* above) flows to the CHM/Bisognano–Wichmann geometric form — *exactly* the object the entanglement-equilibrium derivation consumes ([v323/v358](#)); the bridge’s modular side and the gravity side meet in the continuum limit (a lattice exhibit, not a continuum proof — AP2 stays [\[O\]](#); the natural continuum frame for the seam version is the *four-interval* multilocal modular geometry of the four marks, [v480](#)). *Leg B (the measure):* the single-scale corollary is exact — $d\mu = \delta(t-t_0)dt$ gives $\mu_2/\mu_1^2 = e^{t_0}$, so the one-moment condition forces $t_0 = \ln \frac{72}{17} + 9 \ln(8\pi) = 30.461$, one dimensionless KMS time in closed form; the near-miss to $h(E_8)=30$ (difference 0.461) is explicitly *declined* per the no-free-pattern rule ([v354/v355](#)) — deriving the measure from the entropic structure stays [\[O\]](#). [\[E\]](#)/[\[C\]](#)/[\[O\]](#)

The Simple-Current Extension Theorem — the central G_{net} closing statement [\[E\]](#)/[\[C\]](#) ([[verification/v154_simple_current_theorem.py](#)])

The conditional theorem can now be stated as one clean *simple-current extension*. Let $A = (D_5)_1 \otimes (A_3)_1$ be the carrier net and $L = \langle(1,1)\rangle \subset \text{Disc}(D_5) \oplus \text{Disc}(A_3) = \mathbb{Z}_4 \oplus \mathbb{Z}_4$ the isotropic μ_4 glue line ($q(k(1,1)) = k^2 \in \mathbb{Z}$). Then the local extension $B = A \rtimes L$ has

$$[B : A] = |L| = 4 = |\mu_4|, \quad c(B) = 5+3 = 8, \quad \mu(B) = \frac{\mu(A)}{|L|^2} = \frac{16}{16} = 1$$

so B is *holomorphic* (KLM μ -additivity), and a holomorphic $c=8$ chiral net is the lattice net of the *unique* even-unimodular rank-8 lattice — hence $B \cong (E_8)_1$ (the same- c rival $SO(16)_1$, $\mu=4$, is excluded). Every algebraic ingredient is machine-checked: the $\mathbb{C}[\mathbb{Z}_4]$ Longo Q-system ([[verification/v125_glue_qsystem.py](#)]), the $\mathbb{Z}_4$ -graded Frobenius E_8 hull ([[verification/v143_graded_frobenius.py](#)]), the Ramond glue sectors ([[verification/v148_fock_census.py](#)]). **Consequence:** G_{net} is internally *closed* if the seam–Calderón boundary net is *defined* to be A ; it stays the *one* external identification “the RP seam–Calderón net is A ” if the boundary net is to be *derived* from the seam kernel. Either way this is the single strict-TOE certification theorem still carrying weight — a clean extension statement, not “find a quantum gravity”.

And that one identification reduces to a single shared premise ([[verification/v155_quasifree_boundary.py](#)]). The identification splits ([v110](#)) into (a) “the seam Calderón involution is sheet-odd” — the double-cover deck, the same $|\mathbb{Z}_2|$ that halves $c_3 = \frac{1}{2 \cdot 4\pi}$ (structural) — and (b) “the seam 2-point kernel is quasi-free”. Part (b) is *not independent*: the boundary marginal of a reflection-positive *Gaussian* bulk is the compression *PTP* of the bulk covariance, and the compression of a contraction ($i\Gamma$ with spectrum in $[-1,1]$) is a contraction — so the boundary/seam state is automatically quasi-free, with Wick inherited (every boundary correlator a Pfaffian of Γ_d). Hence (b) *follows* from “the seam bulk is a reflection-positive free field” — the *same* premise already load-bearing in the area law ([[verification/v59_area_law_evidence.py](#)]) and the EH mechanism ([[verification/v150_replica_eh_model.py](#)]/[[verification/v151_bfk_split.py](#)]; since exercised on the discretized collar with the seam’s own kernel and real replica sheets, [v471](#) — the R3 kernel premise is discharged at the finite level, its residual retyping to the continuum scaling limit shared with SEAM.EQUIV.01, [v336](#), plus the [v152](#) anchor). **So the entire seam-side residual is one physical premise** — the seam bulk is a free RP field — shared by G_{net} , the area law and R3, and given it the chain Gaussian bulk $\Rightarrow$ quasi-free boundary $\Rightarrow$ one kernel = net ([v113](#)) $\Rightarrow$ carrier $A \Rightarrow (E_8)_1$ is exact. It is not reducible by a finite computation — it is the one place TFPT meets functional analysis.

The construction, exhibited and verified ([[verification/v156_seam_net_construction.py](#)]). The “ $B \cong (E_8)_1$ ” step is no longer an assertion — the chiral net is built explicitly and the identification *checked*. The Dirichlet-to-Neumann (Calderón) operator of the 2d Laplacian acts on a boundary mode e^{ikx} by the harmonic extension $e^{ikx-|k|y}$: $-\partial_n u|_0 = |k|u$, so for a free bulk the seam operator is the massless chiral dispersion $\omega_k = |k|$ that defines the free-fermion net on S^1 . Sixteen Majoranas give $c = 8 = 5+3$; the weight-1 currents are $248 = 120_{\text{NS}} + 128_{\text{R}}$ (the $SO(16)$ bilinears plus the Ramond

spinor, Frenkel–Kac); and the holomorphic character is

$$\chi = \frac{E_4}{\eta^8} = q^{-1/3}(1 + 248q + 4124q^2 + 34752q^3 + 213126q^4 + \dots) = j^{1/3}$$

(verified by q -series), the 248 at level 1 being $\dim E_8$ — so the net *is* $(E_8)_1$. *What remains* is exactly (i) the free-bulk premise above and (ii) the operator-algebraic existence / uniqueness of the net generated by the seam Calderón kernel (a Kawahigashi–Longo-type net-construction theorem). Route 3 thus *exhibits and verifies* the target; the irreducible residual is the one premise + the standard net-existence step — where TFPT meets constructive QFT, now sharply located.

And the free-bulk premise itself reframes: freeness is a rigid fixed point, not a postulate ([`verification/v157_rigid_fixed_point.py`]). Three facts click together. (i) The DtN/Calderón symbol is *universally* $|k|$: the boundary operator of *any* second-order elliptic bulk is an order-1 Ψ DO with principal symbol $|\xi|_g$ (Lee–Uhlmann), so the free chiral dispersion is bulk-detail-independent and any interaction lives in the lower-order (irrelevant) symbol. (ii) A holomorphic $c=8$ theory is *rigid*: every field has weight $(h, \bar{h}=0)$, an exactly-marginal deformation needs $(1, 1)$, and with $\bar{h}=0$ *no* $(1, 1)$ field exists — the fixed point is isolated, there is no continuous knob along which a non-Gaussian interaction could turn on. (iii) The *same* $|\mathbb{Z}_2|$ that halves the one-sided S^2 Gauss–Bonnet ($c_3 = \frac{1}{|\mathbb{Z}_2| 2\pi \chi(S^2)} = \frac{1}{8\pi}$) labels the sheet *and* is the Ramond/GSO projection giving $\mu=1$ ($248=120_{\text{NS}}+128_{\text{R}}$) — one object, three roles, so the geometry fixing c_3 is the projection making the net holomorphic. **Consequence:** premise (A) shrinks from “prove the seam is Gaussian” to “the gapped ($\Delta=6 \log \frac{3}{2} > 0$) boundary flow reaches the holomorphic fixed point” — given which, freeness is *forced* by rigidity + uniqueness ($(E_8)_1$, v83), not fine-tuned. A strictly smaller and simpler residual.

And the destination is a proven stable fixed point ([`verification/v158_fixed_point_stable.py`]). An exact operator-dimension count settles it: the chiral Majorana has $h=\frac{1}{2}$, the bilinear current $h=1$, the quartic $h=2$. The relevant window $0 < h < 1$ for bosonic deformations is *empty* (bosonic chiral operators have integer h , lowest 1) — nothing relevant flows the free theory away; the $h=1$ currents are chiral (v157), not $(1, 1)$ -marginal; and the lowest genuinely-interacting operator (the quartic, $h=2>1$) is *irrelevant*, as are all higher $2k$ -fermi operators ($h=k$). So the free chiral $c=8$ fixed point is *isolated and stable* against all interactions. Premise (A) thereby reduces to a *basin-of-attraction* statement with a proven stable target — the gapped ($\Delta=6 \log \frac{3}{2}$) flow lands here; the convergence is conformal-perturbation analysis, the destination is settled.

And the closure can be stated as one fixed-point theorem — which relocates (A) to a single internal premise ([`verification/v160_seam_gaussianity_from_pf.py`]). Phrase the argument internally: (i) a quasi-free state has *all* higher connected cumulants $\kappa_{2n}=0$ (the signed set-partition/Pfaffian inversion of $\langle \psi_1 \cdots \psi_{2n} \rangle = \text{Pf}(\Gamma)$ leaves only the degree-2 kernel — the v113 “one kernel = the net” property at the cumulant level); (ii) a kernel-preserving transport fixes the whole quasi-free state by Wick functoriality; (iii) a gapped primitive Perron–Frobenius transport has a *unique* fixed ray, so a non-Gaussian piece is impossible — if κ_{2n} were fixed it would be a *second* fixed ray (killing uniqueness), and if not fixed it sits in the gapped complement and decays to 0 (any gap > 0 suffices; the value $(2/3)^6$ only sets the rate). *The honest residual:* v56’s gapped PF uniqueness is established on the *three-dimensional* scalar transfer spectrum $\{1, (2/3)^6, (1/3)^6\}$ (the flavour gap / horizon Page-recovery rates), whereas the 16-Majorana seam state already carries $\binom{16}{4}=1820$ and $\binom{16}{6}=8008$ higher-cumulant directions — the Schwinger cone the transport must contract has dimension $\gg 3$, and uniqueness must hold there (not merely on the quasi-free sub-cone) for the argument to be non-circular. **So (A) is equivalent to one strictly smaller, internal, exact-testable statement:** the admissible TFPT seam transport is a *primitive, spectrally gapped, RP-cone-preserving* map on the *full* Schwinger cone, not only on the scalar readout spectrum of v56 — which the suite does *not* yet establish [E]/[O].

And that full-cone premise is now discharged to a single finite one-particle statement ([`verification/v161_one_particle_reduction.py`]). The lever is Bogoliubov second quantisation: a quasi-free transport is $T_\Sigma = \Gamma(t)$, the second quantisation of a one-particle contraction t on the 16-dimensional Majorana space H_Σ , and the Fock space grades by correlator degree, so $\Gamma(t)$ acts on the degree- m sector as the exterior power $\Lambda^m(t)$. Three exact consequences follow (all machine-checked): (i) $\text{spec}(\Gamma(t)|_{\deg m})$ is the set of products of m one-particle eigenvalues — the whole correlation spectrum is generated by the finite one-particle spectrum; (ii) **the cone gap equals the one-particle gap:** the sub-leading eigenvalue of $\Gamma(t)$ over *every* degree is $(2/3)^6$ (pad the largest contracted mode with fixed factors), so premise 5 on the cumulant space *is* premise 5 on the 16-dim space, the same number; (iii) the eigenvalue-1 sector is Λ^* of the fixed subspace — the disconnected Wick powers of K_Σ — so with a strict

one-particle gap no independent connected cumulant is fixed, and an independent fixed cumulant *on* the fixed subspace is excluded by [v113](#) (one polarisation $\Leftrightarrow$ one vacuum). Premise 1 (quasi-freeness) is itself *forced*, not assumed: the $120 = \dim \mathfrak{so}(16)$ weight-1 chiral currents *are* the Bogoliubov generators on the 16 Majoranas, and [v158](#) makes every non-Gaussian deformation irrelevant. **So the three full-cone premises collapse to one finite, 16-dimensional statement:** the seam RG transport’s one-particle symbol is the gapped Calderón Bogoliubov map — fixed eigenspace = the rank-8 K_Σ polarisation ([v113](#)), sub-leading gap = $(2/3)^6$ ([v56](#)). Every *number* in it is already supplied by the suite; the one irreducible step is the physical identification that the actual seam RG generator equals this Bogoliubov map — the same operator-algebraic (Kawahigashi–Longo-class) statement as the net-existence residual of (A_2) , now with the infinite-dimensional cone eliminated [\[E\]/\[O\]](#).

And that finite statement carries no new open content: it factors into the two already-open items ([\[verification/v162_seam_transport_identification.py\]](#)). Take the one-particle symbol apart into its two data, the gap value (β) and the fixed space (α) . (β) *is forced*: a μ_4 -equivariant operator is diagonal in the cusp basis — the deck average is the diagonal projection, $\sum_{k=0}^3 U^k X U^{-k} = |\mu_4| \text{diag}(X)$ for $U = \text{diag}(1, i, -i)$, and the commutant of U is exactly the diagonal algebra — so its spectrum is the cusp weights $\text{spec}(A_0^*) = \{0, \frac{1}{3}, \frac{2}{3}\}$ (the parabolic residue, [v115](#)), and the transfer eigenvalue $(1-\alpha)^{p_2}$ with $p_2=6=|R_+(A_3)|$ gives $\{1, (2/3)^6, (1/3)^6\}$, gap $6 \log \frac{3}{2}$ — every entry carrier data, no free knob; the lift to the 16 Majoranas is the established A_3 subnet grading ([v137](#)). (α) is the rank-8 Calderón polarisation of the $\omega_k=|k|$ free dispersion ([v113/v156](#)), the IR conformal free-fermion vacuum. **What is left** is only that the seam RG generator *is* μ_4 -equivariant and admissible — which is the R1 “seam clock” identification (the transfer operator as the seam’s gapped generator) — together with the A_2 net-existence theorem. Both are pre-existing open items, so premise (A) introduces *no independent open content*: it factors exactly into $A_2 + R_1$ with every number forced. The “free-bulk cloud” is gone; what remains is the same two identifications the programme already names — (A) is closed *conditionally* on them (a unification, honestly not an unconditional proof) [\[E\]/\[O\]](#).

And both residual identifications are pinned to their floor, so (A) has no independent open gate left ([\[verification/v165_premise_a_closure.py\]](#)). A_2 (net existence) is not open research but *assembly* of established mathematics: the 16 free Majoranas define a free-fermion conformal net (Buchholz–Mack–Todorov), the holomorphic $c=8$ lattice net *is* $(E_8)_1$ (Frenkel–Kac–Segal / Dong–Xu / Kawahigashi–Longo), the μ_4 extension is a Longo Q-system of index 4 ([v125](#)); the one TFPT-specific input — that the seam Calderón kernel is the free-fermion two-point function — is done because the DtN symbol is universally $|k|$ ([v156/v157](#)). So A_2 re-types [\[O\]](#) $\rightarrow$ [\[C\]](#) (closed by assembly). R_1 (the seam clock) reduces to GATE.QGEO: a μ_4 -equivariant flat gapped generator on the carrier is *unique* ($= A_0^*$, [v115/v116](#)), so its gap $(2/3)^6$ is forced once the seam generator lies in that class, and the residual is exactly one geometric identification — the seam collar boundary is the μ_4 -punctured sphere $\mathbb{P}^1 \setminus \mu_4$ — whose cohomological half is already established ([v137](#), $b_1=3=N_{\text{fam}}$). That is GATE.QGEO, an already-open realisation premise, *not* a new gate. By the No-Unit Theorem ([v153](#)) the irreducibles stay exactly $\{\pi, v_{\text{geo}}\}$ — closing (A) adds none. **Final accounting:** premise (A) = (A_2 assembly, [\[C\]](#)) + (R_1 =GATE.QGEO, already open) + the $\{\pi, v_{\text{geo}}\}$ irreducible core (a theorem) — it carries *zero* independent open gates and ceases to be its own gate [\[E\]/\[O\]](#).

And the two structural residuals — net existence and full-cone reflection positivity — are now discharged to a theorem, leaving the seam realisation as the one irreducible open premise ([\[verification/v175_net_existence_full_cone.py\]](#)). The CAR second-quantisation functor $\Gamma(t) = \bigoplus_m \Lambda^m(t)$ lifts the one-particle seam contraction t to the *complete* Fock space $\Lambda^\bullet(\mathbb{R}^{16})$ ($\dim 2^{16}=65536$); by the exterior-power spectrum theorem every eigenvalue of $\Gamma(t)$ is a subset product of $\text{spec}(t)$, so checking all 2^{16} products shows $\Gamma(t)$ is a *contraction* (full-cone RP), with eigenvalue-1 multiplicity $2^8=256$ (the wedges of the rank-8 K_Σ polarisation) and sub-leading eigenvalue exactly the one-particle gap $(2/3)^6$ — so full-cone RP/gap reduce *for every* m (not only $m \leq 6$, [v161](#)) to the one-particle contraction (Shale–Stinespring), *discharging the v160 scope premise by a theorem*. And A_2 is now an *assembled and verified* certificate: the $(E_8)_1$ character $E_4/\eta^8 = 1 + 248q + 4124q^2 + 34752q^3 + 213126q^4 + 1057504q^5$ (level-1=248), the embedding $248=120+128$, and the E_8 Cartan matrix even with $\det 1$ (the unique rank-8 even unimodular lattice, [v83](#)) — the net *exists* by cited theorems. **Both residuals then collapse to the same input:** that the actual seam Calderón kernel is this gapped one-particle contraction on $\mathbb{P}^1 \setminus \mu_4$, i.e. QGEO.REALIZE.01, whose finite half is proven ([v168](#)). That realisation half is a constructive-geometry/AQFT statement, *not* closeable by a finite computation, and is left honestly [\[O\]](#) — the single irreducible structural premise; net existence and full-cone RP are [\[E\]](#).

The No-Unit Theorem — v_{geo} is an irreducible metrology primitive [E]/[O]
 ([verification/v153_no_unit_theorem.py])

v_{geo} is not an open gap. Every load-bearing compiler datum — $g_{\text{car}}, |\mu_4|, N_{\text{fam}}, \text{rank } E_8, \alpha^{-1}$, the determinant ladder $(3, 4, 8, 20)$, $c_3 = \frac{1}{8\pi}$, φ_0^{ret} — is mass-dimension 0 and *invariant* under a unit rescaling $L \mapsto \lambda L$, whereas a mass / energy / $1/G$ carries nonzero dimension. A map from dimension-0 inputs to a dimension- $d \neq 0$ output is simultaneously invariant and covariant: impossible unless the unit is introduced. **So a dimensionless boundary compiler provably cannot select an absolute scale** — v_{geo} is a *forced input*. The three apparent residuals then collapse onto the one unit, differing only in mass-dimension:

$$U_{\text{point}} \sim v_{\text{geo}}, \quad \frac{1}{G} \sim v_{\text{geo}}^2, \quad m/\mu = e^{3/4} \text{ (a pure ratio),}$$

so the irreducibles stay exactly $\{\text{one dimensionful unit}, \pi\}$. *Re-typing*: v_{geo} is an *irreducible metrology primitive* (a declared unit, like choosing a metre), not a missing derivation.

Sharpened across gravity, and the final tally ([verification/v364_vgeo_sharpen.py]). After the parameter-free gravity (v358/v359/v361) the *gravity* sector adds no new scale either: the Einstein coefficient $1/c_3 = 8\pi$ and the Λ prefactor $(8\pi)^2 \cdot 48c_3^4 = 3/(4\pi^2)$ are dimensionless, so $1/G \sim v_{\text{geo}}^2$ and $\Lambda \sim (\text{dimensionless}) v_{\text{geo}}^4$ both reduce to v_{geo} times atoms. Hence v_{geo} is the *single* dimensionful input of the *complete* theory (matter and gravity), and the final tally is $\{v_{\text{geo}}, \pi\}$ with *zero* dimensionless dials — a $\sim 26 \rightarrow 1$ reduction against the Standard Model. ($1/G$ stays UV-sensitive (Sakharov, v68), so it is a metrology readout, not a clean prediction — consistent with v_{geo} being the anchor.)

Theorem G — the closure statement [O] (target)

The relative spectral boundary kernel over the diffeomorphism-quotiented metric sector admits a regulator-independent, reflection-positive projective-limit measure μ_{QG} . Its Ward identities imply covariant conservation, and its low-curvature expansion is the TFPT $R + R^2$ seam action

$$f_R R_{\mu\nu} - \frac{1}{2} f(R) g_{\mu\nu} + (g_{\mu\nu} \square - \nabla_\mu \nabla_\nu) f_R = \bar{M}_{\text{Pl}}^{-2} T_{\mu\nu} + O(R^3/M^4),$$

so the closed covariant equation becomes the controlled IR shadow of the full measure. *Status*: the local covariant field equation is parameter-free (v358/v359); Theorem G remains open only for the projective-limit ambient measure μ_{QG} (G6, the deepest item). G5 certified; the regime equation closed [verification/v28_gravity_fR.py].

Certifiability and recommended order

Building block	Machine	Hardness
W_{wall} enumeration ($C_U^{(1)}$)	Lean / Python / Wolfram	easy (<i>done</i>)
D_4 fixed-locus polynomial ($C_U^{(2)}$)	Sage / Mathematica	medium
selector uniqueness ($C_U^{(3)}$)	interval arithmetic	medium
SNF, determinants, spectra	Lean	easy-medium
principal-symbol ellipticity (G1)	xAct / Cadabra / Lean	medium
gap inequality (G5)	Lean	easy (<i>done</i>)
operator-norm bound (G5)	analysis + numerics	hard
projective limit (G6)	not primarily machine	very hard
Ward identities (G7)	BV/BRST formal + checks	hard

Recommended order.

Selector-triangle pairings $\rightarrow v_{\text{geo}}$ scale anchor $\rightarrow G_{\text{net}}$ inclusion theorem

The flavor interface (historically U_{wall}) is now reduced to the *selector-triangle pairings* (the dual normal pair (d, n) forcing R columnwise, v136/v139); it is finite, falsifiable and the cheapest visible upgrade. The v_{geo} scale anchor is the one dimensionful input (same nature as $1/G$). G_{net} — the index-4 seam-net inclusion certifying the metric sector — is more fundamental but a *programme*,

not a single proof: heavy machinery, with the projective limit the genuine summit. F_{transfer} (source→pole / relic / cosmology) is the remaining downstream interface.

F_{transfer} is a typed functor, not a bag of open topics. It is standard physics fed TFPT *source* data, $F_{\text{transfer}} = F_{\text{observable}} \circ F_{\text{threshold}} \circ F_{\text{RG}}$, with exactly four interfaces:

Interface	Map	Status	Source data → observable
F_{pole}	source→pole masses	[E] / [C]	Koide 53/54 = $a^T(R+Q)\mathbf{1}/(2\mathbf{1}^T R a)$ is an exact operator factor [E] (the missing Sheet-Diamond corner); the pole-transfer interpretation is [C] ([verification/v183_koide_f_corner_transfer.py])
$F_{\text{Boltzmann}}$	CP source→ η_B	[C]	$\tilde{m}_1=m_3/A_\Lambda$ anchors the washout, but $M_1=M_R\varphi_0^4$ only relocates the free input (M_R is the seesaw scale, not a compiler power), so η_B stays a <i>sharper scenario</i> , not a closure ([verification/v184_etaB_anchored_boltzmann.py])
F_{relic}	$(f_a, m_a, \theta_i, N_{\text{DW}}) \rightarrow \Omega_a$	[C]	the determinant-line axion can be all of DM, but only on the $\theta_i \approx 170^\circ$ hilltop where the relic is exponentially sensitive — a viable scenario ([verification/v185_axion_relic_solver.py])
F_{QCD}	$(\alpha_s, b_3, \Lambda_{\text{QCD}}, \text{lattice}) \rightarrow m_p/m_e$	[O]	a QCD matching contract ($b_3=-7$ a carrier output), <i>not</i> a compiler number — the rejected 1836 near-formula is the discipline working

A machine guard enforces the typing: the four families must always carry a [C]/[O] marker and must never be promoted to a primitive [E] compiler prediction ([verification/v187_fttransfer_laws.py]). Exact *algebraic* sub-parts (the Koide factor, $b_3=-7$) may be [E]; the *physical* prediction never is.

The functor contract CONTRACT.F.01, four axioms ([verification/v213_fttransfer_functor.py]). Beyond the typing guard, the functor itself satisfies four checkable structural axioms, each instance a *discrete* compiler kernel [E] composed with an *external* continuous solver [C]: (1) **μ_4 -deck equivariance** — the discrete kernels are carrier/deck invariants $\{|\mu_4|, N_{\text{fam}}, g_{\text{car}}, A_\Lambda\}$, and the Koide multiplier $\lambda_2=(2/3)^6=64/729$ is the μ_4 -deck transfer eigenvalue (v82/v54); (2) **Plücker preservation** — the dimensionless readouts are Plücker readings ($53=a^T(R+Q)\mathbf{1}$, $\text{Pl}(F)=(1, 22, 30)$; $\|\text{Pl}(K)\|_1=11$); (3) **positivity/stochasticity** — $\text{spec } T=\{1, (2/3)^6, (1/3)^6\}$ is positive (RP/OS), $\kappa_f \in (0, 1)$; (4) **external modules explicit** — the continuous content (pole transfer, the flavored Boltzmann network, the finite- T $\chi(T)$ ODE, the lattice C_N) is named and external, never a compiler number. So CONTRACT.F.01 is the third research contract alongside (U_{wall}) and (G_{metric}): a typed functor, *not* a bag of open topics — but a consolidation, *not* a closure.

The discrete→dynamic lens: F_{transfer} is the readout end of the one principle ([verification/v303_fttransfer_dynamics.py], FR.DYNAMICS.01). The typing above is sharpened by a *dynamical* reading. The main-branch mechanism is a gapped, positivity-preserving relaxation to a *unique* attractor (the local-averaging update $m \leftarrow (A+2I)/4m \rightarrow$ the E_8 marks, gap $6 \ln(3/2)$). All four F_{transfer} instances share that *same shape*: [E] **F_{pole}** (Koide) is the seam dynamics — the source→pole/branch relaxation is the Möbius map with fixed-point multiplier $F'=(2/3)^6=\lambda_2$, the seam-clock subleading eigenvalue, so $|F'| < 1 \Rightarrow$ a unique attractor at the *seam* rate (v82); [C] **$F_{\text{Boltzmann}}$** (η_B) is a gapped positive contraction (washout $\kappa_f \in (0, 1)$ toward the unique balance, an H-theorem); [C] **F_{relic}** (axion) freezes at an adiabatic fixed point (the comoving number relaxes to a constant); [E] **F_{QCD}** (m_p/m_e) is the 1-loop RG flow with the carrier $b_3=-7$, whose Gaussian UV fixed point is the attractor (asymptotic freedom, Λ_{QCD} generated). *Honest boundary*: only F_{pole} has the seam rate $(2/3)^6$ exactly; the other three share the *shape* with external rates (thermal, cosmological, RG). So F_{transfer} is not a bolt-on but the downstream manifestation of the one discrete→dynamic principle — the compiler fixes the discrete seed + integer kernel, a gapped flow carries it to the observable — and the theory's simplicity is preserved precisely because the external rates are honestly fenced (v187), not because they reduce to the seam. (Even F_{pole} 's flow time is non-integer, v93: the relaxation is the *continuous* semigroup

$e^{t \log F}$ generated by the discrete operator — exactly the diffusion-semigroup reading of the main branch.)

The single-flow reduction (DYN.TRANSFER.UNIVERSAL.01). The dynamic lens is sharpened to a precise reduction ([`verification/v425_dyn_transfer_universal.py`]): the shared “shape” is literally the *one* native flow of the theory — the seam modular/recovery semigroup (v238/v240, gap $6 \ln(3/2)$) — restricted to each sector, extending the static universal-gap principle (v383) to a dynamic one. [E] the v99 Koide generator $\dot{q}=(\Delta/N_{\text{fam}})(q-2)(q-5)$ linearised at the attractor $q=2$ has rate $-\Delta$, so the time-1 contraction is $(2/3)^6$ (the recovery eigenvalue, $q=5$ unstable); [E] the Boltzmann washout is that same contraction, with the CP phase $4\pi/3$ native (v320); [E] the QCD rate is $b_3 = -7$. So the transfer *dynamics* is native; only the *anchors* stay external — the IR pole scale and M_R are v_{geo} -class, C_p is the lattice $O(1)$, and F_{relic} ’s cosmological initial condition θ_i is the one genuine wall. A reduction, *not* internalisation: no new premise, and the firewall (v187) holds — the observable stays [C]. So the visible residual is one geometric postulate (QGE0.SYM.01), one unit gauge (v_{geo}), and this one typed transfer functor — everything else is out-of-sample data.

External works built upon

The TFPT discrete core (the two axioms, the E_8 glue, the integer skeleton) is self-contained and machine-checked. Its *bridges* to continuum/AQFT physics, to gravity and to the lattice/singularity picture build on established mathematics, listed here by topic; the body refers to these by author name.

Modular theory and algebraic quantum field theory.

- M. Takesaki, *Tomita’s Theory of Modular Hilbert Algebras and its Applications*, Lecture Notes in Math. **128**, Springer (1970).
- J. J. Bisognano, E. H. Wichmann, *On the duality condition for a Hermitian scalar field / for quantum fields*, J. Math. Phys. **16** (1975) 985–1007; **17** (1976) 303–321.
- K. Osterwalder, R. Schrader, *Axioms for Euclidean Green’s functions* I, II, Comm. Math. Phys. **31** (1973) 83–112; **42** (1975) 281–305.
- S. Doplicher, R. Haag, J. E. Roberts, *Local observables and particle statistics* I, II, Comm. Math. Phys. **23** (1971) 199–230; **35** (1974) 49–85.
- R. Haag, *Quantum field theories with composite particles and asymptotic conditions*, Phys. Rev. **112** (1958) 669–673; D. Ruelle, *On the asymptotic condition in quantum field theory*, Helv. Phys. Acta **35** (1962) 147–163.
- R. Longo, *Index of subfactors and statistics of quantum fields* I, II, Comm. Math. Phys. **126** (1989) 217–247; **130** (1990) 285–309.
- Y. Kawahigashi, R. Longo, M. Müger, *Multi-interval subfactors and modularity of representations in conformal field theory*, Comm. Math. Phys. **219** (2001) 631–669.

Conformal nets, vertex operator algebras and lattice constructions.

- I. B. Frenkel, V. G. Kac, *Basic representations of affine Lie algebras and dual resonance models*, Invent. Math. **62** (1980) 23–66; G. Segal, *Unitary representations of some infinite-dimensional groups*, Comm. Math. Phys. **80** (1981) 301–342.
- T. J. Osborne, A. Stottmeister, *Conformal field theory from lattice fermions*, Comm. Math. Phys. **398** (2023) 219–289, [arXiv:2107.13834](https://arxiv.org/abs/2107.13834), [doi:10.1007/s00220-022-04521-8](https://doi.org/10.1007/s00220-022-04521-8). (This is the seam scaling-limit theorem the TFPT continuum leg cites — the Koo–Saleur $\rightarrow$ Virasoro convergence, $c = \frac{1}{2}$, and the smeared WZW-current Theorem 5.1. The internal TFPT code-token “MMST” and the legacy claim-ids SEAM.MMST.* / filenames v*_mmst_* denote *this* Osborne–Stottmeister theorem.)
- V. Morinelli, G. Morsella, A. Stottmeister, Y. Tanimoto, *Scaling limits of lattice quantum fields by wavelets*, Comm. Math. Phys. **387** (2021) 299–360, [arXiv:2010.11121](https://arxiv.org/abs/2010.11121), [doi:10.1007/s00220-021-04152-5](https://doi.org/10.1007/s00220-021-04152-5); and *Operator-algebraic renormalization and wavelets*, Phys. Rev. Lett. **127** (2021) 230601. (The operator-algebraic (wavelet/Wilson–Kadanoff) renormalization *framework* the Osborne–Stottmeister theorem above builds on.)

- M. S. Adamo, Y. Moriwaki, Y. Tanimoto, *Osterwalder–Schrader axioms for unitary full vertex operator algebras*, [arXiv:2407.18222](#) (2024).
- M. S. Adamo, L. Giorgetti, Y. Tanimoto, *Rational and non-rational two-dimensional conformal field theories arising from lattices*, [arXiv:2506.01008](#) (2025); *Representations of conformal nets associated with infinite-dimensional groups*, [arXiv:2508.07109](#) (2025).
- P. Naaijken, D. Penneys, B. Wallick, *(commuting-projector reflection positivity / LTO-RP)*, [arXiv:2605.10693](#).

Lattices and quadratic forms.

- J. H. Conway, N. J. A. Sloane, *Sphere Packings, Lattices and Groups*, 3rd ed., Grundlehren der math. Wiss. **290**, Springer (1999).
- C. L. Siegel, *Über die analytische Theorie der quadratischen Formen* (the Minkowski–Siegel mass formula), *Ann. of Math.* **36** (1935) 527–606.

Singularities and geometry.

- E. Brieskorn, *Beispiele zur Differentialtopologie von Singularitäten*, *Invent. Math.* **2** (1966) 1–14.
- J. Milnor, *Singular Points of Complex Hypersurfaces*, *Ann. of Math. Studies* **61**, Princeton Univ. Press (1968).

Spectral geometry, gravity and thermal time.

- A. H. Chamseddine, A. Connes, *The spectral action principle*, *Comm. Math. Phys.* **186** (1997) 731–750.
- A. Connes, C. Rovelli, *Von Neumann algebra automorphisms and time–thermodynamics relation in general covariant quantum theories*, *Class. Quantum Grav.* **11** (1994) 2899–2918.
- R. M. Wald, *Black hole entropy is the Noether charge*, *Phys. Rev. D* **48** (1993) R3427–R3431.
- H. Casini, M. Huerta, R. C. Myers, *Towards a derivation of holographic entanglement entropy*, *JHEP* **05** (2011) 036, [arXiv:1102.0440](#).




Figure 3: The structural-residual reduction chain ($v175 \rightarrow$ the certification round). The naive “build a quantum-gravity measure” shrinks, one machine-checked step at a time, from a vague measure to one falsifiable physical statement — *is the raw seam $(E_8)_1$ at $\tau = i$?* (SEAM.EQUIV.01; its conformal-deck face QGEO.SYM.01 is a corollary, v335). The closing arc ($v276$ – $v302$) pins the shared fact and derives the gap $6\ln(3/2) > 0$; the certification round discharges the extension leg to the crossed-product package (v469), the R3 graph→geometry bridge to Kronheimer 1989 (v479) and exhibits the four-interval μ_4 mechanism (v480). Every step above the bedrock is a theorem, a Lean proof or an established citation; the residual — the cited continuum scaling-limit existence (v336) — stays honestly [O].

F_{transfer} : one shape, four rates — a gapped relaxation to a **UNIQUE** attractor




Figure 4: **The discrete→dynamic lens on F_{transfer}** ([verification/v303_ftransfer_dynamics.py]). All four transfer instances share *one* shape — a gapped, positivity-preserving relaxation to a *unique* attractor, the same Perron-Frobenius / H-theorem / RG-fixed-point structure as the main-branch update $m \leftarrow (A+2I)/4m \rightarrow$ the E_8 marks. *Only* F_{pole} (Koide, gold) runs it at the *seam* rate: its Möbius multiplier is $F'=(2/3)^6=\lambda_2$, the seam-clock subleading eigenvalue. $F_{\text{Boltzmann}}$ (η_B , thermal washout), F_{relic} (axion, cosmological freeze) and F_{QCD} (m_p/m_e , the RG flow to the Gaussian UV fixed point with the carrier $b_3 = -7$) share the *shape* with *external* rates. So **F_{transfer}** is the readout end of the one principle, not a bolt-on; the theory's simplicity is preserved because the external rates are honestly fenced (v187), not reduced to the seam.