

TFPT — E_8 Audit, Cascade Bridge and Bootstrap

The seven E_8 slices as an audit raster, the old cascade as the same even-integer spine, and the Möbius self-consistency loop

Stefan Hamann

Alessandro Rizzo

June 22, 2026 · v5.2 (rev 236)

What this document is about

E_8 as an *audit container*, not a mystery: the seven maximal slices of 248 as a falsification raster (every load-bearing number must appear in at least one projection), the bridge showing the old E_8 orbit cascade $D=60-2n$ is the same even-integer spine as the compiler, and the Möbius bootstrap in which $g_{\text{car}} = 5$ and the “8” in c_3 are overdetermined E_8 -closure fixed points (only π irreducible).

What E_8 is *not* (and why the known no-go results do not apply)

E_8 here is the unimodular *audit / compiler hull* that classifies the admissible discrete charge and residue structures — **not** an unbroken physical gauge group. The Standard Model is a *readout* after projection (not a direct E_8 gauge theory), the Lorentz signature appears only after projection, and fermions are not “everything in the adjoint”. Consequently the standard objections to literal E_8 world-formulas do *not* bite: Distler–Garibaldi (no chiral fermions / representation obstructions for E_8 as a 4d gauge group) and Coleman–Mandula (no nontrivial mixing of spacetime and internal symmetry) both constrain E_8 as a *spacetime gauge symmetry*, which TFPT does not claim. The defensible statement is: *TFPT is the unique parabolic compilation of the anchor $a = (1, 1, 2)$ into the E_8 audit hull* — a discrete classifier, attacked on its own (combinatorial) terms, not on gauge-theoretic ones.

The TFPT document set — what is treated where

Plain language: TFPT (Topological Fixed-Point Theory) is a small discrete compiler. Two inputs — the seam constant $c_3 = \frac{1}{8\pi}$ and the carrier rank $g_{\text{car}} = 5$ — build $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ and read off the Standard-Model skeleton, the constants and the scale grammar. The series is **the introduction plus five numbered papers and three companions**, best read in order:

0. **introduction** — reading guide, compiler closure, predictions, the dependency DAG and the single proof ledger (`verification/status_ledger.csv`). (*entry point*)
1. **tfpt_1_architecture_e8** — the two axioms, the derivation map, the EM fixed point α^{-1} , the $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ construction, the QBL theorem chain.
2. **tfpt_2_standard_model** — the SM in one φ_0 -ladder formula, flavor from parabolic transport (sheet diamond, dual anchor, Q_+ cohomology), the worked closures.
3. **tfpt_3_e8_audit_bootstrap** — the seven E_8 slices as an audit raster, the cascade bridge, the Möbius bootstrap.
4. **tfpt_4_frontier** — honest status of η_B , m_p/m_e , Koide, dark matter, full QG.
5. **tfpt_5_redteam** — the adversarial audit: declared attacks, what survives, what each reduces to, and the kill tests.

Companions: **R. tfpt_research_contracts** — the open interfaces ($v_{\text{geo}}, G_{\text{net}}, F_{\text{transfer}}$) as numbered contracts; **H. tfpt_horizon_readouts** — Appendix H (unit-system reframe): one seam constant as the universal horizon thermal code, the Nariai/anchor surface, the resummed seam clock; **O. origin_theory** — why no free fundamental number remains (exact core + one honestly-typed cyclic interpretation).

You are reading document 3 (E_8 Audit & Bootstrap).

TFPT status at a glance — read this card the same way in every document

Two inputs, one compiler. From the boundary constant $c_3 = \frac{1}{8\pi}$ (axiom P1) and the five-slot carrier $g_{\text{car}} = 5$ (axiom P2) — jointly the elementary symmetric data of one anchor $a = (1, 1, 2)$ — the discrete compiler $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ is *built*, and the Standard-Model skeleton (gauge group, three families, hypercharges, the flavor matrix), the dimensionless constants ($\alpha^{-1} = 137.0359992$) and the scale grammar are read off as fixed points. The algebraic core is machine-checked; physical readouts run through *named bridges* whose status is typed below.

What is closed, and what genuinely remains. The discrete/algebraic layer is closed ([E]); the everything-open residual reduces to *three named interface problems*:

Interface	Question	Status
v_{geo}	the one metrology unit ($= 1/\sqrt{G} = m/\mu$); No-Unit Thm: no compiler scale	primitive [O]
G_{net}	simple-current extension $(D_5)_1 \otimes (A_3)_1 \rtimes \langle (1, 1) \rangle \cong (E_8)_1$	[E]/[C]
F_{transfer}	one functor, four typed interfaces (Koide, η_B , axion, m_p/m_e)	[C]

The single open *structural theorem* is G_{net} (= SEAM.EQUIV.01: the raw seam *is* the holomorphic $(E_8)_1$ net), a constructive-AQFT theorem whose only open input is the continuum Osterwalder–Schrader reconstruction of the gapped quasi-free collar. QGEO.SYM.01 (the μ_4 deck acting geometrically) is now its *corollary*: a conformal net’s vacuum is rotation-invariant by axiom, so the deck follows with no extra premise (v335, Lean FORM.QGEO.BW.01). The full nonperturbative quantum-gravity measure QG.AMB.01 is the *general* Euclidean-QG conformal-factor problem (Gibbons–Hawking–Perry 1978), from which TFPT is rigorously *gap-decoupled* (margin $1.648 > 0$, v76/v330) — an inherited, decoupled problem, *not* a TFPT-specific open item. So the honest residual is *one* solvable theorem (G_{net}) plus the typed v_{geo} and F_{transfer} interfaces. Everything carries a claim ID resolving to a row of `verification/status_ledger.csv` (the single source of truth); **if the text and the ledger ever disagree, the ledger wins**. The historical labels (U_{wall} , G_{metric} , F_{frontier}) are retained only for ledger continuity — the live residual is $v_{\text{geo}} \oplus G_{\text{net}} \oplus F_{\text{transfer}}$.


Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in `verification/status_ledger.csv`.

[E] exact / machine-proven

[C] conditional

[O] open / axiom

[X] kill test



Contents

I The E_8 secondary-branching atlas

5

1	The compiler core (where this note sits)	5
2	The $(1, 1, 2)$ anchor microcode	5
3	The residue and length matrices as spectral objects	7
4	The secondary branching atlas	9
5	Group-theoretic audits (optional)	11
6	Further audit lemmas (machine-verified this round)	12
7	What this changes, and what stays open	12
II The E_8 cascade bridge		14
8	The old cascade, recalled	14
9	The spine: every rung is a current number	14
10	Three views of one spine	17
11	Direct φ_0^{ret} relations: old vs current	18
12	Sector inventory: have all the old-cascade sectors been computed?	18
13	What is new, what is old, what stays open	19
III The self-consistency loop		20
14	The loop for g_{car} — an overdetermined fixed point	20
15	The loop for c_3 — the “8” is the E_8 rank	21
16	The Möbius origin of the irreducible π	22
17	Honest status: bootstrap, not creation from nothing	23
IV Open items and the E_8 -slice scan		25
18	The honest open-items list	25
19	The E_8 -slice scan: results (honest hit / miss)	26
20	Following the $SU(9)$ lead: cross-slice consistency	27
Appendix: computational verification		27


Part I

The E_8 secondary-branching atlas

Purpose and discipline. This note does *not* introduce new physics. It builds an *audit raster*: the claim that *every load-bearing TFPT number should appear in at least one E_8 branching projection*. A number that appears in none is suspect; a number appearing in several independent projections becomes structurally strong. All arithmetic identities below are exact [E]; the *physical readings* are audit-level [O] unless tied to a derivation. (All 77 numeric claims were machine-checked with exact integer/rational arithmetic.)

1 The compiler core (where this note sits)

The strongest current form of the theory is a five-layer compiler:



E_8 is now a *closure* object, not an input: the carrier is the $D_5 = \mathfrak{so}_{10}$ half-spinor, the four punctures $\mathbb{P}^1 \setminus \mu_4$ give $A_3 = \mathfrak{su}_4$, and the common discriminant group $\mathbb{Z}_4$ glues them, with the norm test $q(D_5) + q(A_3) = \frac{5}{4} + \frac{3}{4} = 2$. This note charts the *secondary* maximal subalgebras of E_8 as audit windows onto the TFPT modules.

2 The (1, 1, 2) anchor microcode

The parabolic anchor $a = (1, 1, 2)$ (splitting type of $E_\bullet = \mathcal{O}(-2) \oplus \mathcal{O}(-1)^2$) is a miniature compiler. Its power sums $p_n(a) = 1^n + 1^n + 2^n = 2 + 2^n$ generate the discrete budget:

n	$p_n = 2 + 2^n$	TFPT reading (verified)
1	4	$ \mu_4 = -\deg E_\bullet$ (glue index)
2	6	$ R^+(A_3) $ (positive A_3 roots) = $\mathbb{Z}_6$ cusp order
3	10	$\mathcal{A}_\Lambda = \binom{5}{2} = \dim V_{D_5}$ (action ladder / pair sector)
4	18	$N_{\text{fam}} R^+(A_3) = 3 \cdot 6$
5	34	$2h(D_5) + 3 R^+(A_3) = 16 + 18$; and $34 - R(A_3) = 22 = \sum R$

and the budget compresses to products of power sums:

$$q(A_3) = \frac{p_2}{2p_1} = \frac{3}{4}, \quad q(D_5) = \frac{p_3}{2p_1} = \frac{5}{4}, \quad \sum L = p_1 p_3 = 40,$$

$$\Omega_{\text{adm}} = 2p_1 p_2 = 48, \quad D_{\text{start}} = p_2 p_3 = 60, \quad |R(E_8)| = 4p_2 p_3 = 240.$$

The first three power sums alone reach E_8 : $p_1 p_2 p_3 = 4 \cdot 6 \cdot 10 = 240 = |R(E_8)|$, and $p_1 p_2 p_3 + (p_4 - p_3) = 240 + 8 = 248 = \dim E_8$.

The anchor difference ladder: the binary carrier spine [E]

While the *products* of the power sums give the large budgets, the *differences* give the binary carrier spine in one line:

$$p_n(a) = 2 + 2^n \implies \Delta p_n = p_{n+1} - p_n = 2^n : \quad (2, 4, 8, 16) = (|\mathbb{Z}_2|, |\mu_4|, h(D_5), \dim S^+).$$

So the single anchor $a = (1, 1, 2)$ generates the big numbers multiplicatively and the doubling spine 2, 4, 8, 16 additively — one vector, two readings.

The anchor reconstructs the carrier rank — and the 41-chain

The characteristic polynomial of the anchor vector is $(t-1)^2(t-2) = t^3 - 4t^2 + 5t - 2$, with elementary symmetric data

$$e_1(a) = 4 = |\mu_4|, \quad e_2(a) = 5 = g_{\text{car}}, \quad e_3(a) = 2 = |\mathbb{Z}_2|_{\text{sheet}}.$$

So the parabolic flavor geometry reads the carrier rank $g_{\text{car}} = 5$ *backwards*: $D_5 \Rightarrow a$ but also $a \Rightarrow g_{\text{car}}$. Moreover $3^2 + 4^2 = 5^2$ sits as $(p_0, e_1, e_2) = (3, 4, 5)$, giving the *anchor form* of the abelian index

$$10b_1 = e_1(a)^2 + e_2(a)^2 = 4^2 + 5^2 = 41$$

alongside the known $41 = 40 + 1 = \sum L + N_\Phi$. [E]

Pascal closure: $g_{\text{car}} = 5$ is the unique solution (kills alternatives) [E]

The carrier rank is not merely a good fit — it is the *unique* root of an exact closure. The even-exterior carrier dimension 2^{g-1} must equal the truncated Pascal sum $\binom{g}{0} + \binom{g}{1} + \binom{g}{2}$ (the $1+g+\binom{g}{2}$ that builds $\dim S^+$):

$$2^{g-1} = \binom{g}{0} + \binom{g}{1} + \binom{g}{2} \iff g = 5$$

($g=3 : 4 \neq 7$; $g=4 : 8 \neq 11$; **$g=5 : 16=16$** ; $g=6 : 32 \neq 22$; $g=7 : 64 \neq 29$). At $g = 5$ the row is $\binom{5}{0}, \binom{5}{1}, \binom{5}{2} = 1, 5, 10$, the same Pascal triple that reads family, action ladder and the E_8 root count $16 \cdot 5 \cdot 3 = 240$. This is the fourth independent lock on $g_{\text{car}} = 5$ (with rank-fill, Coxeter-match, and integer-glue/norm), and unlike those it explicitly *kills* the neighbours $g = 4, 6$.

The Pascal Ladder Theorem ([[verification/v108_pascal_ladder.py](#)]). The closure generalises exactly: for *every* truncation degree K ,

$$2^{g-1} = \sum_{k \leq K} \binom{g}{k} \iff g = 2K + 1$$

(odd g solves by the binomial midpoint identity; even g is excluded by the strict straddle $\sum_{k < g/2} < 2^{g-1} < \sum_{k \leq g/2}$). **Consequence (zero slack): the carrier selection $g = 5$ is *equivalent to the truncation degree* $K = 2$ — the Pascal-selection residual of the red team (Target B) is therefore exactly the one named hypothesis *Quadratic Boundary Locality* (the seam certifies code degrees ≤ 2 ; the architecture document). Neighbour worlds on the ladder: $K=1 \rightarrow g=3$ (a *one-family* world), $K=2 \rightarrow g=5$ (three families), $K=3 \rightarrow g=7$ (*nine* families), $K=4 \rightarrow g=9$ *inconsistent* ($N_{\text{fam}} = 255/9 \notin \mathbb{Z}$). Overdetermination: among the ladder worlds, only $K = 2$ also satisfies the independent rank closure $g + N_{\text{fam}} = 8 = \text{rank } E_8$ — two independent selections agree on the same world. The hypothesis itself stays [C].**

The seed is an anchor expression

The retained seed is exactly the two-term anchor combination of c_3 :

$$\varphi_0^{\text{ret}} = \frac{2p_1}{p_2} c_3 + 2p_1 p_2 c_3^4 = \frac{4}{3} c_3 + 48 c_3^4 = \frac{1}{6\pi} + \frac{3}{256\pi^4}$$

(verified to 30 digits). The Cabibbo base is $\lambda_Y = \sqrt{\varphi_0^{\text{ret}}(1 - \varphi_0^{\text{ret}})}$, the hierarchy engine of the whole fermion spectrum. [E]

3 The residue and length matrices as spectral objects

With the residue matrix R and the full word-length matrix $L = R + 6W$ (W the rank-one winding),

$$R = \begin{pmatrix} 1 & 3 & 0 \\ 1 & 5 & 2 \\ 2 & 5 & 3 \end{pmatrix}, \quad L = \begin{pmatrix} 7 & 3 & 0 \\ 7 & 5 & 2 \\ 8 & 5 & 3 \end{pmatrix}, \quad W = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{pmatrix},$$

the spectral invariants are pure compiler numbers (all verified):

Invariant of R	value	Invariant of L	value
$\text{tr } R$	$9 = N_{\text{fam}}^2$	$\text{tr } L$	15
$\det R$	$8 = h(D_5)$	$\det L$	$20 = R(A_4) = 2\mathcal{A}_\Lambda$
$\chi_R(t)$	$t^3 - 9t^2 + 10t - 8$	$\chi_L(t)$	$t^3 - 15t^2 + 40t - 20$
$\text{PrinMin}_2(R)$	$(2, 3, 5)$	$\text{PrinMin}_2(L)$	$(5, 14, 21)$
$\ R\ _F^2$	$78 = \dim E_6$	—	—

The three principal 2×2 minors of R are $(2, 3, 5)$ — sheet, family, carrier — with product $2 \cdot 3 \cdot 5 = 30 = h(E_8)$ and sum $10 = \mathcal{A}_\Lambda$. The wrong down-branch $\{1, 3, 4\}$ fails exactly these invariants.

Bilinear anchor forms. With $\mathbf{1} = (1, 1, 1)^\top$ and $a = (1, 1, 2)^\top$ (all verified):

R	$\mathbf{1}$	a	L	$\mathbf{1}$	a
$\mathbf{1}^\top$	$22 = \sum R$	$27 = 3^3$	$\mathbf{1}^\top$	$40 = R(D_5) $	$45 = \dim D_5$
$a^\top$	$32 = 2^5$	$40 = R(D_5) $	$a^\top$	$56 = \dim \mathbf{56}_{E_7}$	$64 = \dim S^+ \cdot \mu_4 $

The corner value $a^\top L a = 64 = (16, 4)$ reads exactly one spinor-gluon sheet of the E_8 branching. And the rank-one winding update shows up in the inverses:

$$R^{-1} \mathbf{1} = \frac{1}{4}(1, 1, -1)^\top, \quad L^{-1} \mathbf{1} = \frac{1}{10}(1, 1, -1)^\top$$

i.e. the residue response reads the glue denominator $|\mu_4| = 4$, the full-length response reads the decouple denominator $\mathcal{A}_\Lambda = 10$. [E]

Anchor-plane Plücker coordinates: scattered hits become oriented areas [E]

The strongest anti-numerology tool is to read *orientation invariants* of the anchor plane $\langle \mathbf{1}, a \rangle$ rather than loose integers. For $M \in \{R, Q, K, L\}$ the left block $B_M = (\mathbf{1}^\top M; a^\top M)$ (a 2×3) and the right

block $C_M = (M \mathbf{1} \ Ma)$ (a 3×2) have the 2×2 Plücker minors

	$\text{Pl}(\cdot)$ (left)	$\text{Pl}_R(\cdot)$ (right)
R	$(-6, 2, 14)$	$(8, 12, 4) = (h(D_5), R(A_3) , \mu_4)$
Q	$(3, 6, 3)$	$(0, 4, 5) = (0, \mu_4 , g_{\text{car}})$
K	$(-1, 6, 4) = (-N_\Phi, R^+(A_3) , \mu_4)$	$(12, 12, -2) = (R(A_3) , R(A_3) , - \mathbb{Z}_2)$
L	$(6, 26, 14)$	$(20, 30, 10) = (2\mathcal{A}_\Lambda, h(E_8), \mathcal{A}_\Lambda)$

Three load-bearing flavor numbers are now single oriented areas, not coincidences:

$$\|\text{Pl}(K)\|_1 = 11, \quad \|\text{Pl}_R(K)\|_1 = \text{Pl}(L)_{13} = 26, \quad \sum \text{Pl}_R(L) = 60 = D_{\text{start}}$$

(11 is the mass-power anchor-plane norm; 26 the c_t/c_b denominator; 60 the E_8 cascade start). The left-null-space response complements the right one above: the column sums of $\det L \cdot L^{-1}$ and $\det R \cdot R^{-1}$ are $(-5, 1, 6)$ and $(1, -5, 6)$, the same magnitudes $\{N_\Phi, g_{\text{car}}, |R^+(A_3)|\} = \{1, 5, 6\}$. Finally the sector rows themselves are anchor-shifted: $K_e - K_u = (1, 1, 2) = a$ and $L_e - L_u = (1, 2, 3) = \text{Exp}(A_3)$ — leptons are up-quarks displaced by the family roots. All entries are exact [\[E\]](#); this is an audit layer, not a closure (the physical quark amplitudes stay [\[C\]](#)). [\[verification/v37_plucker_anchor.py\]](#)




Figure 1: The E_8 slice atlas: seven maximal-subalgebra cuts of 248. Gold = spinor-glue ($64=(16, 4)$, $112=7 \cdot 16$, $84=7 \cdot 12$, 128); red = flavor-residue ($78=\|R\|_F^2$, $8=\det R$); green = family/action blocks; blue = Lie adjoints.

The numerology null test: the look-elsewhere burden quantified [E]

The raster above disciplines *which* integers may appear; the complementary statistical question — *could a random “theory” of equal formula complexity land this many data hits?* — is answered by an explicit, fully declared null model. For each of the 13 scored frozen-registry observables the *entire* grammar class is enumerated exactly (no sampling): one-term expressions $\frac{p}{q} (\varphi_0^{\text{ret}})^a \lambda_C^c (1+\lambda_C)^d \pi^b e^{u/w}$ with $p, q \leq 130$, $|a| \leq 2$, $|c| \leq 3$, $|d|, |b| \leq 1$, $|u| \leq 6$, $w \leq 6$, and two-term sums with coefficients ≤ 12 — a class that provably contains every TFPT formula it scores (machine-checked fairness condition). Data windows are conservative: the *maximum* of the achieved deviation, the documented experimental σ , the quote resolution and the 5% quark scheme spread, so the known tension observables ($\sin^2 \theta_{13}$, s_{23} , δ_{CKM} , m_μ/m_τ , β) enter with their honest *large* windows and contribute almost nothing. Result of the census: joint formula-fishing probability $\prod_i p_i = 10^{-25.8}$, stable to < 2 orders of magnitude when the coefficient bound is varied $80 \dots 200$ and the plausibility window $5 \times \dots 20 \times$. Monte-Carlo cross-check (deterministic seeds): $2 \times 2 \cdot 10^5$ pseudo-theories — random rational dials inside TFPT’s own formula skeletons, with the seed φ_0^{ret} fixed resp. drawn log-uniformly — reach at most 5/13 resp. 4/13 hits versus TFPT’s 13/13. The test has *power*: perturbing φ_0^{ret} by 1%/10%/30% collapses TFPT’s own scorecard to 9/6/4 hits, and shuffling which measured value is assigned to which formula collapses it to a mean of 1.2. Separately, all 94500 complexity-matched variants of the $F_{U(1)}$ fixed-point equation (coefficients, M , transport exponent, resolvent power) are root-solved: exactly *one* — the TFPT instance $(2, \frac{4}{5}, M=41, e^{-2\alpha}, -\frac{5}{4})$ — lands inside the achieved $4 \cdot 10^{-8}$ CODATA window, so $p_\alpha \leq 1.1 \cdot 10^{-5}$. Headline:

$$P(\text{null reproduces the scorecard incl. } \alpha) \leq 10^{-30.7} \quad (\approx 102 \text{ bits})$$

conditional on the declared grammar — a look-elsewhere-corrected null-model rejection, never “certainty”. Complementary to the frozen registry: the freeze (v84) prevents *future* look-elsewhere drift, this test quantifies the *retrospective* look-elsewhere burden. [verification/v100_numerology_null_mc.py]

4 The secondary branching atlas

Each maximal subalgebra of E_8 mirrors a TFPT module. The dimension identities are exact [E]; the readings are audit-level [O]. The seven *slices* of $\dim E_8 = 248$ are shown below: the same total cut seven ways, each cut a different TFPT module.

Subalgebra	248 =	TFPT reading
D_8	120 + 128	family-transition adjoint + two spinor-gluon sheets (128=2 dim $S^+ \mu_4 $)
$D_5 \times A_3$	45 + 15 + 60 + 64 + 64	main split: dim D_5 , dim A_3 , $\mathcal{A}_\Lambda R^+(A_3) $, two (16, 4) glue sheets
$E_6 \times A_2$	78 + 8 + 81 + 81	flavor residue: $\ R\ _F^2$, det R =dim A_2 , cubic family blocks
$E_7 \times A_1$	133 + 3 + 112	scalaron: 112=7 · 16 (exponent × one family)
$F_4 \times G_2$	52 + 14 + 182	occupancy/gauge: $ R(F_4) $ =48= Ω_{adm} , $ R(G_2) $ =12=dim $\mathfrak{g}_{\text{SM}}$
$A_4 \times A_4$	24 + 24 + 4 · 50	action ladder: 24+24=48= Ω_{adm} ; off-diagonal (5, 10) reps $\mathcal{A}_H, \mathcal{A}_\Lambda$
A_8	80 + 84 + 84	rank-decuplet + scalaron: 80=8 · 10, 84=7 · 12

The D_8 split is a magnitude/phase channel [E]/[C]/[O] ([verification/v227_degree_exponent_channel_split.py])
The D_8 row 248 = 120 + 128 is the canonical <i>exponent/degree</i> split of E_8 : the exponents sum to 120 = $ R^+(E_8) $ (a magnitude channel — positive roots, root transitions, the dimensionless ratio/product readouts), the fundamental degrees {2, 8, 12, 14, 18, 20, 24, 30} sum to 128 = $2^{\text{rank}-1}$ = rank E_8 · dim S^+ (a phase/gluon channel — the spinor/Ramond sheets). This gives the red-team Target D CP residual a clean home: the (ratios, product) $\Rightarrow v_{\text{geo}}$ magnitude bijection lives in the 120 channel, the un-covered CP phases in the 128 channel (cf. the hexagonal phase fiber, red team Target D). Whether 128 also hosts a dark sector is Frontier [O] — this is the honest replacement for the over-strong “ S^- is dark matter” reading, not an assertion of it.

E_8 Slice Compression Lemma [E] (audit)
The seven slices are not seven separate narratives — they are <i>one</i> projection of two small alphabets. The anchor power sums $p_n(a) = 2 + 2^n$ give $P = (p_0, p_1, p_2, p_3) = (3, 4, 6, 10)$, whose six K_4 edge products are exactly the carrier blocks $(p_0 p_1, \dots, p_2 p_3) = (12, 18, 30, 24, 40, 60) = (R(A_3) , N_{\text{fam}} R^+(A_3) , h(E_8), W(A_3) , R(D_5) , D_{\text{start}})$; the six Sheet-Diamond operators carry determinants (det Q , det K , det R , det C , det L , det F) = (3, 4, 8, 14, 20, 32) with $\sum \det = 81 = N_{\text{fam}}^4$, the total determinant charge of the flavour space. All seven 248-slices, the row-budget cross row($C + sU + tV$) = (7, 11, 13) + $s(3, 3, 3)$ + $t(1, 2, 3)$ (democratic winding + A_3 exponents), and dim $E_6 = 78 = p_2(p_0 + p_3)$ then follow from P , those determinants and $\Delta = p_0 + p_3 = 13$. The atlas is therefore a closed grammar over admissible invariant classes (power sums, determinants, branching dimensions), not a shelf of trophies ([verification/v170_diamond_slice_compression.py]).

Theorem 1 ($E_6 \times A_2$ flavor-residue shadow **[O]**). The principal flavor branching reads the residue matrix:

$$\|R\|_F^2 = 78 = \dim E_6, \quad \det R = 8 = \dim A_2 = h(D_5), \quad \mathbf{1}^\top R a = 27,$$

and the full $E_6 \times A_2$ dimension count is the residue identity

$$248 = \|R\|_F^2 + \det R + 2(\mathbf{1}^\top R a) N_{\text{fam}} = 78 + 8 + 2 \cdot 27 \cdot 3.$$

Here E_6 reads the Frobenius norm of the residue matrix, A_2 the pure three-family symmetry, and the **27** the cubic family block. (Strongest new E_8 flavor fingerprint.)

Theorem 2 ($E_7 \times A_1$ scalaron shadow **[O]**). With the scalaron exponent $7 = \Omega_{\text{adm}} - 10b_1 = -\deg E + \text{rk } E$,

$$248 = 133 + 3 + 112, \quad 112 = 7 \cdot \dim S^+ = 7 \cdot 16, \quad |R(E_7)| = 126 = 7 \cdot 18,$$

where $18 = N_{\text{fam}}|R^+(A_3)|$. The scalaron exponent 7 (the seam power $M_{\text{scal}}^2/\bar{M}_{\text{P1}}^2 = c_3^7$) is thus exactly the E_7 off-block coefficient. *Cross-check:* the centred E_8 exponents give $\sum_{m \in \text{Exp}(E_8)} |m - 15| = 56 = \dim \mathbf{56}_{E_7}$ (exponents {1, 7, 11, 13, 17, 19, 23, 29}, the primitive residues mod 30).




Figure 2: The same seven 248-slices, now coloured by their *alphabet* source rather than by physics module: blue = anchor power-sum blocks ($P = (3, 4, 6, 10)$), gold = Sheet-Diamond determinant blocks ($D = (3, 4, 8, 14, 20, 32)$, $\sum \det = 81 = N_{\text{fam}}^4$), green = symmetric/glue data. Every slice is one projection of the two alphabets (`[verification/v170_diamond_slice_compression.py]`).

Proposition 1 ($A_4 \times A_4$ action-ladder and $F_4 \times G_2$ occupancy [O]). $248 = 24 + 24 + 4(5 \cdot 10)$ with $24 + 24 = 48 = \Omega_{\text{adm}}$ and off-diagonal reps $(\mathbf{5}, \mathbf{10}) = (\mathcal{A}_H, \mathcal{A}_\Lambda)$ — the action ladder appears inside E_8 (an audit, *not* an $SU(5)$ GUT). And $F_4 \times G_2$: $|R(F_4)| = 48 = \Omega_{\text{adm}}$, $|R(G_2)| = 12 = \dim \mathfrak{g}_{\text{SM}}$, $\text{rank } F_4 = 4 = |\mu_4|$, $\text{rank } G_2 = 2 = |\mathbb{Z}_2|$.

Flavor-matrix and prime-table fingerprints (audit-level, machine-verified). Five exact linear-algebra readouts of R, L and the prime atoms, kept here under the no-free-pattern rule.

- **Row norms ladder.** $\|R_u\|^2, \|R_d\|^2, \|R_e\|^2 = (10, 30, 38) = (\mathcal{A}_\Lambda, h(E_8), 2 \cdot 19)$ with $10+30+38 = 78 = \dim E_6$ ($19 \in \text{Exp}(E_8)$).
- **Residue inventory.** The entries of R on C_6 are $m = (1, 2, 2, 2, 0, 2)$: the *only* gap is $4 = |\mu_4| = \sum a$. Moments $\sum R = 22$, $\sum R^2 = 78 = \dim E_6$ — so $\|R\|_F^2 = 78$ is the second moment of the C_6 inventory with the μ_4 hole.
- **First-column norm** $\rightarrow E_6 \times A_2$ **off-blocks.** $L_{\cdot 1} = (7, 7, 8)$, $\|L_{\cdot 1}\|^2 = 162 = 81 + 81$ — the two **81** off-blocks of the $E_6 \times A_2$ slice are exactly the wound first-generation column.
- **Lepton exponent polynomial.** The φ_0^{ret} -ladder lepton powers $(k_e, k_\mu, k_\tau) = (5, 3, 2)$ are the Coxeter primes $30 = 2 \cdot 3 \cdot 5$ in hierarchy order: $(t-5)(t-3)(t-2) = t^3 - 10t^2 + 31t - 30$ ($e_1 = \mathcal{A}_\Lambda$, $e_3 = h(E_8)$). The lepton c -ratios obey $\frac{(16/7)(7/6)}{4/3} = 2 = |\mathbb{Z}_2|$ and $\frac{(16/7)(7/6)}{(4/3)^2} = \frac{3}{2}$ (the near-Koide factor).
- **The 2, 3, 5 multiplication table** (the shortest integer-landscape compression):

$$(2, 3, 5) \xrightarrow{\times} (6, 10, 15, 30) \rightarrow (8=3+5, 16, 240=8 \cdot 30, 248=240+8).$$

5 Group-theoretic audits (optional)

Weyl completion index. Over the main stabiliser,

$$\frac{|W(E_8)|}{|W(D_5)| |W(A_3)|} = \frac{696729600}{1920 \cdot 24} = 15120 = 63 \cdot 240 = (2^6 - 1) |R(E_8)|.$$

A group-theory machine-check, not physics: it quantifies the Weyl symmetry created beyond the $D_5 \times A_3$ stabiliser. [E] (arithmetic) / [O] (reading)

E_8 **theta series as a higher-shell raster.** $\Theta_{E_8} = E_4 = 1 + 240 \sum_{n \geq 1} \sigma_3(n) q^n$ has shell degeneracies 240, 2160, 6720, 17520, ... The first is $|R(E_8)| = 240$; the second is $2160 = 9 \cdot 240 = N_{\text{fam}}^2 |R(E_8)|$. If later one-loop or CMB shell degeneracies touch these multiplicities, the theta series is the right place to type them — as an audit raster, not a claim. [O]

6 Further audit lemmas (machine-verified this round)

Five more exact identities tie the flavor matrices and the E_8 exponents to the compiler. All checked with exact integer arithmetic.

(i) **Full-length Frobenius norm is three E_6 blocks.** $\|L\|_F^2 = 234 = N_{\text{fam}} \|R\|_F^2 = 3 \dim E_6$, and $\|L\|_F^2 = \|R\|_F^2 + 12\langle R, W \rangle + 36\|W\|_F^2 = 78 + 48 + 108$, with $\langle R, W \rangle = 4 = |\mu_4|$ and $\|W\|_F^2 = 3 = N_{\text{fam}}$. The full word-length load is exactly the residue E_6 norm + the A_3 -winding over μ_4 + the family-winding square. [E]

(ii) **A_3 -winding determinant lemma.** The rank-one winding $L = R + 6\mathbf{1}e_1^\top$ and the matrix-determinant lemma with $e_1^\top R^{-1}\mathbf{1} = \frac{1}{4}$ give $\det L = \det R(1 + \frac{|R^+(A_3)|}{|\mu_4|}) = 8(1 + \frac{6}{4}) = 20 = 2\mathcal{A}_\Lambda$.

So $h(D_5) = 8 \xrightarrow{A_3^+/\mu_4} 20 = |R(A_4)|$: the winding injects the carrier factor 5 into the cokernel ($\text{coker } R \simeq \mathbb{Z}_8$, $\text{coker } L \simeq \mathbb{Z}_{20}$). [E]

(iii) **The scalaron 7 is the E_8 exponent variance per decuple.** Centring the exponents at $h/2 = 15$, $\sum_m (m - 15)^2 = 560$ and $\text{Var} = \frac{560}{8} = 70 = 7 \cdot \mathcal{A}_\Lambda = (\text{scalaron } 7) \cdot (\text{decuple } 10)$. The halved positive offsets $\{1, 2, 4, 7\}$ satisfy $\sum = 14 = 2 \cdot 7$, $\sum^2 = 70$, $\prod = 56 = \dim \mathbf{56}_{E_7}$ — a micro-code carrying the scalaron, the decuple and the E_7 minuscule at once. Also $\sum_{\text{pairs}} m(30 - m) = 620 = \frac{h \dim E_8}{12} = 2 \cdot 10 \cdot 31$. [E]

(iv) **Exponents + degrees give the 120+128 split.** The invariant degrees $d_i = m_i + 1 = \{2, 8, 12, 14, 18, 20, 24, 30\}$ obey $\sum_i m_i = 120$, $\sum_i d_i = 128$, $248 = 120 + 128$, $\prod_i d_i = |W(E_8)|$, matching the D_8 branching $248 = 120 + 128$; the degrees are a TFPT dictionary (2, 8, 12, 14, 18, 20, 24, 30 = sheet, $h(D_5)$, $|R(A_3)|$, $\dim G_2$, p_4 , $\det L$, $\dim A_4$, $h(E_8)$). [E]

(v) **Root shell = local $E_7 \times A_1$.** Around any E_8 root the inner-product multiplicities are $\{+2:1, +1:56, 0:126, -1:56, -2:1\}$, i.e. the local $248 = (133, 1) + (1, 3) + (56, 2)$. A distinguished root (seam/Higgs direction) leaves an orthogonal E_7 and an active coupling shell $56+56$. [E]

7 What this changes, and what stays open

E_8 is one container whose maximal subalgebras each project a different TFPT module:

Net. The atlas turns isolated fingerprints (e.g. $\|R\|_F^2 = 78 = \dim E_6$, the scalaron 7) into *branching projections* of one container. The discipline rule — “every load-bearing number must appear in ≥ 1 projection” — makes the structure falsifiable: a number in no projection is suspect. This is a *program*, not a proof of new physics: all readings are [O]. The two genuinely load-bearing, derived results remain the μ_4 glue ($E_8 = (D_5 \oplus A_3) + \mu_4$, [E] [verification/v1_e8_glue.py]) and the spectral selector of the flavor matrix ($\det R = h(D_5)$, $\text{PrinMin}_2 = (2, 3, 5)$, [E] [verification/v4_flavor_matrix.py]). The remaining flavor-geometry step (historically “H2”) is sharpened in the parabolic note to: *the D_4 -equivariant stable parabolic structure on $\mathcal{O}(-2) \oplus \mathcal{O}(-1)^2$ realises the unique branch with $\det R = 8$ and $\text{PrinMin}_2(R) = (2, 3, 5)$ — no longer “find the matrix”, but “prove this one geometric realisation”.* (Status: since v118–v122 the dictionary *values* are exact theorems — hexagon = sign-twisted family spectrum, addresses pinned, margins forced; what remains is exactly the realisation statement above, the selector establishment $R4'$.)




Figure 3: E_8 as an audit container: each maximal subalgebra projects one TFPT module. The discipline rule — every load-bearing number appears in ≥ 1 projection — makes the structure falsifiable.

Part II

The E_8 cascade bridge

The connection in one sentence. The old E_8 cascade (Paper V1.06: the nilpotent-orbit chain $D_n = 60 - 2n$, $n = 0 \dots 26$) and the new compiler/atlas are *the same even-integer E_8 spine*, read three ways: the cascade *orders* the numbers $\{8, 10, \dots, 60\}$ as a scale ladder, the atlas *reads* them as E_8 subalgebra branchings, and the compiler *generates* them from the $(1, 1, 2)$ anchor. The cascade *starts* at $D_{\text{start}} = 60 = p_2 p_3$ and *ends* at $8 = h(D_5) = \det R$ — the flavor selector.

8 The old cascade, recalled

Paper V1.06 builds a deterministic scale ladder from the nilpotent orbits of E_8 . With centralizer dimension $D = 248 - \dim \mathcal{O}$ and the Hasse chain restricted to $\Delta D = 2$, a unique maximal chain runs from $A_4 + A_1$ ($D_0 = 60$) down to the regular orbit ($D_{26} = 8$):

$$D_n = 60 - 2n, \quad n = 0, \dots, 26; \quad \varphi_{n+1} = \varphi_n e^{-\gamma(n)}, \quad \gamma(n) = \lambda [\ln D_n - \ln D_{n+1}]$$

The cascade arithmetic (endpoints, exponent rungs, IR-tail product, variance) is machine-checked. [\[verification/v5_e8_cascade.py\]](#) with the single structural normalisation $\lambda = \gamma(0)/(\ln 248 - \ln 60) = 0.58770$, fixed by a curvature-extremum on the chain (no fit). The E -windows of the two-loop flow anchor it: $\alpha_3 = 1/(8\pi) = c_3$ (the E_8 window) and $\alpha_3 = 1/(6\pi)$ (the E_6 window, near φ_0^{ret}).

9 The spine: every rung is a current number

Reading $D = 60 - 2n$ against today's compiler atoms, the secondary-branching atlas and the $(1, 1, 2)$ anchor power sums $p_n = 2 + 2^n$, **25 of the 27 rungs are load-bearing numbers** (machine-checked):

n	D	current meaning	n	D	current meaning
0	60	$D_{\text{start}} = p_2 p_3 = \mathcal{A}_\Lambda R^+(A_3) $ (start)	13	34	p_5 anchor power = $2 \cdot 17$
1	58	$2 \cdot 29$ (E_8 exponent)	14	32	$2^5 = 2 \dim S^+$
2	56	$\dim \mathbf{56}_{E_7} = a^\top L \mathbf{1} = \sum m-15 $	15	30	$h(E_8) = 2 \cdot 3 \cdot 5$ (Coxeter rung)
3	54	$2 \cdot 27$, $27 = \mathbf{1}^\top R a = 3^3$	16	28	$4 \cdot 7$ (half of 56)
4	52	$\dim F_4 = R(D_5) + R(A_3) $	17	26	$2 \cdot 13$ (old hadronic block)
5	50	$A_4 \times A_4$ block = $5 \cdot 10$	18	24	$\dim A_4$; $24+24=48$
6	48	$\Omega_{\text{adm}} = R(F_4) = 2p_1 p_2$	19	22	$\sum R = \mathbf{1}^\top R \mathbf{1}$ (residue sum)
7	46	$2 \cdot 23$ (E_8 exponent)	20	20	$\det L = R(A_4) = 2\mathcal{A}_\Lambda$
8	44	$4 \cdot 11$ (E_8 exponent)	21	18	$p_4 = N_{\text{fam}} R^+(A_3) $
9	42	$6 \cdot 7 = R^+(A_3) \cdot 7$	22	16	$\dim S^+$ (one generation)
10	40	$ R(D_5) = \sum L = \mathbf{1}^\top L \mathbf{1}$	23	14	$\dim G_2 = \sum L_d = 8+6$
11	38	$2 \cdot 19$ (E_8 exponent)	24	12	$ R(A_3) = \dim \mathfrak{g}_{\text{SM}}$
12	36	6^2 (old EW block)	25	10	$\mathcal{A}_\Lambda = \binom{5}{2}$ (old cosmo block)
			26	8	$h(D_5) = \det R = 2p_1$ (end: flavor selector)

Three structural facts jump out (all [\[E\]](#) arithmetic):

- **Endpoints are the compiler boundary.** The cascade *starts* at $60 = p_2 p_3 = D_{\text{start}}$ and *ends* at $8 = h(D_5) = \det R$ — the flavor spectral selector is literally the deepest orbit of E_8 .
- **Coxeter rung is $h(E_8)$.** At $n = 15$, $D = 30 = h(E_8)$ — this is the *Coxeter rung*, not the arithmetic centre; the median node is $n = 13$, $D = 34 = p_5(a)$ (below). It matches the exponent-centering audit $\sum_m |m - 15| = 56 = \dim \mathbf{56}_{E_7}$ (which is itself rung $n = 2$).
- **Doubled E_8 exponents are rungs.** $\{58, 46, 38, 34, 26, 22, 14\} = 2 \times \{29, 23, 19, 17, 13, 11, 7\}$ are exactly the cascade rungs at $n = 1, 7, 11, 13, 17, 19, 23$; three of them coincide with anchor

numbers ($34 = p_5$, $22 = \sum R$, $14 = \dim G_2$). Only $n = 8$ ($D = 44 = 4 \cdot 11$) is not independently labelled.

Cascade arithmetic — endpoints, centre and length (verified)

The 27 rungs encode E_8 in their endpoints, and their centre is the anchor:

$$\begin{aligned} \frac{D_{\text{start}} D_{\text{end}}}{2} &= \frac{60 \cdot 8}{2} = 240 = |R(E_8)|, & 240 + D_{\text{end}} &= 248 = \dim E_8, & D_{\text{start}} &= 2h(E_8) = 60; \\ \text{arithmetic median} &= \frac{60+8}{2} = 34 = p_5(a) = h(E_8) + |\mu_4|, & \sum_{n=0}^{26} (D_n - 30) &= 27 \cdot 4 = 27|\mu_4|; \\ \text{width} &= D_{\text{start}} - D_{\text{end}} = 52 = \dim F_4, & 60 \rightarrow 30 &= h(E_8), & 30 \rightarrow 8 &= \sum R = 22; \\ \# \text{nodes} &= 27 = 3^3 = \mathbf{1}^\top R a \text{ (} E_6 \text{ cubic block)}, & \# \text{steps} &= 26 = \dim \mathbf{26}_{F_4}. \end{aligned}$$

So the cascade is an E_8 Coxeter ladder centred on $p_5(a)$ with a per-node $|\mu_4|$ glue shift; its length (27 nodes, 26 steps) carries the E_6 cube and the F_4 fundamental. [E]

The IR tail $n = 24 \dots 30$: the compiler atoms

Continuing $D = 60 - 2n$ past the last orbit ($D=8$, $n=26$) is no longer a nilpotent-orbit chain but the *algebraic compiler tail* of the Coxeter clock:

$$D_{24 \dots 30} = \underbrace{12}_{|R(A_3)|=\dim \mathfrak{g}_{\text{SM}}}, \underbrace{10}_{\mathcal{A}_\Lambda}, \underbrace{8}_{h(D_5)}, \underbrace{6}_{|R^+(A_3)|}, \underbrace{4}_{|\mu_4|}, \underbrace{2}_{\text{sheet}}, \underbrace{0}_{\text{closure}}.$$

This is exactly what the old “ $n = 0 \dots 30$ clock” was scanning: $n=0 \dots 26$ the E_8 orbit spine, $n=27 \dots 30$ the compiler atoms. (Mark the tail as the *compiler tail*, not orbits.) [E]

Three further cascade lemmas (all machine-verified)

The clock carries more E_8 data than the endpoints:

- **Primitive rungs sum to the root count.** Evaluating $D_n = 60 - 2n$ at the *primitive* indices $n \in (\mathbb{Z}/30)^\times = \{1, 7, 11, 13, 17, 19, 23, 29\} = \text{Exp}(E_8)$ gives $\{58, 46, 38, 34, 26, 22, 14, 2\}$ with

$$\sum_{n \in \text{Exp}(E_8)} (60 - 2n) = 240 = |R(E_8)|,$$

and since $30 - n$ is primitive when n is, the clock indexes its own E_8 exponents.

- **The IR tail product is the main Weyl stabiliser.** The nonzero tail $\{12, 10, 8, 6, 4, 2\}$ has


$$\prod = 46080 = |W(D_5)| |W(A_3)| = 1920 \cdot 24, \quad \sum = 42 = |R^+(A_3)| \cdot 7.$$

- **Cascade variance = E_6 flavour norm times the scalaron-gauge block.** The 27 rungs $60, \dots, 8$ have mean $34 = p_5(a)$ and

$$\sum_{n=0}^{26} (D_n - 34)^2 = 6552 = 78 \cdot 84 = \dim E_6 \cdot (7 \dim \mathfrak{g}_{\text{SM}}),$$

tying the cascade simultaneously to the $E_6 \times A_2$ flavour shadow and the A_8 scalaron-gauge shadow.
[E] (audit level)

10 Three views of one spine



The unification. The same numbers $\{8, 10, \dots, 60\}$ are: (a) the *scale ladder* φ_n of the old cascade (orbit dimensions, ordered by $e^{-\gamma}$); (b) the *branching blocks* of the secondary-atlas ($\dim F_4=52$, $\dim \mathbf{56}_{E_7}$, $5 \cdot 10, \dots$); (c) the *anchor power sums and matrix invariants* of the compiler ($p_n = 2 + 2^n$, $\det R=8$, $\sum L=40$). The cascade is the *UV*→*IR* ordering; the compiler is its *IR endpoint*, where the small numbers 8, 10, 12, 16 are the flavor/family/carrier data.

The spine tetrahedron: the microkernel as one finite object [E]
([verification/v91_spine_tetrahedron.py])

The spine $T = \{2, 3, 4, 5\} = \{\mathbb{Z}_2, N_{\text{fam}}, |\mu_4|, g_{\text{car}}\}$ equals $\{e_3(a), p_0(a), e_1(a), e_2(a)\}$ for the anchor $a = (1, 1, 2)$ — the tetrahedron on these four nodes is the anchor plus its zeroth moment, not a second structure. Its combinatorics *are* the central integer grammar: edge products $\{6, 8, 10, 12, 15, 20\}$ ($|R^+(A_3)|$; $\text{rank } E_8=h(D_5)=\det R$; $\mathcal{A}_\Lambda=|E(K_5)|$; $|R(A_3)|=\dim \mathfrak{g}_{\text{SM}}$; $\dim A_3$; $\det L$), face products $\{24, 30, 40, 60\}$ ($|W(A_3)|$; $h(E_8)$; $|R(D_5)|=\sum L$; D_{start}), volume $2 \cdot 3 \cdot 4 \cdot 5 = 120 = |R^+(E_8)| = 5!$; and the graph reading $|R(E_8)| = 240 = |\mu_4| \cdot |E(K_4)| \cdot |E(K_5)|$ with $|E(K_4)| = p_2 = 6$, $|E(K_5)| = p_3 = 10$ — including the honest negative control that the identification *breaks* at K_6 ($|E(K_6)| = 15 \neq p_4 = 18$), so it is specific, not generic. Two disciplined caveats: the “dual cuts” (node $\times$ opposite face = 120, etc.) are tautologies of any complementary 4-set partition and are typed as presentation only; and the tetrahedron grammar is a *sub*-grammar — the load-bearing 7, 16, 41, 48, 240, 248 are *not* subset products and still require the

anchor power sums / Pascal / glue layers.

11 Direct φ_0^{ret} relations: old vs current

The old cascade’s “fundamental relations near $n = 0$ ” connect directly to the current note — one is *identical*, two are early forms of current refinements:

Quantity	old (V1.06)	current	relation
Ω_b	$\varphi_0^{\text{ret}}(1 - 2c_3) = 0.04894$	$(1 - \frac{1}{4\pi})\varphi_0^{\text{ret}} = 0.04894$	identical ($2c_3 = \frac{1}{4\pi}$) [E]
V_{us}/V_{ud}	$\sqrt{\varphi_0^{\text{ret}}} = 0.2306$	$\lambda_C = \sqrt{\varphi_0^{\text{ret}}(1 - \varphi_0^{\text{ret}})} = 0.2244$	refined (the $(1 - \varphi_0^{\text{ret}})$ factor) [E]
r (tensor)	$(\varphi_0^{\text{ret}})^2 = 0.00283$	$12/N_\star^2 = 0.00397$	two branches, same order [C]

Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in `verification/status_ledger.csv`.

So $\Omega_b = \varphi_0^{\text{ret}}(1 - 2c_3)$ was already the exact current formula in V1.06; the Cabibbo ratio gained its $(1 - \varphi_0^{\text{ret}})$ factor; and the old direct $r = (\varphi_0^{\text{ret}})^2$ and the current Starobinsky $r = 12/N_\star^2$ agree in order (~ 0.003) but are distinct predictions — a clean, testable discriminant, stated honestly.

12 Sector inventory: have all the old-cascade sectors been computed?

The old cascade also carried *calibrated absolute scales*: per block, one unit ζ_B was fixed on a single reference and every other scale in the block followed fit-free from the ratio law $\varphi_m/\varphi_n = (D_m/D_n)^\lambda$ (Paper V1.06, App. C). These block assignments are the place where the proton mass, the EW masses and the CMB temperature lived. The honest question is: *has the new compiler re-derived each of these, or only some?* The full inventory, with the current status spelled out:

Old block (rung)	Old observable & calibration	Current-model status	Mark
near $n=0$	$\Omega_b = \varphi_0^{\text{ret}}(1 - 2c_3),$ $V_{us}/V_{ud} = \sqrt{\varphi_0^{\text{ret}}},$ $r = (\varphi_0^{\text{ret}})^2$	covered — dimensionless, derived from $\{c_3, \varphi_0^{\text{ret}}\}$; see the φ_0^{ret} -table above. Ω_b is <i>identical</i> .	[E]
EW, $n=12$	$v_H = \zeta_{\text{EW}} M_{\text{P}1} \varphi_{12},$ $M_W = \frac{1}{2} g_2 v_H, M_Z$	scale covered, absolute scheme-level — the EW scale enters the current note via the action-grammar exponent and the two-loop flow; M_W, M_Z are dimensionful and remain RG/scheme outputs (one ζ_{EW}), <i>not</i> new compiler powers.	[C]
hadronic, $n=15, 17$	$m_p = \zeta_p M_{\text{P}1} \varphi_{15},$ $m_b = \zeta_b M_{\text{P}1} \varphi_{15},$ $m_u = \zeta_u M_{\text{P}1} \varphi_{17}$	ratios closed, absolute scale an anchor. The quark mass <i>ratios</i> sit on the φ_0^{ret} -word-length ladder and are closed (integer Plücker readouts, <code>v49/v71</code>); the <i>absolute</i> values need the per-sector $\Lambda/U_f^\star$ normalisation, an anchor ([O]). The proton mass is QCD-confinement, $m_p \sim \Lambda_{\text{QCD}}$ by dimensional transmutation — a <i>cross-sector</i> ratio, listed as a genuine frontier item (m_p/m_e , <code>tfpt_4</code>).	[O]
cosmo, $n=25$	$T_{\gamma 0} = \zeta_\gamma M_{\text{P}1} \varphi_{25},$ $T_\nu = (\frac{4}{11})^{1/3} T_{\gamma 0}, \rho_\Lambda$	partly. ρ_Λ/Λ is <i>typed</i> by the action grammar and the seam (Λ -typing, <code>tfpt_2</code>); $T_\nu/T_{\gamma 0}$ is standard QFT (not cascade). The absolute CMB temperature $T_{\gamma 0}$ is a downstream readout of the reheating pipeline (the old cosmology/CMB layers; today the scalaron-reheating chain, <code>v86</code>), conditional — not re-derived as a compiler number.	[C]

Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in `verification/status_ledger.csv`.

Direct answer: which sectors are closed, which are not

Closed (derived from $\{c_3, g_{\text{car}}, \varphi_0^{\text{ret}}\}$): the dimensionless near- $n=0$ anchors (Ω_b exactly, Cabibbo, tensor branch) and all the *dimension/rung* structure of the spine. **Scale-level only (one ζ_B each, scheme/RG):** the EW masses M_W, M_Z and the cosmological absolute scales $T_{\gamma 0}$ — exactly as in the current SM/cosmology notes; these were *never* pure predictions, not even in V1.06 (each needed a block unit). **Genuinely open:** the **proton mass** m_p (and m_p/m_e) is QCD-confinement across two sectors and is *not* claimed as a compiler power — it stays a frontier item ([O], tfpt_4); the quark mass *ratios* are closed (integer Plücker, $\sqrt{49}/\sqrt{71}$) and only the *absolute* quark scale m_b, m_u needs the per-sector Λ/U_f^* normalisation, an anchor ([O]). So: yes, the *spine* is fully audited and the *dimensionless* sectors are derived; the *dimensionful* hadronic/cosmo absolute scales are honestly still scheme-level, with the proton mass the one explicitly open cross-sector number — nothing is silently claimed.

13 What is new, what is old, what stays open

New structures found via the bridge

1. The flavor selector $\det R = h(D_5) = 8$ is the *terminal orbit* of the E_8 cascade; the compiler is the cascade's IR endpoint. [E]
2. The atlas branching blocks (52, 56, 60) are the cascade's *UV head* ($n = 0, 2, 4$); the secondary atlas and the orbit cascade are the same spine read top-down vs. as subalgebras. [E]
3. The anchor power sums sit on doubled-exponent rungs: $p_5 = 34$ ($n=13$), $\sum R = 22$ ($n=19$), $\dim G_2 = 14$ ($n=23$). [E]
4. The cascade Coxeter rung $D = 30 = h(E_8)$ ($n=15$) ties to the exponent-centering audit $\sum |m-15| = 56$. [E]

Honest separation

The *dimension/rung coincidences* are exact arithmetic [E] — they show the old and new constructions live on one E_8 spine. The old cascade's *scale assignments* φ_n (EW at $n=12$, hadronic at $n=15, 17$, cosmo at $n=25$) are a *different* mechanism from the current φ_0^{ret} -power mass ladder: the cascade orders *orders of magnitude* (which sector sits where) via $e^{-\gamma}$, while the φ_0^{ret} -ladder gives *values within* a sector. Both are real, complementary orderings; neither is forced onto the other. The per-sector ζ -calibrated absolute scales (proton mass, EW masses, CMB temperature) remain scheme-level — audited one by one in §12, with the proton mass the single explicitly open cross-sector number.

Bottom line. The old E_8 cascade was an early, correct intuition: the physically relevant numbers are exactly the even-integer E_8 orbit spine $\{8, \dots, 60\}$. The current work *explains why* — they are the $(D_5 \times A_3) + \mu_4$ compiler atoms and their E_8 branching blocks — and identifies the spine's two ends as the compiler boundary ($D_{\text{start}} = 60$, $\det R = 8$). Old and new are the same structure at different resolution.

Part III

The self-consistency loop

The question, and the honest answer. The theory rests on two inputs, $c_3 = \frac{1}{8\pi}$ and $g_{\text{car}} = 5$. The old theory had a Möbius seam and the seed was the *start* of the E_8 cascade. Run it *backwards*: can the cascade/ E_8 *return* c_3 and g_{car} , closing a self-consistent loop? **Answer:** the *discrete* content is a closed, overdetermined fixed point — $g_{\text{car}} = 5$ and the “8” in c_3 are forced by the E_8 closure. The *only* irreducible primitive left is the geometric π (the Möbius/boundary Gauss–Bonnet). So “two axioms” collapses to “*one* continuous primitive π + *one* discrete self-consistency fixed point.”

14 The loop for g_{car} — an overdetermined fixed point

Two *independent* structural conditions, both pure Lie-theory, force $g_{\text{car}} = 5$ once E_8 is the closure object and A_3 is the four-puncture family geometry:

$g_{\text{car}} = 5$ is forced twice over

(A) rank fill.

The glue $D_{g_{\text{car}}} \oplus A_3 \hookrightarrow E_8$ must fill the rank: $\text{rank } D_{g_{\text{car}}} + \text{rank } A_3 = \text{rank } E_8$, i.e. $g_{\text{car}} + 3 = 8 \Rightarrow g_{\text{car}} = 5$.

(B) Coxeter match.

The carrier Coxeter number must equal the E_8 rank, $h(D_{g_{\text{car}}}) = 2g_{\text{car}} - 2 = 8 \Rightarrow g_{\text{car}} = 5$.

Both give $g_{\text{car}} = 5$, and the loop *closes*: $g_{\text{car}} = 5 \Rightarrow D_5 \oplus A_3 \xrightarrow{\mu_4} E_8 \Rightarrow h(E_8) = 30 = 2 \cdot 3 \cdot 5 \Rightarrow \text{largest prime} = 5 = g_{\text{car}}$. [E]

Reverse glue selection — D_5 and A_3 from unimodular closure

The closure even *selects the two factors*, not just the rank. The discriminant of $A_{\mu-1}$ is $\mathbb{Z}_\mu$ and of D_g (odd g) is $\mathbb{Z}_4$; a common cyclic glue forces $\mu = 4$, hence $A_{\mu-1} = A_3$. The even-glue norm $q(D_g) + q(A_3) = \frac{g}{4} + \frac{3}{4} = 2$ then forces $g = 5$, hence D_5 . A second, independent confirmation of $\mu = 4$ is the *mirror identity* $\dim D_{\mu+1} + \dim A_{\mu-1} = \dim V_{D_{\mu+1}} \cdot |R^+(A_{\mu-1})|$, i.e. $3\mu(\mu+1) = \mu(\mu+1)(\mu-1) \Rightarrow \mu-1 = 3 \Rightarrow \mu = 4$ (the $45 + 15 = 10 \cdot 6$ identity). So

$$\text{common } \mathbb{Z}_4 \text{ glue} + \text{root norm } 2 \implies D_5 + A_3$$

— $g_{\text{car}} = 5$ is recovered from the closure, not just posited. [E]

Familyful E_8 -closure fixed-point theorem

Among rank-filling pairs $D_g \oplus A_m$ ($g+m=8$) admitting an integer glue index $\sqrt{\det D_g \det A_m} = \sqrt{4(m+1)}$, exactly two close to an even unimodular lattice: $D_8 \oplus A_0$ (no family) and $D_5 \oplus A_3$ (familyful). The unique familyful solution is $D_5 \oplus A_3$, with lattice $\cong E_8$ and $g_{\text{car}} = 5$.

$D_g + A_m$	$g+m$	glue idx ² =4(m+1)	integer?	status
$D_8 + A_0$	8	4	2	valid, <i>no family</i>
$D_7 + A_1$	8	8	—	no integer glue
$D_6 + A_2$	8	12	—	no integer glue
$D_5 + A_3$	8	16	4= μ_4 	valid, familyful, E_8
$D_4 + A_4$	8	20	—	no integer glue

Equivalently, with $q(A_{\mu-1}) = \frac{\mu-1}{\mu}$ and $g = 9 - \mu$ the norm closure $q(D_g) + q(A_{\mu-1}) = 2$ becomes $\mu^2 - 5\mu + 4 = 0$, whose nontrivial root is $\mu = 4$ ($A_3, g = 5$). So $g_{\text{car}} = 5$ is forced *three* ways: rank-fill, Coxeter-match, and integer-glue/norm closure. [E] [verification/v6_bootstrap.py]

Bootstrap classification: $D_5 \oplus A_3$ is the *unique* familyful E_8 glue [E]

E_8 is the unique even unimodular rank-8 lattice. Its two-factor root-lattice gluings $L_1 \oplus L_2$ (rank = 8) that admit a *cyclic* glue are exactly those whose discriminant groups are isomorphic and cyclic:

$$D_5 \oplus A_3 (\mathbb{Z}_4), \quad E_6 \oplus A_2 (\mathbb{Z}_3), \quad E_7 \oplus A_1 (\mathbb{Z}_2), \quad A_4 \oplus A_4 (\mathbb{Z}_5).$$

Imposing the two TFPT data — one factor with a 16-dimensional half-spinor ($\dim S_{D_g}^+ = 2^{g-1} = 16 \Leftrightarrow g = 5$) and one factor giving exactly $N_{\text{fam}} = 3$ families (rank $A_m = m = 3$) — selects

$$D_5 \oplus A_3 \text{ (uniquely), glue group } \mathbb{Z}_4 = |\mu_4|.$$

$E_6 \oplus A_2$ gives only 2 families, $E_7 \oplus A_1$ only 1; D_8 has non-cyclic glue $(\mathbb{Z}_2)^2$ and A_8 is a single rank-8 factor. So the carrier+family decomposition is forced. [verification/v15_bootstrap_classification.py]




Figure 4: The closed loop: $g_{\text{car}} = 5$ builds E_8 , whose Coxeter number $30 = 2 \cdot 3 \cdot 5$ returns the carrier prime 5. The value $g_{\text{car}} = 5$ is the unique fixed point (overdetermined by rank-fill *and* Coxeter-match).

The carrier split $3 + 2$ and the one-generation count $\dim S^+ = 2^{g_{\text{car}}-1} = 16$ are then automatic, as is the glue norm $q(D_5) + q(A_3) = \frac{5}{4} + \frac{3}{4} = 2$. The family rank is rank $A_3 = 3 = N_{\text{fam}}$ and the glue index is $h(A_3) = 4 = |\mu_4|$ — the Coxeter numbers of the three players are $(h(A_3), h(D_5), h(E_8)) = (4, 8, 30) = (|\mu_4|, \det R, \text{compiler})$.

15 The loop for c_3 — the “8” is the E_8 rank

The boundary normalisation is $c_3 = 1/(8\pi)$. Its integer is exactly the carrier Coxeter number, which the loop above equals to the E_8 rank:

$$c_3 = \frac{1}{h(D_5) \pi} = \frac{1}{(2g_{\text{car}} - 2) \pi} = \frac{1}{8\pi}, \quad 8 = h(D_5) = \text{rank } E_8 = \det R = \varphi(30)$$

where $\varphi(30) = 8$ is Euler’s totient (the number of primitive residues mod $30 =$ the E_8 exponents) — so the “8” has *five* concordant readings. The tree seed $\varphi_{\text{tree}} = 1/(6\pi)$ has $6 = |R^+(A_3)| = h(E_8)/g_{\text{car}}$, the three Gauss–Bonnet boundary cycles ($6\pi = 3 \times 2\pi$). So the *discrete* parts of both c_3 and φ_{tree} are $E_8/\text{carrier}$ invariants; only the factor π is genuinely continuous.

Honest caveat on 8π

8π is *also* the standard Einstein normalisation ($G_{\mu\nu} = 8\pi G T_{\mu\nu}$). The loop *identifies* this gravitational 8 with rank $E_8 = h(D_5)$ — a non-coincidence *if* E_8 is the closure, but one should be clear that 8π has an independent GR reading. The loop promotes the two 8's to the same number; it does not derive π .

16 The Möbius origin of the irreducible π

What the loop does *not* remove is π . In the old theory this is exactly the Möbius datum: the orientable double cover of the Möbius fibre is a cylinder with two boundaries plus one $\mathbb{Z}_2$ seam Γ ; Gauss–Bonnet gives $\oint k ds = 2\pi + 2\pi + 2\pi = 6\pi$ (hence $\varphi_{\text{tree}} = 1/(6\pi)$), and the $\mathbb{Z}_2$ identification is the sheet “2” of $30 = 2 \cdot 3 \cdot 5$. So the Möbius seam supplies precisely the two things the discrete loop cannot: the continuous π and the sheet parity $\mathbb{Z}_2$.

Why “Möbius” is not incidental: it is the boundary Lorentz structure. The Möbius group of the Riemann sphere is *exactly* the proper orthochronous Lorentz group: $\text{Möb}(\hat{\mathbb{C}}) \cong \text{PSL}(2, \mathbb{C}) \cong \text{SO}^+(3, 1)$, with the sphere realised as the celestial sphere of null directions via stereographic projection (Penrose & Rindler, *Spinors and Space-time*, Vol. 1, CUP 1984, §1.2–1.4). This is the spacetime meaning of the self-consistency loop: the seam’s boundary-2-sphere carries the Lorentz symmetry, the carrier clock $z \mapsto iz$ is an elliptic $\text{PSL}(2, \mathbb{C})$ rotation of that sphere (the order-4 Möbius map, [verification/v180_clock_is_mobius.py]), and the $\text{RP} \rightarrow \text{OS}$ -reconstructed causal cone (the shared $v_\gamma = v_g = c$ of `tfpt_horizon_readouts`) is the same Lorentz structure read off the seam. So “Möbius origin of π ” and “one Lorentzian cone” are two faces of one boundary $\text{PSL}(2, \mathbb{C})$ — not a coincidence of names. [O] (π stays primitive; the identification is a structural reading).

Net: two axioms $\rightarrow$ one primitive + one fixed point.

$$\{c_3, g_{\text{car}}\} \longrightarrow \underbrace{\pi \text{ (Möbius/Gauss–Bonnet)}}_{\text{one continuous primitive}} + \underbrace{E_8 \text{ closure}}_{\text{one discrete fixed point}}$$

The discrete data $\{2, 3, 4, 5, 6, 8, 30\}$ are *all* self-consistent outputs of the E_8 closure ($g_{\text{car}} = 5$, $h(D_5) = 8$, $h(A_3) = 4 = |\mu_4|$, $h(E_8) = 30 = 2 \cdot 3 \cdot 5$, $|R^+(A_3)| = 6$). The only irreducible input on the *dimensionless* axis is π . [E] (discrete loop) / [O] (π is primitive)

On the *dimensionful* axis there is exactly *one* more irreducible: a single mass scale. But even that is no longer a *free* constant — fixing the physical Λ branch pins $\bar{M}_{\text{Pl}}$, so G_N is a Λ -metrology *output* (`origin_theory`, Λ -typing); the one genuinely irreducible scale is the induced-gravity (Sakharov) anchor. So the complete honest reduction is “ π + the E_8 -closure fixed point + *one* dimensionful scale”.

One sharpening of the “8”. Both π -coefficients are 2π times a compiler integer: $\varphi_{\text{tree}} = 1/(6\pi)$ with $6 = 2N_{\text{fam}}$ (three Gauss–Bonnet boundary circles) and $c_3 = 1/(8\pi)$ with $8 = 2|\mu_4|$ (seam winding), so $\varphi_{\text{tree}} = (|\mu_4|/N_{\text{fam}})c_3 = \frac{4}{3}c_3$. The *same* integer 8 is the Hawking/Einstein coefficient ($G_{\mu\nu} = 8\pi T_{\mu\nu}$, Appendix H): the seam normaliser, the E_8 rank and the universal horizon factor are one number, overdetermined ($8 = 2|\mu_4| = \text{rank } E_8 = h(D_5) = \varphi(30) = \det R$). [E] [verification/v54_seam_horizon_keystones.py]

A cyclic element *inside* the hull: the order-30 Coxeter rotation [E] ([verification/v55_coxeter_cycle.py])

The compiler hull carries a *literal* cyclic symmetry. Built from the E_8 Cartan matrix, the Coxeter element $c = s_1 \cdots s_8$ has **order exactly** $h(E_8) = 30 = |\mathbb{Z}_2| \cdot N_{\text{fam}} \cdot g_{\text{car}} = 2 \cdot 3 \cdot 5$ (verified by direct matrix powers), and its eigenvalues are the primitive 30th roots $e^{2\pi i m/30}$ with m the E_8 exponents $\{1, 7, 11, 13, 17, 19, 23, 29\}$ — which are precisely the $\varphi(30) = 8$ *totatives* of 30. Hence

$$\text{rank } E_8 = 8 = \varphi(30) = \#\{\text{exponents}\} = \#\{\text{live phases of the order-30 cycle}\}$$

So the 8 in $c_3 = \frac{1}{8\pi}$ is the count of live phases of a primitive order-30 cyclic rotation of the hull, whose period $30 = |\mathbb{Z}_2| \cdot N_{\text{fam}} \cdot g_{\text{car}}$ is the product of the three core integers. The cyclic structure is *present in the mathematics*, not interposed.

The compiler atoms are the E_8 Casimir degrees [E] ([verification/v66_e8_casimir_degrees.py])

The load-bearing integers are not an ad-hoc list: they are the fundamental *invariant degrees* of E_8 (exponents +1),

$$\deg(E_8) = \{2, 8, 12, 14, 18, 20, 24, 30\},$$

with the clean (non-trivial) readings $2=|\mathbb{Z}_2|$, $8=\text{rank } E_8$, $14=\dim G_2$ (the $[K, Q]$ commutator G_2 , flavor sector), $20=\det L$, $24=|W(A_3)|=4!$, $30=h(E_8)$ (and the fittable $12=|R(A_3)|$, $18=p_4(a)=2+2^4$). Two exact sums close it:

$$\begin{aligned} \sum \deg(E_8) &= 128 = 2^7 \quad (\text{the de Sitter prefactor } 128c_3^4; \ 7=\text{scalaron exp}), \\ \sum \exp(E_8) &= 120 = |R^+(E_8)|. \end{aligned}$$

Since E_8 is the audit hull, its invariant degrees being the theory's integers is *structurally expected* — an organizing simplification that confirms the E_8 choice rather than a coincidence.

Two-level foundation, and the parameter audit

Local foundation: c_3, g_{car} are the two operational inputs (the two axioms). **Global bootstrap foundation:** g_{car} and the integer in c_3 are E_8 -closure fixed points; only π is primitive. The audit:

former input	status now	recovered by
$g_{\text{car}} = 5$	discrete fixed point	rank-fill, Coxeter-match, integer-glue/norm, $h(E_8)$
8 in c_3	discrete fixed point	$h(D_5)$, rank E_8 , $\det R$, $\varphi(30)$, $2 \mu_4 $ (seam winding)
6 in φ_{tree}	discrete fixed point	$ R^+(A_3) = h(E_8)/g_{\text{car}}$
π	geometric primitive	Möbius / Gauss–Bonnet (<i>not</i> derived)
φ_0^{ret}	derived readout	$\frac{4}{3}c_3 + 48c_3^4$

Correct one-liner: *TFPT has no freely tunable fundamental numbers left; the only remaining continuous origin is π from the Möbius boundary geometry.* (Not “zero axioms”.) [E]/[O]

17 Honest status: bootstrap, not creation from nothing

- **What is genuinely shown [E]:** $g_{\text{car}} = 5$ is an *overdetermined* fixed point of the E_8 closure (two independent Lie-theory conditions); the integer “8” in c_3 is the carrier Coxeter = E_8 rank; the discrete data closes a loop. The earlier landscape scan confirms $g_{\text{car}} = 5$ is the *unique* value at which all hierarchies/family/gauge close — so the fixed point is unique, not one of many. **This uniqueness is now machine-formalised [E] in Lean 4** (Theorem A, `lean4-carrier-rigidity`): the division-free Pascal condition $2^g = g^2 + g + 2$ has the *unique* solution $g = 5$ (`carrier_rank_pascal_unique`; growth lemma for $g \geq 6$; $N_{\text{fam}} = (2^4 - 1)/5 = 3$, $g + N_{\text{fam}} = 8$), audited to use only the standard kernel axioms. The 2026-06-11 hardening round adds two formal modules: `AnchorLadder.lean` (the anchor power-sum law $p_n(a) = 2 + 2^n$ for

all n , the ladder atoms $(3, 4, 6, 10)$, $240/8/248$, and the binary ladder $p_{n+1}-p_n = 2^n$ as general theorems) and `PascalLadder.lean` (the odd-midpoint ladder identity $\sum_{k \leq K} \binom{2K+1}{k} = 2^{2K}$ for all K via Mathlib, plus the kernel-decidable bounded uniqueness of the $K=2$ carrier case) — the solves-side of the Pascal Ladder Theorem (`[verification/v108_pascal_ladder.py]`) is thereby [E].

- **What this is *not*:** a derivation from zero inputs. A self-consistency loop is consistent *at* a value; uniqueness is what makes it predictive (and that is supplied by the landscape scan). The continuous π is not produced by the loop — it is the Möbius/boundary geometry, the one irreducible primitive.
- **Conceptual gain:** the theory’s foundation tightens from “two free axioms” to “*the* E_8 -closure fixed point + the geometric π .” That is the Möbius self-consistency you intuited: the seed is the start of the cascade, and the cascade returns the seed’s discrete content.

Part IV

Open items and the E_8 -slice scan

18 The honest open-items list

This is the complete current ledger of what is *not* fully closed, graded by how hard the gap is.

#	Open item	Grade	Hardening path / what closes it
<i>1. SM core — finite computations, no new physics</i>			
1a	absolute quark c -scale (ratios closed)	[O]	quark <i>ratios</i> are integer Plücker readouts on the derived selector stratum ($\sqrt{49}/\sqrt{69}/\sqrt{71}$); only the <i>absolute</i> U_f^* normalisation remains — an anchor
1b	θ_{12} solar angle	[E]	texture closed (below): explicit $\mu\tau$ -symmetric texture verified; only $\varepsilon = \frac{3}{4}\varphi_0^{\text{ret}} \approx c_3$ (0.23%) stays conditional
<i>2. Scheme layer — standard RG, no new input</i>			
2a	$m_W, m_Z, m_H, \sin^2 \theta_W, \alpha_s$	[C]	RG threshold matching R_{SM} ; m_H on the closed UV quartic $1/16\pi^2$
2b	hadron pole spectrum (singlet)	[C]	singlet algebra $\mathcal{A}_{\text{sing}}$ explicit; dynamics already closed (gap $6 \log \frac{3}{2}$)
<i>3. Conditional — named analytic hypotheses</i>			
3a	gravity (Hodge/ R^2 branch)	[C]	APS-Fredholm + truncation written out; full spectral UV completion remains
3b	strong CP / QFT closure	[C]	reflection positivity (RP) is the one remaining input
3c	Λ branch label (rec vs car)	[C]	one discrete choice; typed via $(8\pi)^2$, value not fixed
<i>4. Axioms — not eliminable, only hardenable</i>			
4	P1 (c_3), P2 ($g_{\text{car}}=5$)	[O]	more Lean; P2 algebra already machine-verified; remain inputs
<i>5. Frontier — honestly not on a clean ladder (see frontier doc)</i>			
5	$\eta_B, m_p/m_e$, Koide, dark matter, full QG	[O]	genuine handles only; not forced (scan results below)
<i>6. Cosmology late-time pipeline — not simplified by the compiler</i>			
6	C_ℓ transfer, $a_{\ell m}$ seam, baryogenesis	[C]	Paper-7 machinery; $a_{\ell m}$ “good CMB world, not yet this CMB world” (SMICA test)

Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in *verification/status_ledger.csv*.

Net. Items 1 are finite arithmetic; 2 are standard RG; 3 hinge on named hypotheses (RP, the Λ branch, the UV completion); 4 are the two irreducible axioms; 5–6 are honestly open. The inflation sector (the only cosmology piece the compiler *does* sharpen) is now a parameter-free prediction ($M=c_3^{7/2}\bar{M}_{\text{P1}}$, document 2).

Item 1b closed: the explicit solar-angle Majorana texture [E]

The last open SM angle is now realised by an *explicit* texture, not a fitted number. Take the $\mu\tau$ -symmetric Majorana matrix (tri-bimaximal at $\varepsilon=0$)

$$M_\nu(\varepsilon) = \begin{pmatrix} -\eta(\varepsilon) & 1 & 1 \\ 1 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix}, \quad \eta(\varepsilon) = \frac{2\sqrt{2}}{\tan 2\theta_{12}} - 1, \quad \sin^2 \theta_{12} = \frac{1}{3}(1 - 2\varepsilon). \quad (1)$$

$\mu\tau$ -symmetry forces $\theta_{23} = 45^\circ$ and $\theta_{13} = 0$ *exactly*; diagonalisation at the TFPT solar deviation $\varepsilon = \frac{3}{4}\varphi_0^{\text{ret}}$ gives

$$\sin^2 \theta_{12} = \frac{1}{3} - \frac{\varphi_0^{\text{ret}}}{2} = 0.306747, \quad \theta_{23} = 45^\circ, \quad \theta_{13} = 0 \quad [\text{E}] \quad (2)$$

Honest residual ([`verification/v9_neutrino_texture.py`]). The texture and the angle are exact. The seam identification is $\varepsilon = \frac{3}{4}\varphi_0^{\text{ret}} = \frac{3}{4}\left(\frac{1}{6\pi} + \frac{3}{256\pi^4}\right) = c_3 + \frac{9}{1024\pi^4}$: the *leading geometric term is exactly* $c_3 = \frac{1}{8\pi}$, and the 0.23% residual is purely the φ_0^{ret} seed tail (document 2). So $\varepsilon = c_3$ holds at leading order; the full $\varepsilon = \frac{3}{4}\varphi_0^{\text{ret}}$ is what enters the texture. The reactor angle θ_{13} ($\sin^2 \theta_{13} = e^{-5/6}\varphi_0^{\text{ret}}$, document 2) enters through a separate $\mu\tau$ -breaking term and is not part of this minimal solar texture. **Scope:** this realises the TFPT solar shift in the $\mu\tau$ -symmetric *limit*; it is *not* the full PMNS matrix. The $\theta_{23}=45^\circ$ value is the symmetric-limit centre, *not* an octant prediction — current global fits (NuFIT 6.0) leave the octant unresolved — and the reactor angle plus the combined matrix require the separately stated $\mu\tau$ -breaking channel.

19 The E_8 -slice scan: results (honest hit / miss)

All seven maximal slices of 248 were scanned for the unassigned quantities. The rule was a meaningful (low-complexity) compiler/slice expression matching to $< 0.5\%$; matches that need a free fitted factor are reported as *misses*.

Slice / target	Result	Verdict
$A_4 \times A_4 = SU(5)^2$	$\sin^2 \theta_W _{\text{GUT}} = 3/8$ sits in the $SU(5)$ hypercharge embedding; the 12 coset states are the proton-decay X, Y bosons	[E] (known)
$F_4 \times G_2$ / Koide	Koide $Q = 2/3 = \mathbb{Z}_2 /N_{\text{fam}}$ (sheet over families); $J_3(\mathbb{O})$ has $\dim 27=26+1$, 3×3 octonionic “generations” on the diagonal	[C] structural
$D_8 = SO(16)$	$128 = (16, 4) \oplus (\overline{16}, \overline{4})$ — <i>same</i> spinor as $D_5 \times A_3$; ν_R (seesaw) lives in the 16; no new dark-matter state (only the 4th μ_4 direction)	[E]
$A_8 = SU(9)$, $9=3 \times 3$	$\Lambda^3(3, 3) = (10, 1) \oplus (8, 8) \oplus (1, 10)$: the $(8, 8)=64$ dominates — naive family $\times$ colour fails	miss
$A_8 \supset SU(5) \times SU(4)$	$\Lambda^3(5+4)$ contains the SM 10 : a genuine chiral-family-unification <i>lead</i> , but not TFPT-forced	lead
$E_6 \times A_2$ / Jarlskog	cross term $2(\mathbf{1}^\top Ra)N_{\text{fam}} = 162 = 248 - 78 - 8$ exact; $J \sim \lambda_Y^6$ only to right order, no clean hit	weak
E_8 theta shells	root shells 240, 2160, 6720, $\dots$; $31=2^{g_{\text{car}}}-1$ is <i>not</i> a clean shell ratio	inconcl.
$m_p/m_e = 1836.15$	no clean E_8 /compiler expression (a QCD-confinement ratio); apparent hits need a free factor	miss
$\eta_B \approx 6.1 \times 10^{-10}$	sits <i>between</i> $c_3^6=4 \times 10^{-9}$ and $c_3^7=1.6 \times 10^{-10}$; right order, no integer power	order only

Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable
kill test. Per-claim fine type in `verification/status_ledger.csv`.

Genuine new readings the scan produced

- **Koide** = $|\mathbb{Z}_2|/N_{\text{fam}}$. The empirical Koide value $2/3$ is exactly the sheet over the family count — a clean (if simple) compiler reading, promoting Koide from “near $2/3$ ” to “democratic $|\mathbb{Z}_2|/N_{\text{fam}}$ ”
- **E_8 as the optimal 8-bit code.** E_8 is the optimal 8D lattice (kissing number 240). In byte framing $256-248 = 8 = \text{rank } E_8$ and $256-240 = 16 = \dim S^+$ — a clean information-theoretic reading of the two load-bearing E_8 numbers.
- **Dark matter = no new state.** The $SO(16)$ scan shows the 128 is the *same* matter spinor; there is no spare E_8 singlet to be a WIMP. Any DM candidate must be the determinant-line axion (frontier doc), not a new E_8 rep — a useful *negative*.
- **$SU(9) \supset SU(5) \times SU(4)$ chiral lead.** Λ^3 puts the SM **10** in the antisymmetric; the cleanest unexploited route to a family-unification reading, flagged for future work.

Honest negatives

m_p/m_e has *no* clean E_8 /compiler derivation (it is a confinement ratio — the apparent “hits” all need a free fitted factor); η_B matches only at the order-of-magnitude level ($\sim c_3^{6.x}$); the Λ rec/car branch is *not* selected by the theta series; the Jarlskog invariant is right only to order λ_Y^6 . These are reported as open, not forced.

20 Following the $SU(9)$ lead: cross-slice consistency

The one scan lead worth following is $A_8 = SU(9) \supset SU(5) \times SU(4)$. Splitting $9 = (5, 1) \oplus (1, 4)$ the E_8 branching $248 = 80 \oplus 84 \oplus \overline{84}$ decomposes *exactly* as:

piece	$SU(5) \times SU(4)$ content	reading
80 (adj)	$(24, 1) + (1, 15) + (5, \overline{4}) + (\overline{5}, 4) + (1, 1)$	$SU(5)_{\text{GUT}} + SU(4)_{\text{fam}}$ gauge
$84 = \Lambda^3 9$	$(10, 1) + (10, 4) + (5, 6) + (1, 4)$	the SM 10 as family-singlet + family- 4
$\overline{84}$	$(\overline{10}, 1) + (\overline{10}, \overline{4}) + (\overline{5}, 6) + (1, \overline{4})$	the conjugate matter

(Dimension checks: $24+15+20+20+1 = 80$; $10+40+30+4 = 84$; total 248.)

The genuine content: A_8 and $D_5 \times A_3$ agree on the family [E]

The $SU(4)$ here is *the same* A_3 family group as in the construction slice $D_5 \times A_3$ — the four punctures μ_4 of $\mathbb{P}^1 \setminus \mu_4$. The $SU(5)$ is the Georgi–Glashow GUT sitting inside the carrier $SO(10) = D_5$ ($SU(5) \subset SO(10)$). So the two maximal slices are *consistent*: both read the family as $SU(4)=A_3$, and A_8 additionally exposes the $SU(5)$ matter content — the SM **10** appears as $(10, 1) \oplus (10, 4)$, i.e. a family-singlet **10** plus a family-**4**, where the **4** of A_3 is exactly TFPT’s “ $4 = N_{\text{fam}} + 1$ ” (three families + the μ_4 direction).

Honest limit: chirality is still the index, not E_8

$84 \oplus \overline{84}$ is *vectorlike* ($10 \oplus \overline{10}$ etc.): E_8 being non-chiral, the slice cannot by itself produce three *chiral* generations. The chirality is the boundary index $\text{Ind}(D_{\text{fam}}^+) = 3$ — the *same* input TFPT already uses. So $SU(9)$ is a consistent re-reading (matter in $SU(5)$, family in A_3), not an independent derivation of $N_{\text{fam}} = 3$. Each $SU(5)$ pair (**10** + **5**) is anomaly-free, so the embedding is clean. [C]

Appendix: computational verification

The audit/bridge identities of this document (marked [E], [E]) are re-derived from $\{c_3, g_{\text{car}}\}$ and machine-checked by the suite in the repository folder `verification/`; the inline tags `[verification/...]` point to the exact script. Relevant here:

- `v1_e8_glue.py` — the μ_4 glue $E_8 = (D_5 \oplus A_3) + \mu_4$ (load-bearing).
- `v4_flavor_matrix.py` — the spectral selector $\det R = h(D_5) = 8$, minors $\{2, 3, 5\}$.
- `v5_e8_cascade.py` — the cascade $D_n = 60 - 2n$: endpoints, exponent rungs, IR tail, variance.
- `v6_bootstrap.py` — the self-consistency loop: $\mu^2 - 5\mu + 4 = 0$, $g_{\text{car}} = 5$ three ways, $8 = \text{rank } E_8$.

The E_8 -slice/atlas readings are explicitly [O] (audit-level), not part of the pass/fail suite. Run `python verification/run_all.py`; full map in `verification/README.md`.