

TFPT — Red Team / Stress Test layer

Five adversarial targets (A–E): try to *break* the reductions, not confirm them

Stefan Hamann

Alessandro Rizzo

June 14, 2026 · v5.1 (rev 109)

The TFPT document set — what is treated where

Plain language: TFPT (Topological Fixed-Point Theory) is a small discrete compiler. Two inputs — the seam constant $c_3 = \frac{1}{8\pi}$ and the carrier rank $g_{\text{car}} = 5$ — build $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ and read off the Standard-Model skeleton, the constants and the scale grammar. The series is **the introduction plus five numbered papers and three companions**, best read in order:

0. **introduction** — reading guide, compiler closure, predictions, the dependency DAG and the single proof ledger (`verification/status_ledger.csv`). (*entry point*)
1. **tfpt_1_architecture_e8** — the two axioms, the derivation map, the EM fixed point α^{-1} , the $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ construction, the QBL theorem chain.
2. **tfpt_2_standard_model** — the SM in one φ_0 -ladder formula, flavor from parabolic transport (sheet diamond, dual anchor, Q_+ cohomology), the worked closures.
3. **tfpt_3_e8_audit_bootstrap** — the seven E_8 slices as an audit raster, the cascade bridge, the Möbius bootstrap.
4. **tfpt_4_frontier** — honest status of η_B , m_p/m_e , Koide, dark matter, full QG.
5. **tfpt_5_redteam** — the adversarial audit: declared attacks, what survives, what each reduces to, and the kill tests.

Companions: **R.** **tfpt_research_contracts** — the open interfaces (v_{geo} , G_{net} , F_{transfer}) as numbered contracts; **H.** **tfpt_horizon_readouts** — Appendix H (unit-system reframe): one seam constant as the universal horizon thermal code, the Nariai/anchor surface, the resummed seam clock; **O.** **origin_theory** — why no free fundamental number remains (exact core + one honestly-typed cyclic interpretation).

You are reading document 5 (Red Team).


TFPT status at a glance — read this card the same way in every document

Two inputs, one compiler. From the boundary constant $c_3 = \frac{1}{8\pi}$ (axiom P1) and the five-slot carrier $g_{\text{car}} = 5$ (axiom P2) — jointly the elementary symmetric data of one anchor $a = (1, 1, 2)$ — the discrete compiler $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ is *built*, and the Standard-Model skeleton (gauge group, three families, hypercharges, the flavor matrix), the dimensionless constants ($\alpha^{-1} = 137.0359992$) and the scale grammar are read off as fixed points. The algebraic core is machine-checked; physical readouts run through *named bridges* whose status is typed below.

What is closed, and what genuinely remains. The discrete/algebraic layer is closed ([E]); the everything-open residual reduces to *three named interface problems*:

Interface	Question	Status
v_{geo}	the one metrology unit ($= 1/\sqrt{G} = m/\mu$); No-Unit Thm: no compiler scale	primitive [O]
G_{net}	simple-current extension $(D_5)_1 \otimes (A_3)_1 \rtimes \langle (1, 1) \rangle \cong (E_8)_1$	[E]/[C]
F_{transfer}	one functor, four typed interfaces (Koide, η_B , axion, m_p/m_e)	[C]

Everything carries a claim ID resolving to a row of `verification/status_ledger.csv` (the single source of truth); **if the text and the ledger ever disagree, the ledger wins**. The historical labels (U_{wall} , G_{metric} , F_{frontier}) are retained only for ledger continuity — the live residual is $v_{\text{geo}} \oplus G_{\text{net}} \oplus F_{\text{transfer}}$.
Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in `verification/status_ledger.csv`.

[E] exact / machine-proven**[C]** conditional**[O]** open / axiom**[X]** kill test**What this note is**

This is the deliberately *adversarial* layer requested in the **v78** review. Its purpose is **not** to confirm TFPT but to attack the five load-bearing reductions at their weakest logical transitions. Each target is run through one fixed protocol — minimal statement, assumptions, logical chain, validity conditions, counterexample search, limiting/degenerate cases, alternative structures, provisional verdict — and produces a deliberately narrow output: where the statement works, where it could fail, and the residual risk. The layer is machine-backed by `[verification/redteam/rt_A_e8net.py]` ... `[verification/redteam/rt_E_vgeo.py]` and `[verification/redteam/run_redteam.py]`; a red-team check asserts an *adversarial* fact (a counterexample really exists, a hidden assumption is really needed, a firewall really holds), and the honest outcome lives in the **status** of each target, never in a green pass.

Target	Minimal claim	Status	Residual risk
A	seam–Calderón measure = $(E8)_1$ net	reduced, not closed	<i>one</i> statement: boundary-net holomorphy + $c=8$ ($\Leftrightarrow$ the index-4 inclusion); E_8 and bulk uniqueness then automatic (v83/v87/v89 ; Lie-level realisation v143). Since reduced to <i>zero new</i> open content: the free-bulk premise is a fixed-point theorem, the cone is eliminated (cone gap = one-particle gap), and the residual factors into the already-open A_2 (net assembly) + GATE.QGEO (v160–v165). The three-residual form below is the historical development.
B	$g_{\text{car}} = 5$ forced by the Pascal rule	survives (narrowed)	boundary derivation of half-spinor exhaustion (degree-2 truncation)
C	$k = c_3/2 = \frac{1}{16\pi}$, $S/A = \frac{1}{4}$	survives	absolute $1/G$ (UV-sensitive) is the one anchor
D	ratios+products $\Rightarrow$ one scale v_{geo}	survives (narrowed)	CP phases ($\delta_{\text{CKM/PMNS}}$) not covered by v_{geo}
E	one scale v_{geo} carries the theory	survives (narrowed)	EW/reheating/leptogenesis scales + seam=Planck identification

Method

The red team treats each reduction as a hostile witness. A confirmatory script that always passes is worthless here: the layer is built so that it *may* downgrade a claim. Three outcomes are allowed — **survives** (stands as worded), **survives (narrowed)** (stands only after a silent assumption is made explicit), and **reduced, not closed** (the conservative wording is the correct one). A fourth, **broken**, is reserved for an actual failure; none occurred.

Target A — the $(E8)_1$ boundary-net identification

Minimal claim. The ambient metric/projective measure is reduced to identifying the seam–Calderón boundary measure with the $(E8)_1$ lattice net, carrier subnet $(D5)_1 \times (A3)_1$. [verification/redteam/rt_A_e8net.py]

Logical chain (established). The Sugawara level-1 central charges are $c(E8)_1 = \frac{248}{31} = 8$, $c(D5)_1 = 5 = g_{\text{car}}$, $c(A3)_1 = 3 = N_{\text{fam}}$, so the conformal-embedding criterion $c_{\text{coset}} = c(E8) - c(D5) - c(A3) = 0$ holds [E]. This is *compatibility*.

Where it breaks: central charge underdetermines the net

Counterexample: $(D8)_1 = SO(16)_1$ also has $c = \frac{120}{15} = 8$, but is *not* holomorphic (four primaries $1, v, s, c$) — a distinct $c = 8$ net. Hence equal central charge is *necessary, not sufficient*: $(E8)_1$ is the unique *holomorphic* $c = 8$ net (even unimodular rank-8 lattice = $E8$), so **holomorphy (single vacuum sector, μ -index 1) is the load-bearing extra assumption**. Dropping chirality is worse: the $c = 8$ Narain family is positive-dimensional. The constructive map seam–Calderón kernel $\rightarrow (E8)_1$ and the uniqueness of bulk reconstruction are not established; the gap $\Delta_{\text{eff}} > 0$ *supports* tightness (clustering) but does not prove it.

Verdict. Keep “reduced to a rigorous boundary-net identification problem”; do *not* write “metric sector closed”. The red team confirms the conservative wording. Status [C]. Residual: holomorphy proof + constructive boundary map + bulk-reconstruction uniqueness.

Target B — carrier rank / the Pascal condition

Minimal formal claim. $2^g = g^2 + g + 2$ has the unique solution $g = 5$ (Lean-verified [E]). This arithmetic is *not* attacked. The attack is upstream: why must the carrier obey the Pascal half-spinor exhaustion $2^{g-1} = \binom{g}{0} + \binom{g}{1} + \binom{g}{2}$? [verification/redteam/rt_B_pascal.py]

Where it narrows: the rule, not the arithmetic, selects 5

Perturbing the constant breaks it: $2^g = g^2 + g + c$ gives $\{5\}$ only at $c = 2$; neighbours give other or empty solution sets. The truncation *degree* is the real choice: with $2^{g-1} = \sum_{k \leq K} \binom{g}{k}$ one finds $K=0 \rightarrow g=1$, $K=1 \rightarrow 3$, $K=2 \rightarrow 5$, $K=3 \rightarrow 7$, i.e. $g = 2K + 1$. So choosing $K = 2$ (to reach $g_{\text{car}} = 5$) is a physical postulate. It is corroborated — the $D5$ half-spinor $\dim S^+ = 2^{g-1} = 16 \Leftrightarrow g = 5$ and the family closure $N_{\text{fam}} = (2^{g-1} - 1)/g$, $g + N_{\text{fam}} = 8$ both pick 5 — but each is itself a postulate, not a theorem.

Verdict. Theorem A’s arithmetic core stands; the Pascal *selection* must be typed [O]/[C], not [E]. Residual: a boundary/seam derivation of half-spinor exhaustion.

Target C — $k = c_3/2$ and the seam-area coefficient

Minimal claim. In reduced units $k = c_3/2 = \frac{1}{16\pi}$ and Fursaev–Solodukhin gives $S/A = \frac{1}{4}$. [verification/redteam/rt_C_kc3.py]

The dimensional firewall holds

The only legal chain is

$$k_{\text{red}} = \frac{c_3}{2} = \frac{1}{16\pi} \text{ (dimensionless)}, \quad k_{\text{phys}} = \frac{k_{\text{red}}}{G}, \quad S = 4\pi k_{\text{phys}} A_{\text{phys}} = \frac{A_{\text{phys}}}{4G}.$$

The forbidden form $k_{\text{phys}} = c_3/2$ is dimensionally false (legal only at $G = 1$). A scan of **v67/v68/v73** finds every $c_3/2$ used as the *reduced* coefficient and **zero** naked dimensionful identities; the firewall's teeth are verified against a synthetic violation. Internally consistent with $1/G$ being UV-sensitive (Sakharov/Connes, **v68**).

Verdict. Safe as worded (reduced coefficient). Status **[E]** (structure) + **[O]** for the residual. Residual: the absolute $4D$ Newton scale $1/G$ ($\propto \Lambda^2 f_2$) stays the one dimensionful anchor; nothing *dimensionless* is open.

Target D — $U_{\text{point}} \rightarrow v_{\text{geo}}$

Minimal claim. Ratios plus sector products determine the dimensionless amplitudes, leaving one common scale v_{geo} . [[verification/redteam/rt_D_upoint.py](#)]

Where it narrows: the bijection needs explicit hypotheses

For positive reals, (ratios, product) $\Leftrightarrow$ (individuals) is a bijection, and the TFPT sectors satisfy it (positive rational c , odd size $3 = N_{\text{fam}}$). But it *fails* off-assumption: even sector size leaves a sign ambiguity ($\{2, 3\}$ vs $\{-2, -3\}$ share ratio $3/2$ and product 6); complex amplitudes leave an n -th-root branch ambiguity ($\{1, 1, 1\}, \{\omega, \omega, \omega\}, \{\omega^2, \omega^2, \omega^2\}$ share all ratios and product 1); a zero amplitude is degenerate. **Genuine residual:** the bijection fixes amplitude *magnitudes* only — CP phases $\delta_{\text{CKM}}, \delta_{\text{PMNS}}$ are complex and are *not* folded into v_{geo} .

Verdict. Holds for TFPT once four hypotheses (real, positive, fixed branch, no zeros / odd sector size) are made explicit; they are silent in **v75** and must be named. Status **[E]** (reduction) + **[O]** (v_{geo}). Residual: the CP-phase sector.

Target E — v_{geo} as the unique dimensional floor

Minimal claim. No pure-number theory can derive an absolute dimensionful scale; all scales are ratios to one unit v_{geo} . [[verification/redteam/rt_E_vgeo.py](#)]

Single-scale audit: exact for the certified tiers, conditional beyond

The closed compiler + protected IR readouts all reduce to (pure number) $\times v_{\text{geo}}$: $M_{\text{scal}} = c_3^{7/2} \bar{M}_{\text{Pl}}$, $f_a = (c_3/|\mu_4|)^{7/2} \bar{M}_{\text{Pl}}$, masses $= (\pi/\sqrt{2}) c_f \varphi_0^k v_{\text{geo}}$, and $1/G$ shares the same anchor (**v68/v75**). The audit *surfaces* candidates that are not pure reductions: the seam cutoff Λ (only an *identification* seam=Planck collapses it to v_{geo}), the electroweak matching scale, the reheating temperature, and the leptogenesis scale — all frontier **[C]/[O]** inputs that must not be silently absorbed into v_{geo} .

Verdict. v_{geo} is the unique dimensional floor of the certified tiers (true by dimensional analysis). Full single-scale-ness is conditional on the flagged identifications staying honestly typed. Status **[E]** (ratios) + **[O]** (one scale). Residual: explicit reduction (or honest frontier typing) of the EW / reheating / leptogenesis scales + the seam=Planck cutoff.

Outcome

What the red team changed. No target broke. The hardest gate (A) is confirmed *open* with the conservative wording; three targets (B, D, E) *survive once a silent assumption is named*; one (C) survives as worded. The net effect is not new physics but sharper typing: the package now states

where each reduction would fail and *which* assumptions are truly necessary, exactly at the weakest logical transitions. Reproduce with `cd verification/redteam && python run_redteam.py`.

Follow-up — the verdicts were then used as a worklist [E]
 ([verification/v83_e8net_holomorphic_uniqueness.py])

Two residuals were attacked directly after the red-team pass; the result is recorded in ledger rows GATE.METRIC.04 and CAR.PASCAL.01.

Target A, residual 1 — closed [E]. At level 1 the number of primaries equals $\det(\text{Cartan}) = |\text{fundamental group}|$: $A_3=4$, $D_5=4$, $D_8=4$, $\mathbf{E}_8=1$. So holomorphy (single vacuum sector, μ -index 1) is *necessary* — it excludes the same- c rival $(D_8)_1 = SO(16)_1$ ($c = 120/15 = 8$ but four primaries $1, v, s, c$) — and *sufficient*: a holomorphic $c = 8$ chiral CFT is the lattice theory of an even unimodular rank-8 lattice, of which there is exactly one, E_8 , by the Minkowski–Siegel mass formula ($\text{mass} = 1/|W(E_8)| = 1/696729600$, and $|\text{Aut}(E_8)|$ saturates it). The $(D_5)_1 \times (A_3)_1$ embedding $c = 5 + 3 = 8$ stays *compatibility* (16 primaries vs 1). **Consequence:** the constructive map need only show the seam–Calderón boundary net is holomorphic with $c = 8$ (then E_8 is automatic), so Target A drops from *three* residuals to *two*: (i) boundary-net holomorphy + $c = 8$ ($\Delta_{\text{eff}} > 0$ supports, not proves), and (ii) bulk-reconstruction uniqueness — both still [C]/[O].

Target B — reduced [E]. The “ $K = 2$ Pascal truncation” is not free: $K = (g - 1)/2$ is the Pascal-row midpoint $\sum_{k \leq (g-1)/2} \binom{g}{k} = 2^{g-1}$ (odd g), i.e. the half-spinor split; the carrier is the even Clifford half-spinor $\Lambda^{\text{even}}(\mathbb{C}^5) = 1 + 10 + 5 = 16 = \dim S^+$. So B’s residual reduces to the single standard input “carrier = half-spinor of $\text{Spin}(10)$ ” (the Lean arithmetic core stays [E]).

Still open, typed honestly: A (i)+(ii); Target D’s CP-phase residual (at $\theta_{13} = 0$ the Dirac phase is undefined, so this needs a texture correction); Target E’s reheating/leptogenesis scales and the seam=Planck identification. Target C is a dimensional *floor*, not a gap. **No target was promoted.**

Second follow-up round [E] ([verification/v87_bulk_uniqueness_reduction.py], [verification/v88_cp_phase_audit.py], [verification/v86_nstar_reheating.py])

Target A, residual (ii) — conditionally closed [E]: Target A = ONE residual. Bulk-reconstruction uniqueness is *not* independent of residual (i). By the algebraic-QFT classification of full 2D CFTs over a chiral net (Longo–Rehren; Kawahigashi–Longo–Müger; Bischoff–Kawahigashi–Longo–Rehren), the possible bulks are the haploid commutative Q-systems $\simeq$ physical modular invariants of $\text{Rep}(A)$. For a *holomorphic* net $\text{Rep}(A) = \text{Vect}$, so the bulk pairing is *unique*. The contrast is machine-checked: the same- c rival $SO(16)_1$ (discriminant category $\mathbb{Z}_2 \times \mathbb{Z}_2$, μ -index 4) admits *six* modular invariants — diagonal, $s \leftrightarrow c$ swap, both E_8 -extension invariants $|\chi_0 + \chi_{s/c}|^2$ (the lattice glue $E_8 = D_8 \cup (D_8 + s)$, $h_s=1$ bosonic), and two twisted pairings — so without holomorphy the $c = 8$ bulk is *multiply* ambiguous; with it, trivially unique. Hence the entire Target A reduces to the single theorem “seam–Calderón boundary net is holomorphic with $c = 8$ ” ([C]/[O]; ledger GATE.METRIC.05). The theorem also has an equivalent, more tractable *index* form ([verification/v89_carrier_index_lemma.py], ledger GATE.METRIC.06): by Kawahigashi–Longo–Müger, $\mu_A = [B:A]^2 \mu_B$, the carrier inclusion has Jones index $[(E_8)_1 : (D_5)_1 \times (A_3)_1] = \sqrt{16} = 4 = |\mu_4|$ — the glue-group order *is* the inclusion index — and the extension is the μ_4 simple-current extension (all three glue sectors are $h = 1$ currents; $248 = 45 + 15 + 64 + 64 + 60$). Since an index-4 simple-current extension of the carrier has $\mu = 16/4^2 = 1$, **holomorphy follows from the index statement:** Target A $\Leftrightarrow$ “the seam net contains the carrier net with Jones index 4 as its μ_4 simple-current extension” — both ends constructed objects, the index controllable by Longo theory, the gap supporting finiteness. And *which* extension carries no freedom ([verification/v92_glue_uniqueness.py], ledger GATE.METRIC.07): exhaustive classification of the carrier discriminant form $(\mathbb{Z}_4 \times \mathbb{Z}_4, q = (5x^2 + 3y^2)/8)$ shows the extension tower is *completely rigid* — exactly two Lagrangian glues (the two chiralities, identified by the sheet $\mathbb{Z}_2$; the same two objects as the E_8 -extension invariants above), and exactly one halfway extension, whose induced form $q = \{0, 0, 0, \frac{1}{2}\}$ *is* the D_8 discriminant form: the $SO(16)_1$ rival is the unique intermediate of the tower, not an arbitrary counter-model. Tower: carrier ($\mu=16$) $\rightarrow SO(16)_1$ ($\mu=4$) $\rightarrow (E_8)_1$ ($\mu=1$, sheet pair) — nothing else exists. Target A is thereby the *bare* index statement.

Target A, final form — the Simple-Current Extension Theorem [E]/[C] ([verification/v154_simple_current_theorem.py]). Assembled into one statement: the carrier net $A = (D_5)_1 \otimes (A_3)_1$ extended by the isotropic glue $L = \langle (1, 1) \rangle$ has $[B : A] = |L| = 4$, $c = 5 + 3 = 8$, $\mu(B) = \mu(A)/|L|^2 = 16/16 = 1 \Rightarrow$ holomorphic $\Rightarrow B \cong (E_8)_1$. The *entire* adversarial residual of Target A is

now the single seam-side identification “the RP seam–Calderón net *is* A ” — the algebra $(B \cong (E_8)_1)$ is exact and machine-checked (v89/v92/v125/v143/v148). That residual is now sharply located: the analytic core (“the kernel is quasi-free”) *follows* from a free RP bulk — the boundary marginal of a Gaussian bulk is the compression of a contraction (v155), the *same* premise the area law and the EH mechanism use — and the net is *explicitly constructed and verified*: DtN = $|k|$, 16 free Majoranas, currents $248=120+128$, holomorphic character $E_4/\eta^8 = j^{1/3}$ with 248 at level 1 (v156). *Final reduction* (v160–v165): the free-bulk premise is recast as a fixed-point theorem (quasi-free $\Rightarrow$ all higher cumulants $\kappa_{2n}=0$, v160); the infinite Schwinger cone is *eliminated* — by Bogoliubov second quantisation the cone gap equals the one-particle gap $(2/3)^6$ (v161), whose value is forced by μ_4 -equivariance (v162); and the remaining 16-dim identification factors into the two *already-open* items — A_2 net existence (an assembly of established net constructions, [C]) and $R_1 = \text{GATE.QGEO}$ (“seam boundary = $\mathbb{P}^1 \setminus \mu_4$ ”) — with the irreducible core $\{\pi, v_{\text{geo}}\}$ a theorem (v165). So Target A carries *no new independent open content*: it is a unification, not an unconditional proof. *The two structural residuals are then discharged to a theorem* (v175): the CAR second-quantisation functor $\Gamma(t) = \bigoplus_m \Lambda^m(t)$ makes full-cone reflection positivity reduce *for every* m to the one-particle contraction (verified on the complete 2^{16} -dim Fock space: contraction, fixed multiplicity 2^8 , sub-leading = $(2/3)^6$), and A_2 is an assembled and verified $(E_8)_1$ certificate (character $1+248q+4124q^2+\dots$, embedding $248=120+128$, E_8 Cartan even unimodular). Both then need the *same* single input — the seam realisation QGEO.REALIZE.01 — which stays the one irreducible [O]; net existence and full-cone RP become [E].

Target D — the CP-phase residual quantified [E] (not closed). With the frozen leading reading $\delta = \pi/3 + 3\lambda_C^2 = 68.65^\circ$: pull vs $\gamma_{\text{PDG}} = 65.7^\circ \pm 3.0^\circ$ is $+0.98\sigma$ (survives). Audit, *not* promoted: the data central value sits 0.07° from the alternative coefficient reading $\pi/3 + 2\lambda_C^2$ ($|\mathbb{Z}_2|$ instead of N_{fam}) — a textbook look-elsewhere trap; the registry keeps coefficient 3 (REG.FREEZE.01). Decision threshold: the readings differ by $\lambda_C^2 = 2.88^\circ$, i.e. 3σ -distinguishable once $\sigma_\gamma \leq 0.96^\circ$. The J -inversion is flagged magnitude-contaminated (source-scale s_{23} vs M_Z). Ledger FLAV.CP.01.

Target E — one flagged scale computed [C]. The reheating temperature is no longer a silent input: $M_{\text{scal}} = c_3^{7/2} \bar{M}_{\text{Pl}}$ [E] + standard physics give $T_{\text{reh}} = 9.6 \times 10^9$ GeV and $N_*(0.05/\text{Mpc}) = 51.4$ [C] (ledger COSMO.NSTAR.01); honestly recorded: this slow Higgs-channel point is A_s -disfavoured (-11.4σ ; the measured A_s requires near-instantaneous reheating), so the point is conditional on the decay channel and the frozen band [50, 60] stays the surface of record; the leptogenesis scale and the seam=Planck identification remain typed inputs. **Again: no target was promoted; one residual was merged (A), one was quantified (D), one input was computed (E).**

Global look-elsewhere closure [E] ([verification/v100_numerology_null_mc.py])

The per-target look-elsewhere traps recorded above (the $\pi/3 + 2\lambda_C^2$ alternative in Target D, the $\ln(46080)$ scale reading destroyed in [verification/v99_koide_flow_time.py]) are local instances of the global adversarial question: *could the whole scorecard be formula-fishing?* The numerology null test answers it with an explicit, declared null model run in adversarial mode: the *entire* complexity-matched formula grammar (which provably contains every scored TFPT formula — the null space includes the theory) is enumerated against conservative data windows. Joint formula-fishing probability $\prod_i p_i = 10^{-25.8}$ over the 13 scored observables; all 94 500 $F_{U(1)}$ variants solved, exactly one hits CODATA ($p_\alpha \leq 1.1 \cdot 10^{-5}$); total $\leq 10^{-30.7}$. The test is falsifiable in both directions: the negative controls (seed perturbation 1%/10% collapses TFPT’s own hit count $13 \rightarrow 9/6$; data shuffling $\rightarrow 1.2$) prove it has power, and the census is stable under budget/window variation (< 2 orders). Honest limit, stated as such everywhere: this is a null-model rejection *conditional on the declared grammar* — no statistical test certifies a theory; the tension observables keep their large windows and contribute nothing to the number. Details and the grammar definition: the audit note (tfpt_3) and the module docstring.